

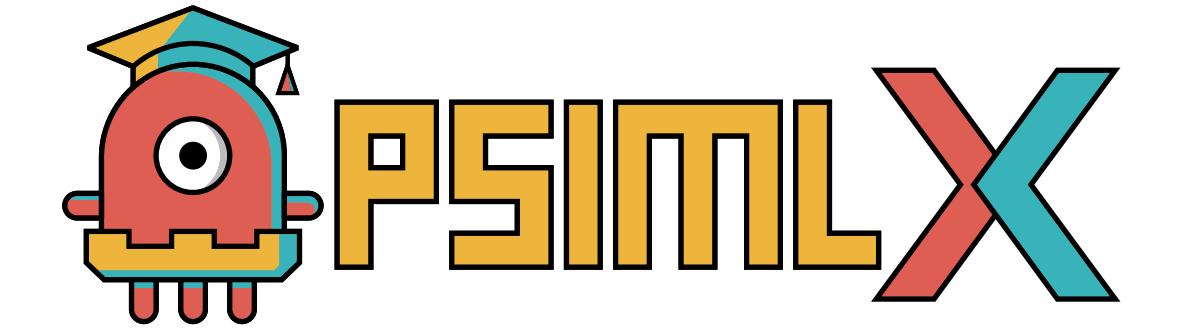
Reinforcement Learning

A gentle introduction

Đorđe Božić
PhD Student
University of Bath

 djordjebozic@gmail.com

Who am I?



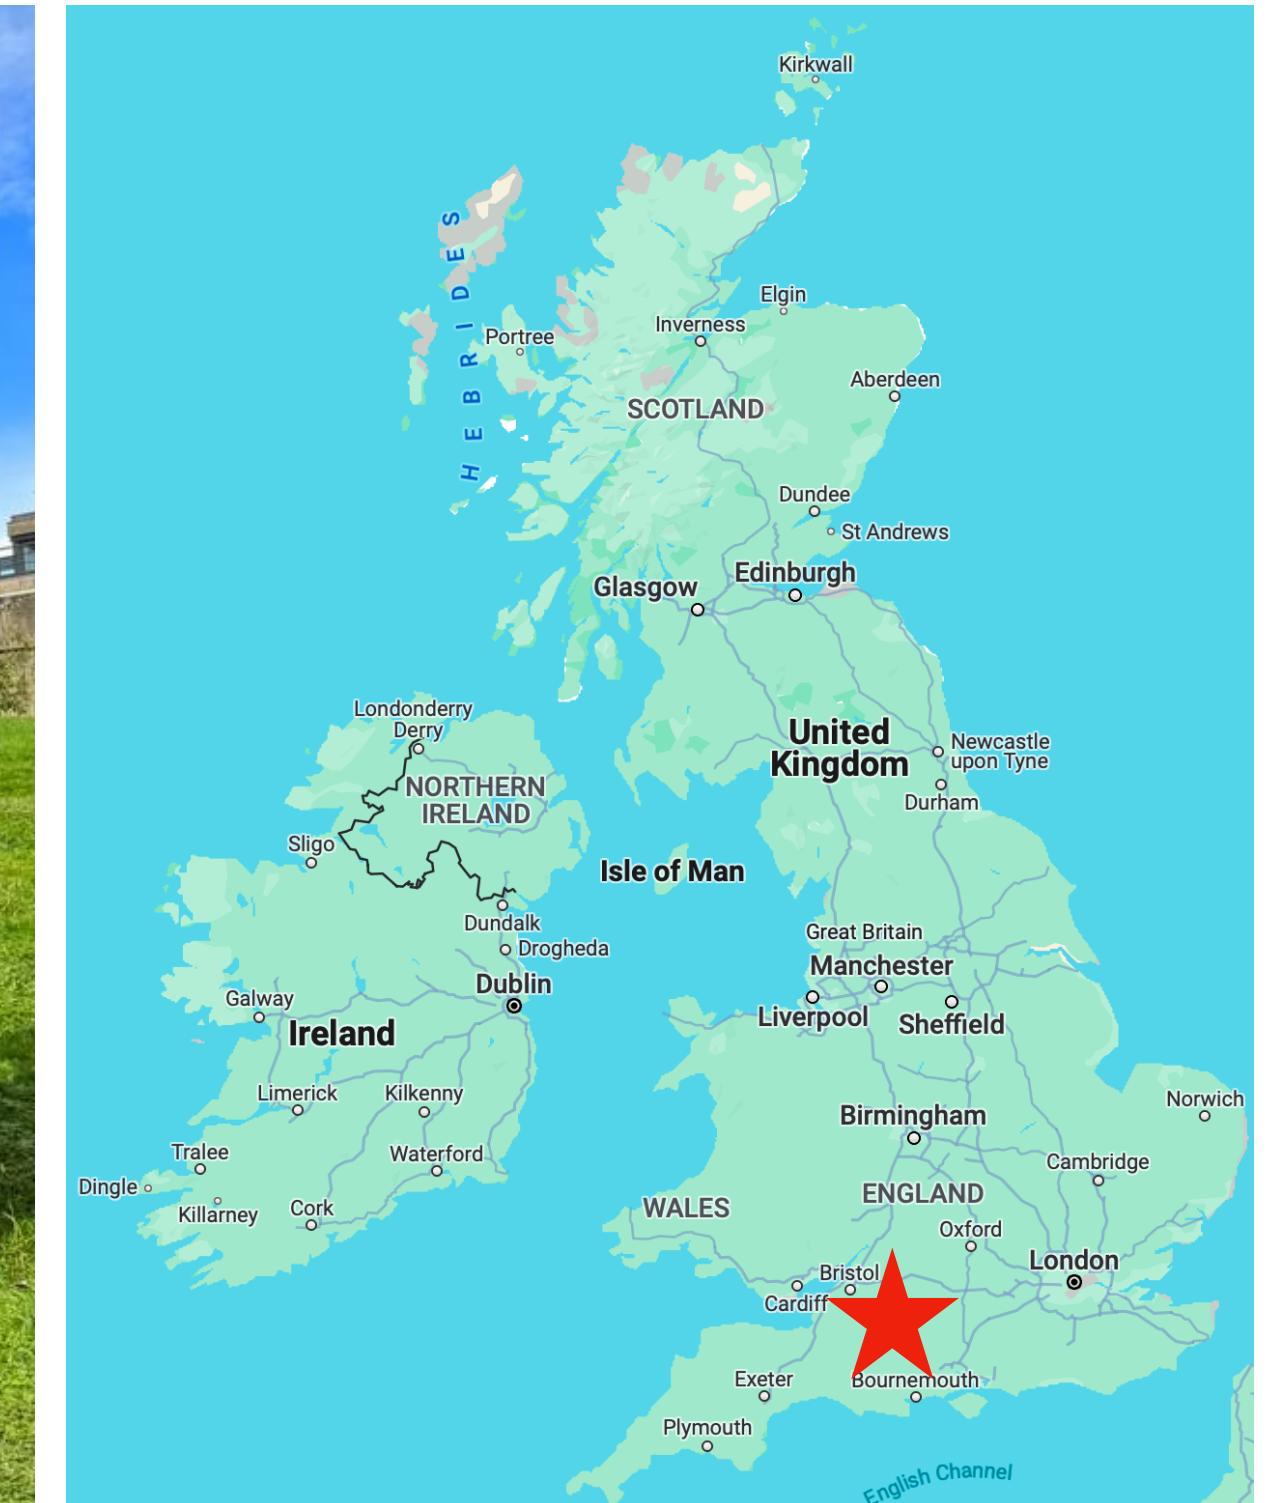
- PhD student at Bath Reinforcement Learning Laboratory under professor Özgür Şimşek
- **Research interests:** transfer learning in reinforcement learning, continual learning, hierarchical reinforcement learning, intrinsically motivation

Đorđe Božić
PhD Student
University of Bath

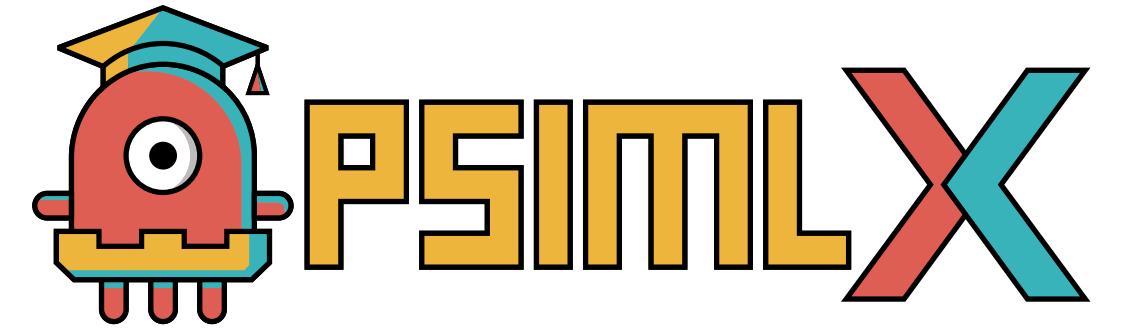
 djordjebozic@gmail.com



UNIVERSITY OF
BATH



Resources

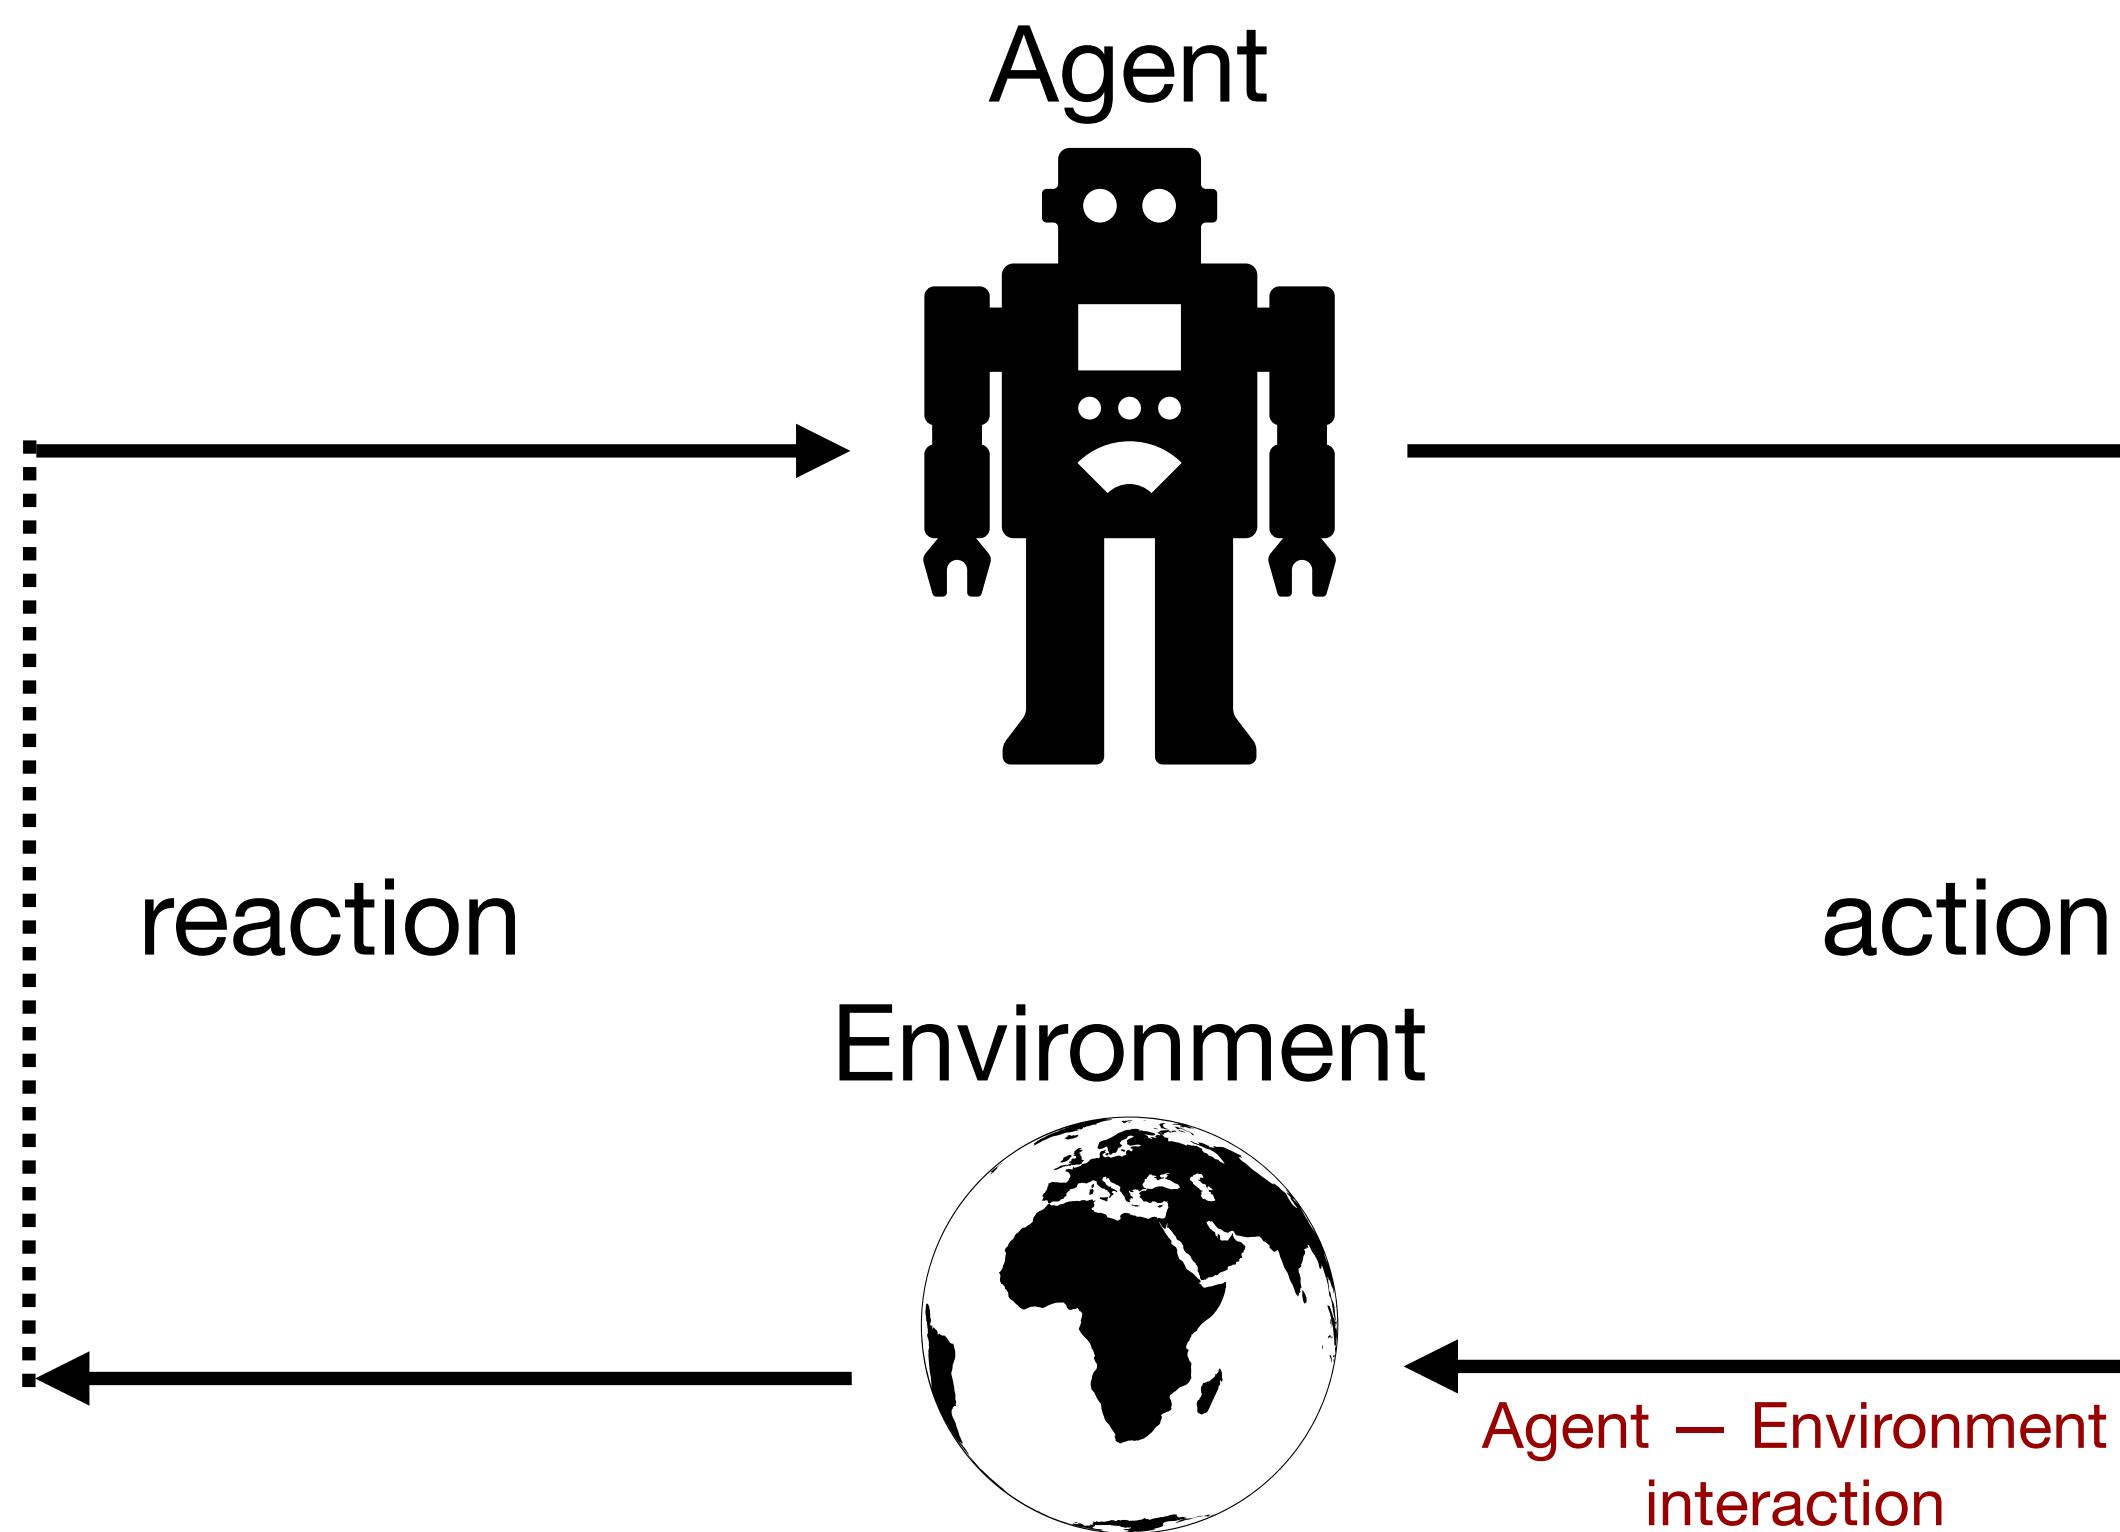


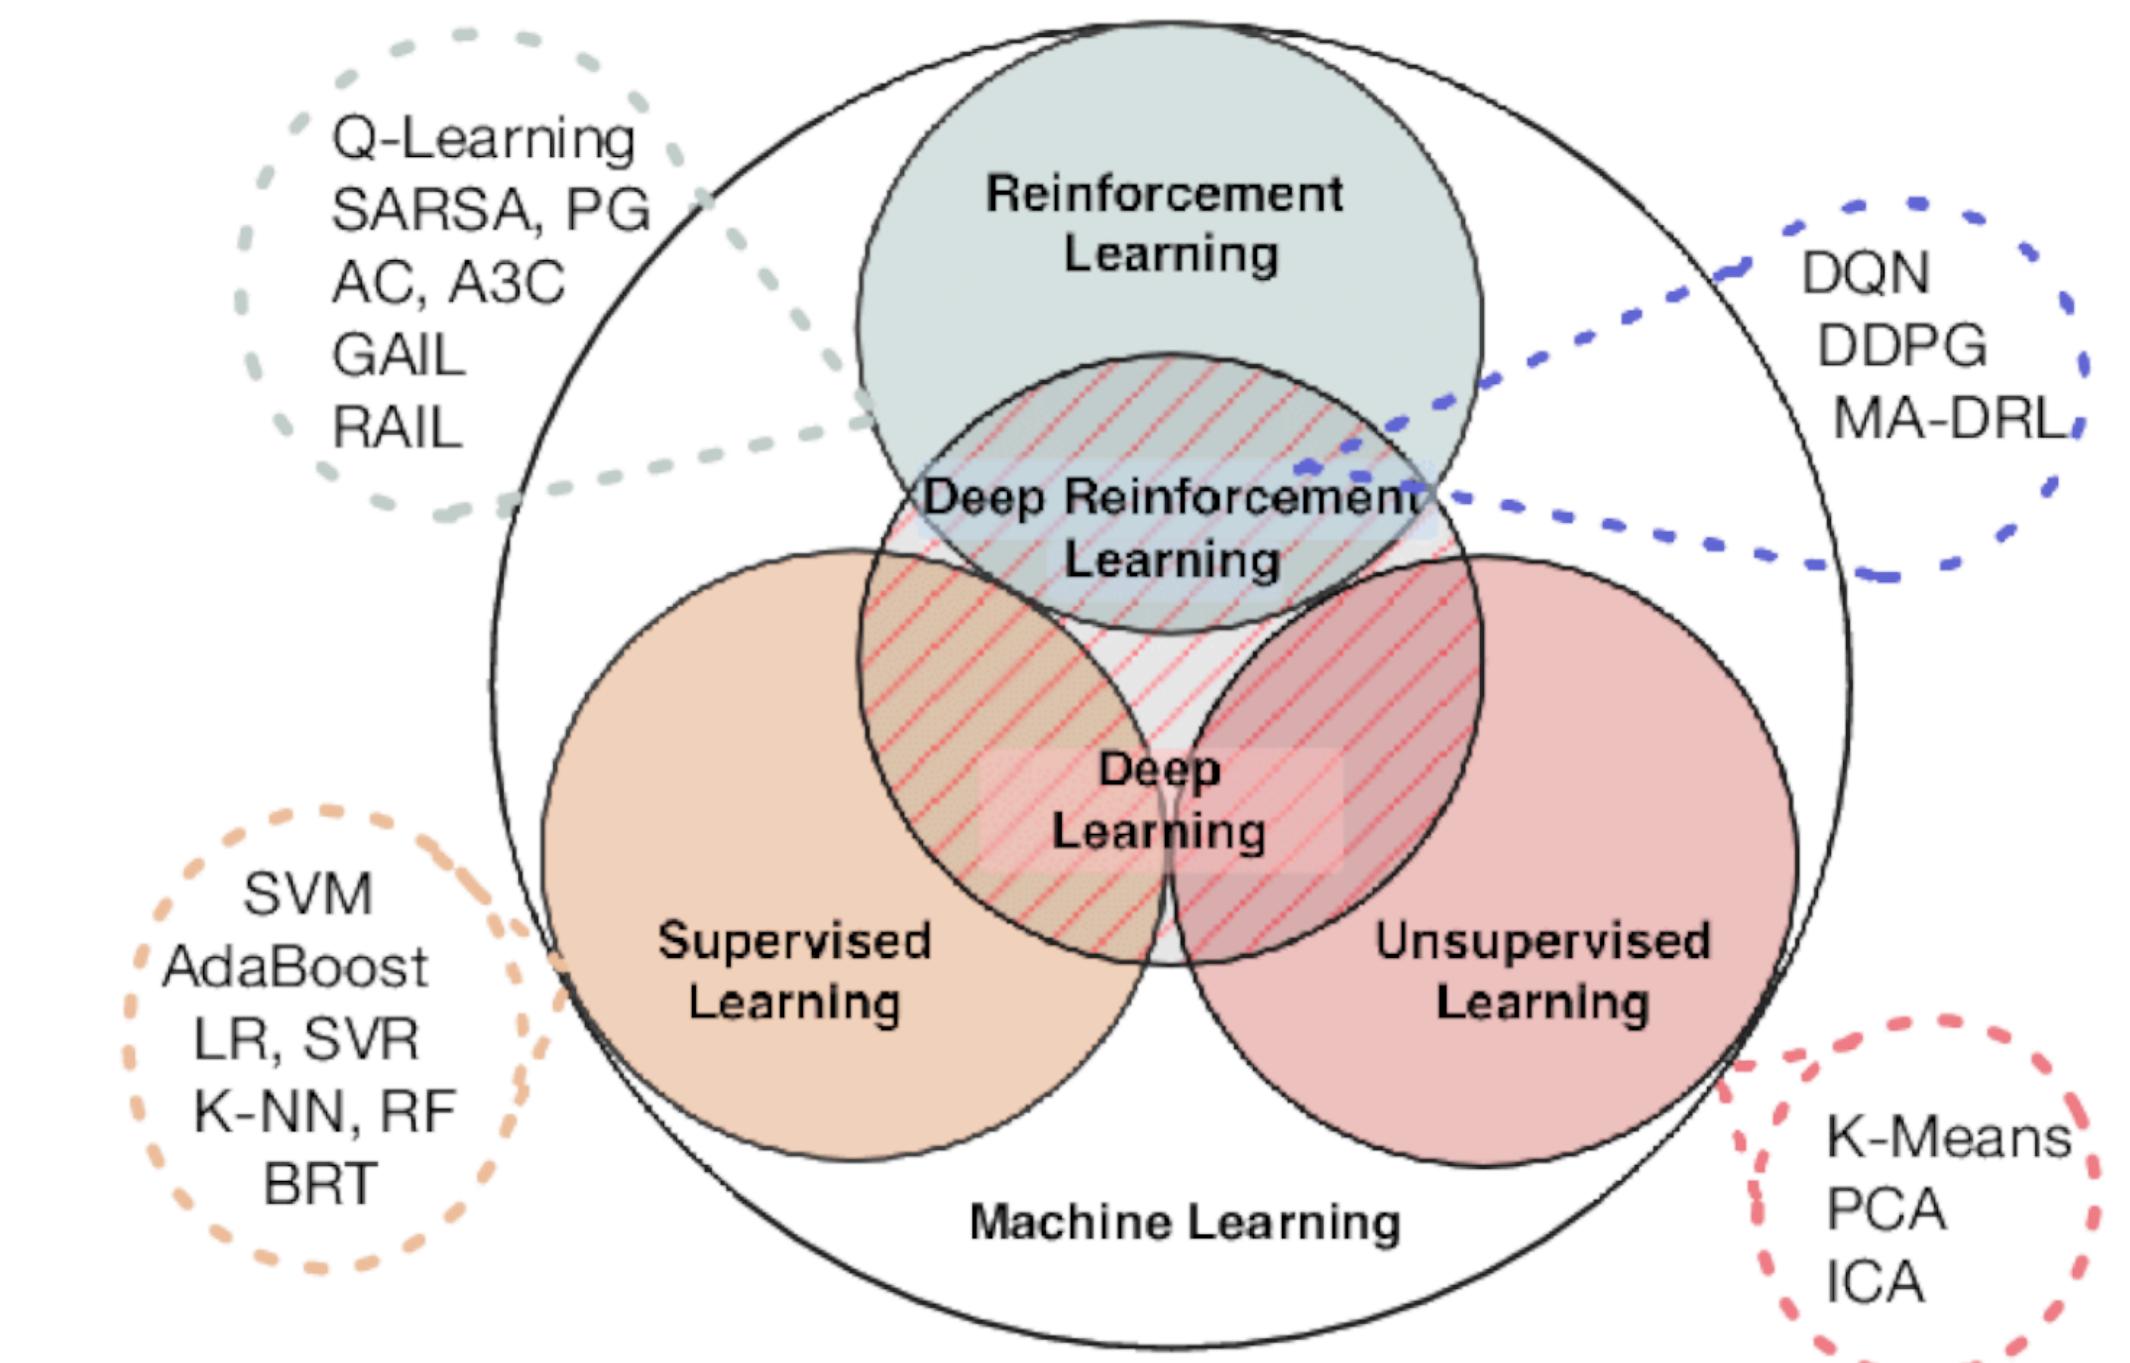
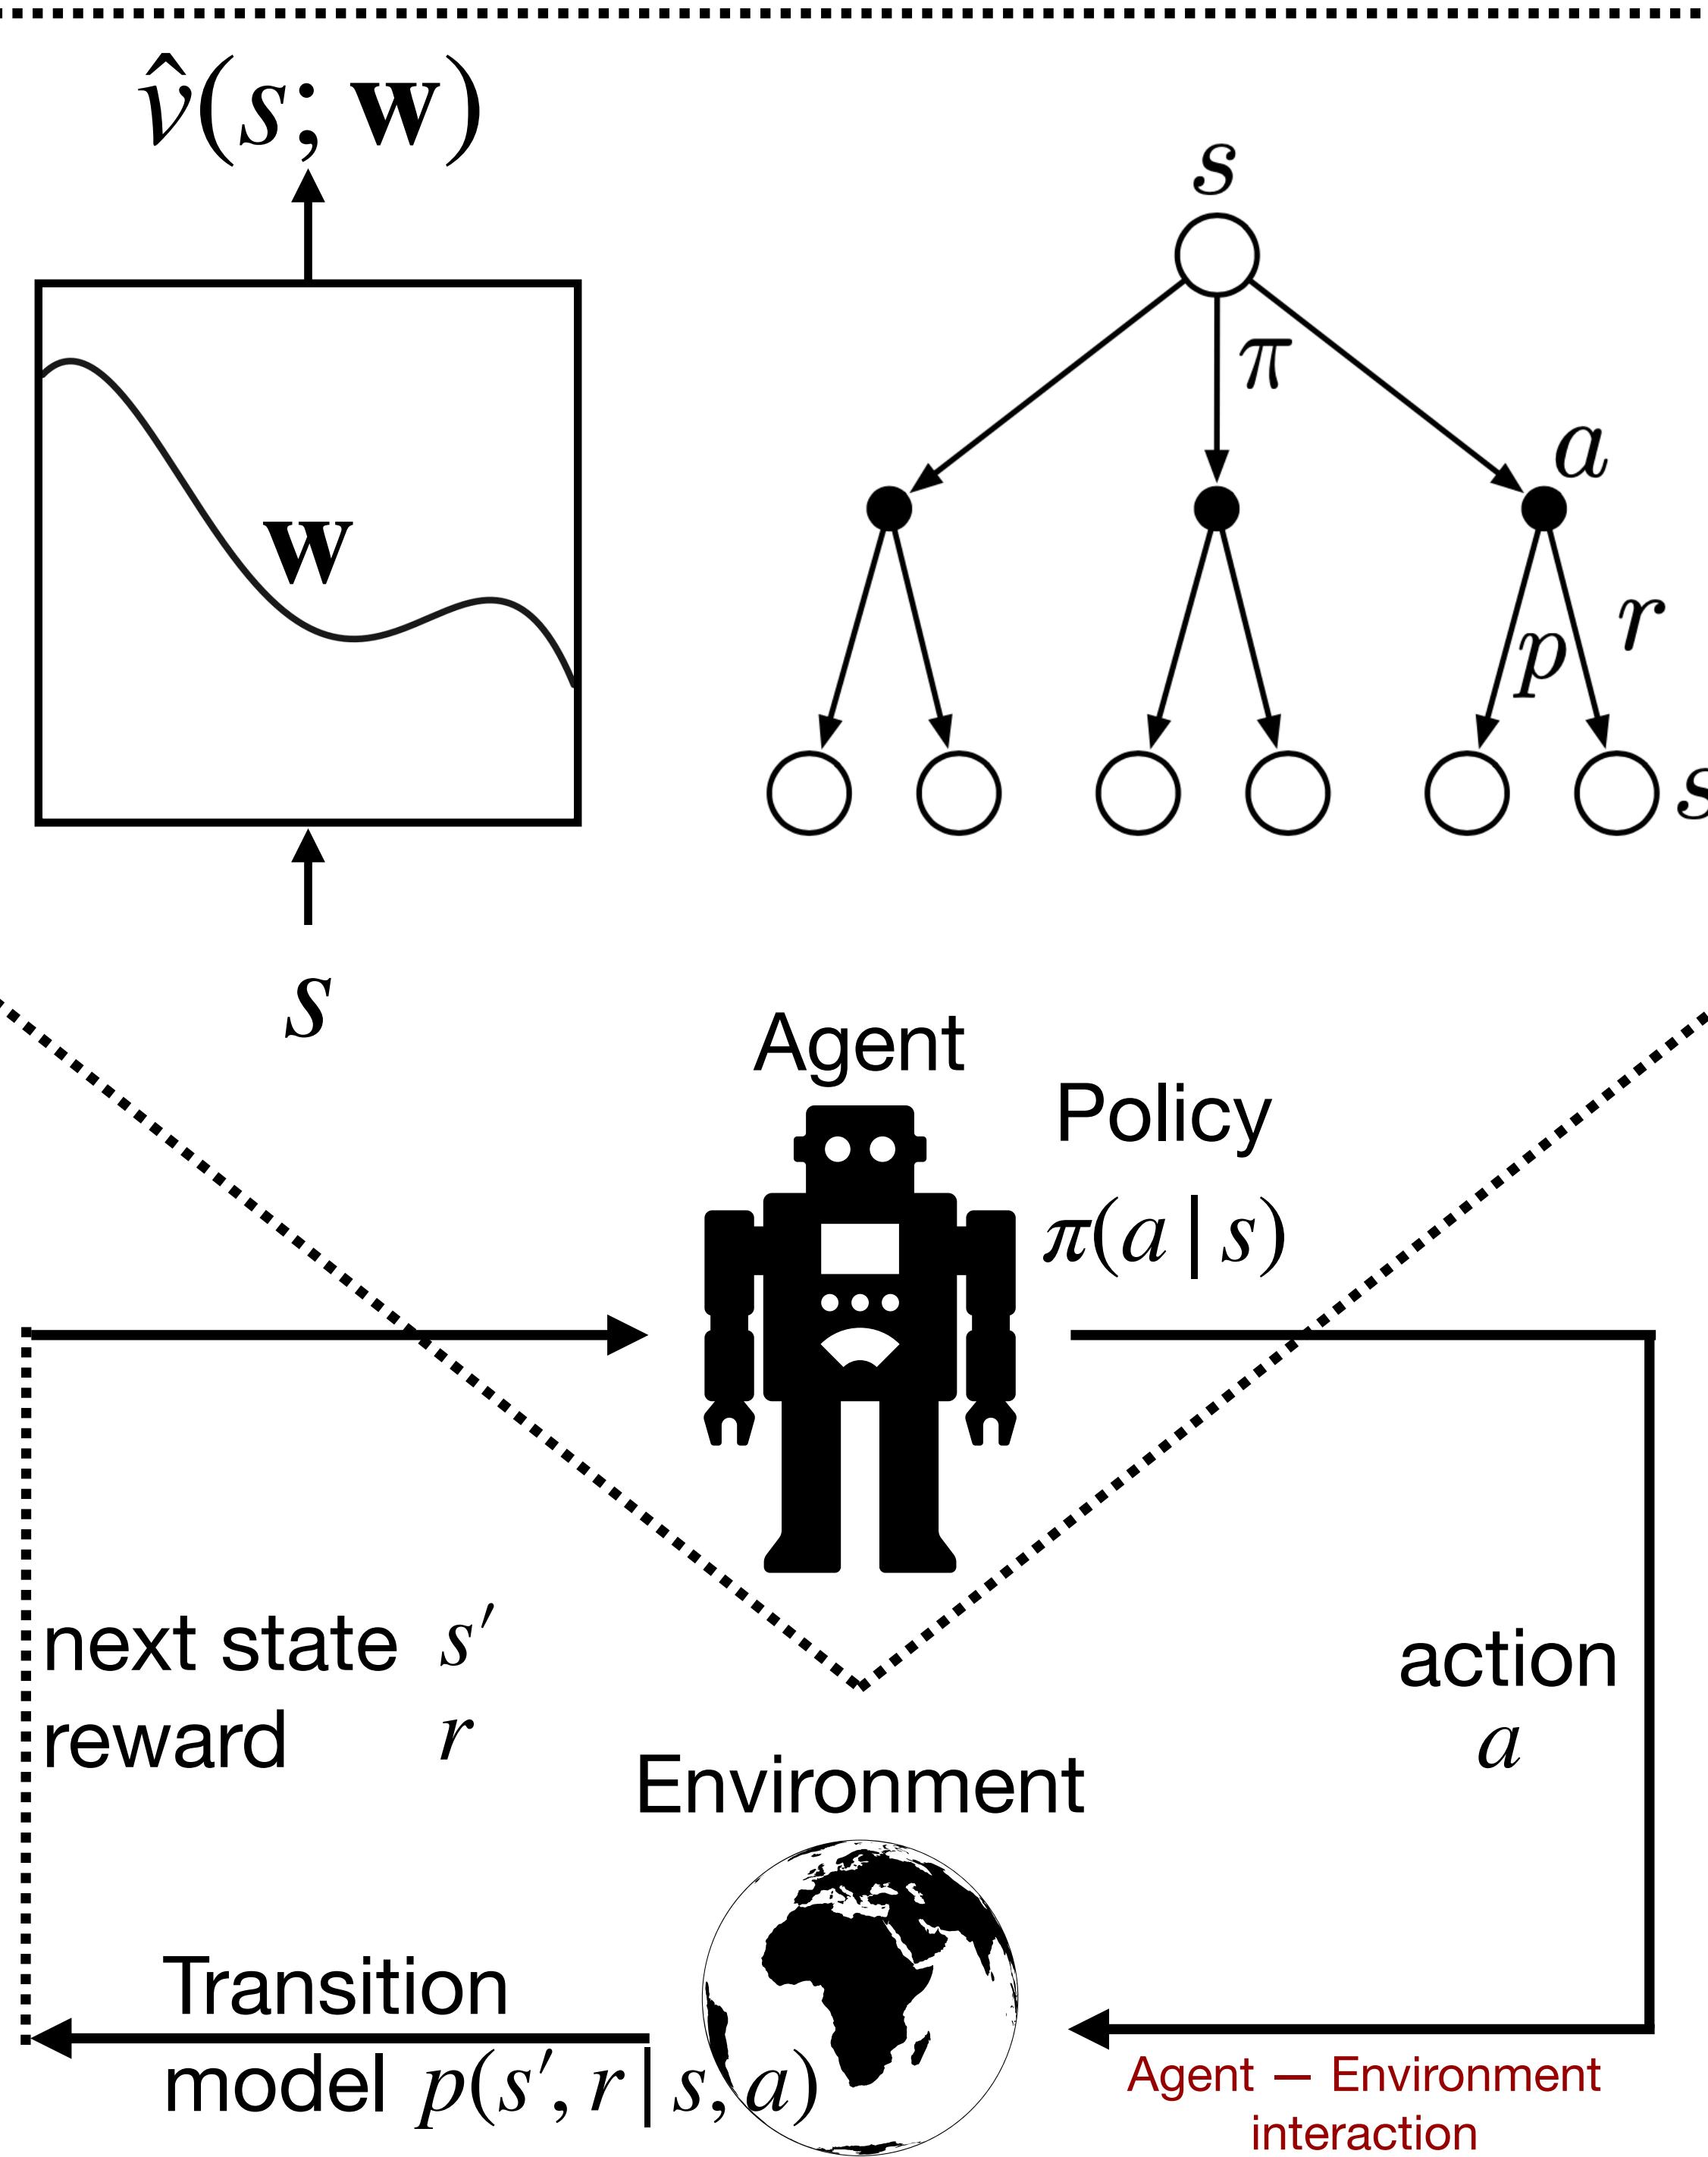
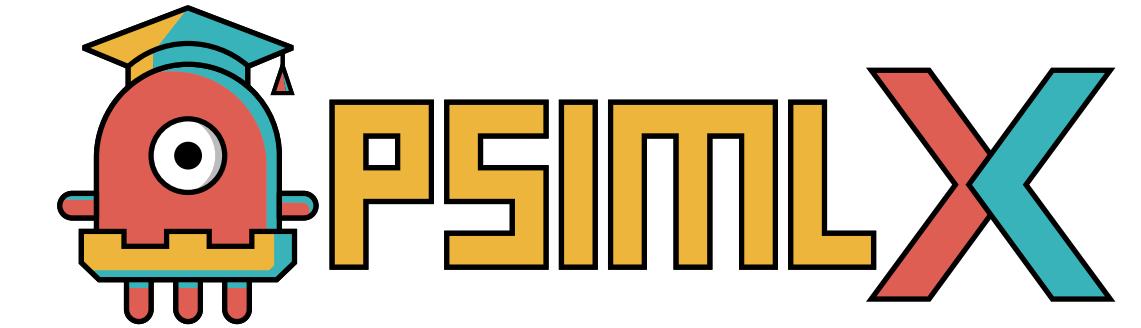
- [Sutton & Barto 2018]: Sutton, Richard S. and Barto, Andrew G.. Reinforcement Learning: An Introduction. Second : The MIT Press, 2018 .
- [Russel & Norvig 2010]: Russell, S. & Norvig, P. (2010), Artificial Intelligence: A Modern Approach , Prentice Hall.

Introduction

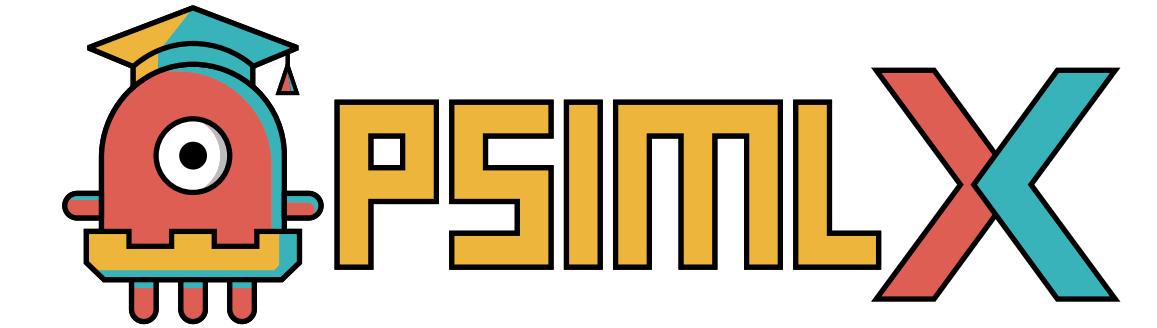


Tunnelling strategy discovered by DQN on Breakout Atari environment





Introduction



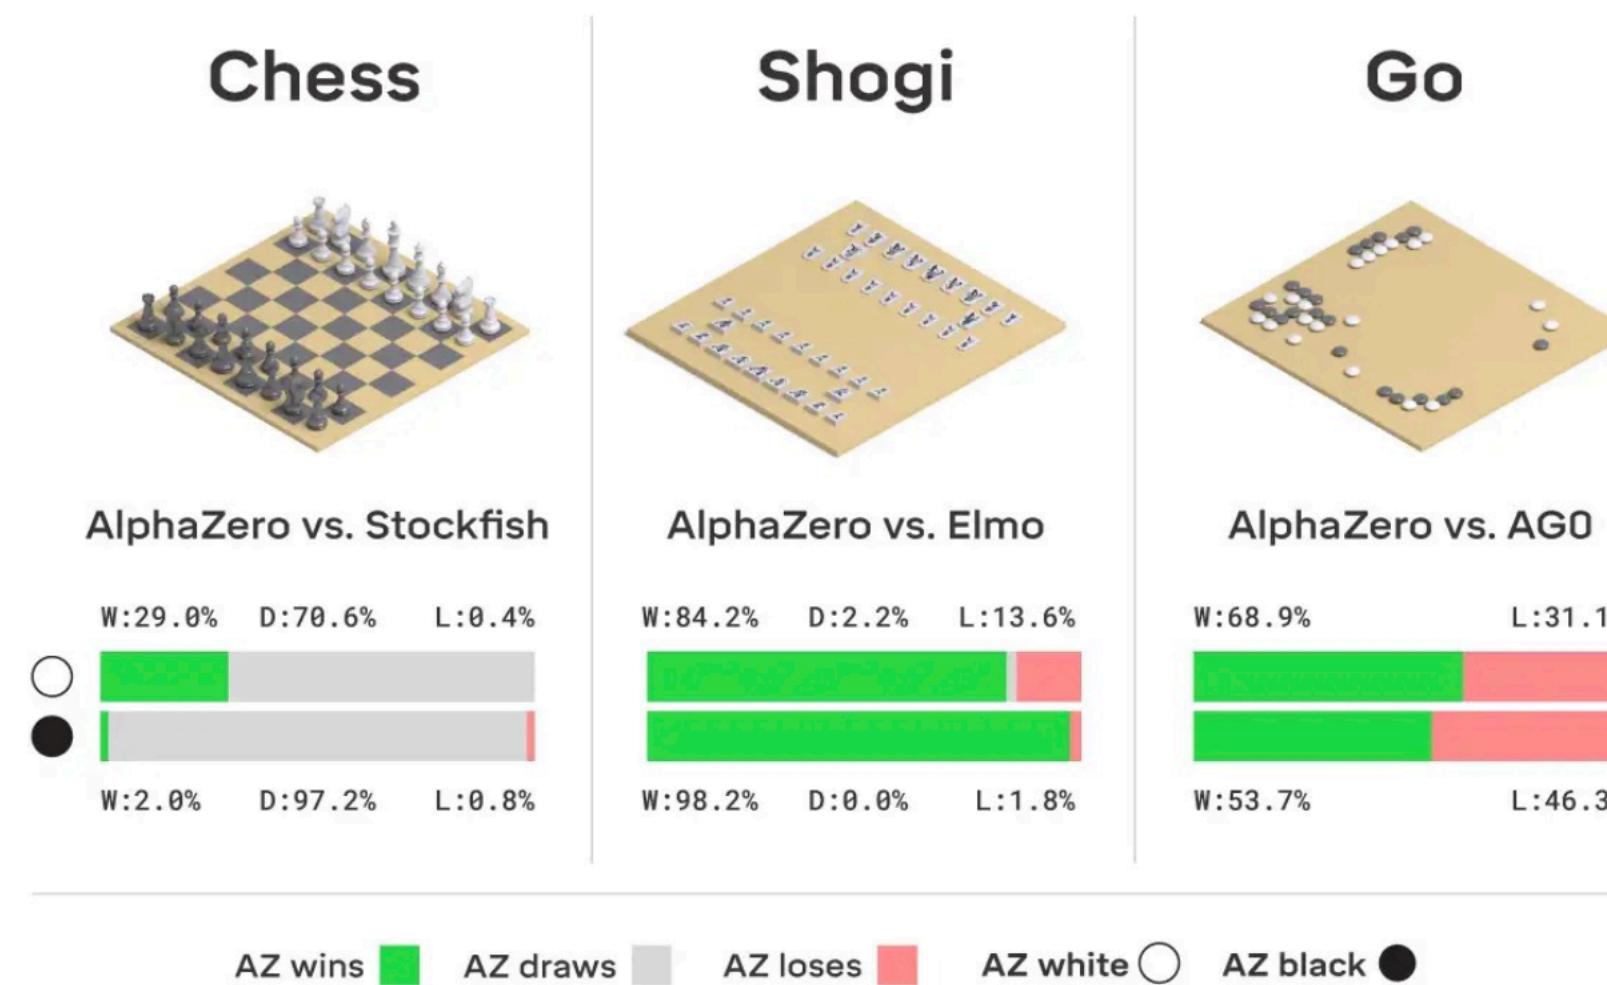
Alpha Go, Silver et al. 2016



Alpha Star, Vinyals et al. 2019



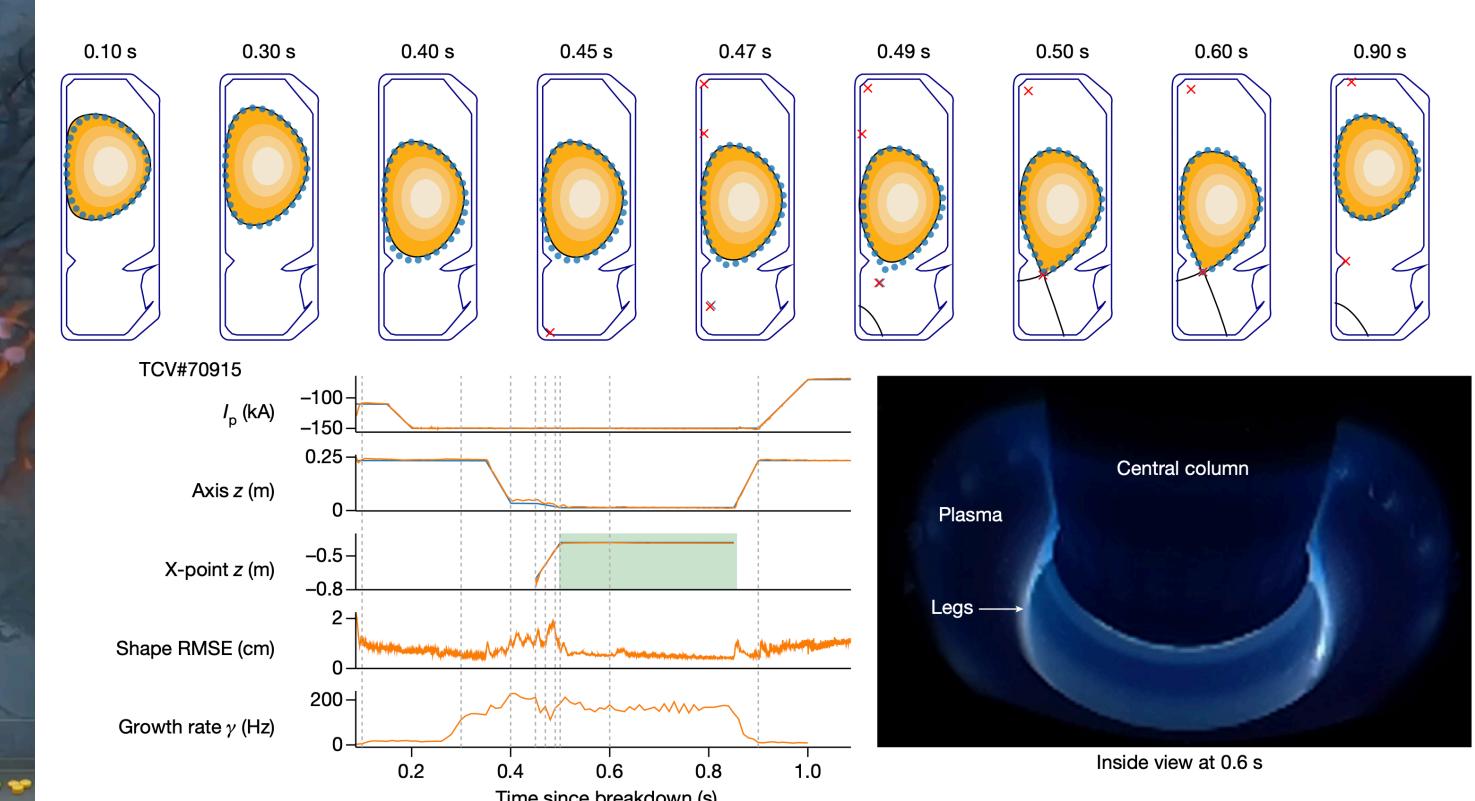
Stratospheric Balloon Navigation, Bellemare et al, 2020



Alpha Zero, Silver et al. 2017

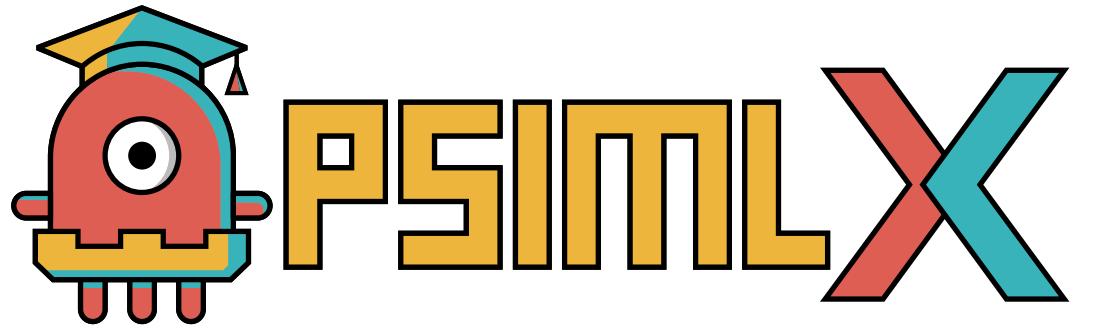


OpenAI Five, Barner et al. 2019



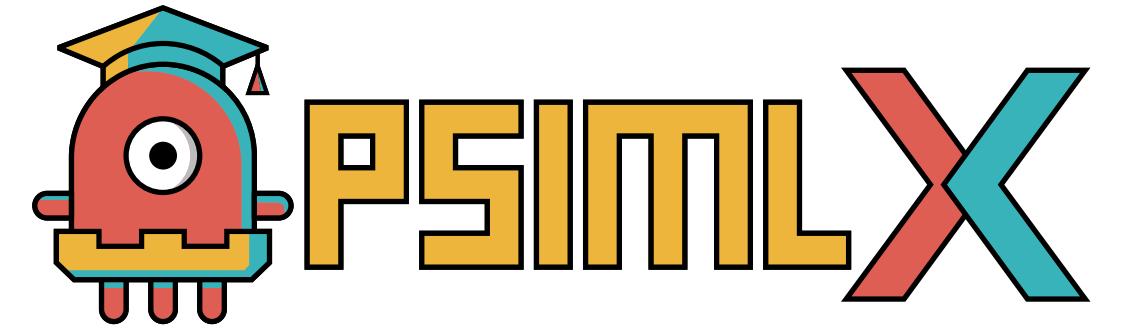
Magnetic Control of Tokamak Plasmas, Degreve et al, 2022

Outline



- Introduction
 - Reinforcement Learning Formalisation
 - Model-Free Reinforcement Learning
 - Value Function Approximation
 - Policy Gradient Methods
- “Deep RL”

Outline

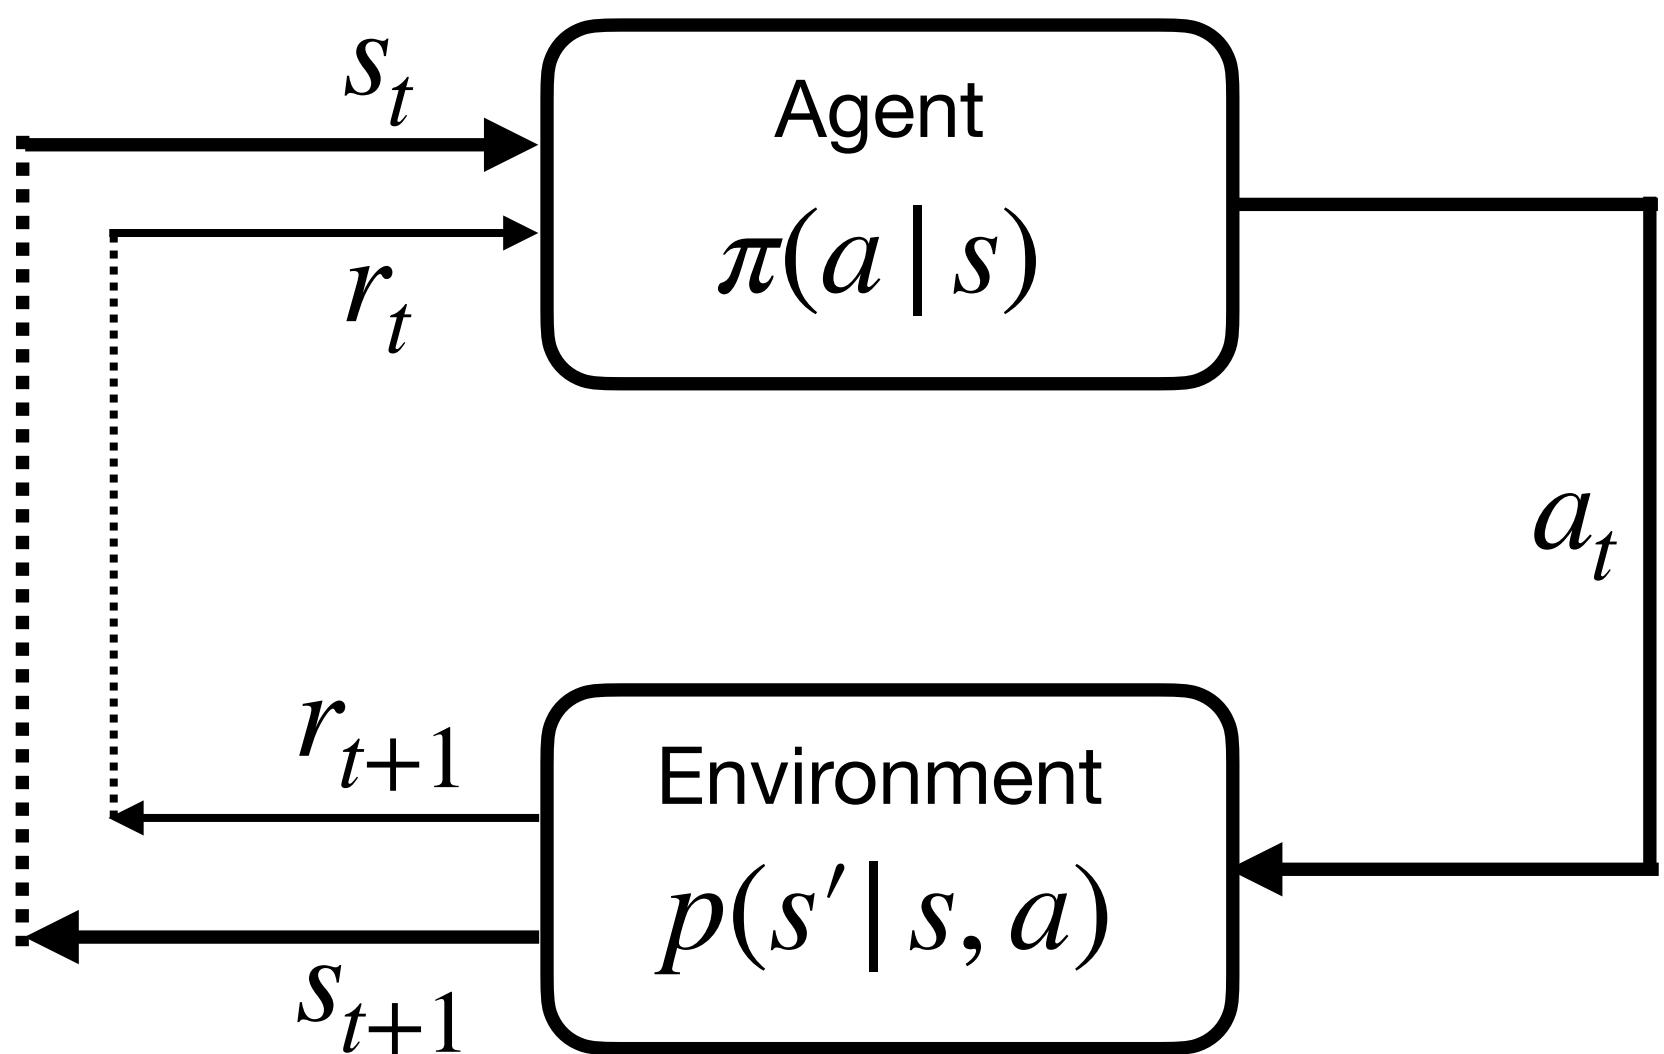


- Introduction
- Reinforcement Learning Formalisation
 - Agents and Environments
 - Markov Decision Process
 - Reward and Return
 - Policy
 - State-Value Function
 - Action-Value Function
 - Optimal Policy
 - Optimal Value Functions
 - Bellman Equations and Planning
 - Generalised Policy Iteration (GPI)
 - Exploration-Exploitation Trade-Off
- ...

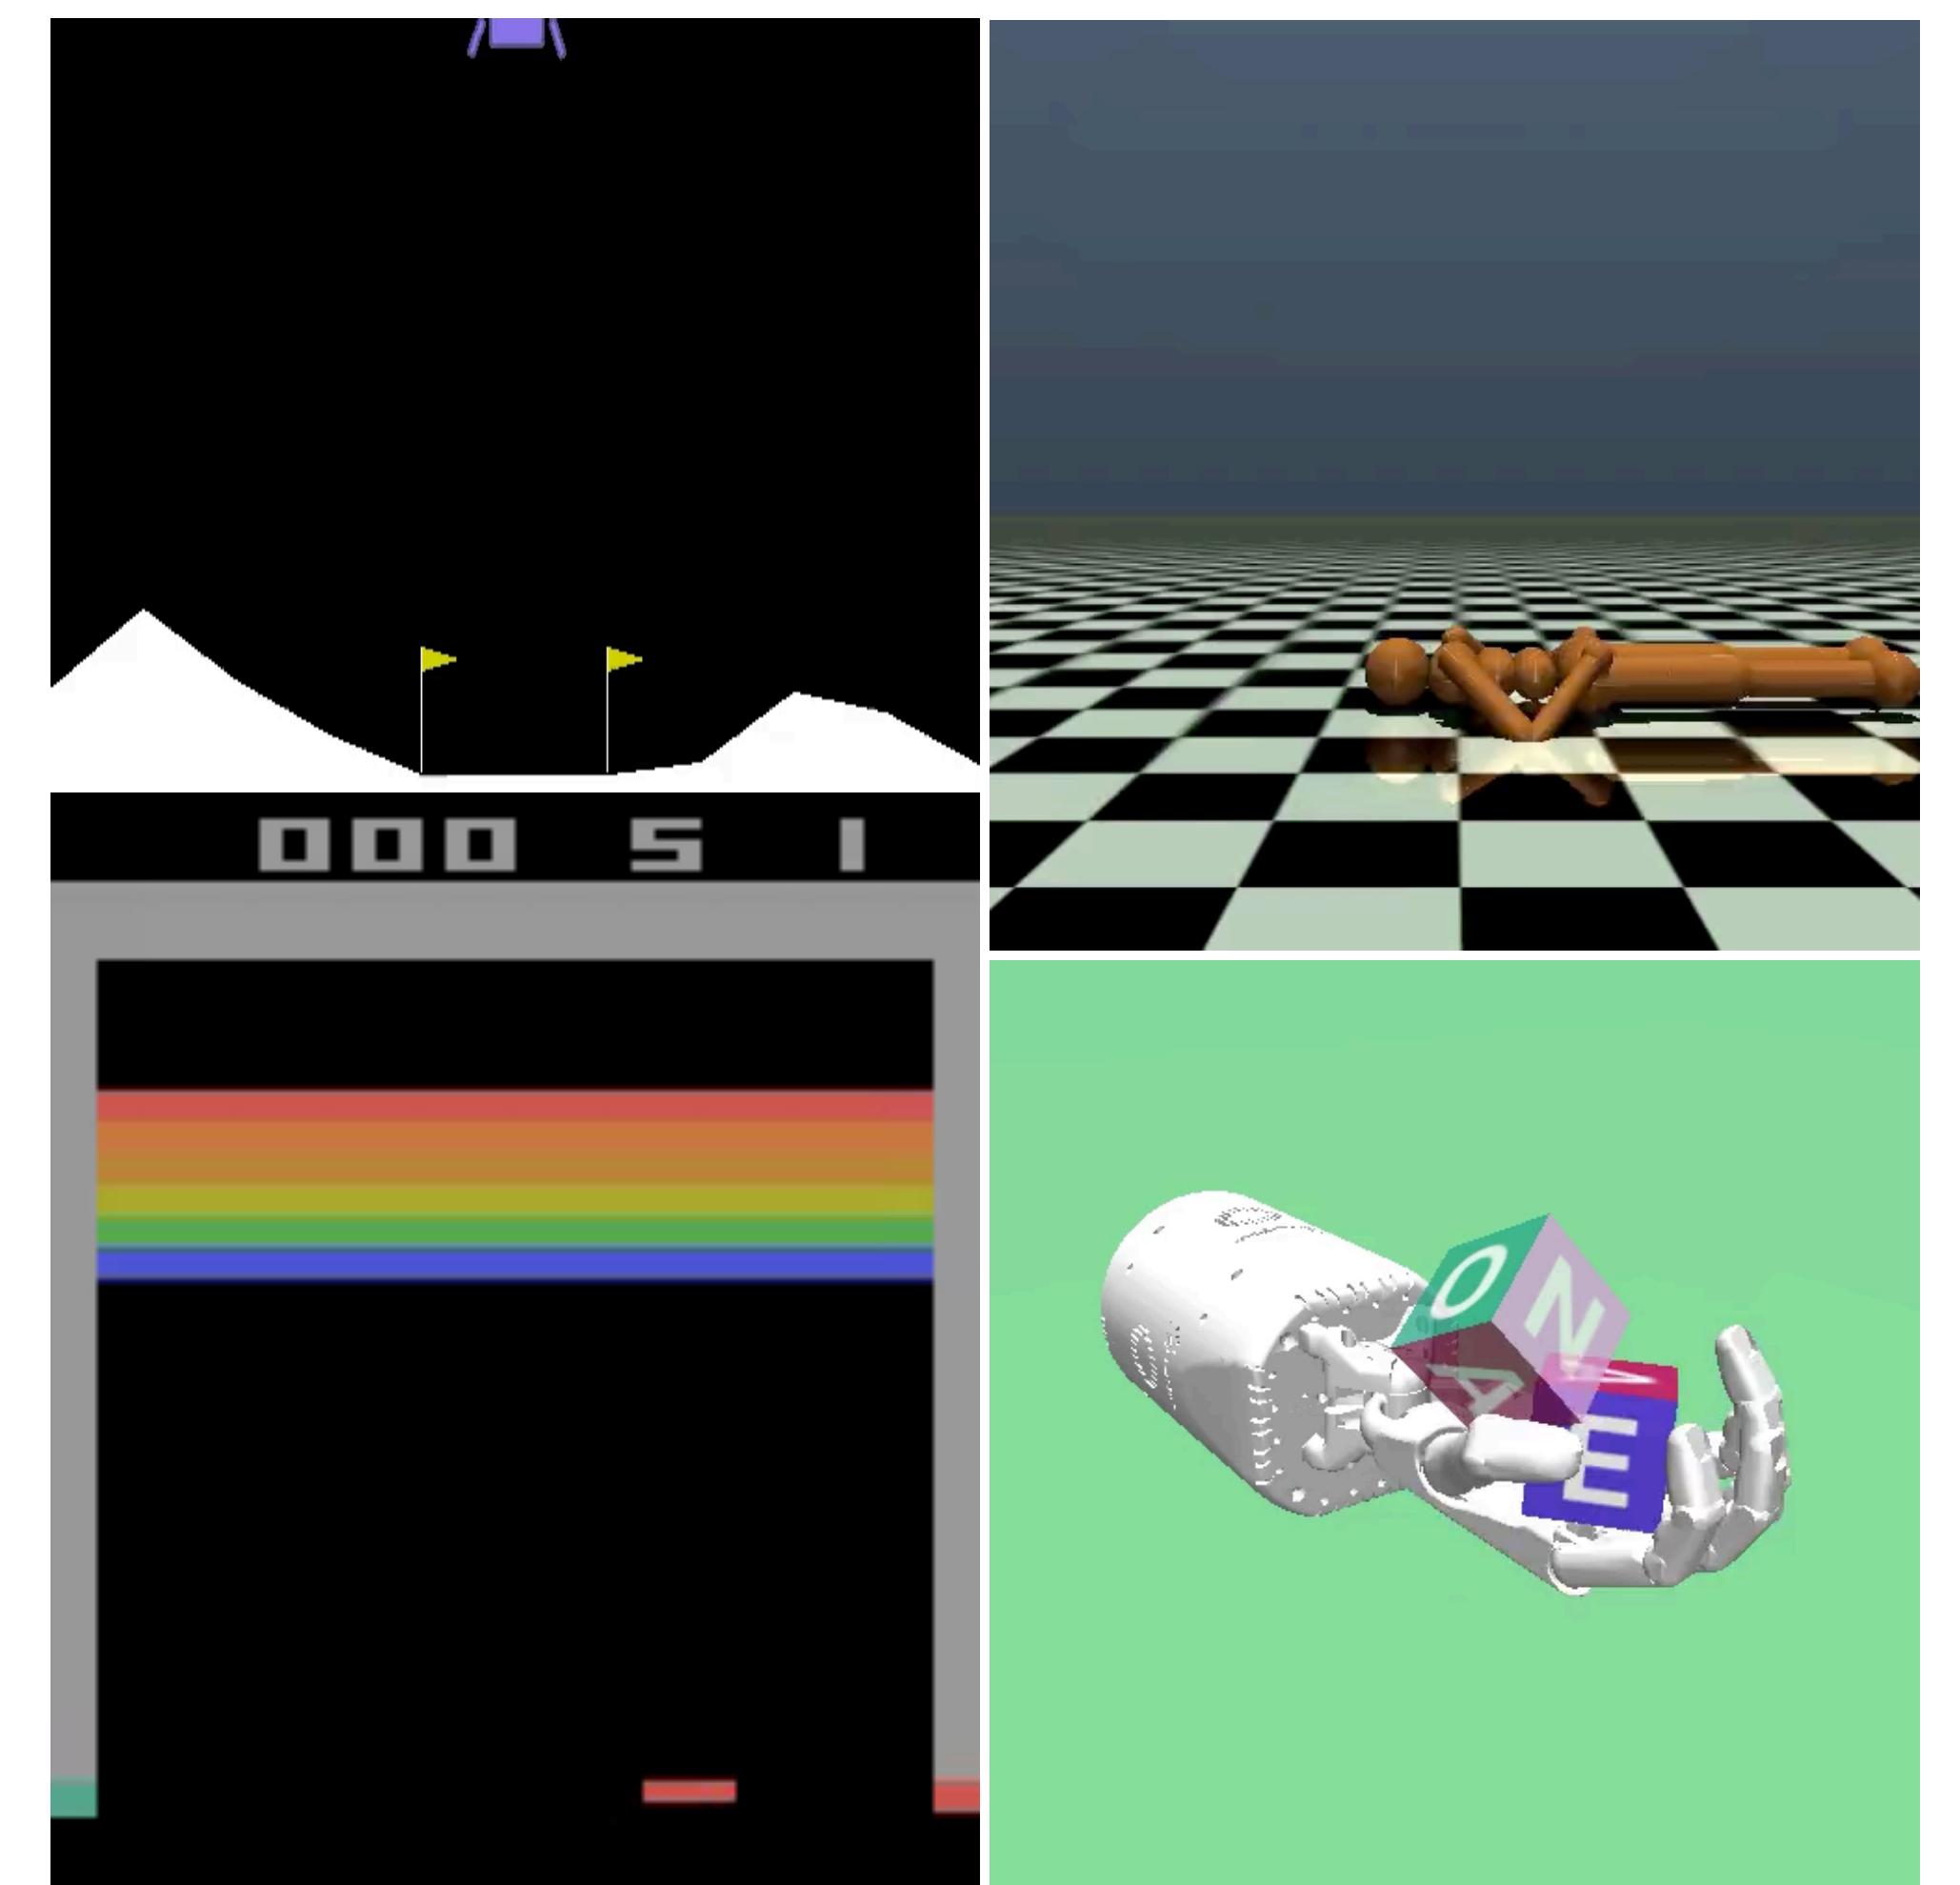
RL Formalisation

Agents and Environments

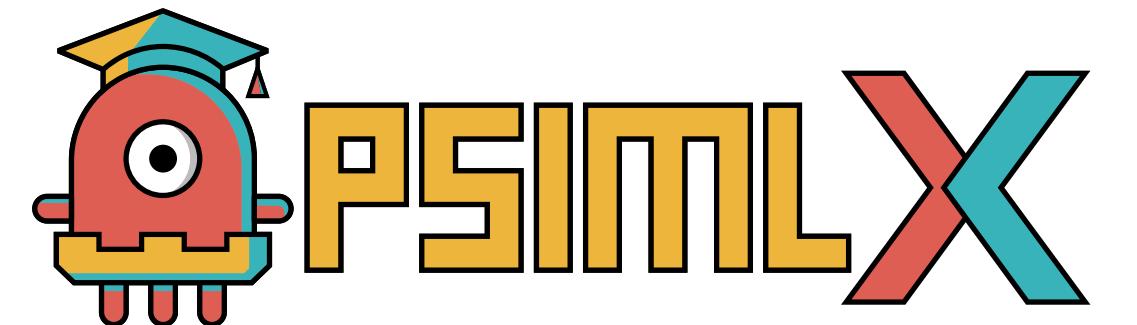
- Agent interacts with the environment
- Rewards given as feedback
- Different kind of supervision:
 - Samples not I.I.D
 - Learning signal may be delayed



Agent-Environment Interaction [Sutton & Barto 2018]



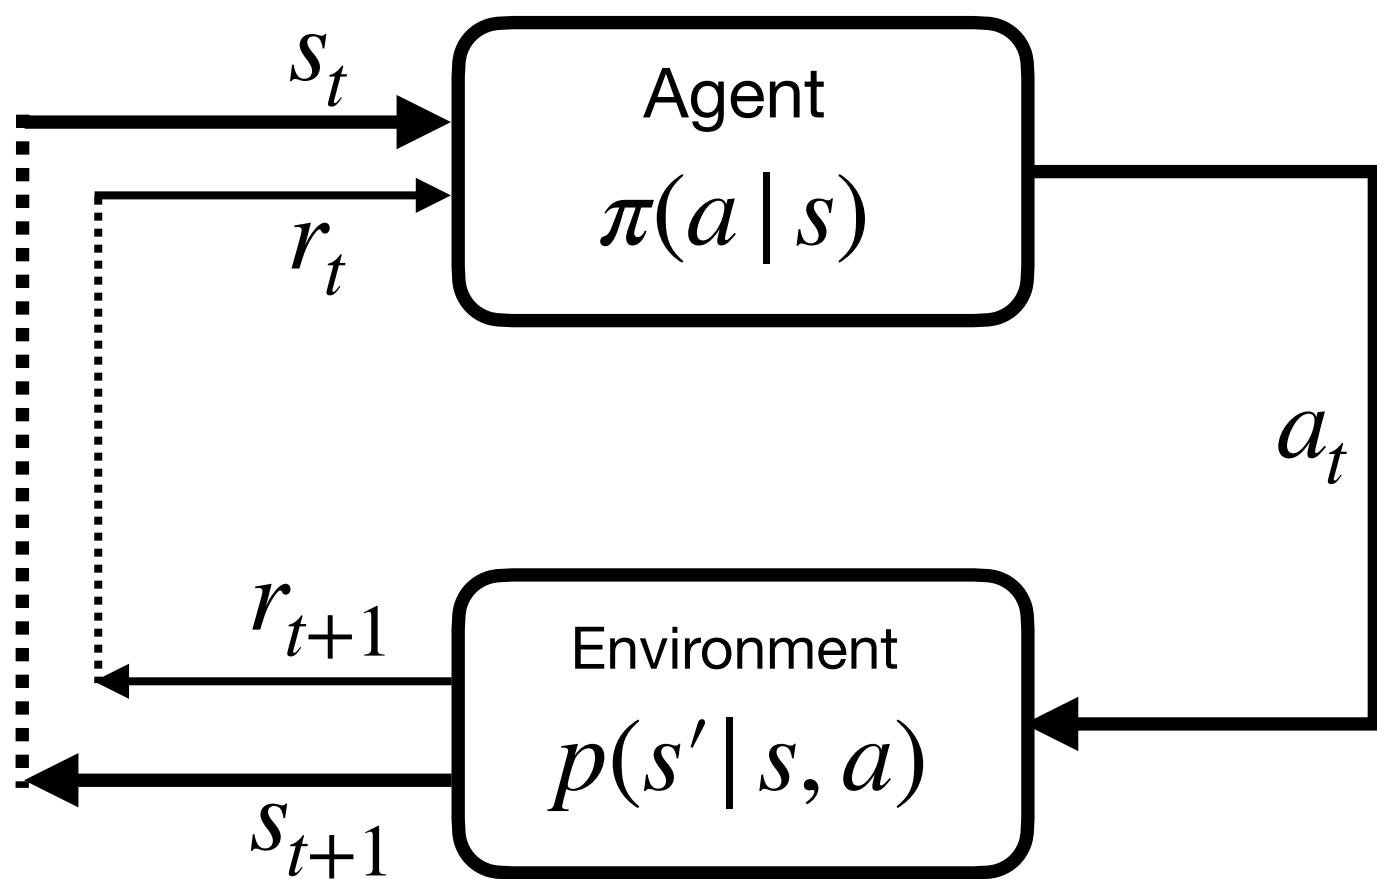
Lunar Lander, Breakout, MuJoCo Humanoid, Hand [Open AI Gym Environment suite]



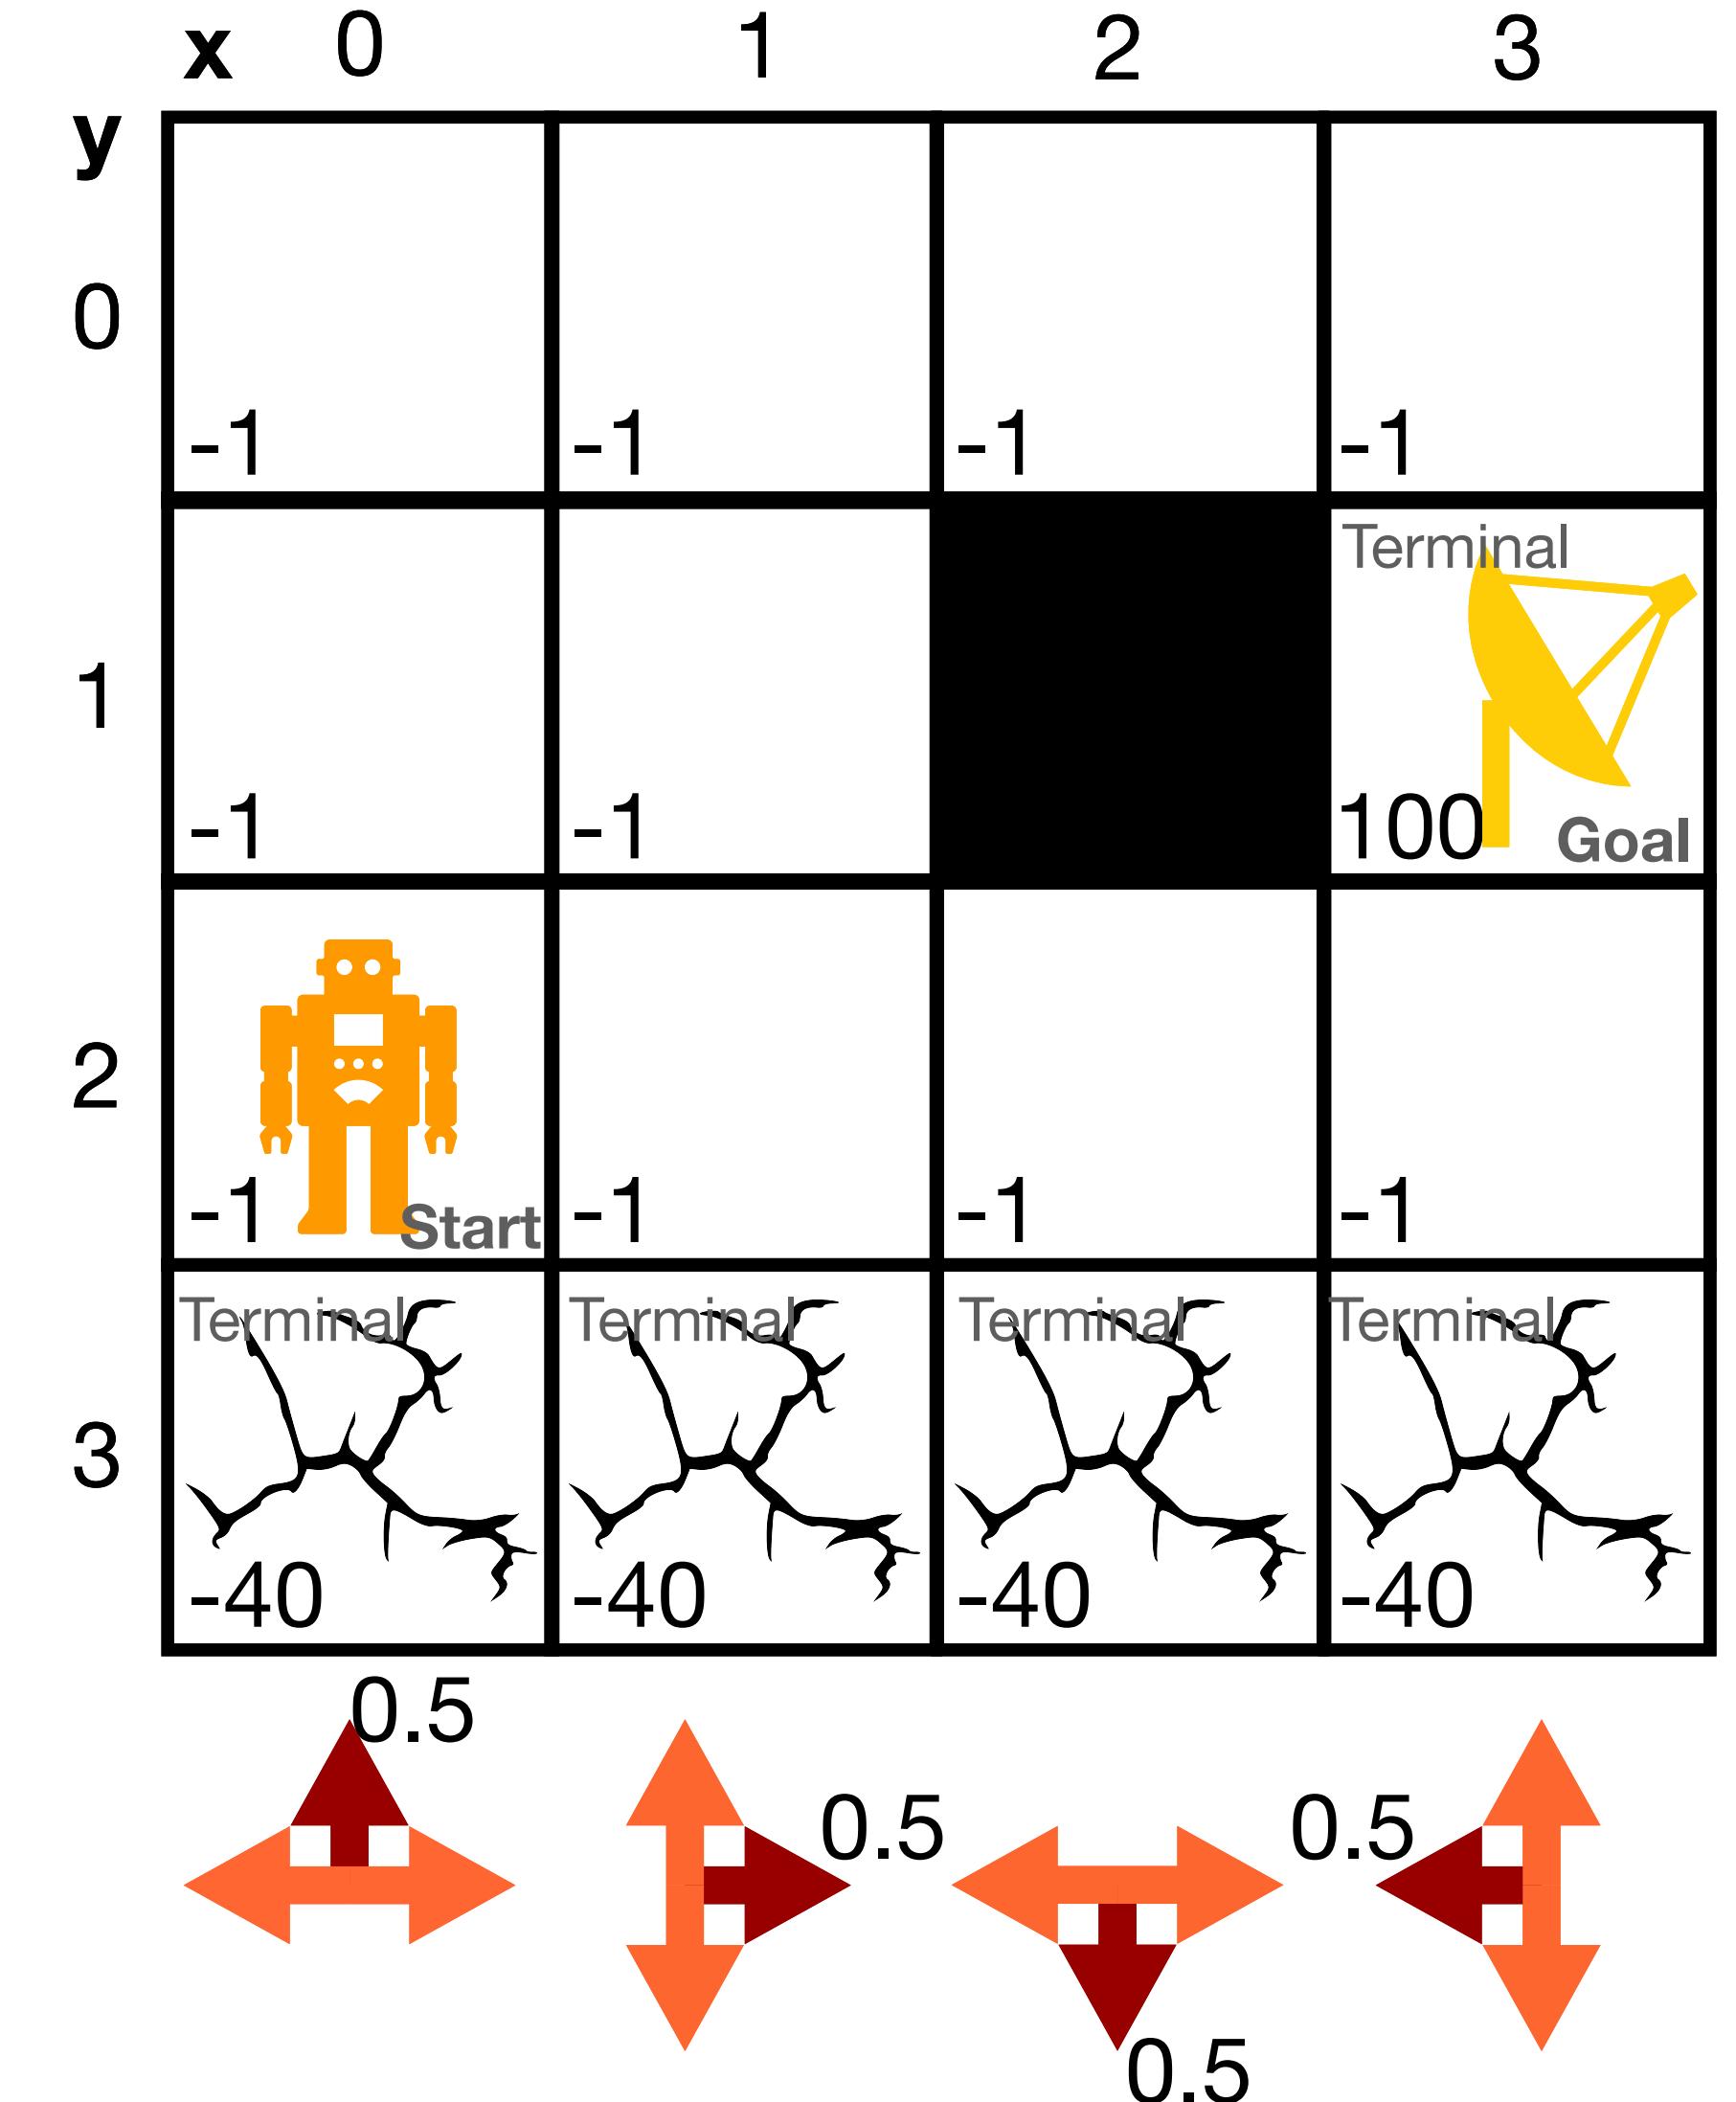
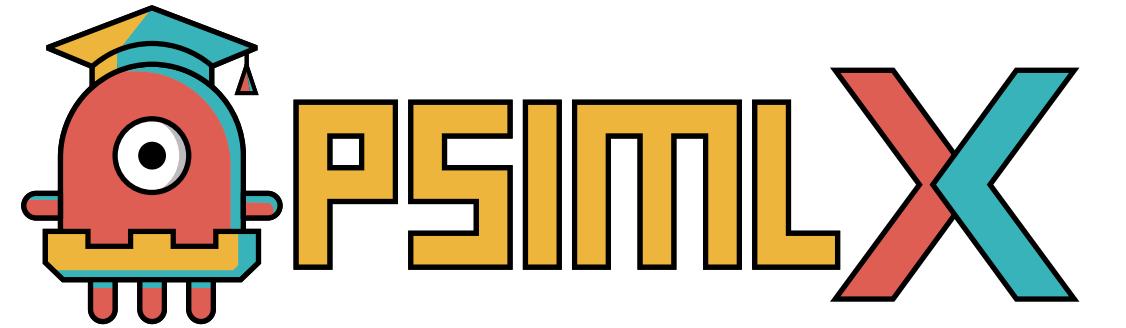
RL Formalisation

Agents and Environments

- Environment characteristics?
 - Episodic (finite-horizon) vs. Continuing (indefinite-horizon)
 - Deterministic vs. Stochastic
 - Fully vs. Partially observable
 - Discrete vs. Continuous



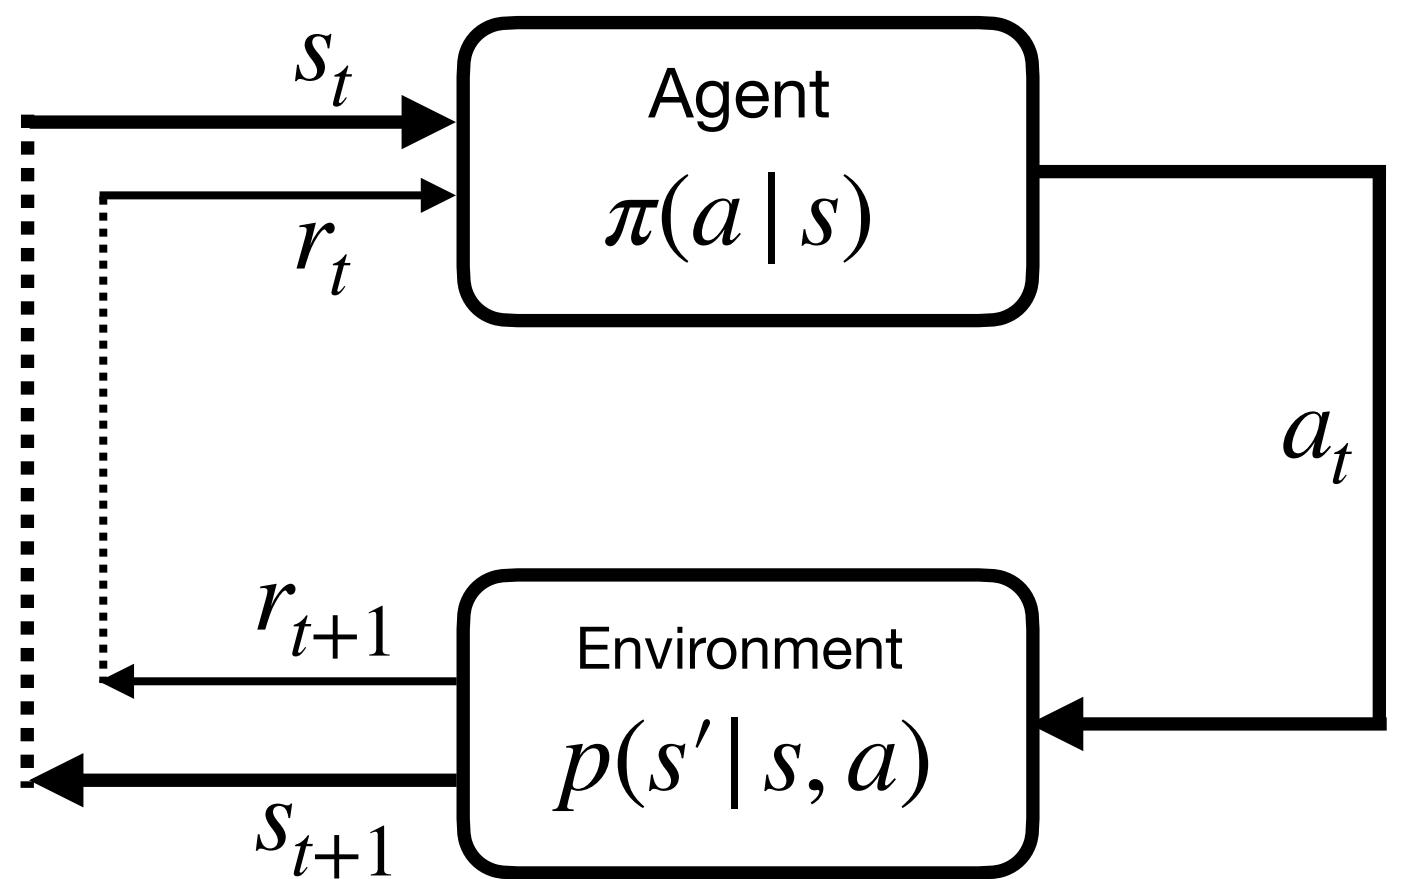
Agent-Environment Interaction [Sutton & Barto 2018]



RL Formalisation

Agents and Environments

- Environment characteristics?
 - Episodic (finite-horizon) vs. Continuing (indefinite-horizon)
 - Deterministic vs. Stochastic
 - Fully vs. Partially observable
 - Discrete vs. Continuous



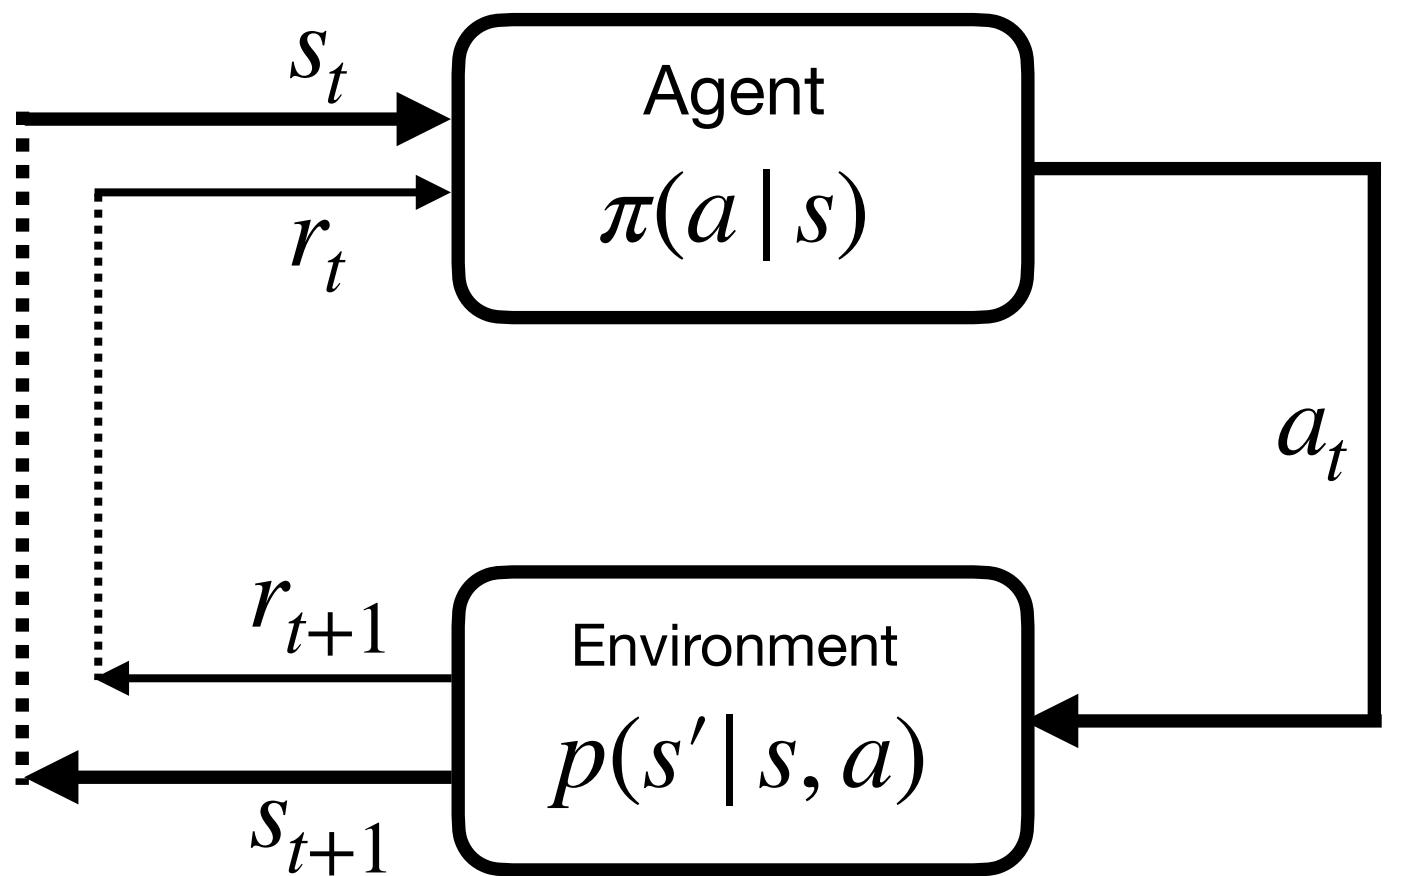
Agent-Environment Interaction [Sutton & Barto
2018]



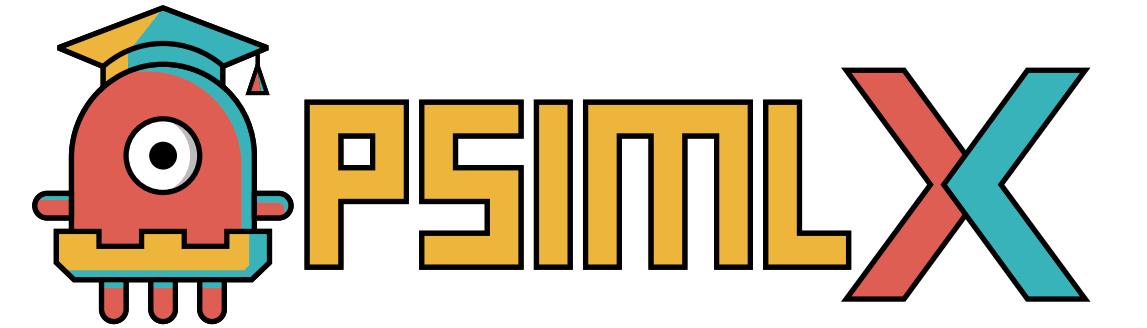
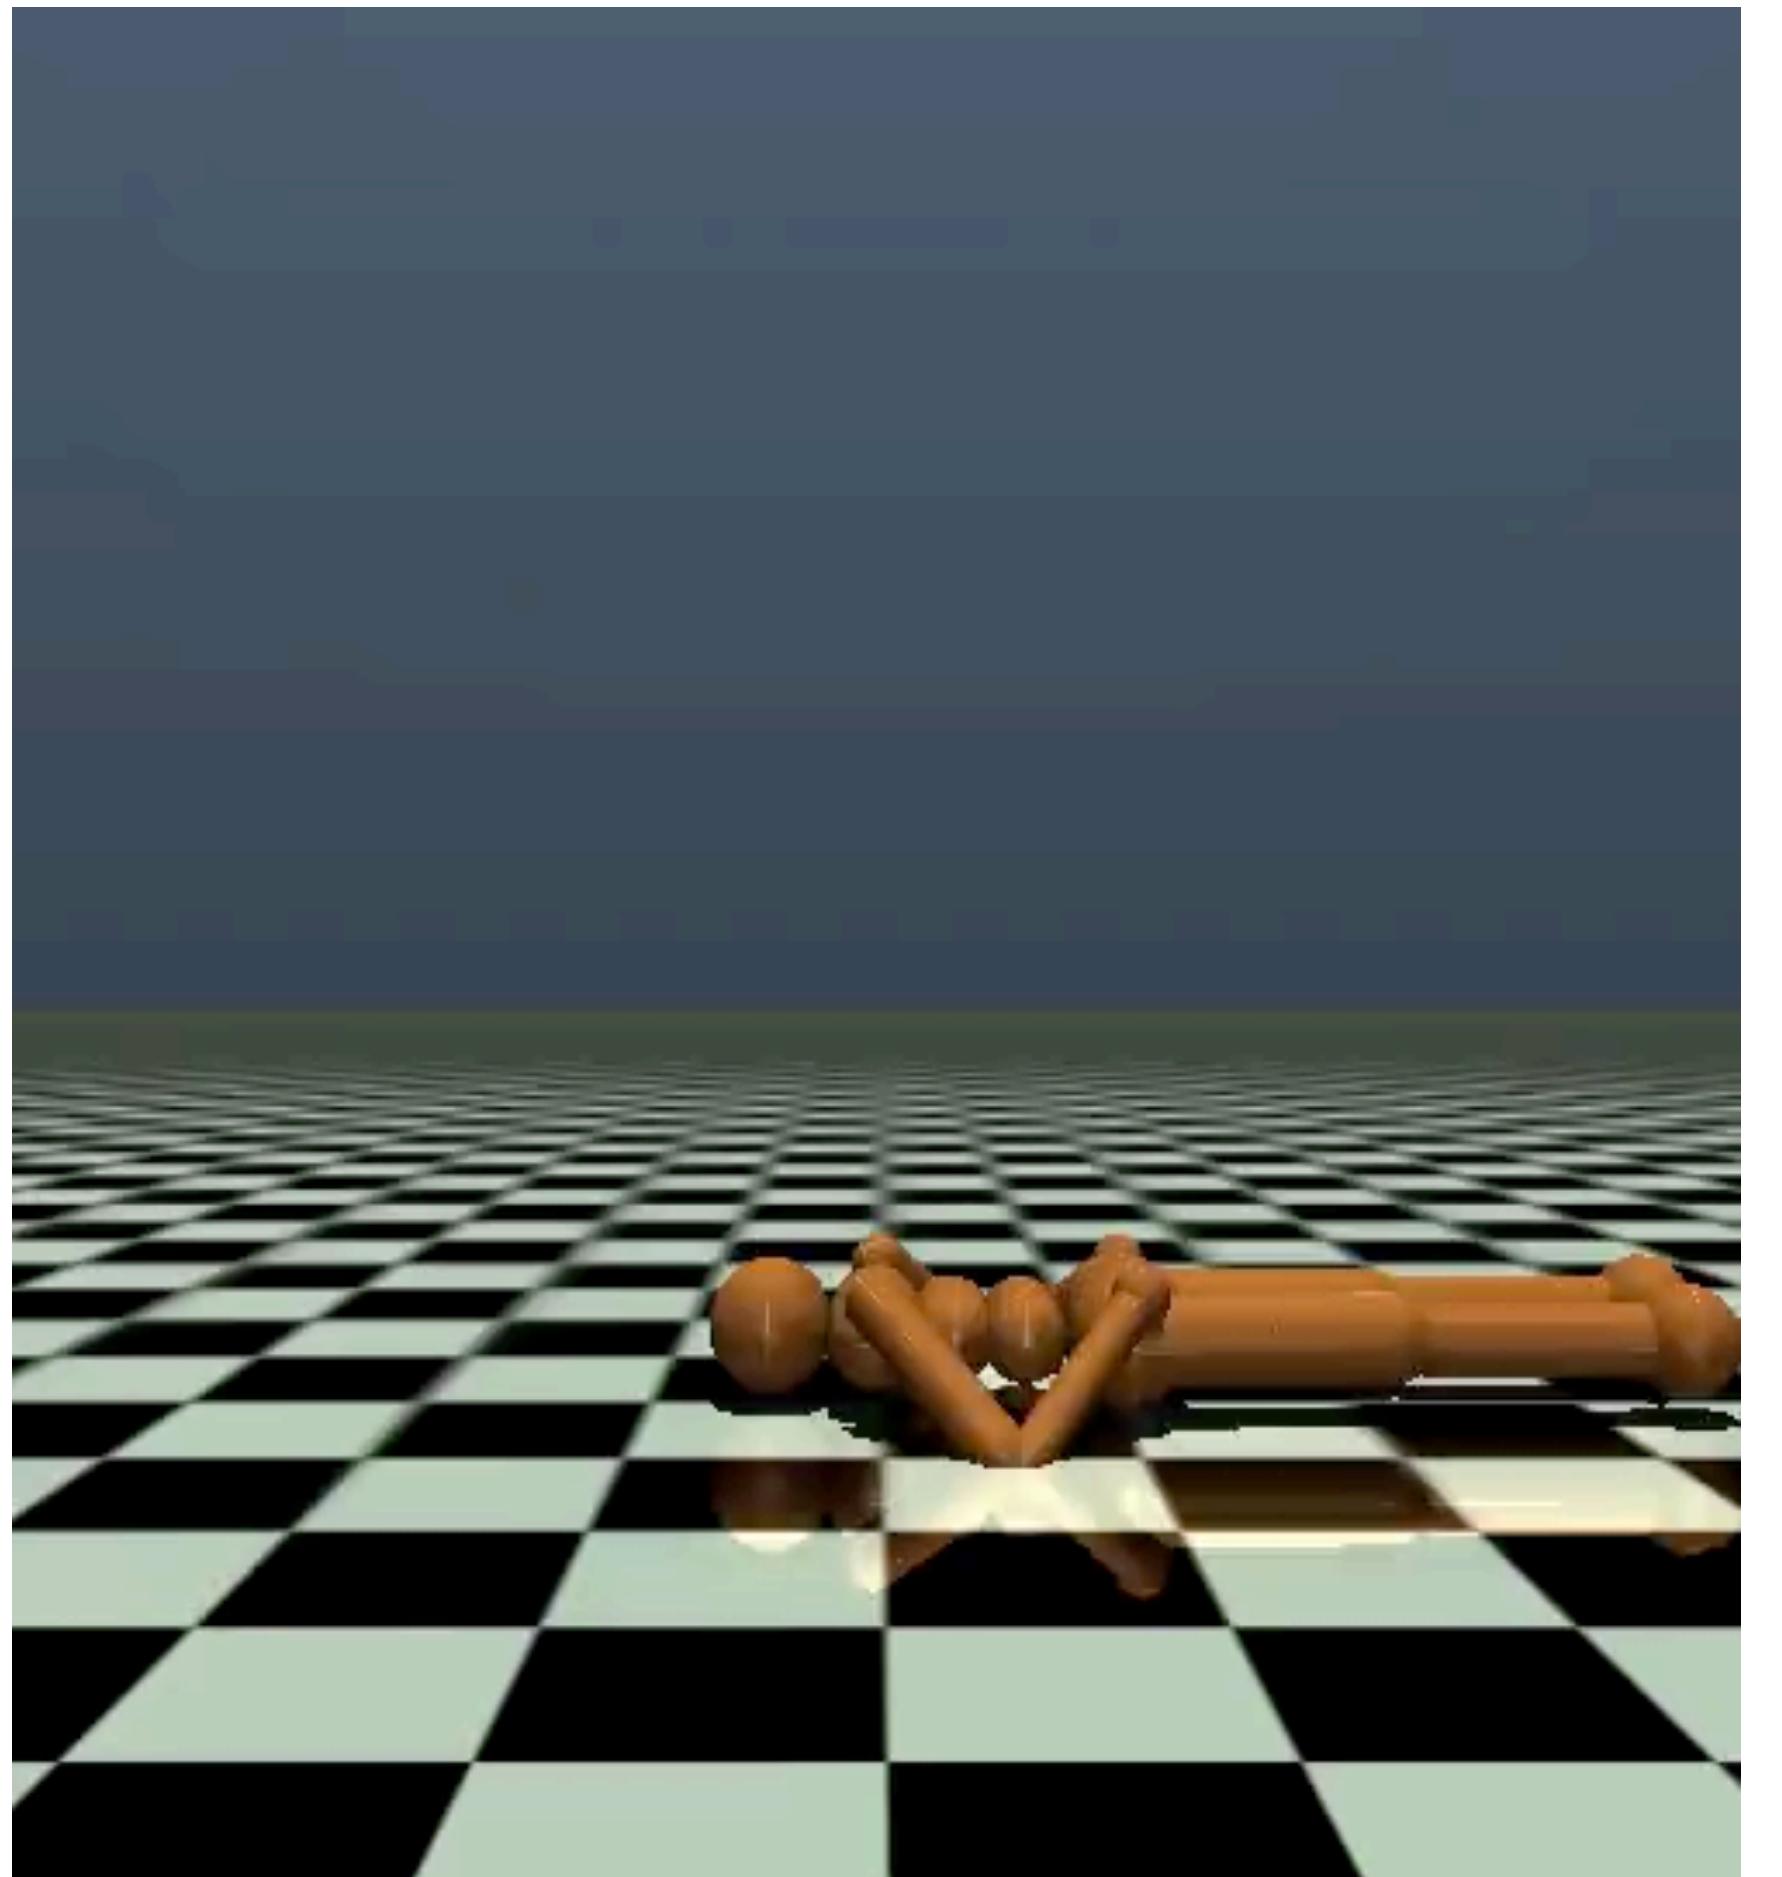
RL Formalisation

Agents and Environments

- Environment characteristics?
 - Episodic (finite-horizon) vs. Continuing (indefinite-horizon)
 - Deterministic vs. Stochastic
 - Fully vs. Partially observable
 - Discrete vs. Continuous



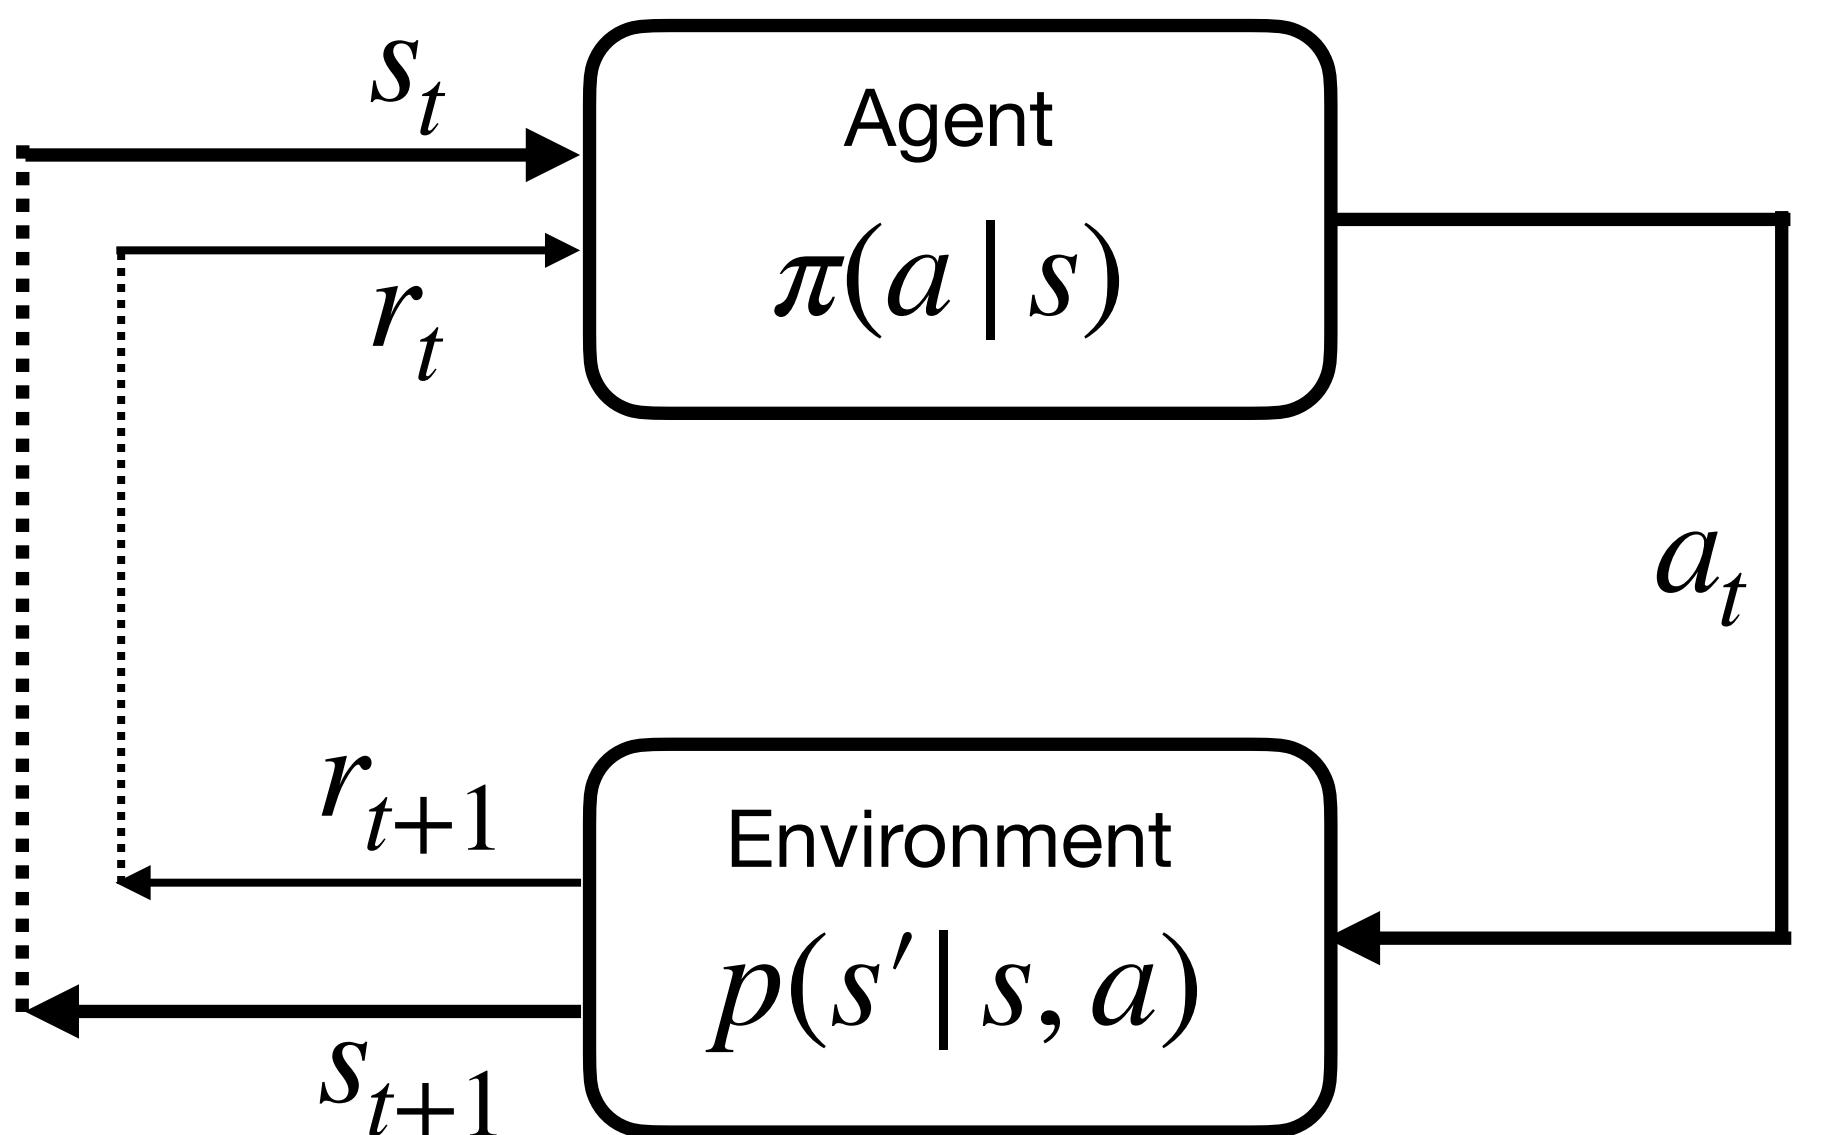
Agent-Environment Interaction [Sutton & Barto 2018]



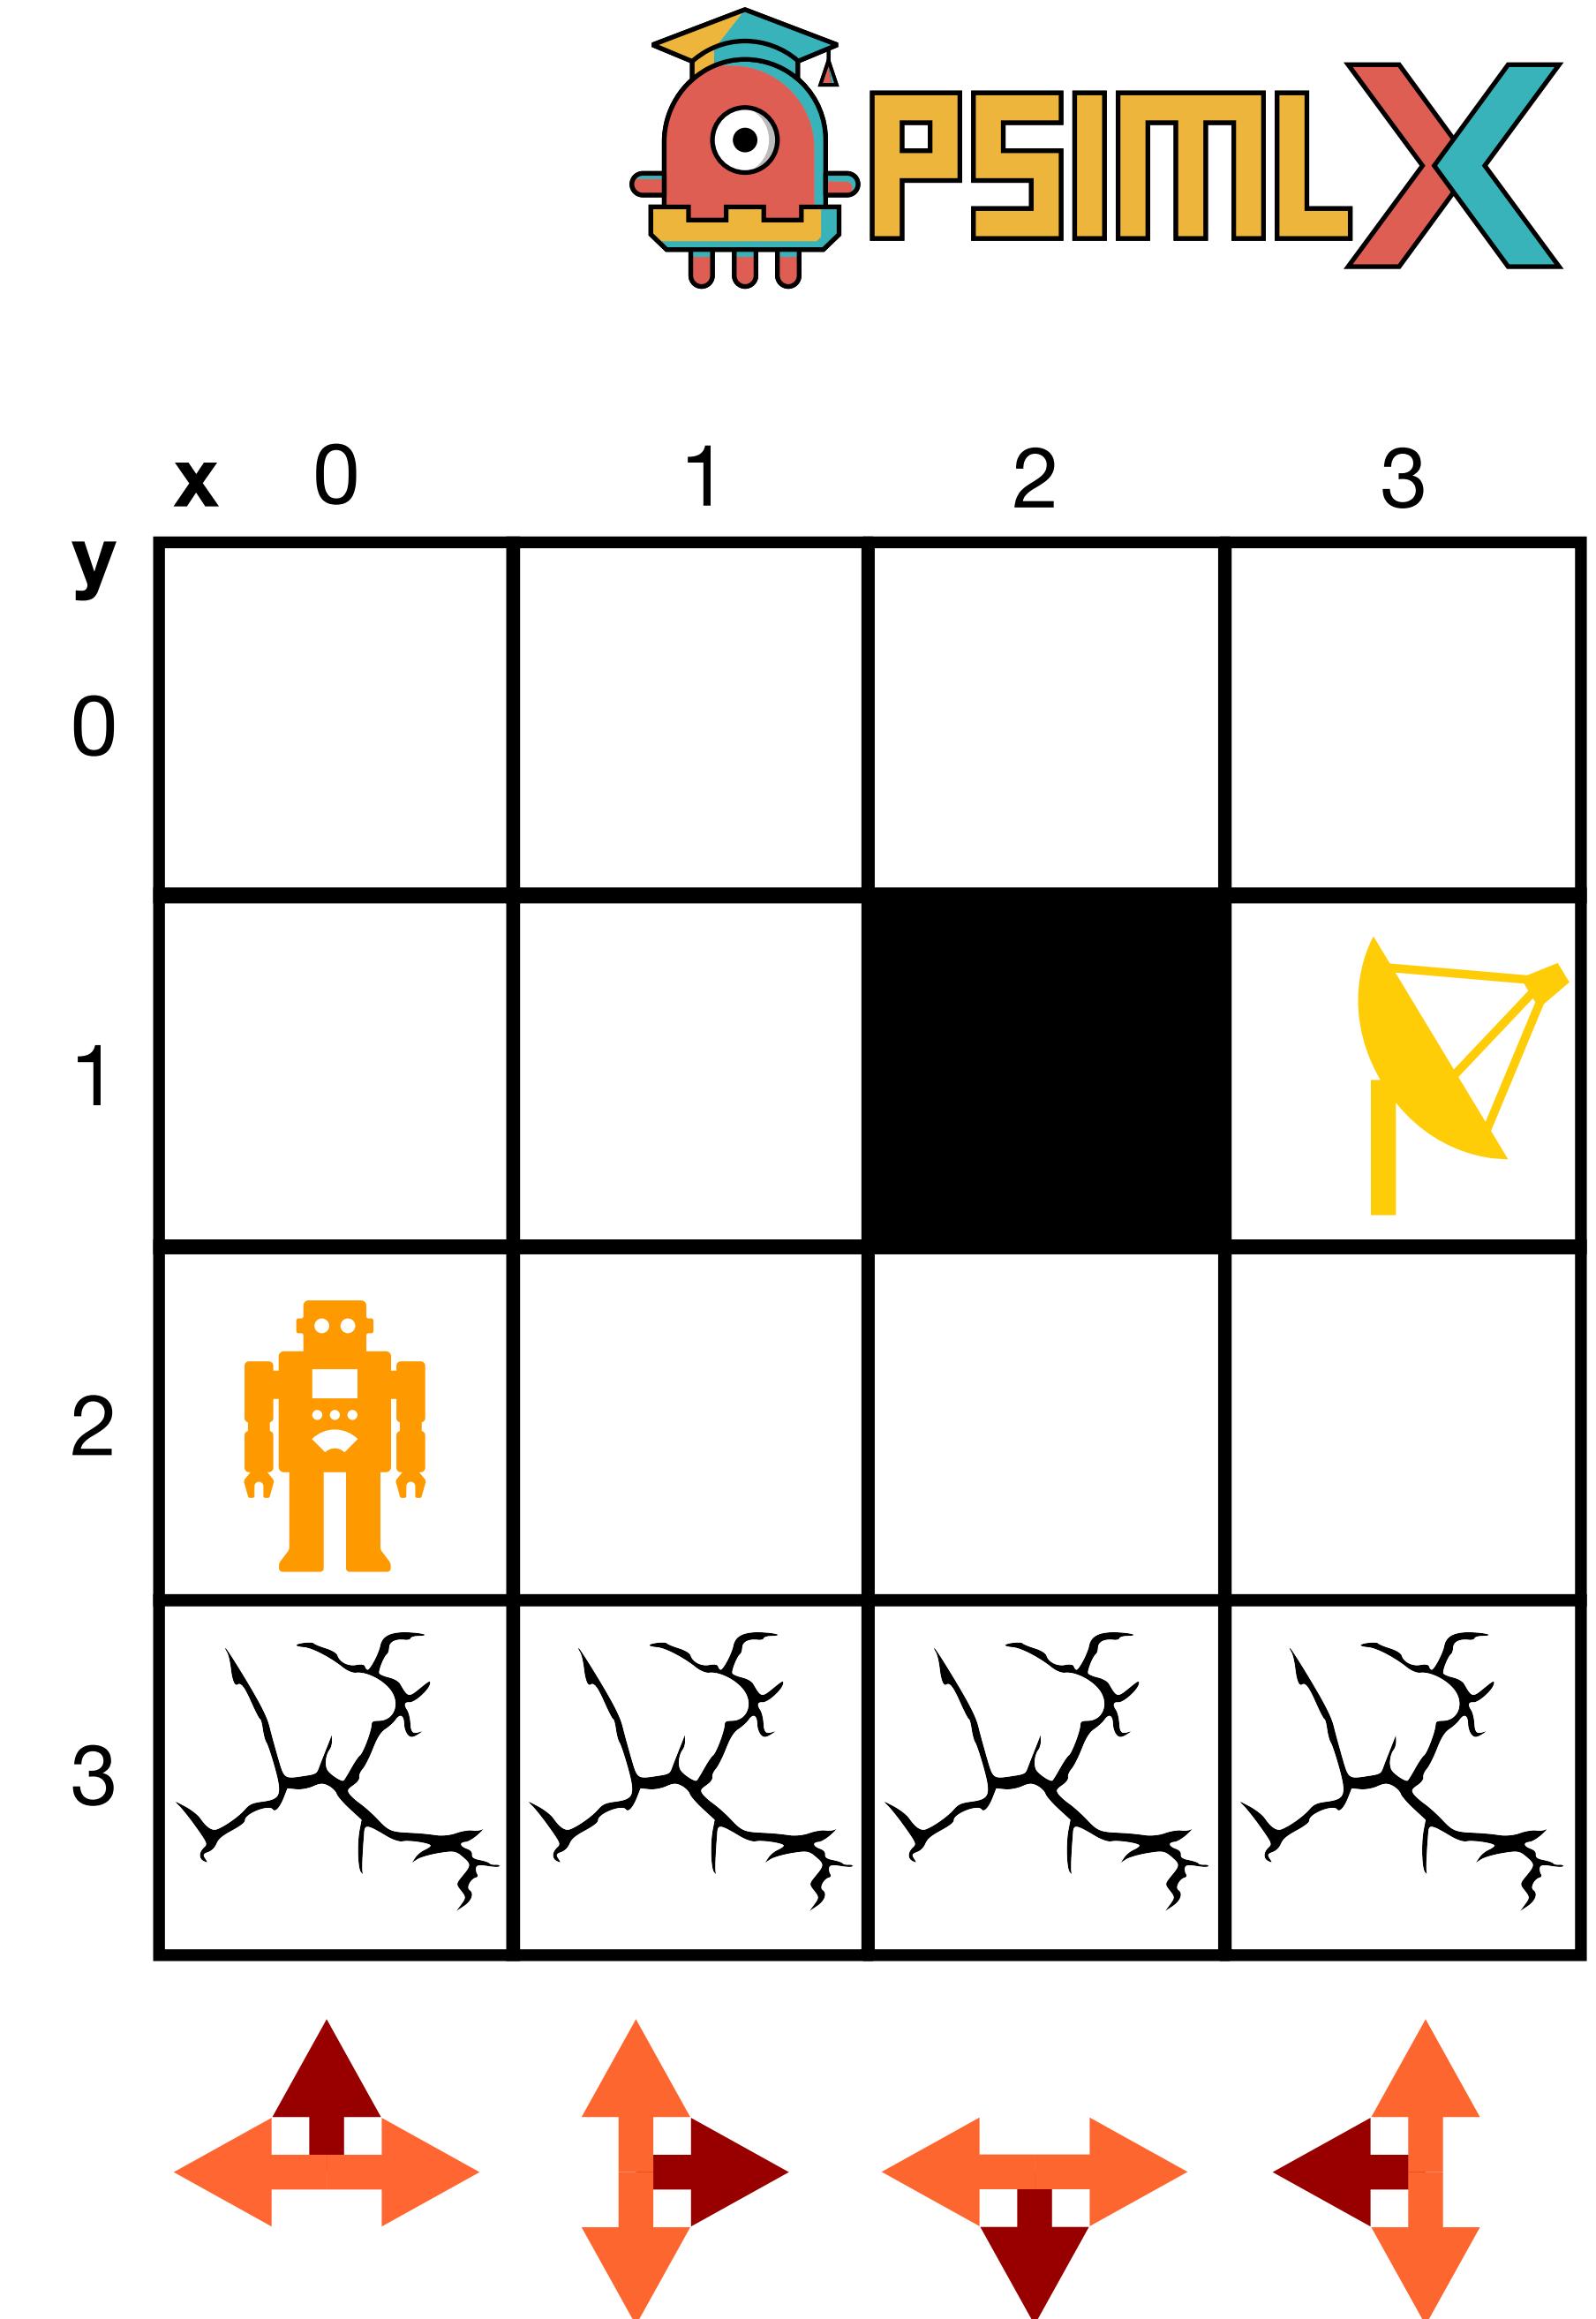
RL Formalisation

Agents and Environments

- Environment characteristics?
 - Stochastic, multiple terminal states
- How to find the optimal strategy?
 - Search algorithms (depth-first, A*, ...) not applicable



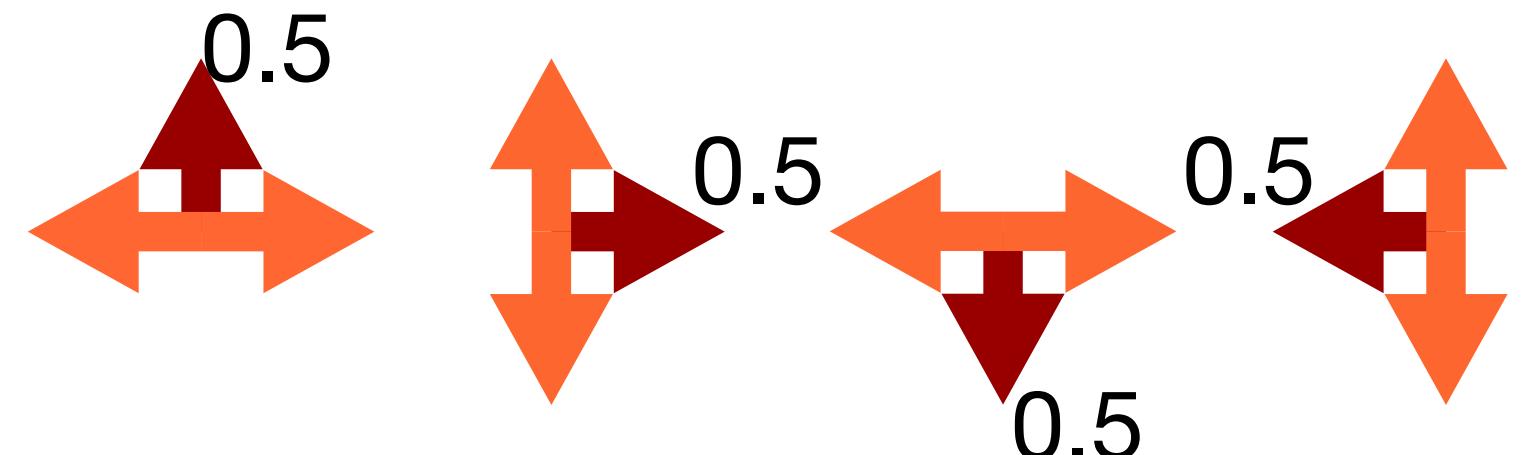
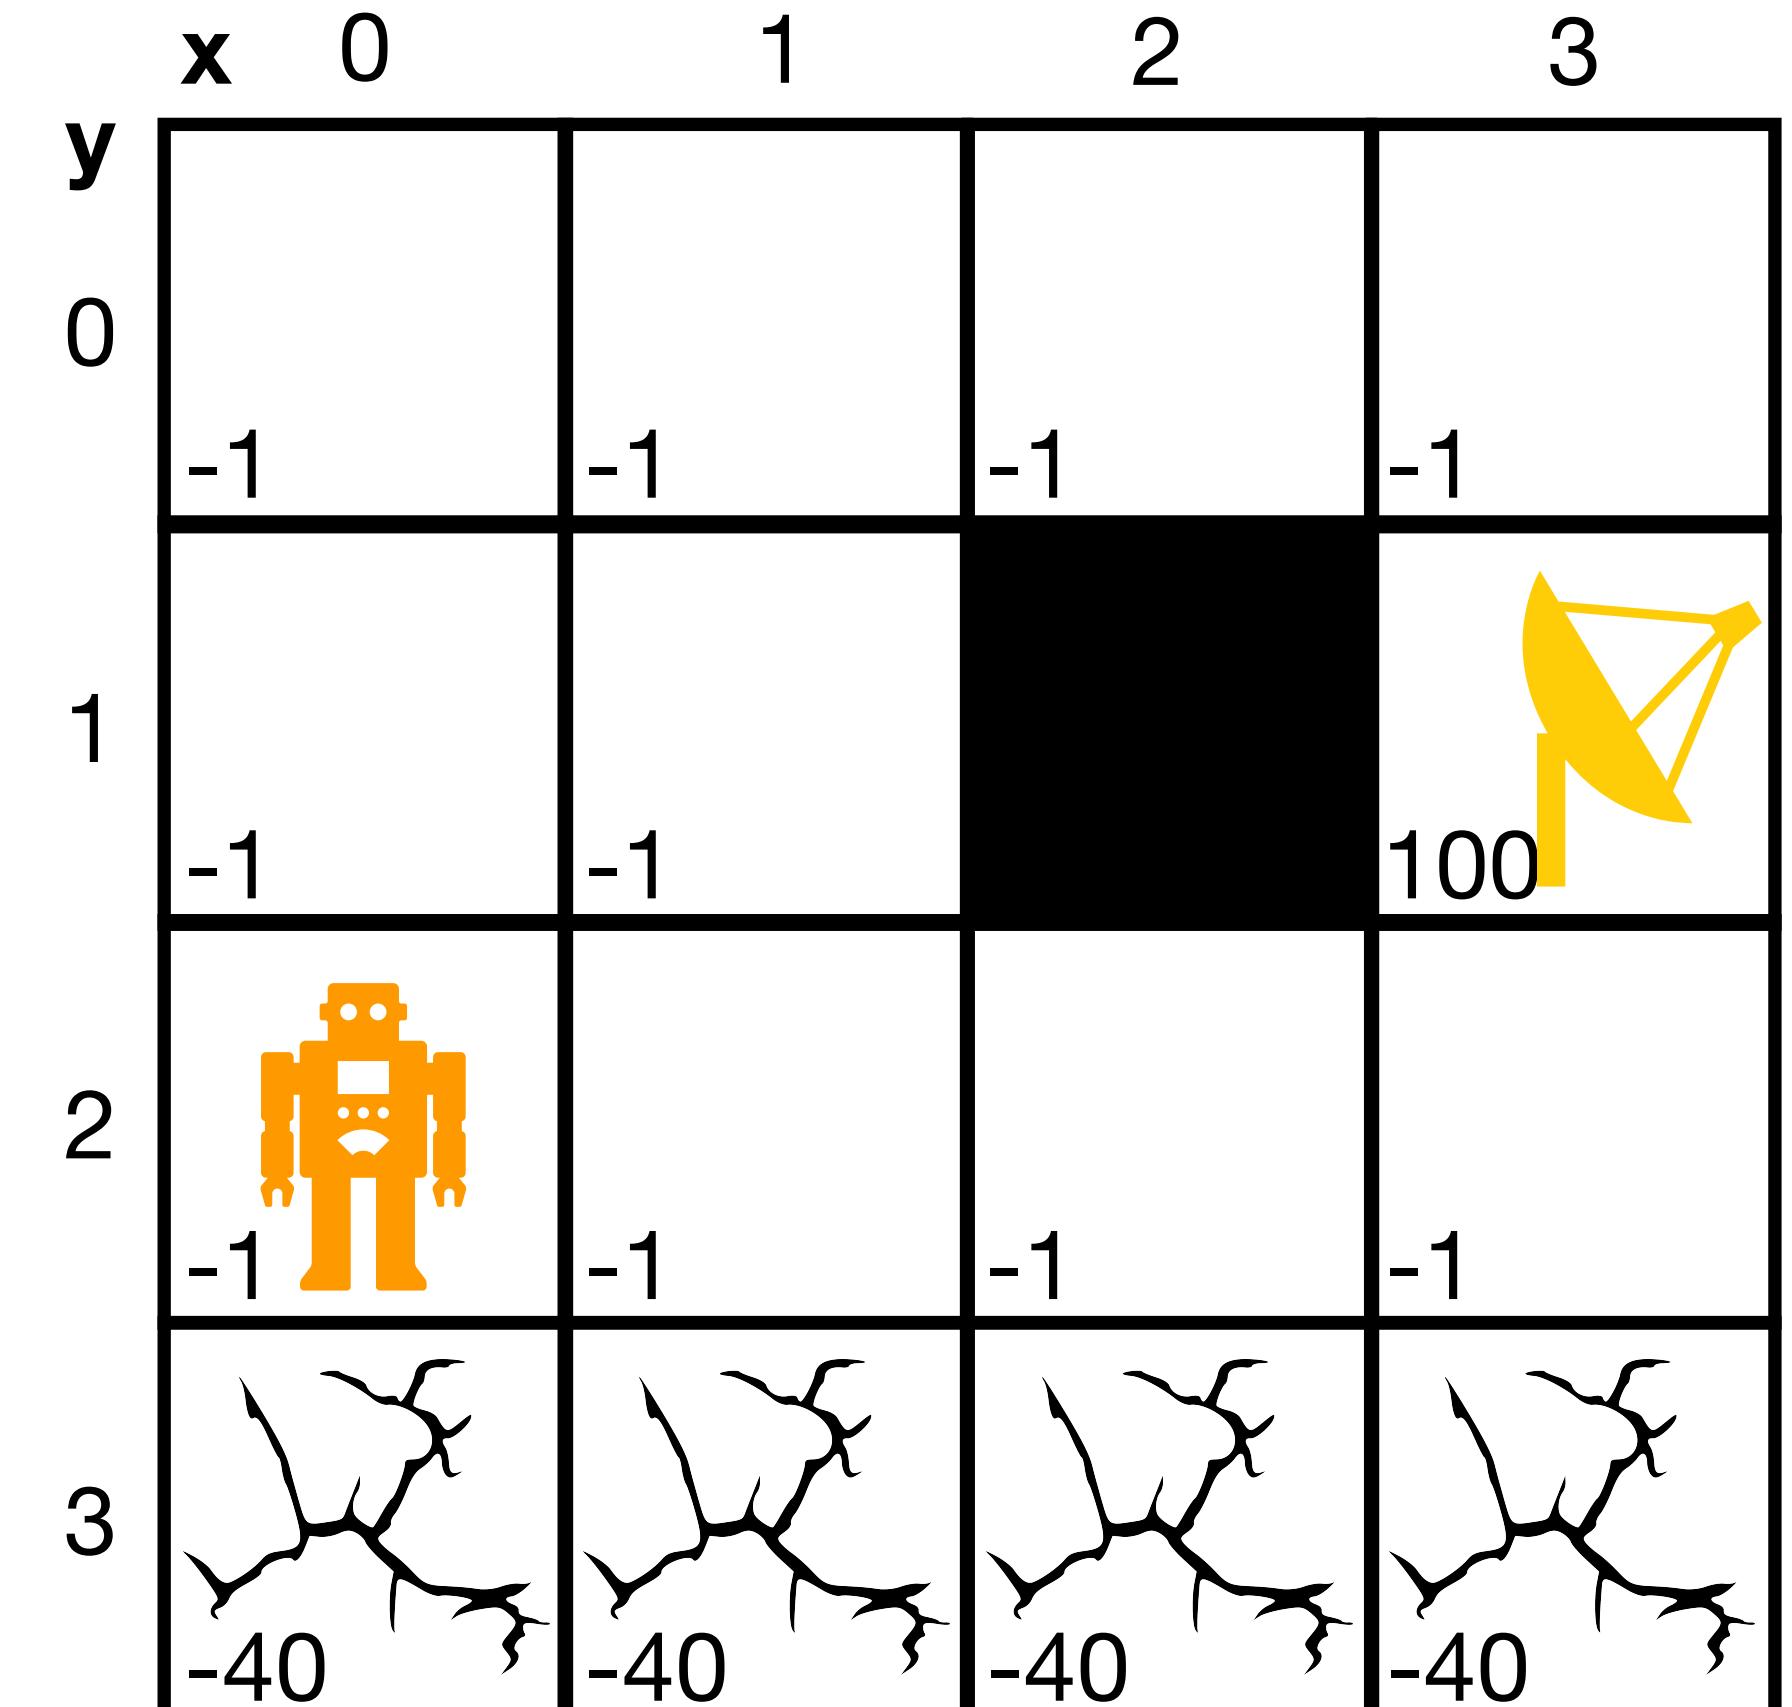
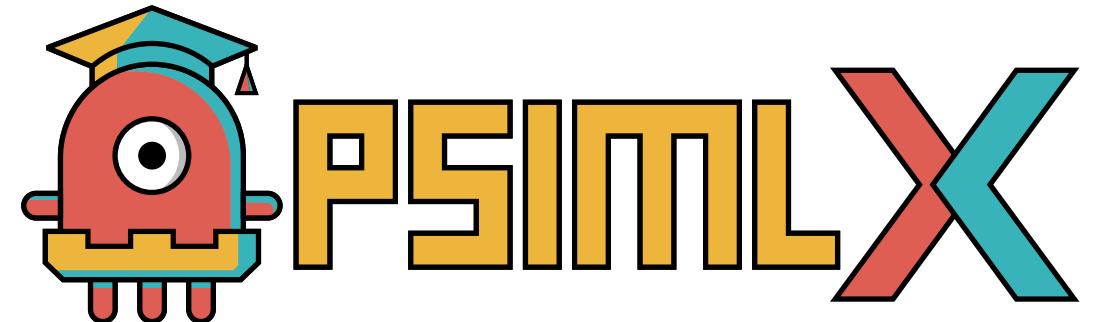
Agent-Environment Interaction [Sutton & Barto 2018]



RL Formalisation

Markov Decision Process

- Markov Decision Process (MDP) is a tuple (S, A, R) :
 - S – finite set of states
 - A – set of actions
 - R – set of rewards
- $p(s' | s, a)$ – MDP transition model
- $r_s^a \doteq \mathbb{E} [R_{t+1} | S_t = s, A_t = a]$
- $= \sum_{s'} p(s' | s, a) r_{ss'}^a$ one step expected reward



RL Formalisation

Markov Decision Process

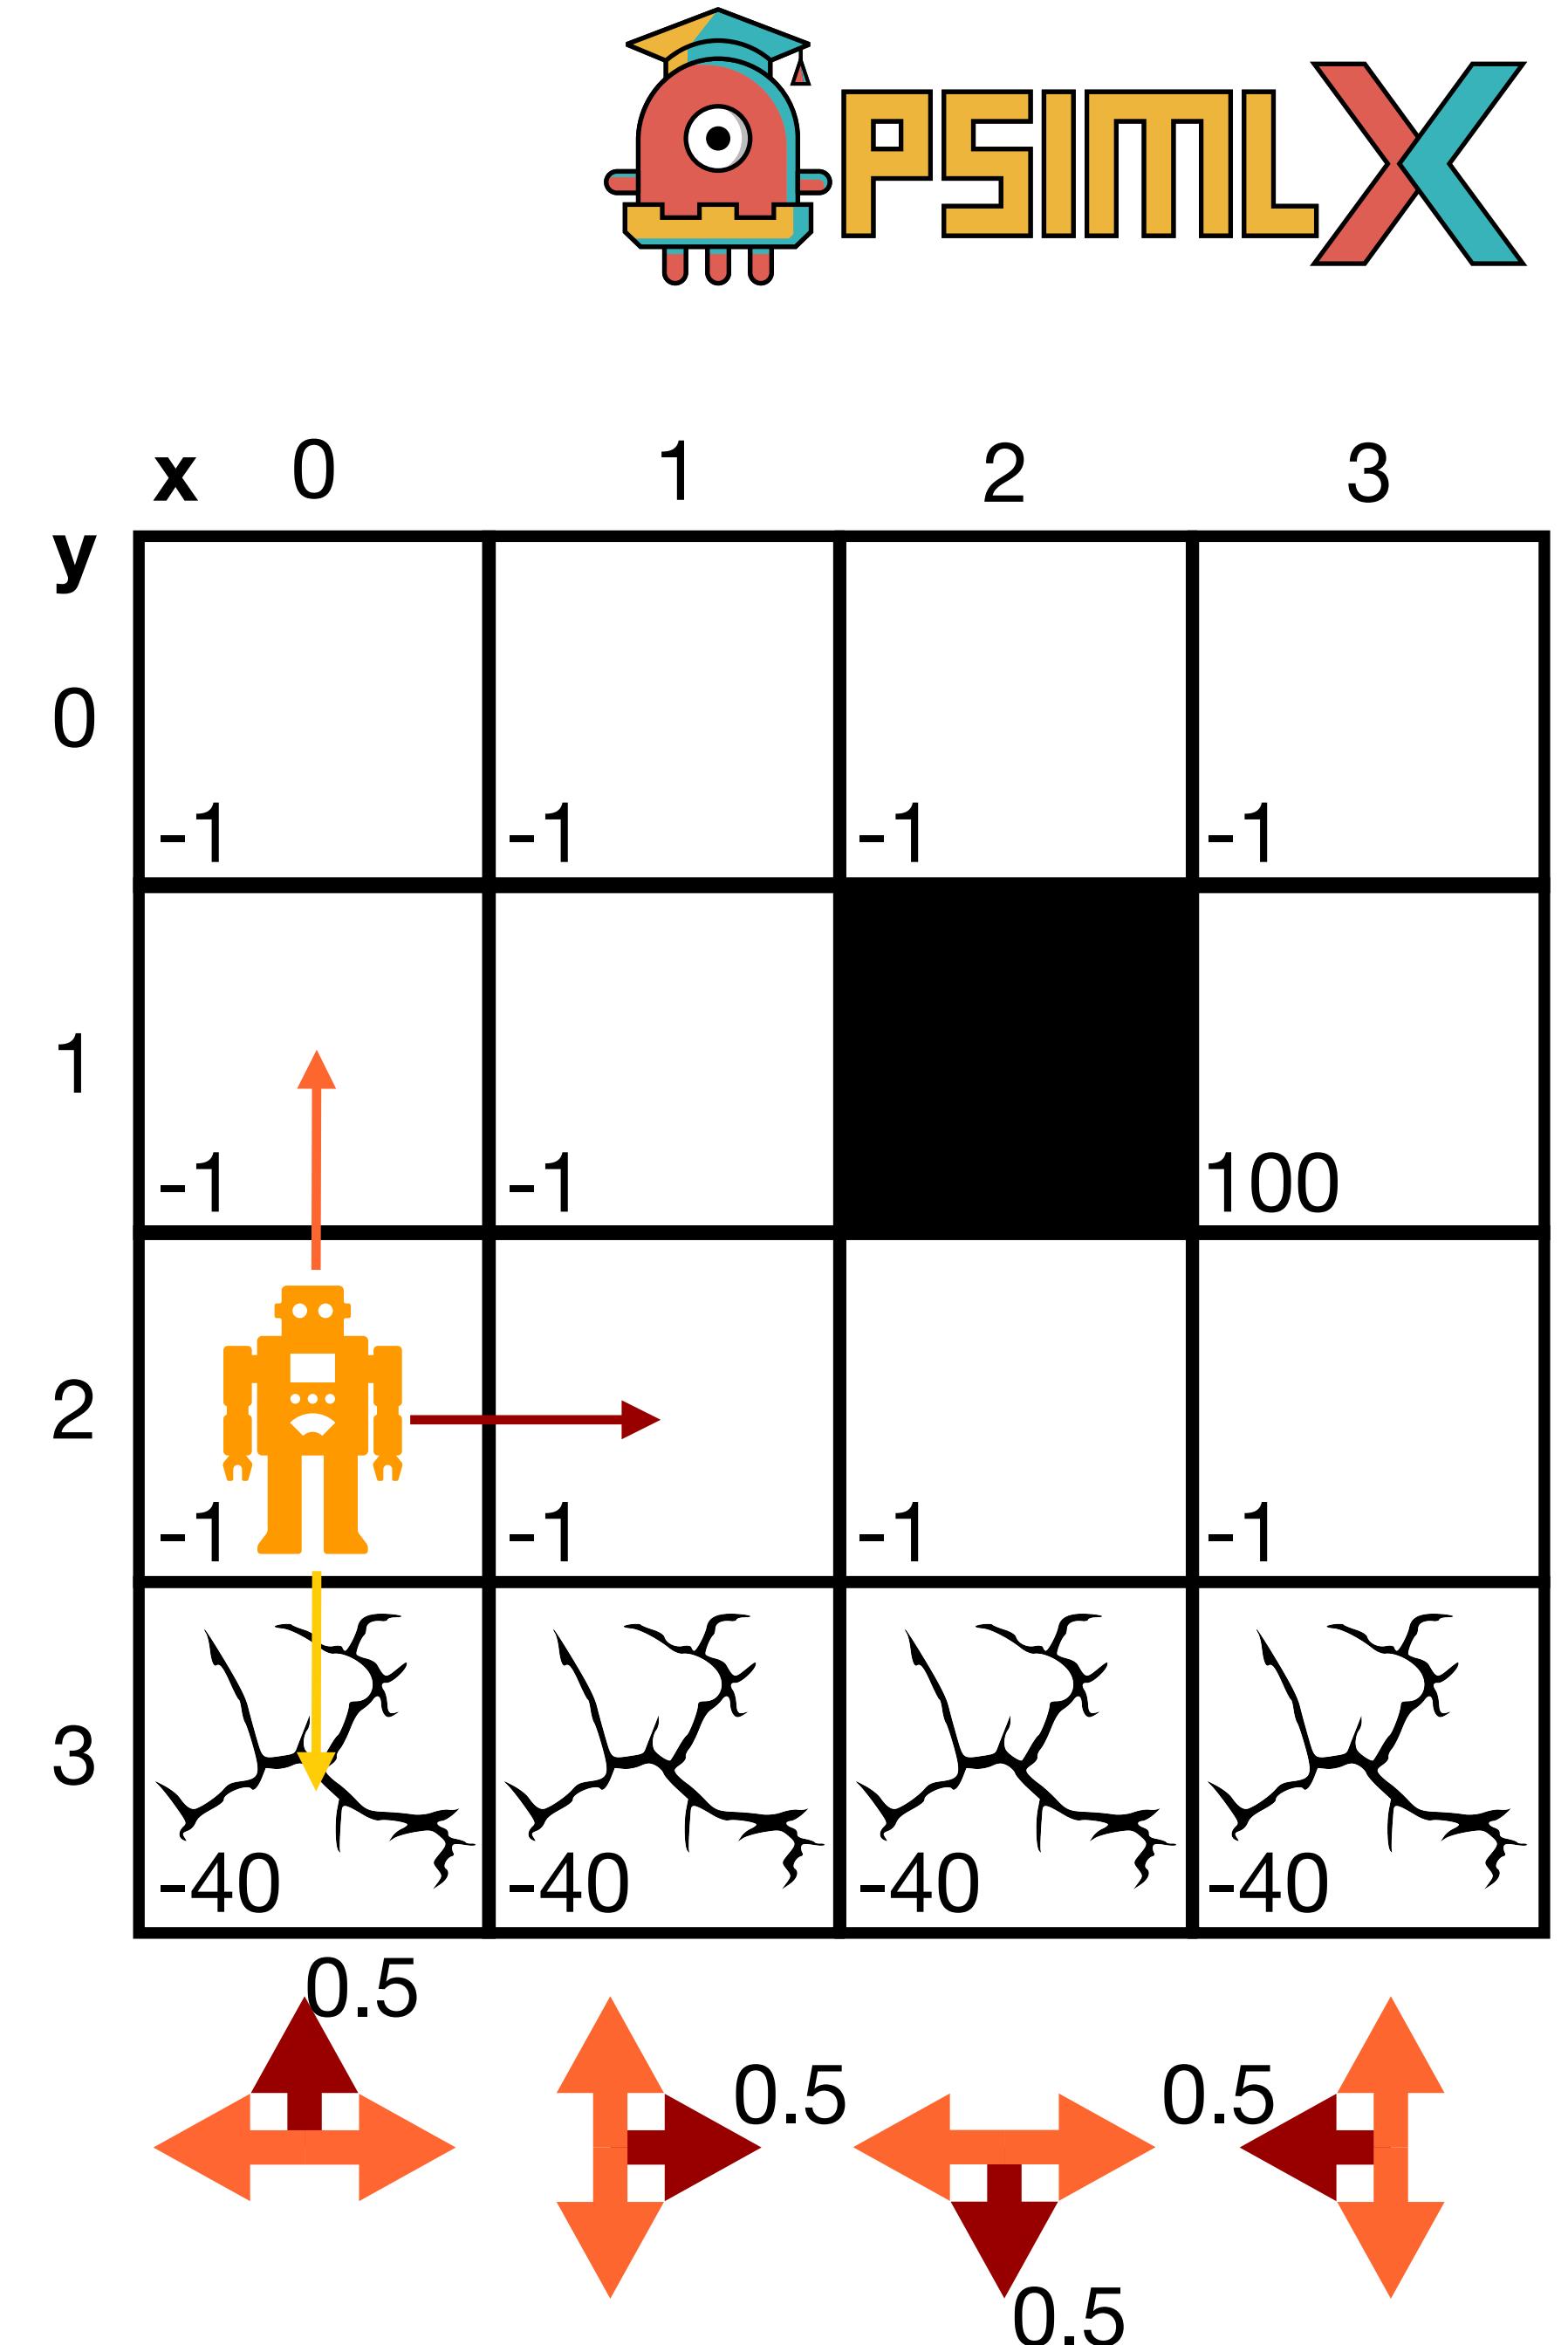
- Markov Decision Process (MDP) is a tuple (S, A, R) :
 - S – finite set of states
 - A – set of actions
 - R – set of rewards
- $p(s' | s, a)$ – MDP transition model
- $r_s^a \doteq \mathbb{E} [R_{t+1} | S_t = s, A_t = a]$
 $= \sum_{s'} p(s' | s, a)r$ one step expected reward

$(S_0 = (0,2), A_0 = \text{RIGHT}, R_1 = -1), (S_1 = (1,2), \dots), \dots$

$(S_0 = (0,2), A_0 = \text{RIGHT}, R_1 = -1), (S_1 = (0,1), \dots), \dots$

$(S_0 = (0,2), A_0 = \text{RIGHT}, R_1 = -40), \quad \text{skull}$

$$r_{(0,2)}^{\text{RIGHT}} = 0.5 \cdot (-1) + 0.25 \cdot (-1) + 0.25 \cdot (-40) = -10.75$$

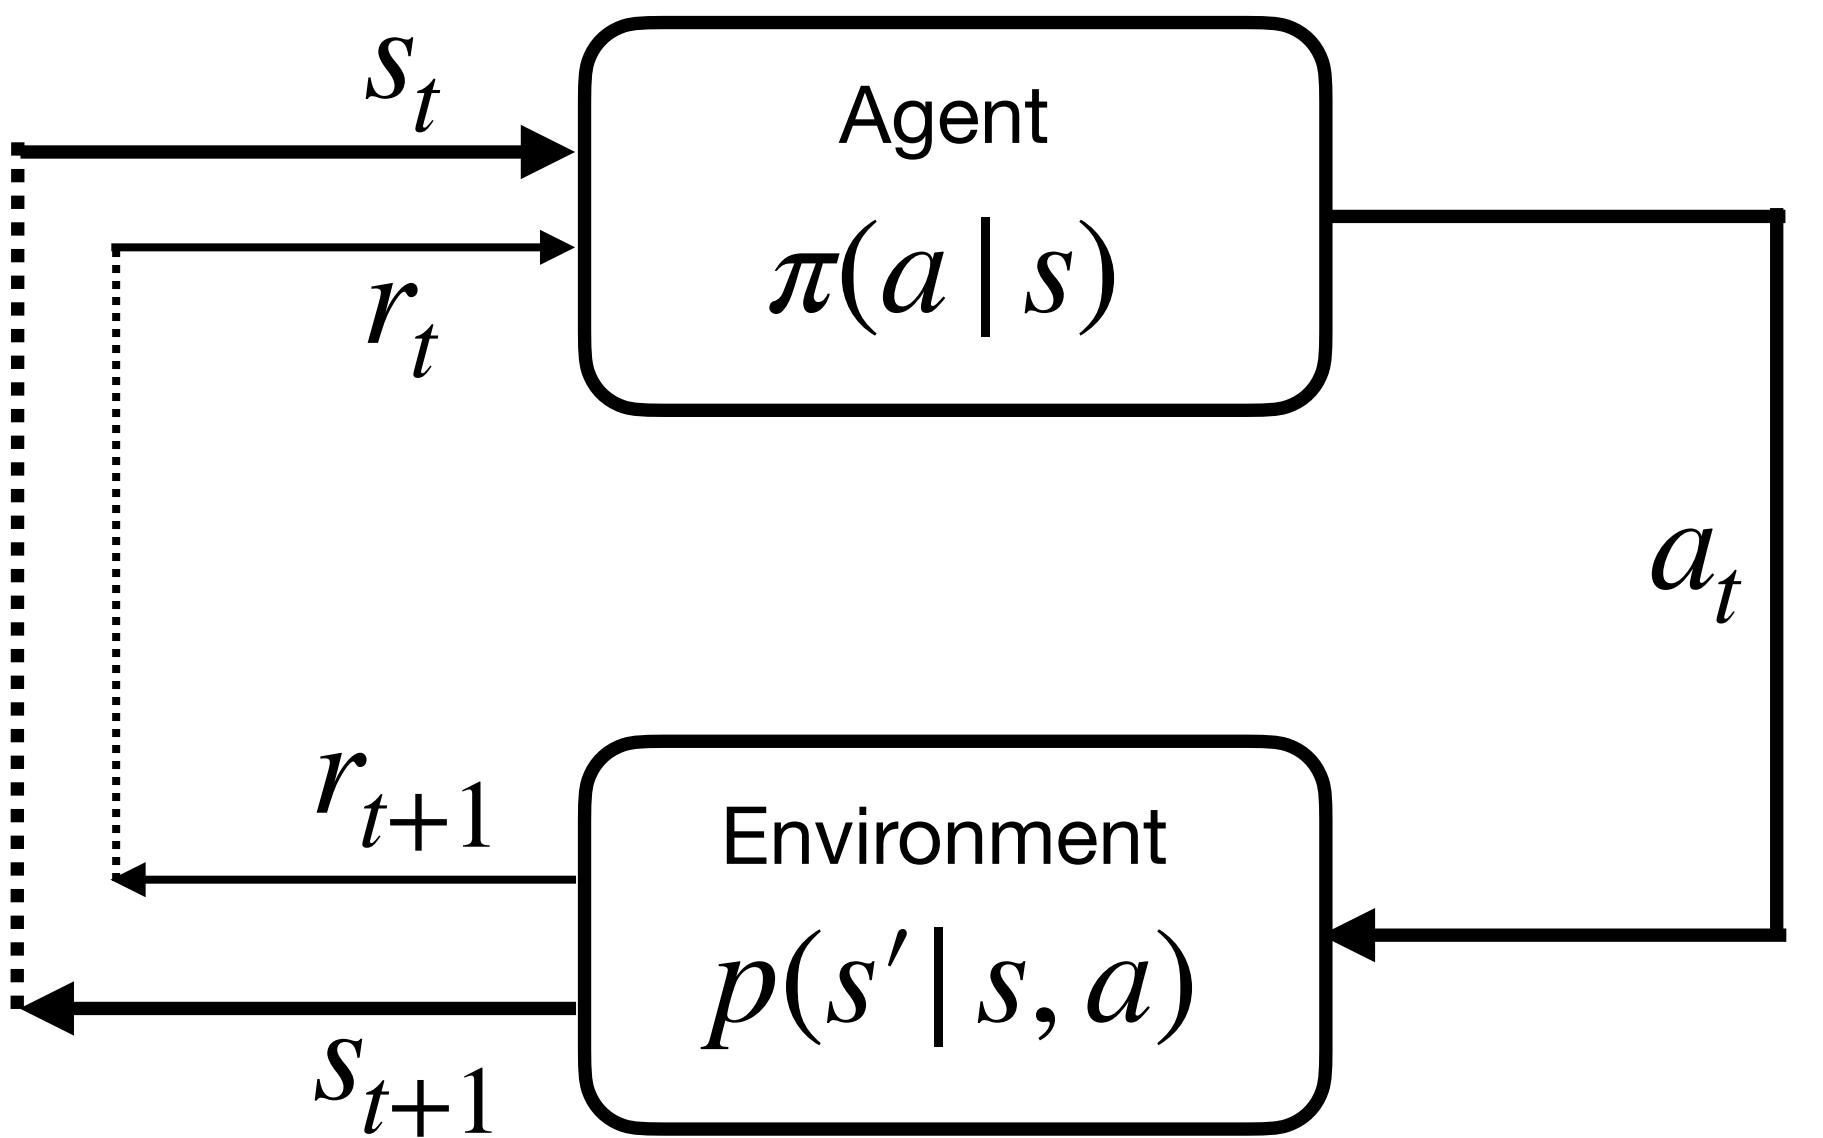


RL Formalisation

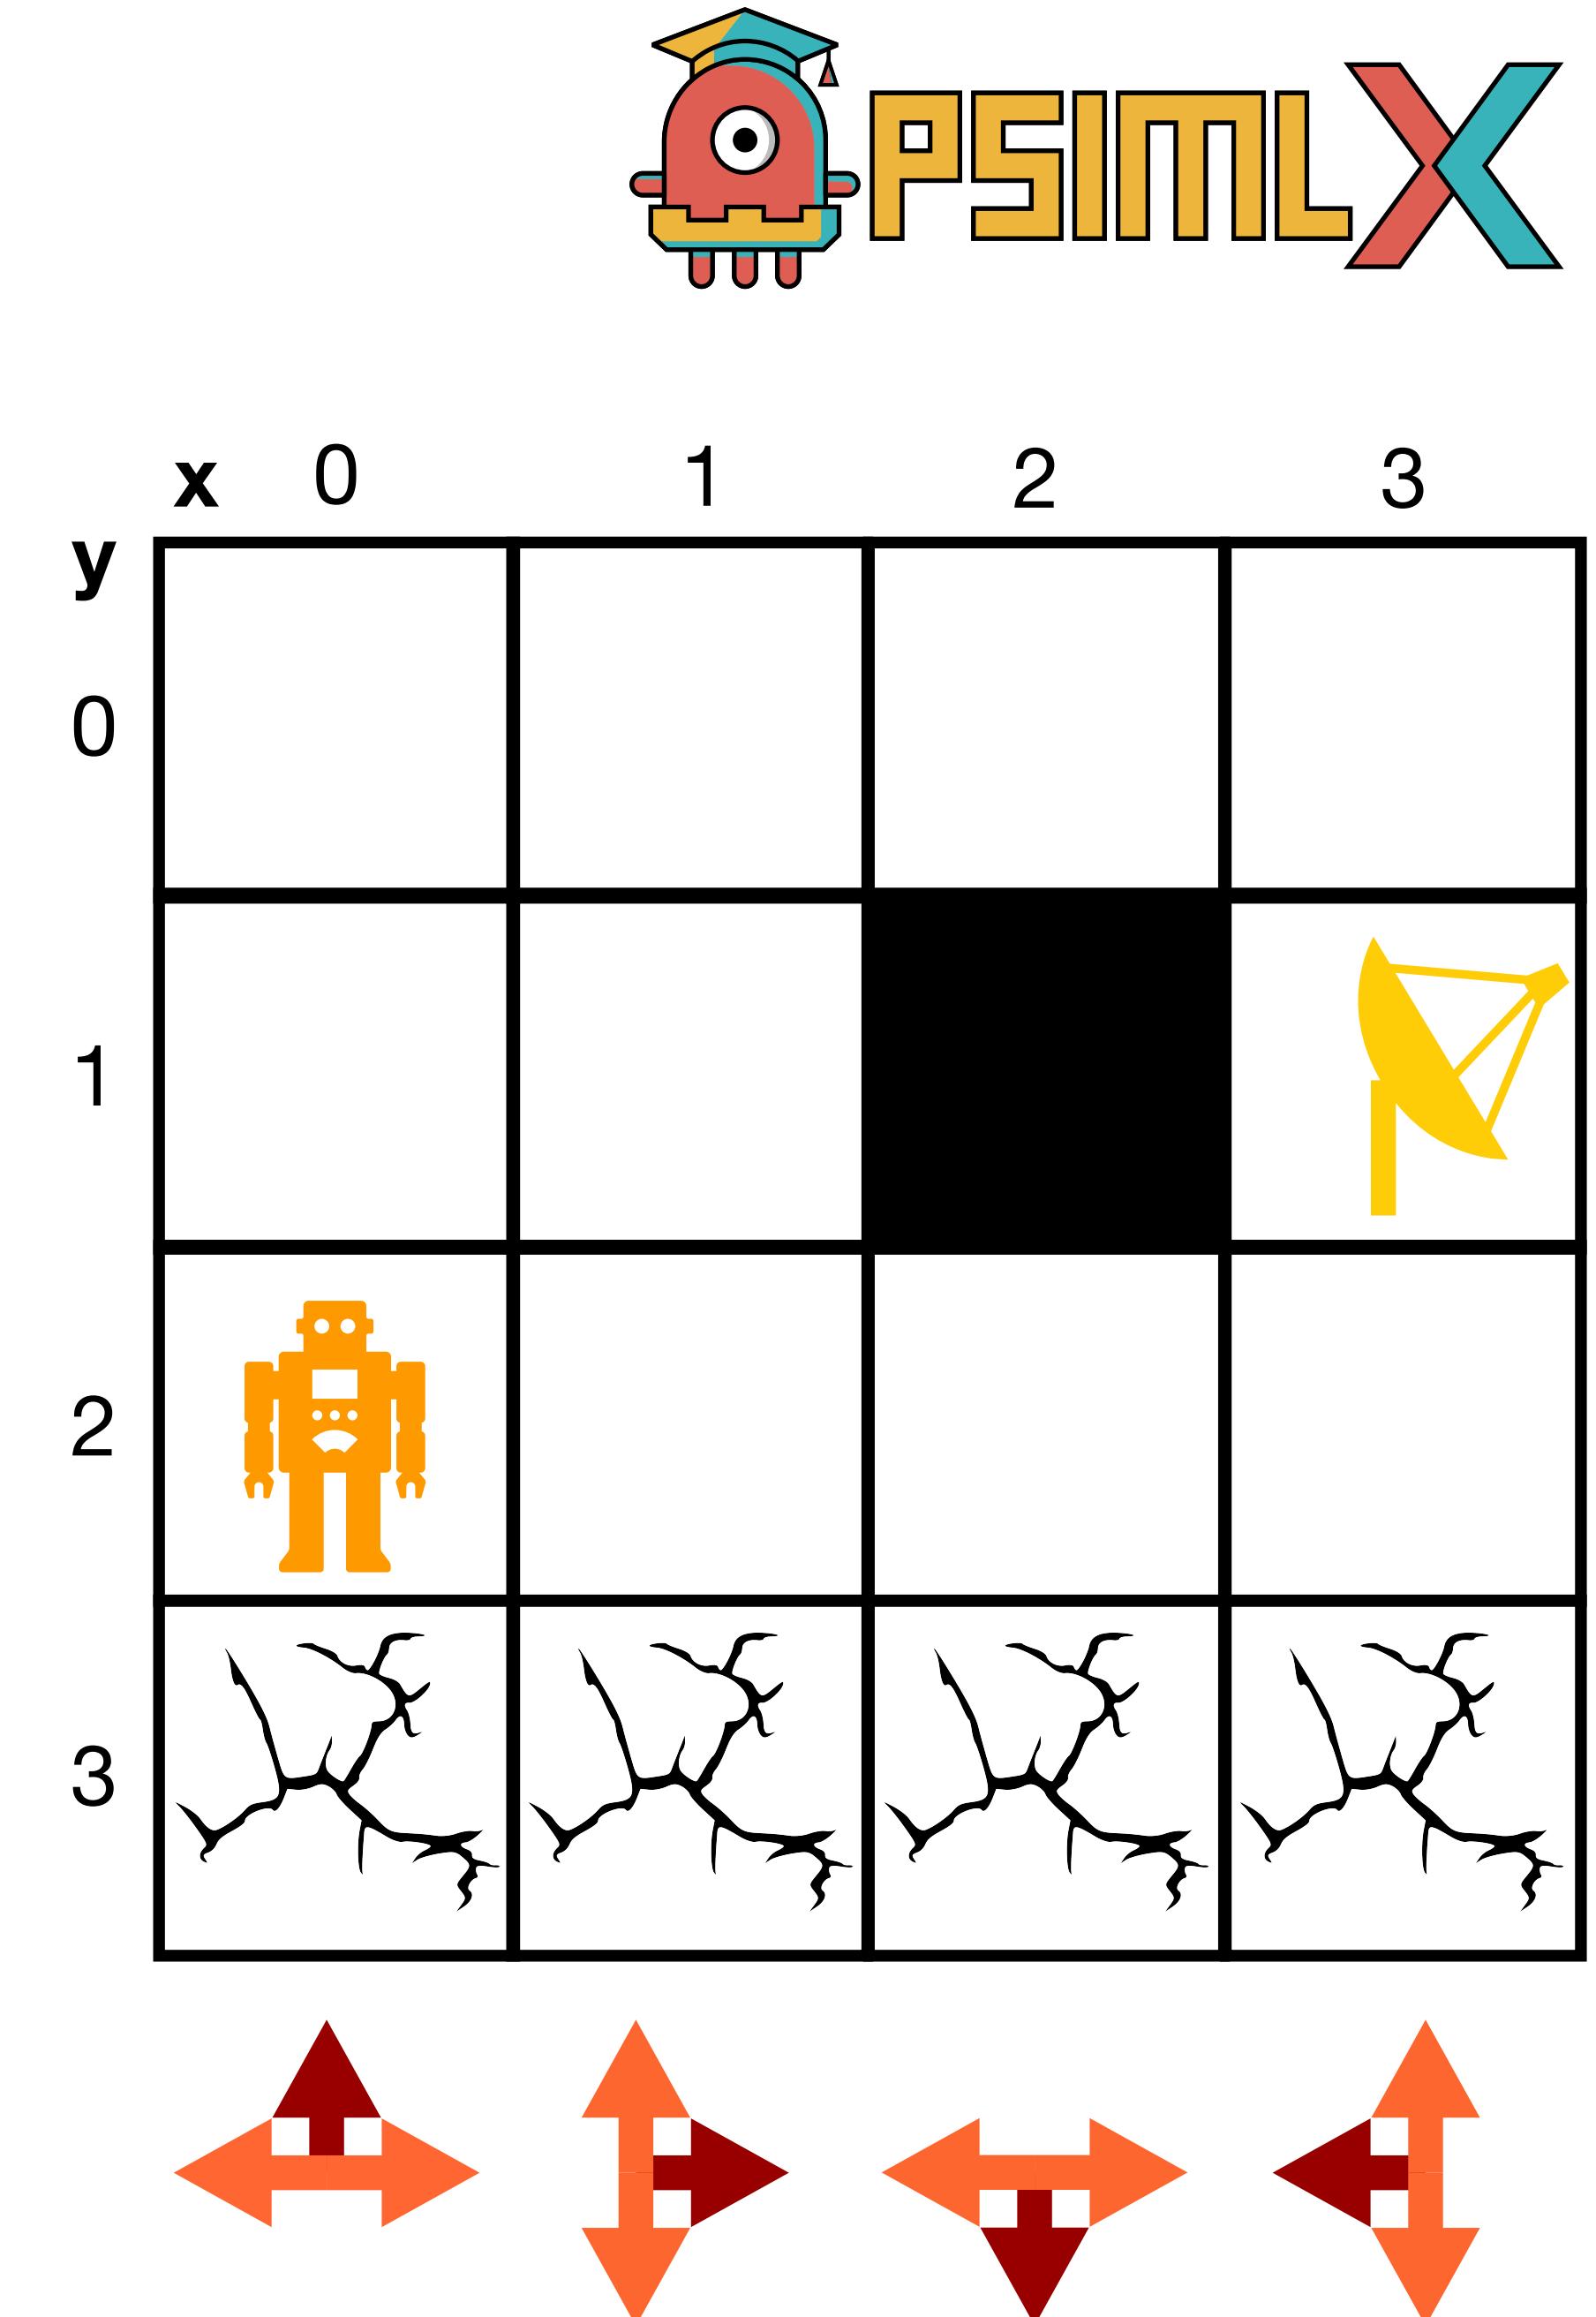
Agents and Environments

- Agent-environment interaction forms a *trajectory*:

$$\tau = (S_0, A_0, R_1), (S_1, A_1, R_2), \dots, (S_t, A_t, R_{t+1}), \dots$$



Agent-Environment Interaction [Sutton & Barto 2018]



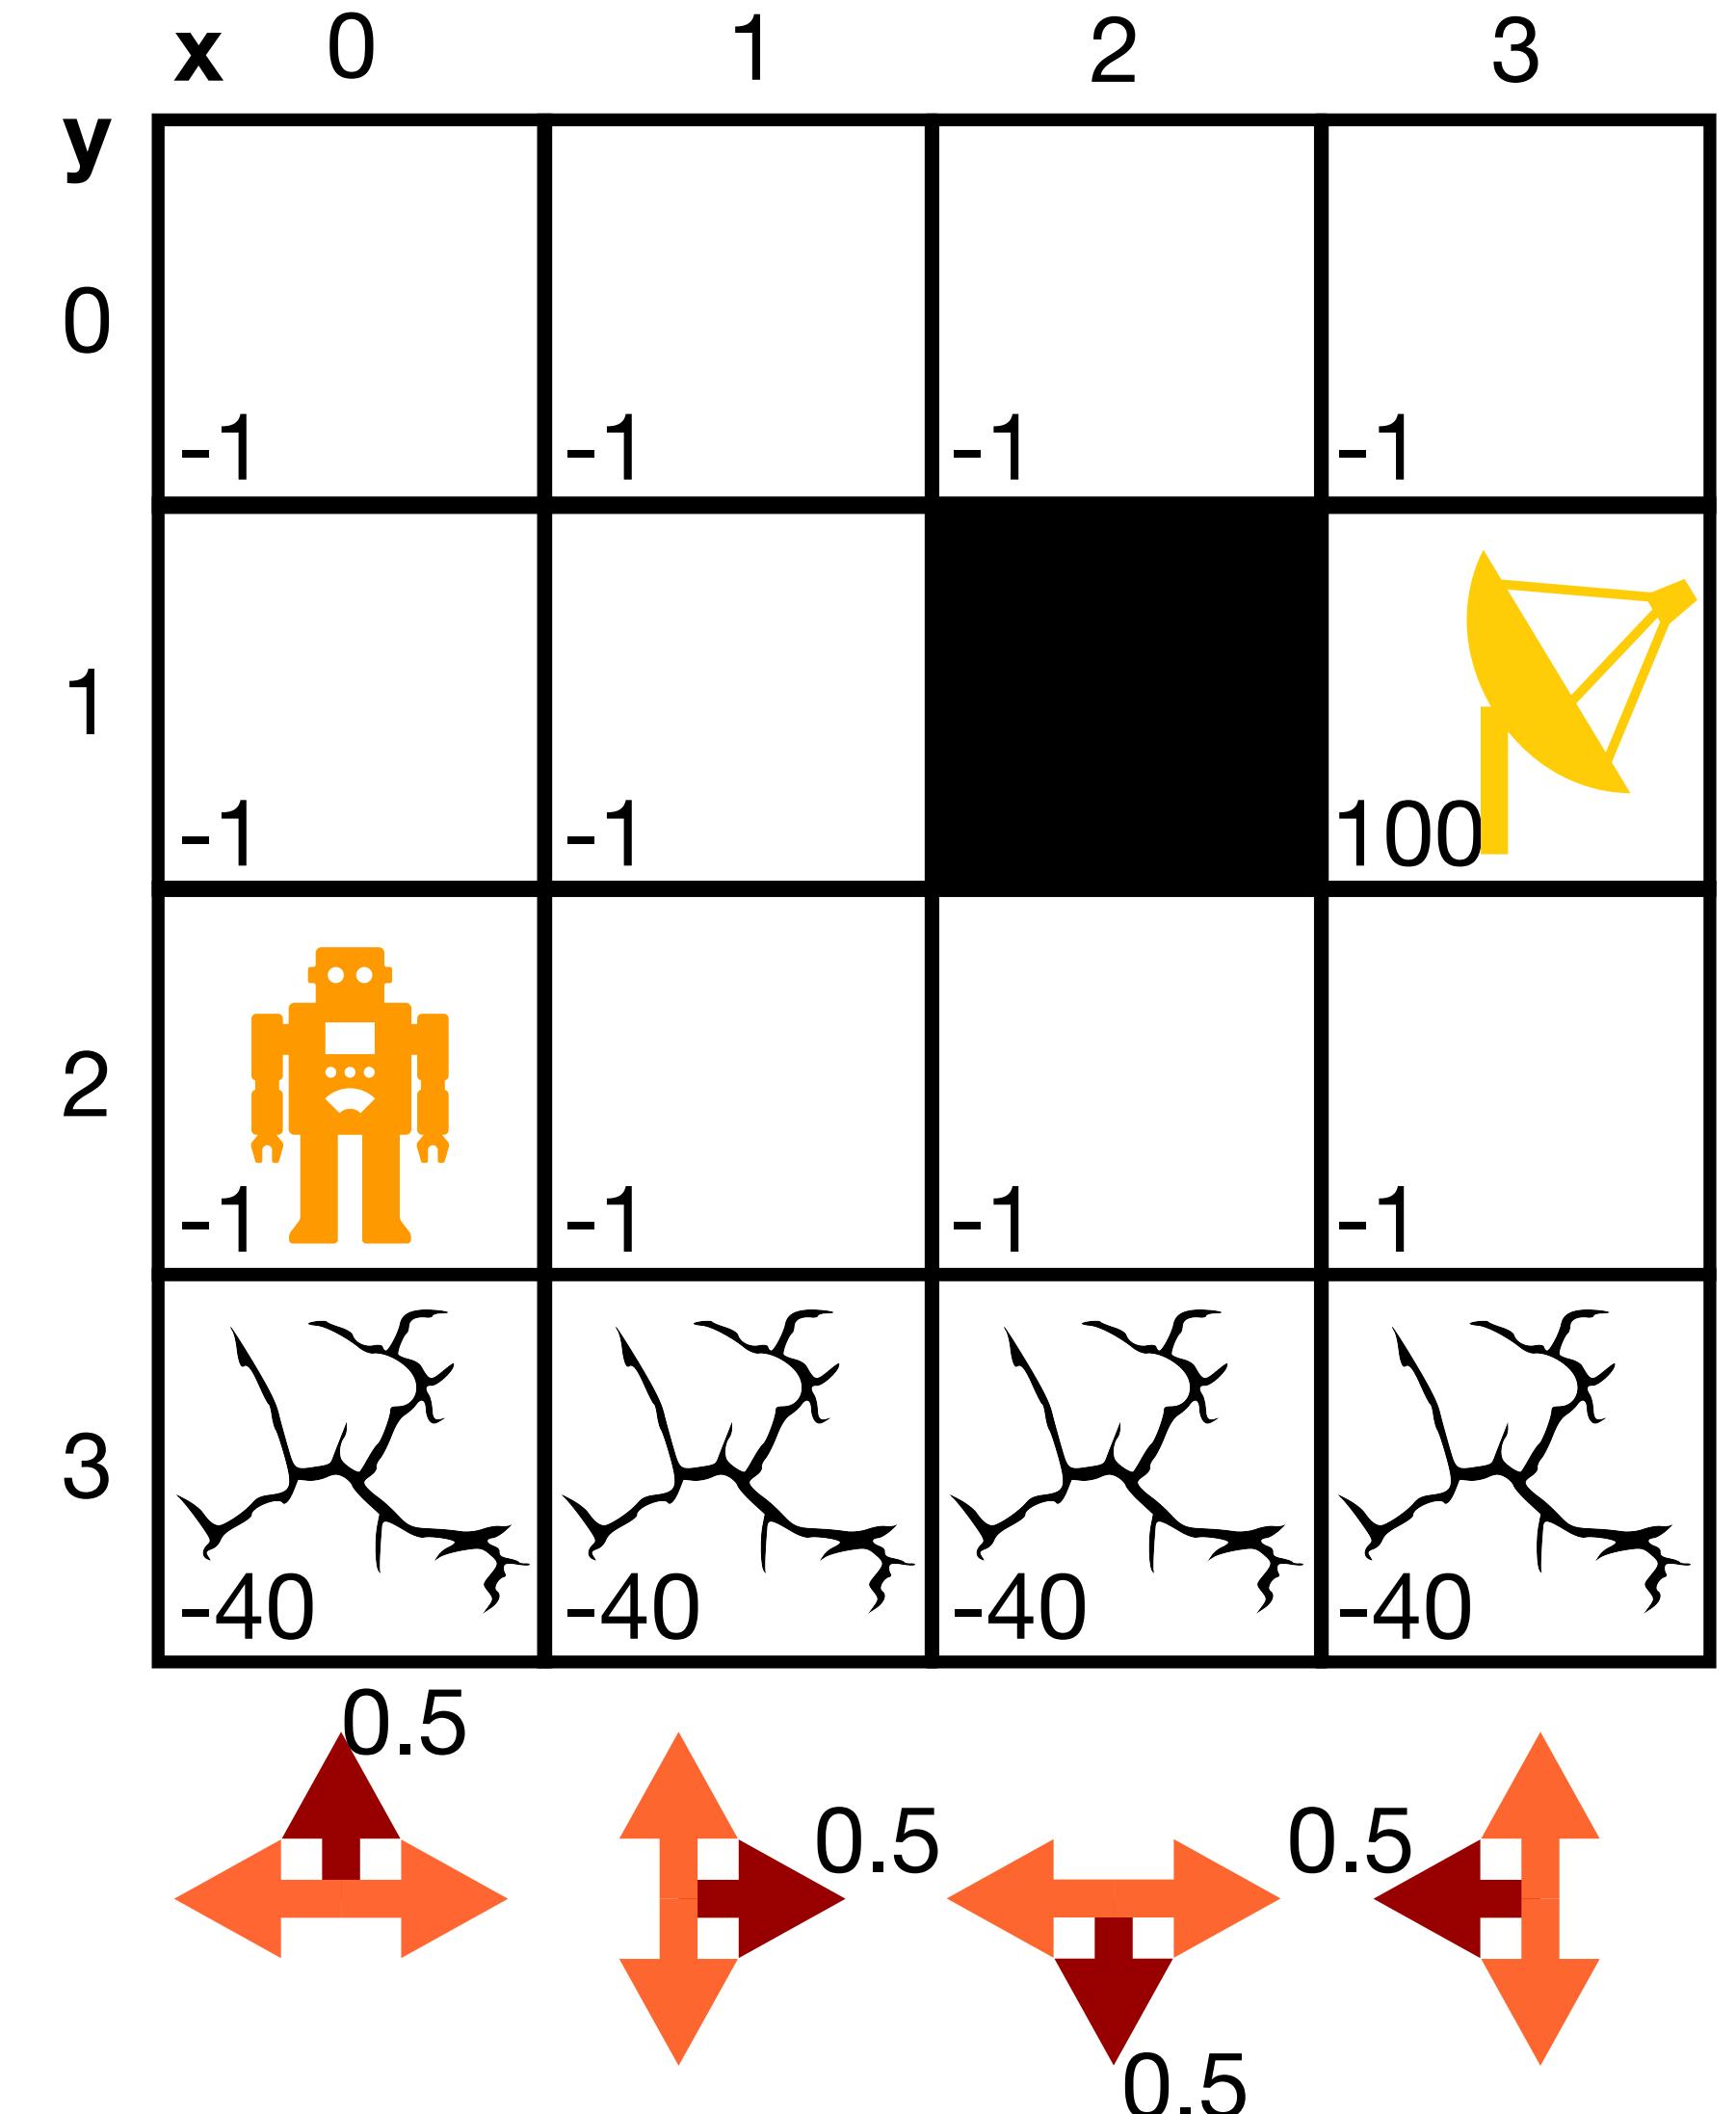
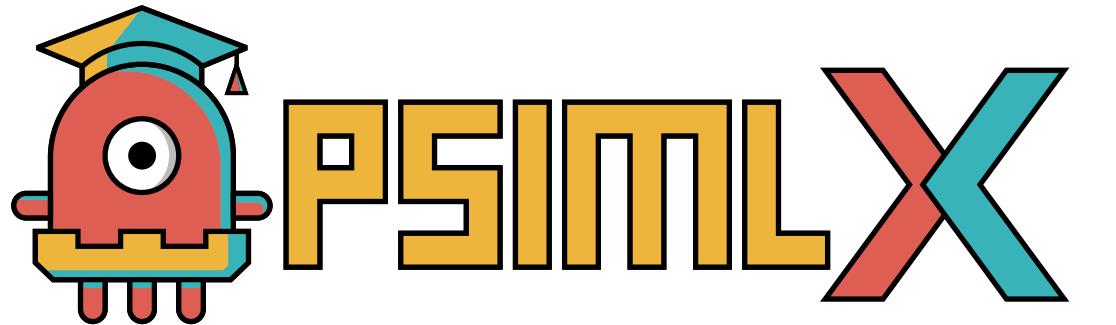
RL Formalisation

Markov Decision Process

- **Markov property:**

$$P[S_{t+1} | S_t] = P[S_{t+1} | S_t, S_{t-1}, \dots, S_0]$$

- The future is independent of the past given the present
- The state is sufficient statistic of the future



RL Formalisation

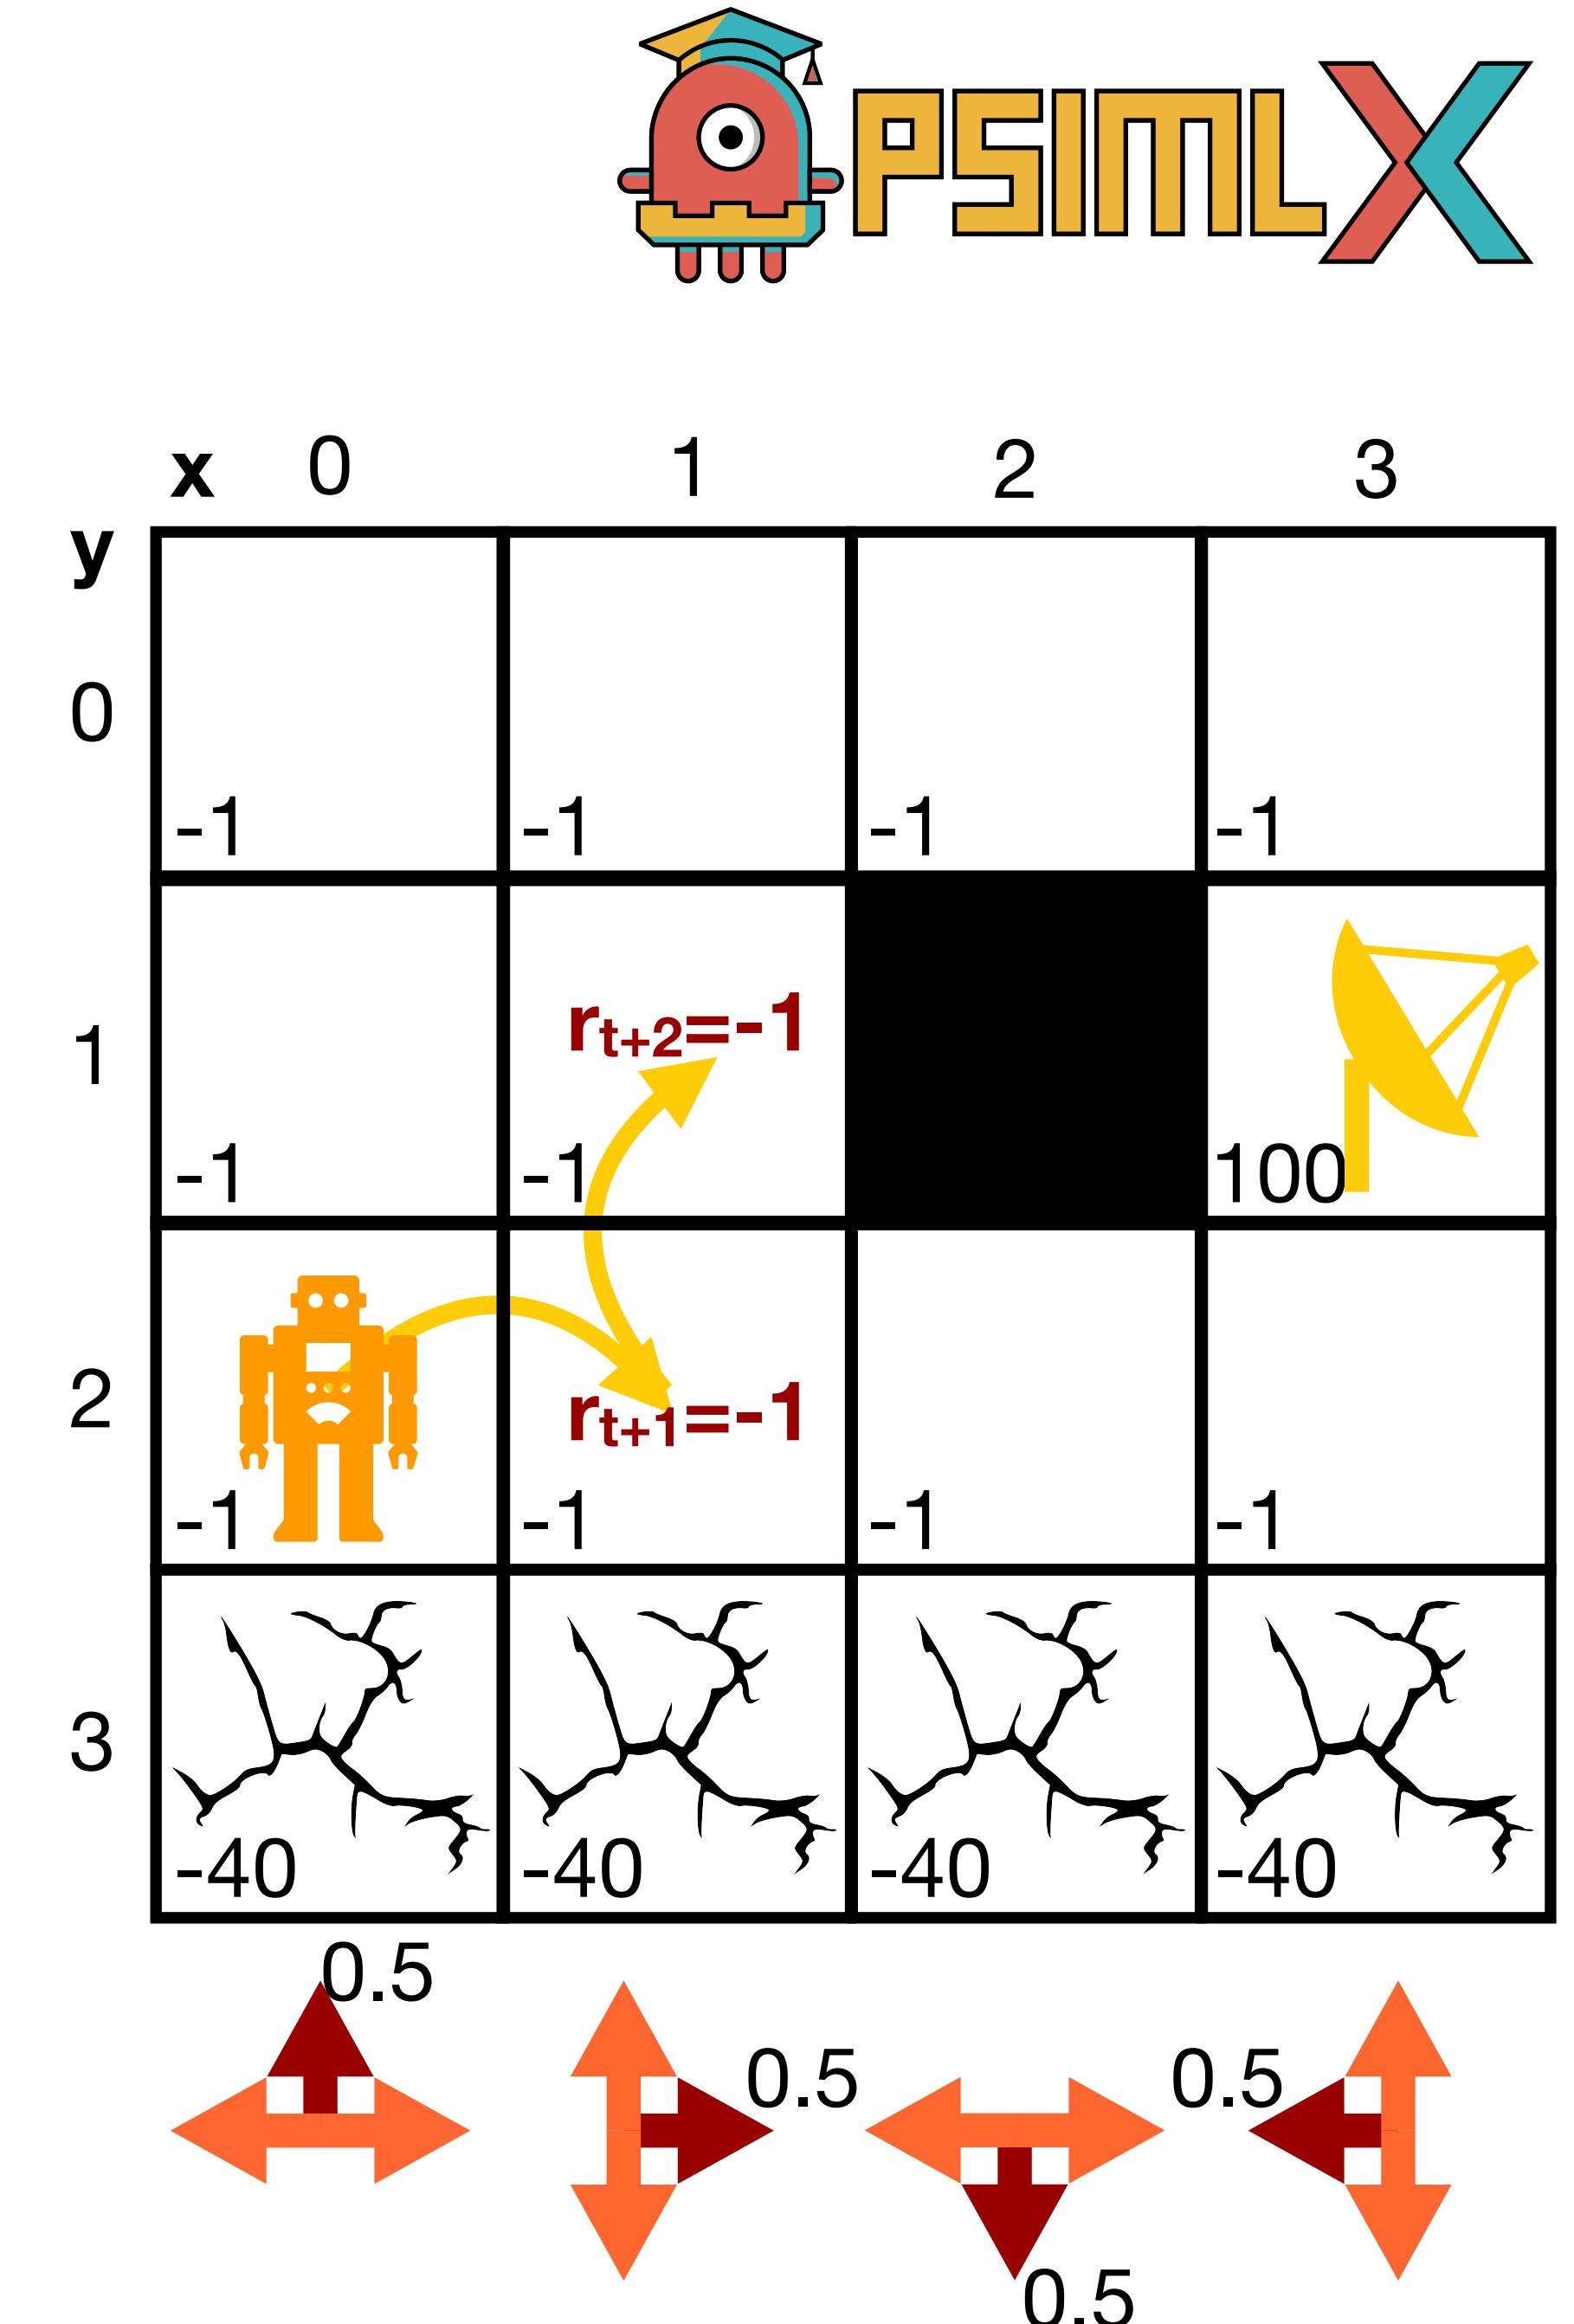
Reward and Return

- Agent's goal is to maximise the *return* G :

$$\begin{aligned} G_t &\doteq r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \\ &= r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} + \dots) \\ &= r_{t+1} + \gamma G_{t+1} \end{aligned}$$

- $\gamma \in [0,1]$ – future reward *discount factor*
 - Varying γ varies the “far-sightedness”
 - Mathematically convenient in continuing problems and cyclic Markov processes:

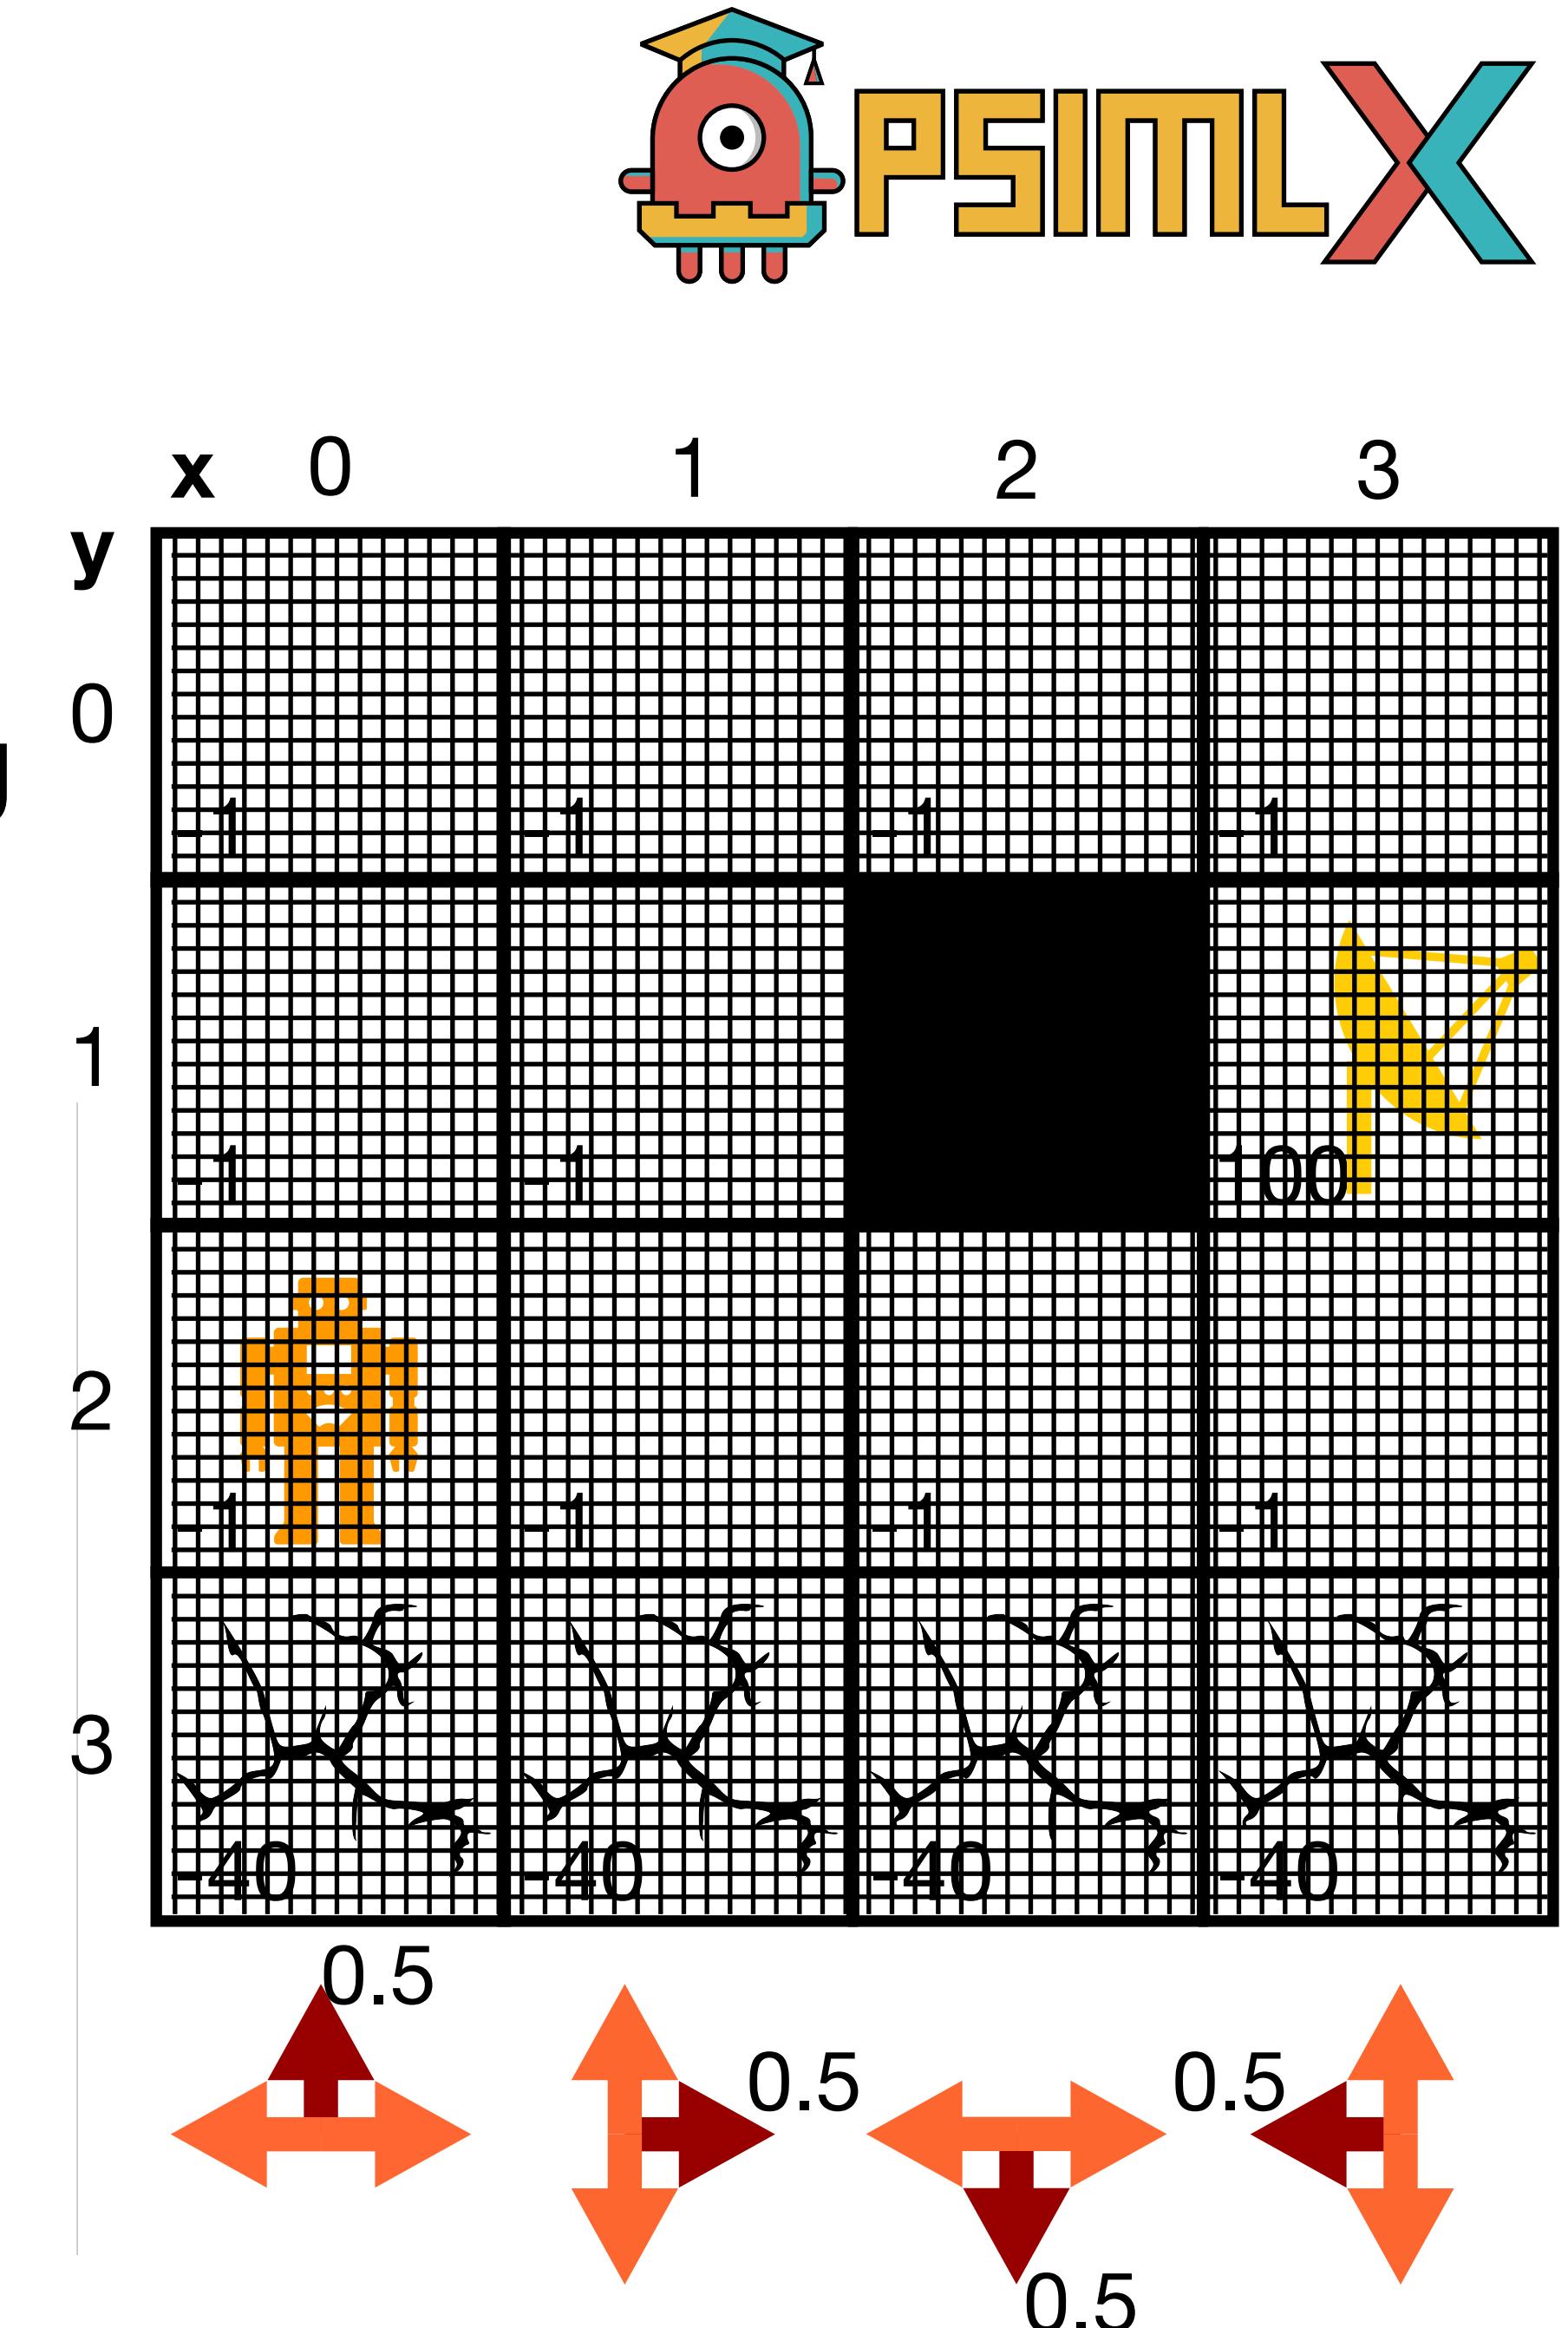
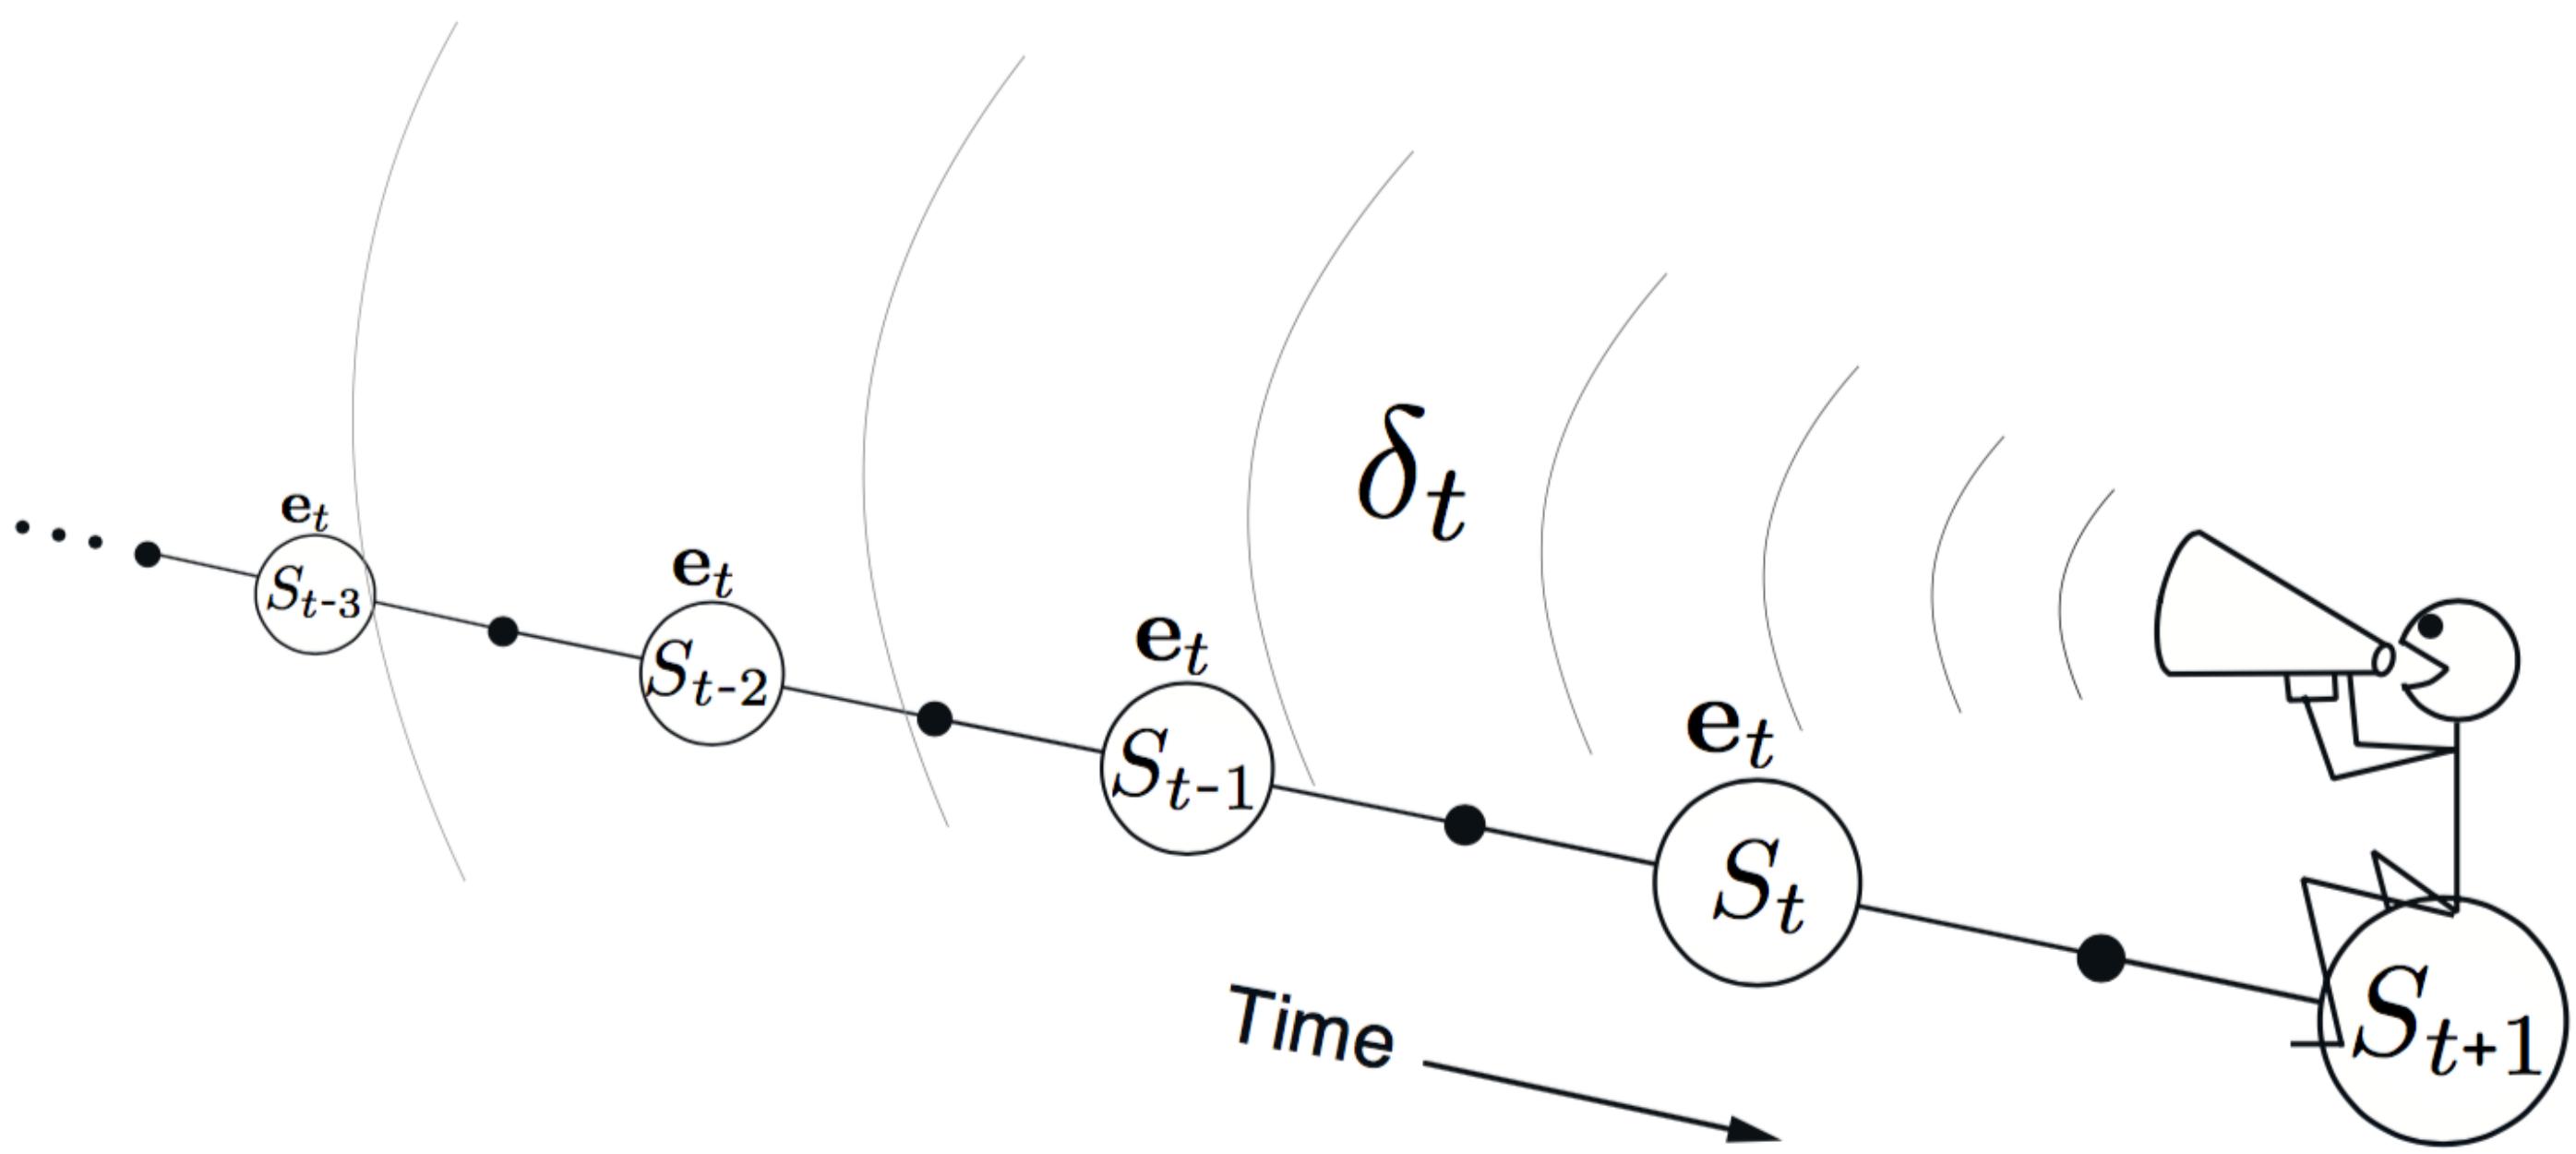
$$G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} = \frac{r}{1 - \gamma}$$



RL Formalisation

Reward and Return

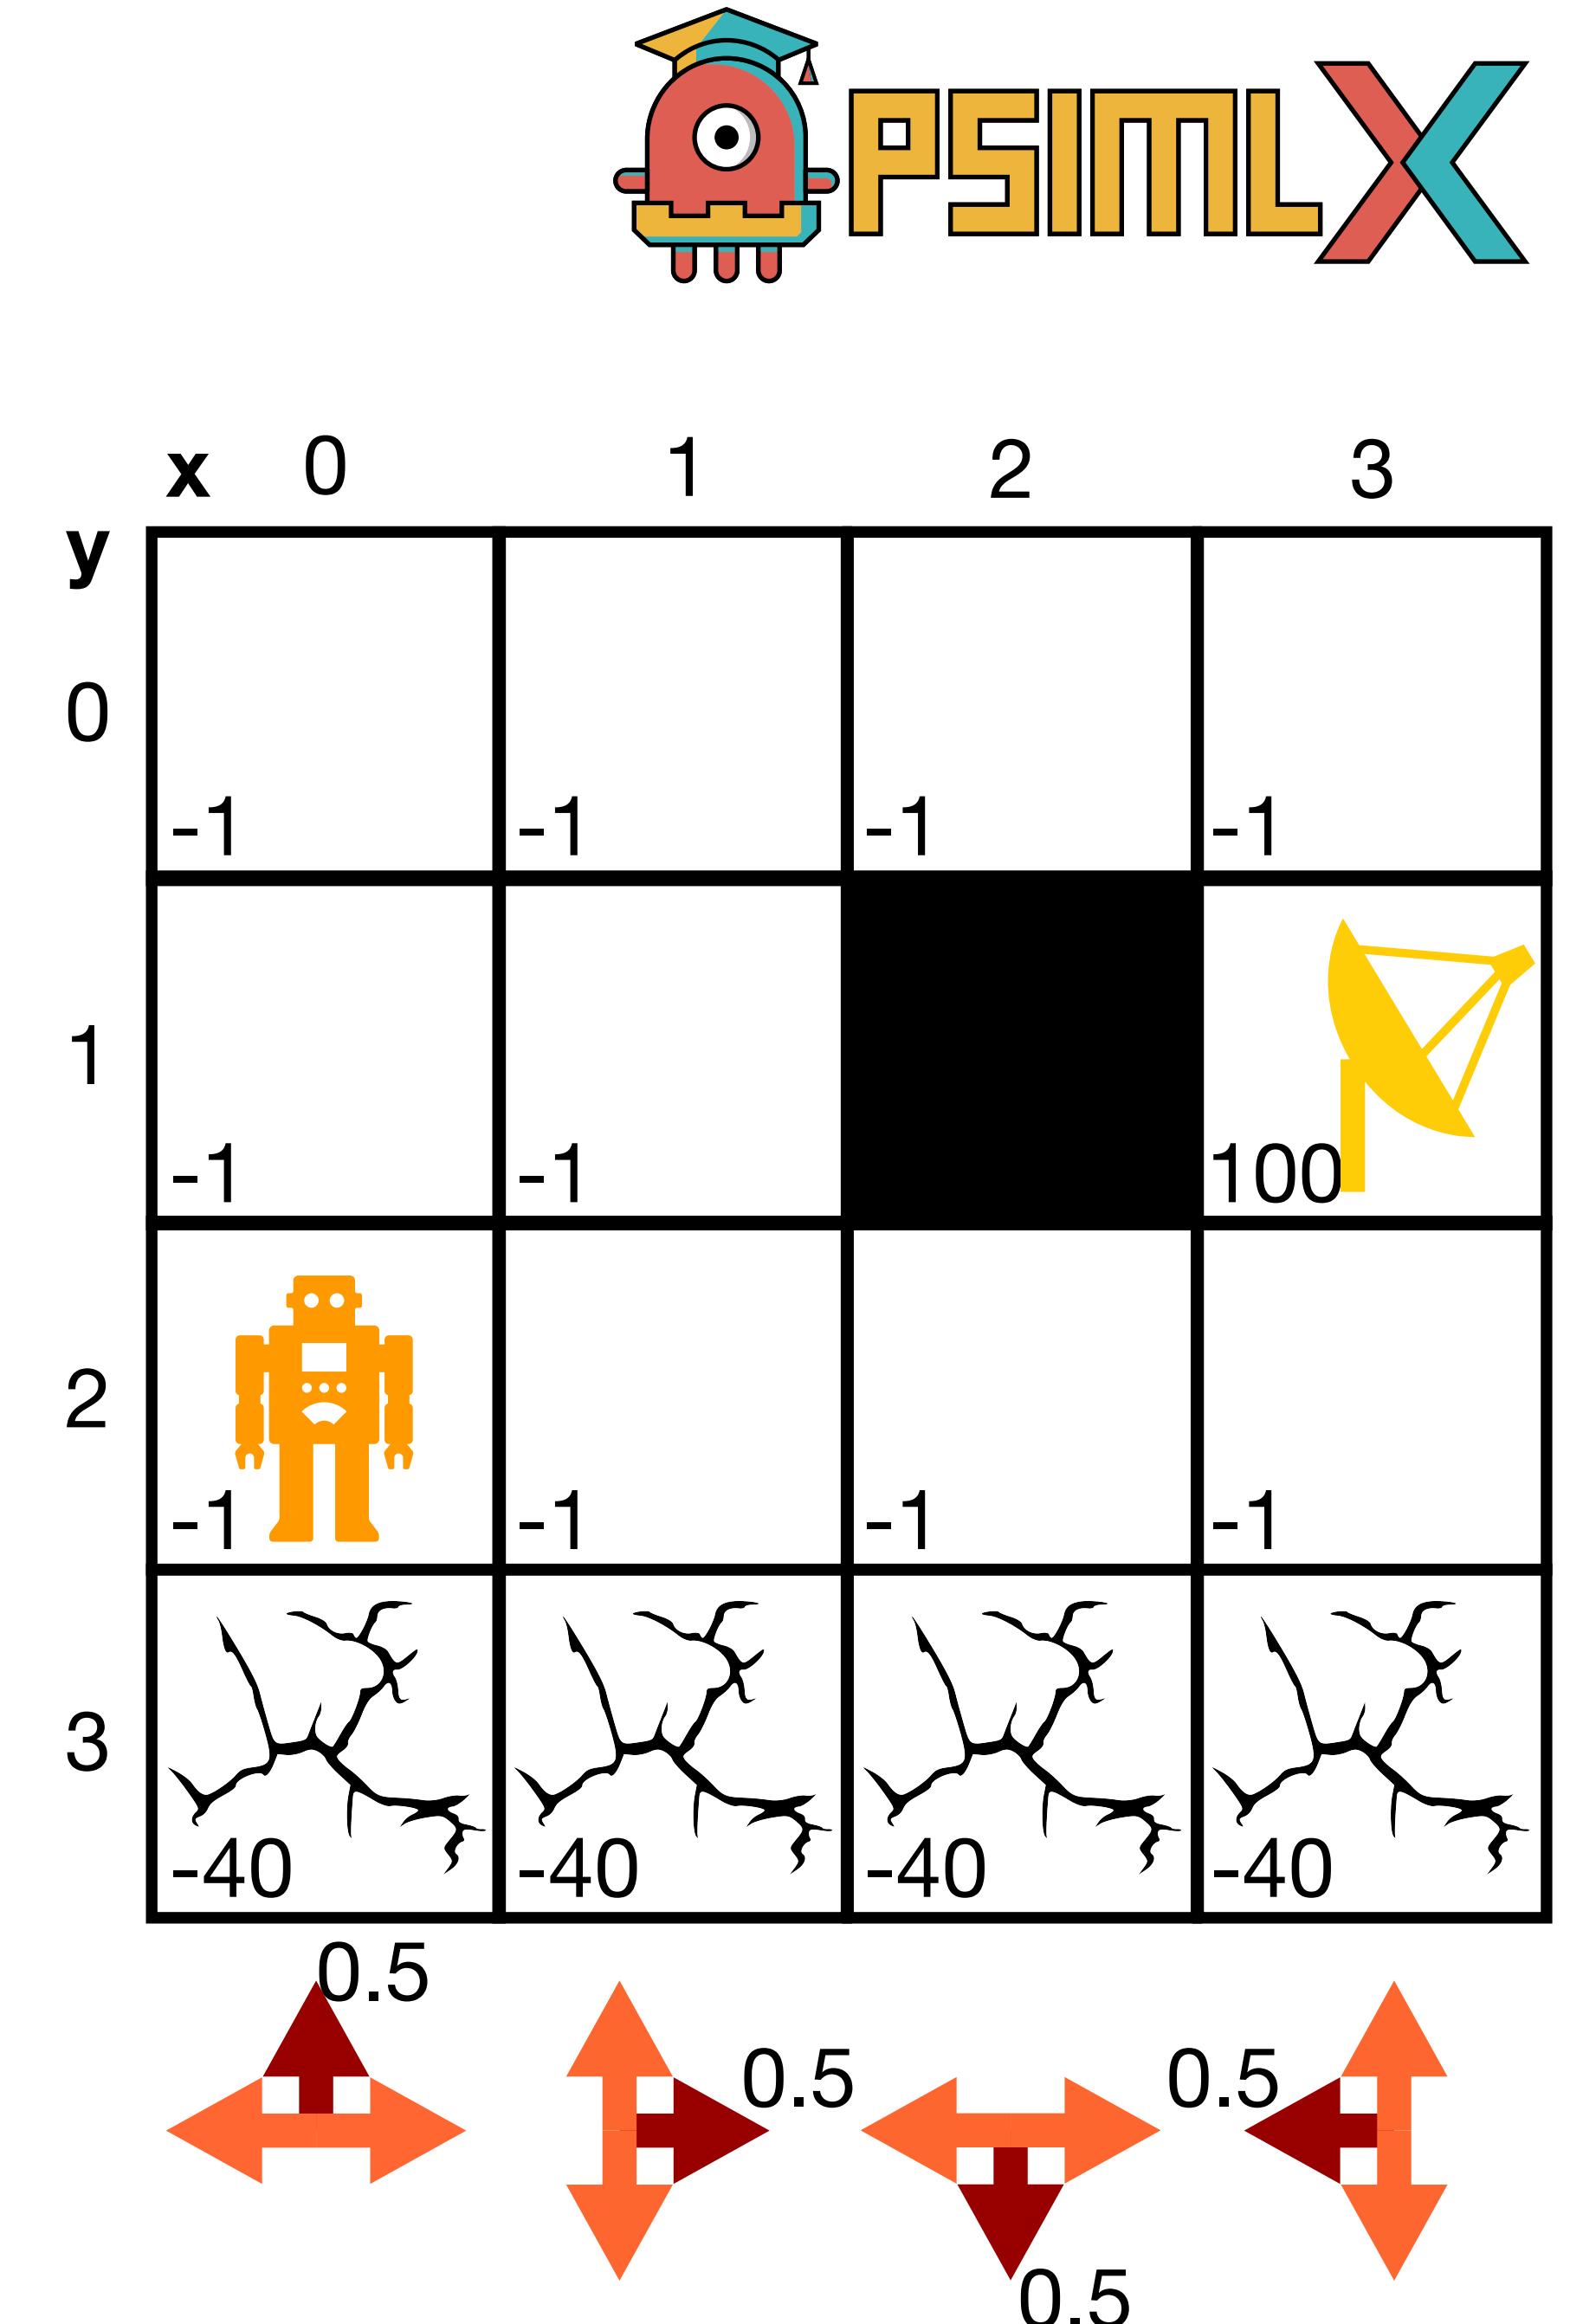
- *Credit assignment problem:*
 - How do you distribute credit for success (or blame for failure) of a decision (action) among the many throughout the episode?



RL Formalisation

Policy

- Policy fully captures agent's reasoning process (agent = policy)
- Is the conditional probability distribution over actions $a \in A$ given states $s \in S$:
 - Deterministic policy: $a = \pi(s)$
 - e.g. *greedy policy*
 - Stochastic policy: $\pi(a | s) = P[A_t = a | S_t = s]$
 - e.g. *exploratory policy*
- *Learning* in RL refers to learning the policy that maximises the return



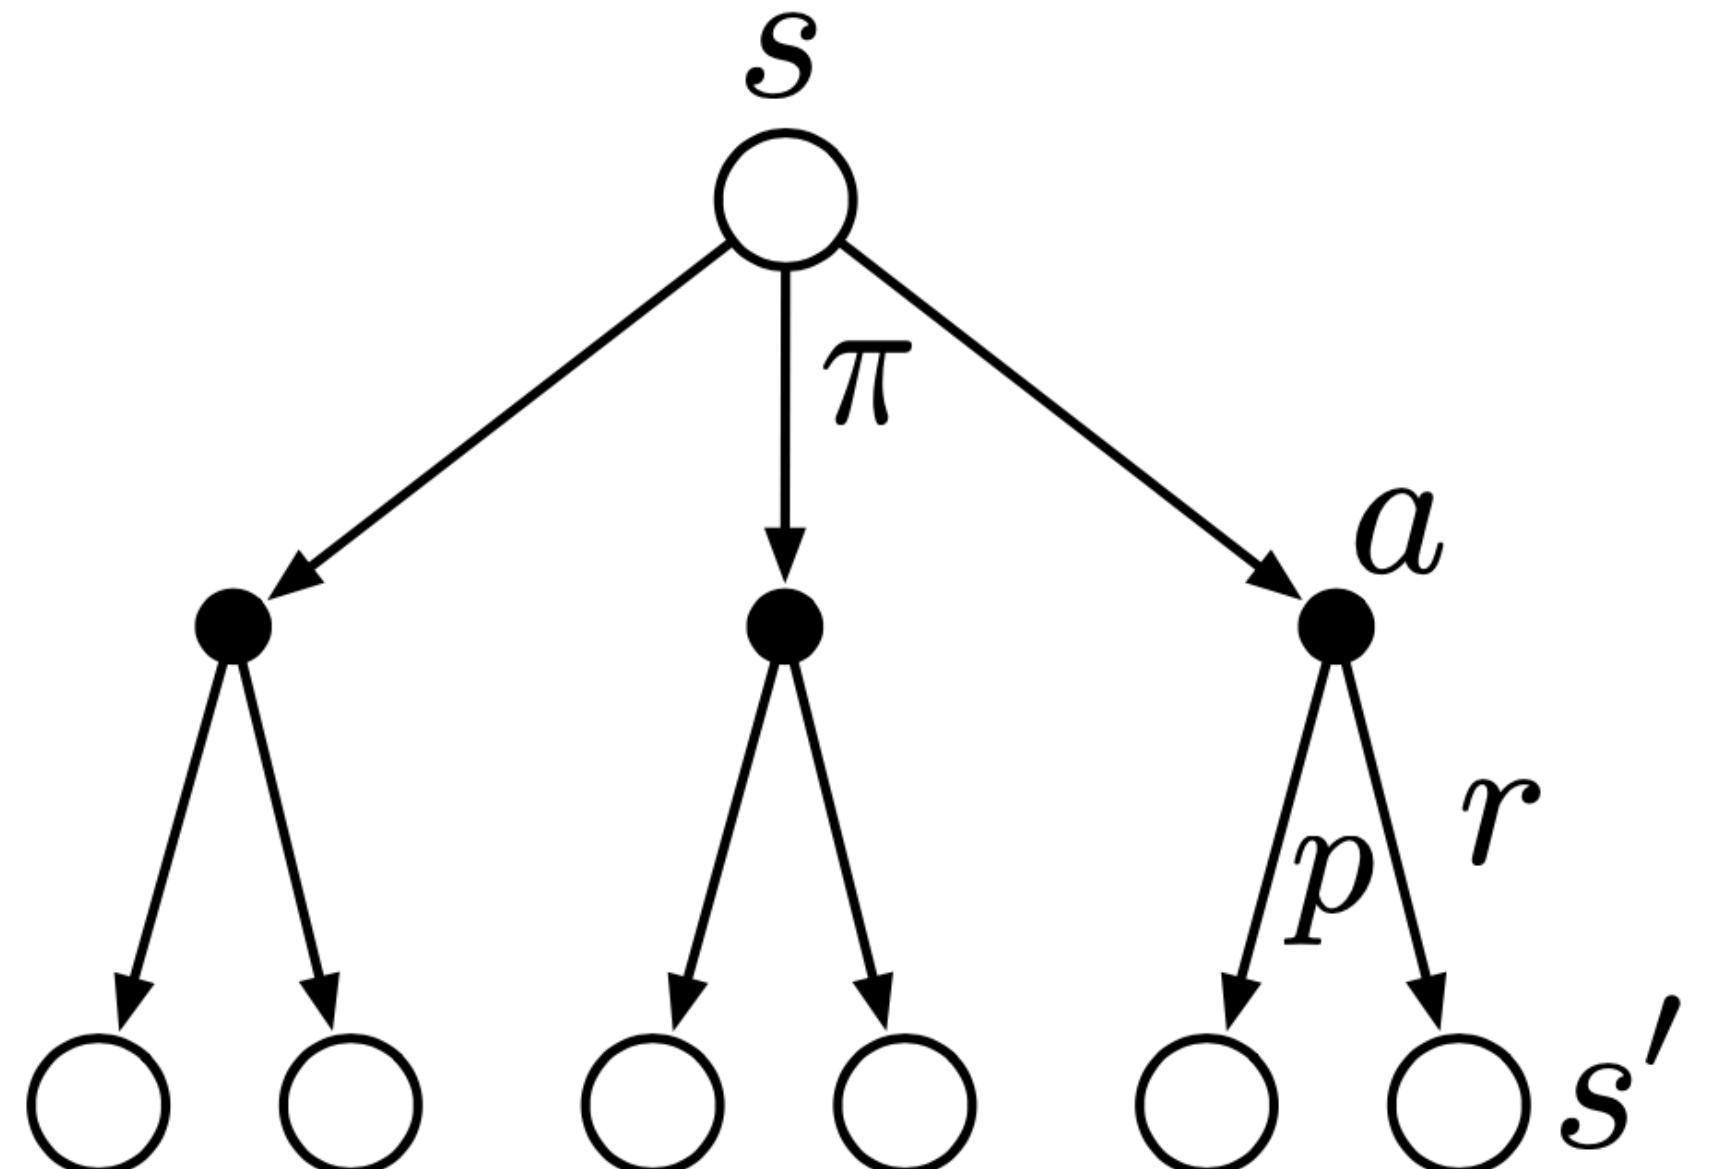
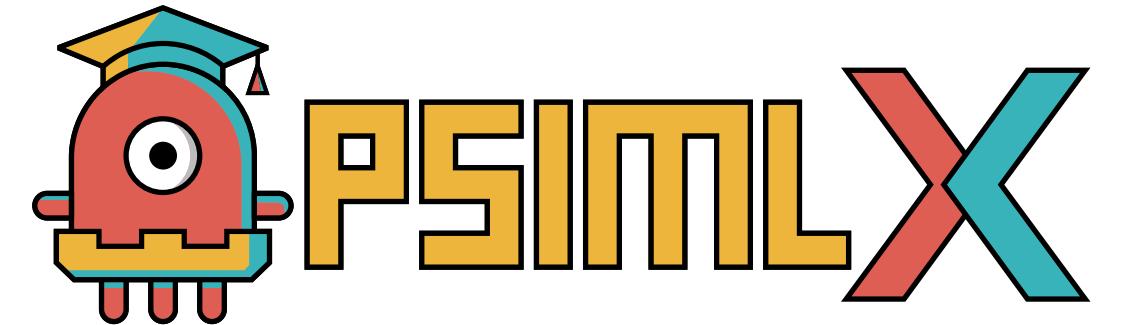
RL Formalisation

State-Value Function

- The state-value function $v_\pi(s)$ of an MDP is the expected return starting from state s , and then following policy π :

$$\begin{aligned} v_\pi(s) &\doteq \mathbb{E}_\pi[G_t | S_t = s] \\ &= \mathbb{E}_\pi[r_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \sum_a \pi(a | s) \left[r_s^a + \gamma \sum_{s'} p(s' | s, a) v_\pi(s') \right] \end{aligned}$$

Between MDPs and SMDPs [Sutton et al. 1999]



Backup diagram for v_π [Sutton & Barto 2018]

- Bellman equation for state-value

RL Formalisation

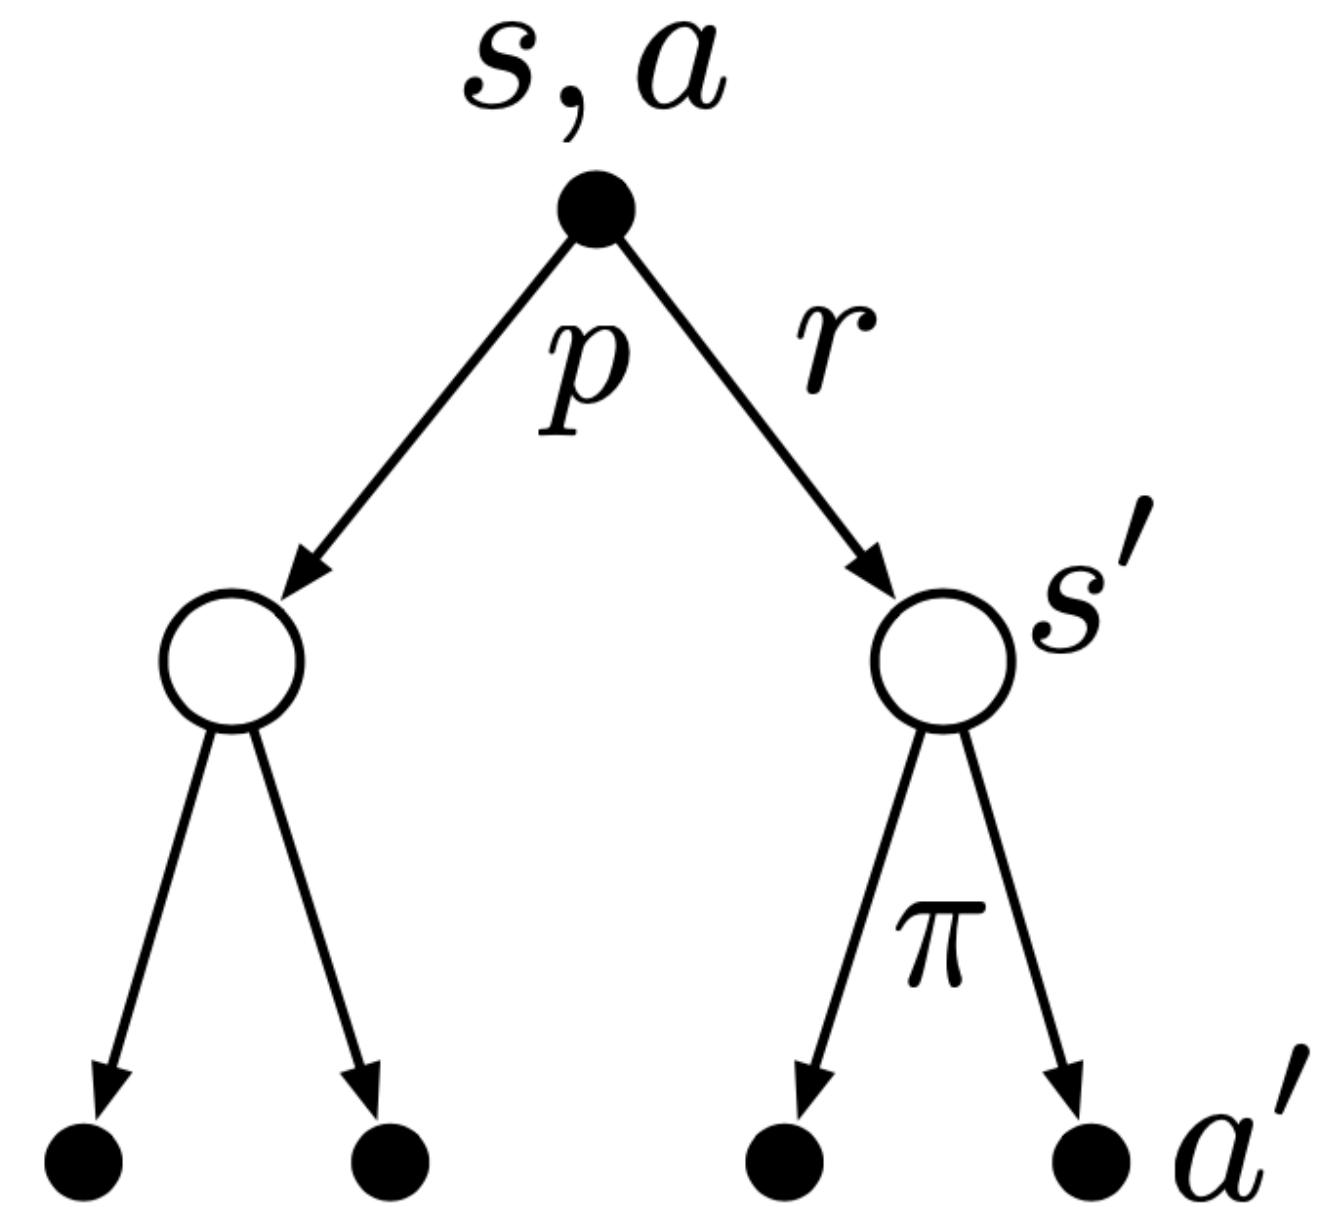
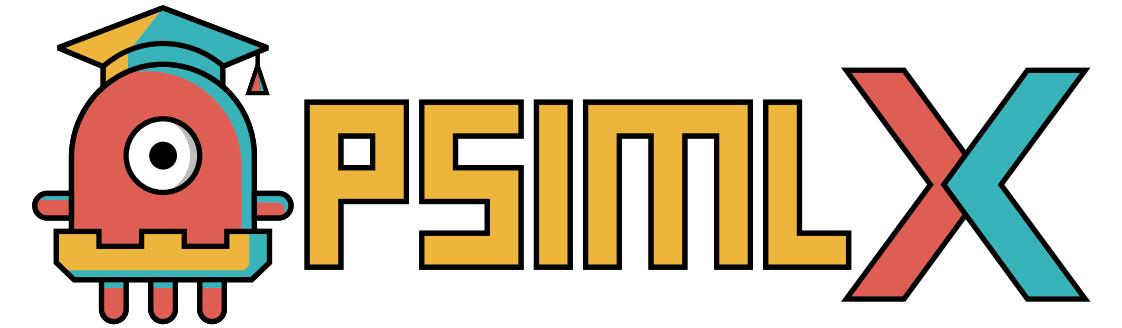
Action-Value Function

- The action-value function $q_\pi(s, a)$ of an MDP is the expected return starting from state s by taking an action a , and then following policy π :

$$\begin{aligned}
 q_\pi(s, a) &\doteq \mathbb{E}_\pi[G_t | S_t = s, A_t = a] \\
 &= \mathbb{E}_\pi[r_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a] \\
 &= r_s^a + \gamma \sum_{s'} p(s' | s, a) v_\pi(s') \\
 &= r_s^a + \gamma \sum_{s'} p(s' | s, a) \sum_{a'} \pi(a' | s') q_\pi(s', a')
 \end{aligned}$$

Between MDPs and SMDPs [Sutton et al. 1999]

- Bellman equation for action-value



Backup diagram for q_π [Sutton & Barto 2018]

RL Formalisation

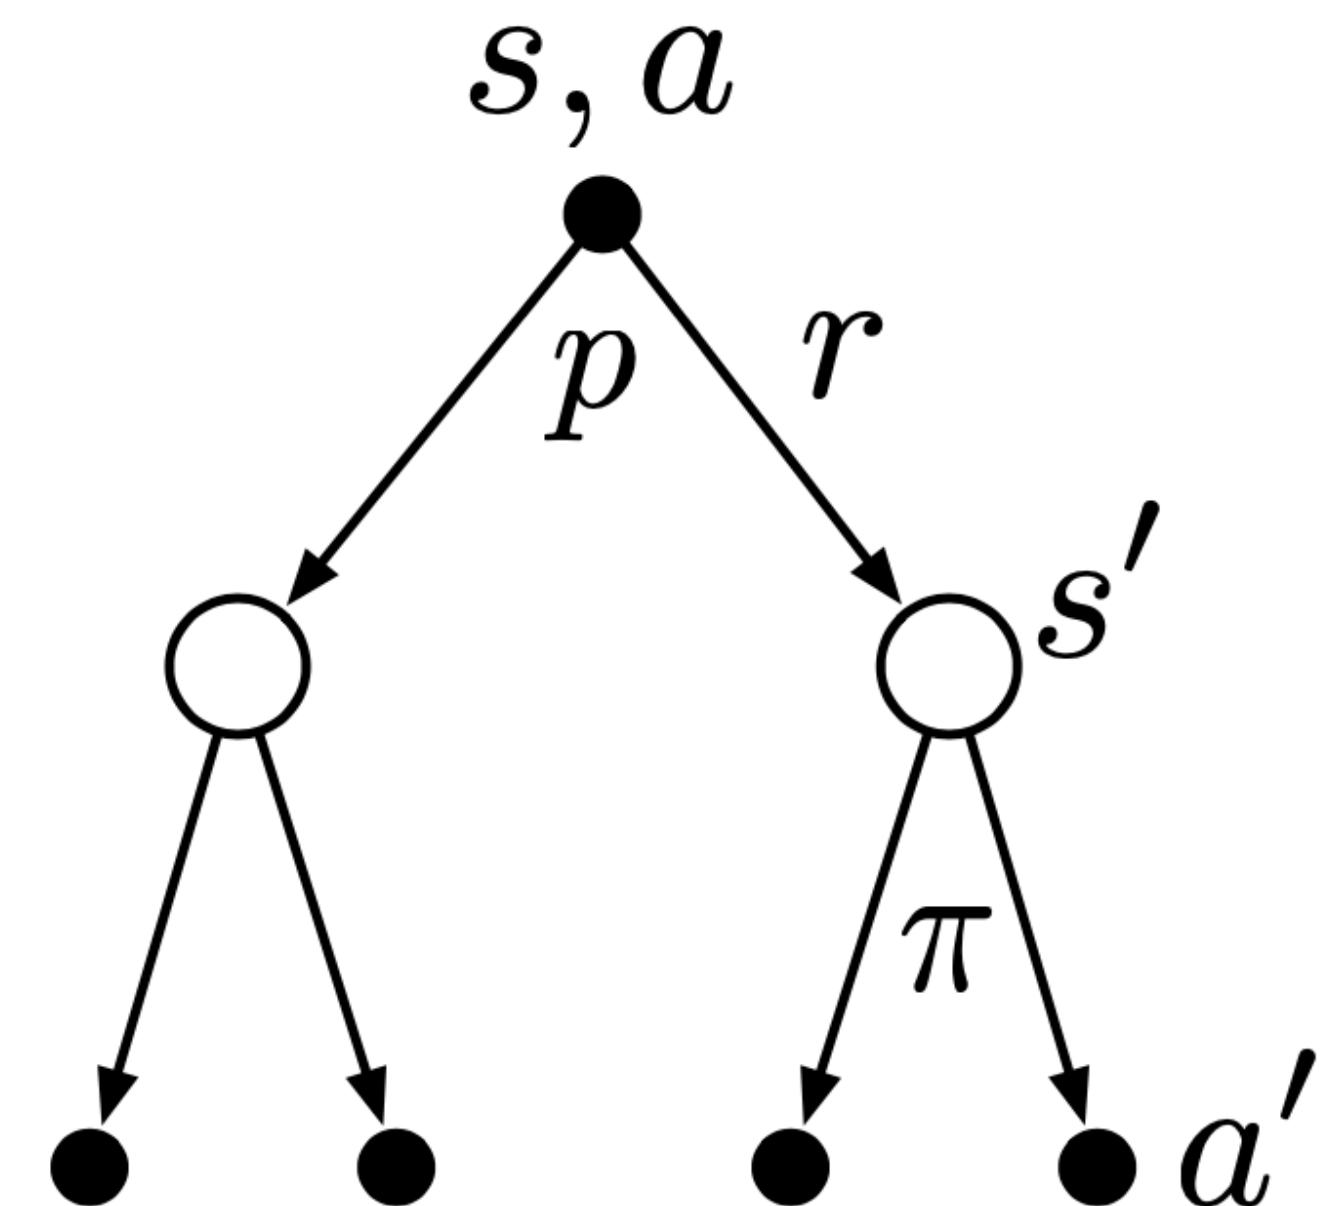
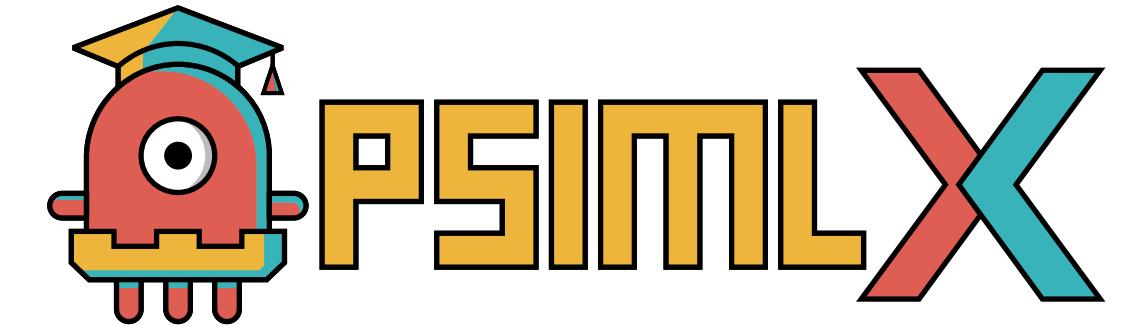
Action-Value Function

- The action-value function $q_\pi(s, a)$ of an MDP is the expected return starting from state s by taking an action a , and then following policy π :

$$\begin{aligned}
 q_\pi(s, a) &\doteq \mathbb{E}_\pi[G_t | S_t = s, A_t = a] \\
 &= \mathbb{E}_\pi[r_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a] \\
 &= r_s^a + \gamma \sum_{s'} p(s' | s, a) v_\pi(s') \\
 &= r_s^a + \gamma \sum_{s'} p(s' | s, a) \sum_{a'} \pi(a' | s') q_\pi(s', a')
 \end{aligned}$$

Between MDPs and SMDPs [Sutton et al. 1999]

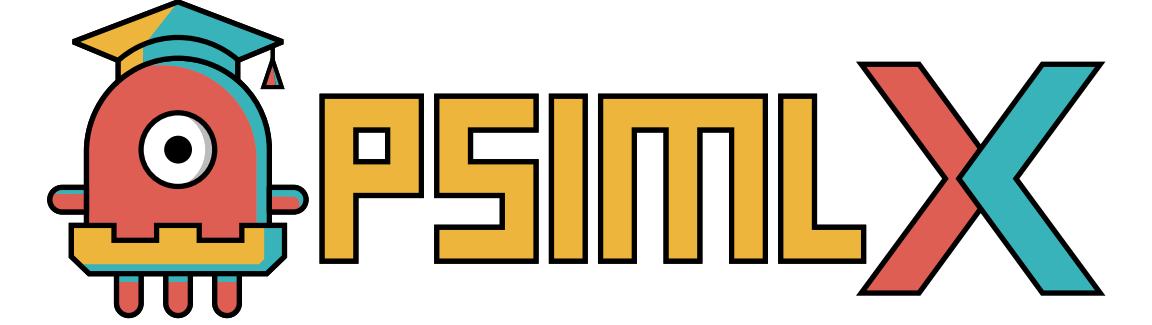
Relationship between v_π and q_π



Backup diagram for q_π [Sutton & Barto 2018]

RL Formalisation

Optimal Policy

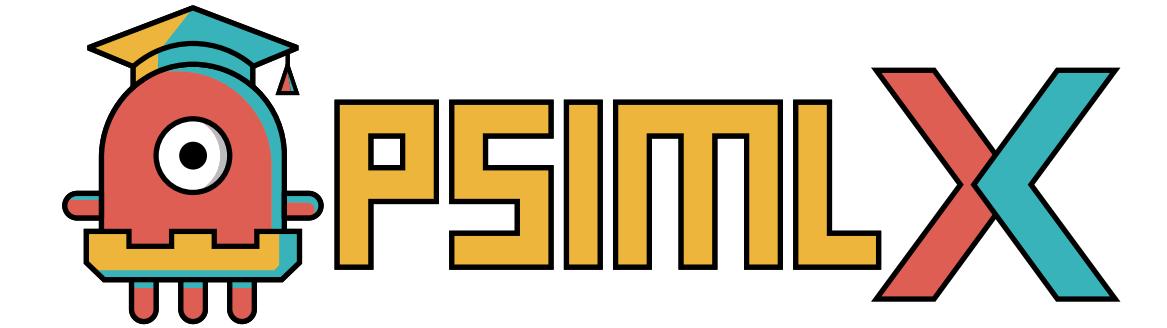


- Value functions define a partial ordering over policies:
 $\pi' \geq \pi$ if and only if $v_{\pi'}(s) \geq v_{\pi}(s) \forall s \in S$
- A policy π' is defined to be better than or equal to a policy π if its expected return is greater than or equal to that of π for all states
- There is always at least one policy that is better than or equal to all other policies, called the *optimal policy*, and denoted π^*
- **Policy Improvement Theorem:** If we have two policies π and π' so that $\pi(s) = \pi'(s)$ for all $s \in S$ except some s' where $\pi'(s') = a \neq \pi(s')$ and $q_{\pi}(s, a) > v_{\pi}(s)$ then $\pi' > \pi$.

Proof: [Sutton & Barto 2018] p. 78

RL Formalisation

Optimal Value Functions



- All optimal policies share the same state-value function $v^*(s)$ and action-value function $q^*(s, a)$
- Correspond to optimal policies, and optimal policies are greedy

$$v^*(s) \doteq \max_{\pi} v_{\pi}(s)$$

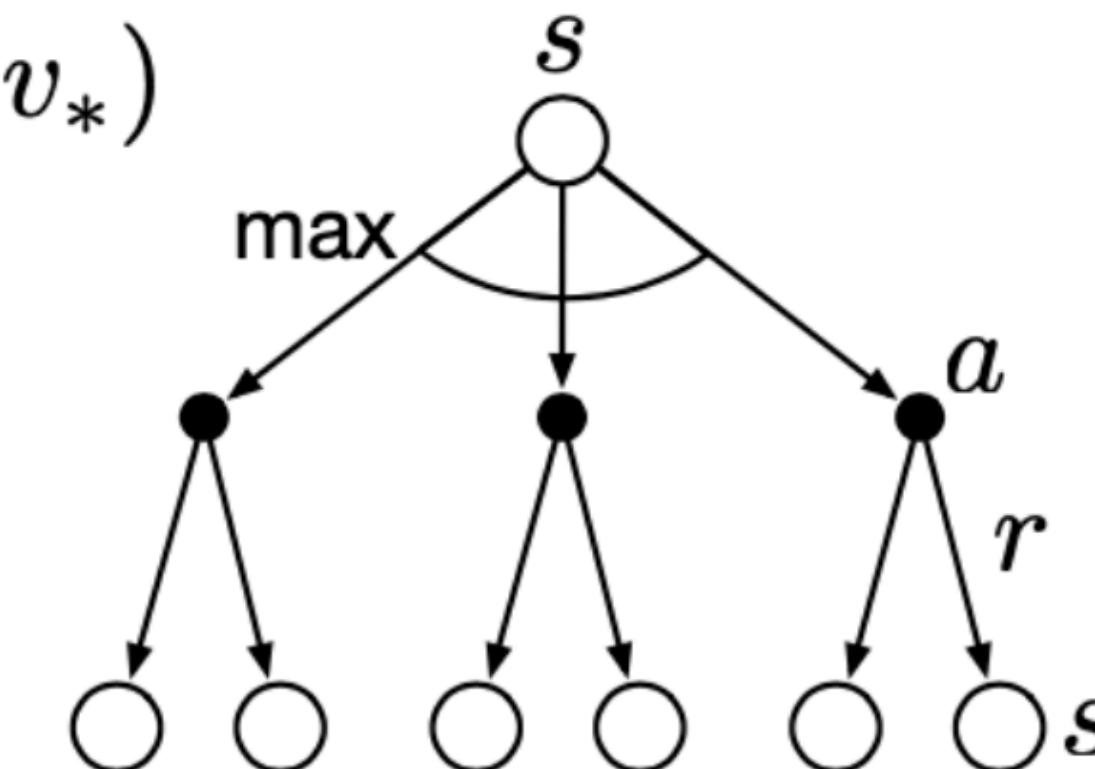
$$= \max_{a \in A} \mathbb{E}_{\pi} [r_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a]$$

$$= \max_{a \in A} \left[r_s^a + \gamma \sum_{s'} p(s' | s, a) v^*(s') \right]$$

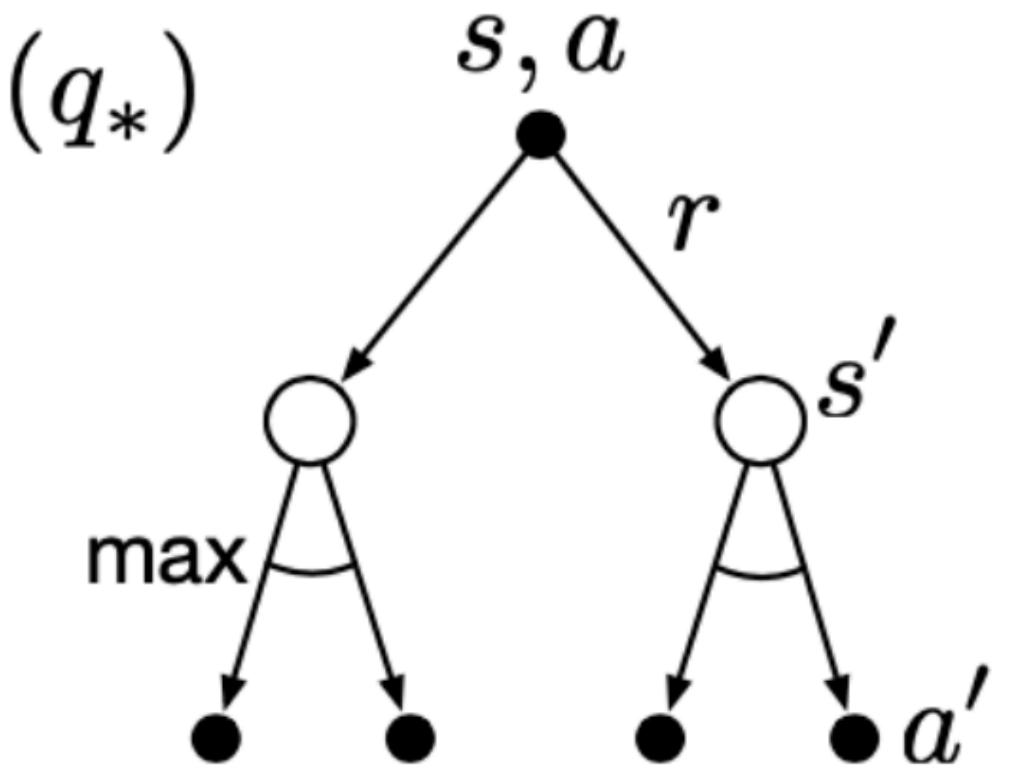
$$q^*(s, a) \doteq \max_{\pi} q_{\pi}(s, a)$$

$$= r_s^a + \gamma \sum_{s'} p(s' | s, a) \max_{a' \in A} q^*(s', a')$$

Between MDPs and SMDPs [Sutton et al. 1999]

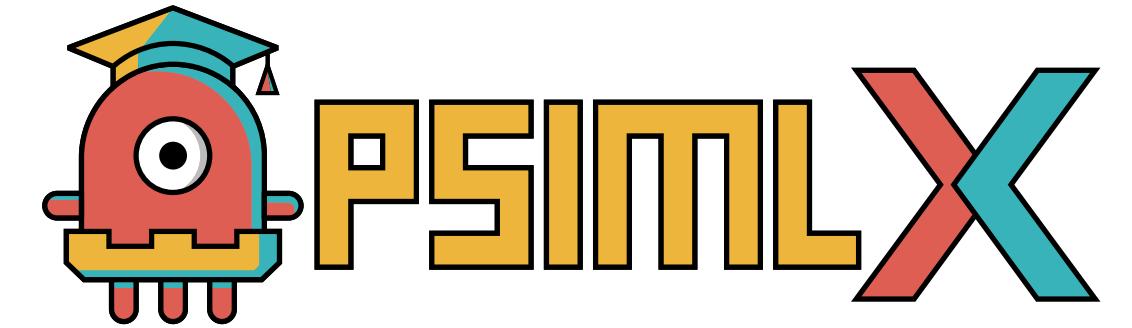


Optimal value backup diagrams [Sutton & Barto 2018]



RL Formalisation

Bellman Equations and Planning



- Equations for v_π and q_π are called *Bellman equations*
 - Set of recursive equations relating states (and actions) to successor states (and actions)
 - In principle, could be solved iteratively, or with a *dynamic programming* methods:
 - *value iteration*, *q-value iteration*, *policy iteration*
- Equations for v^* and q^* are called *Bellman optimality equations*

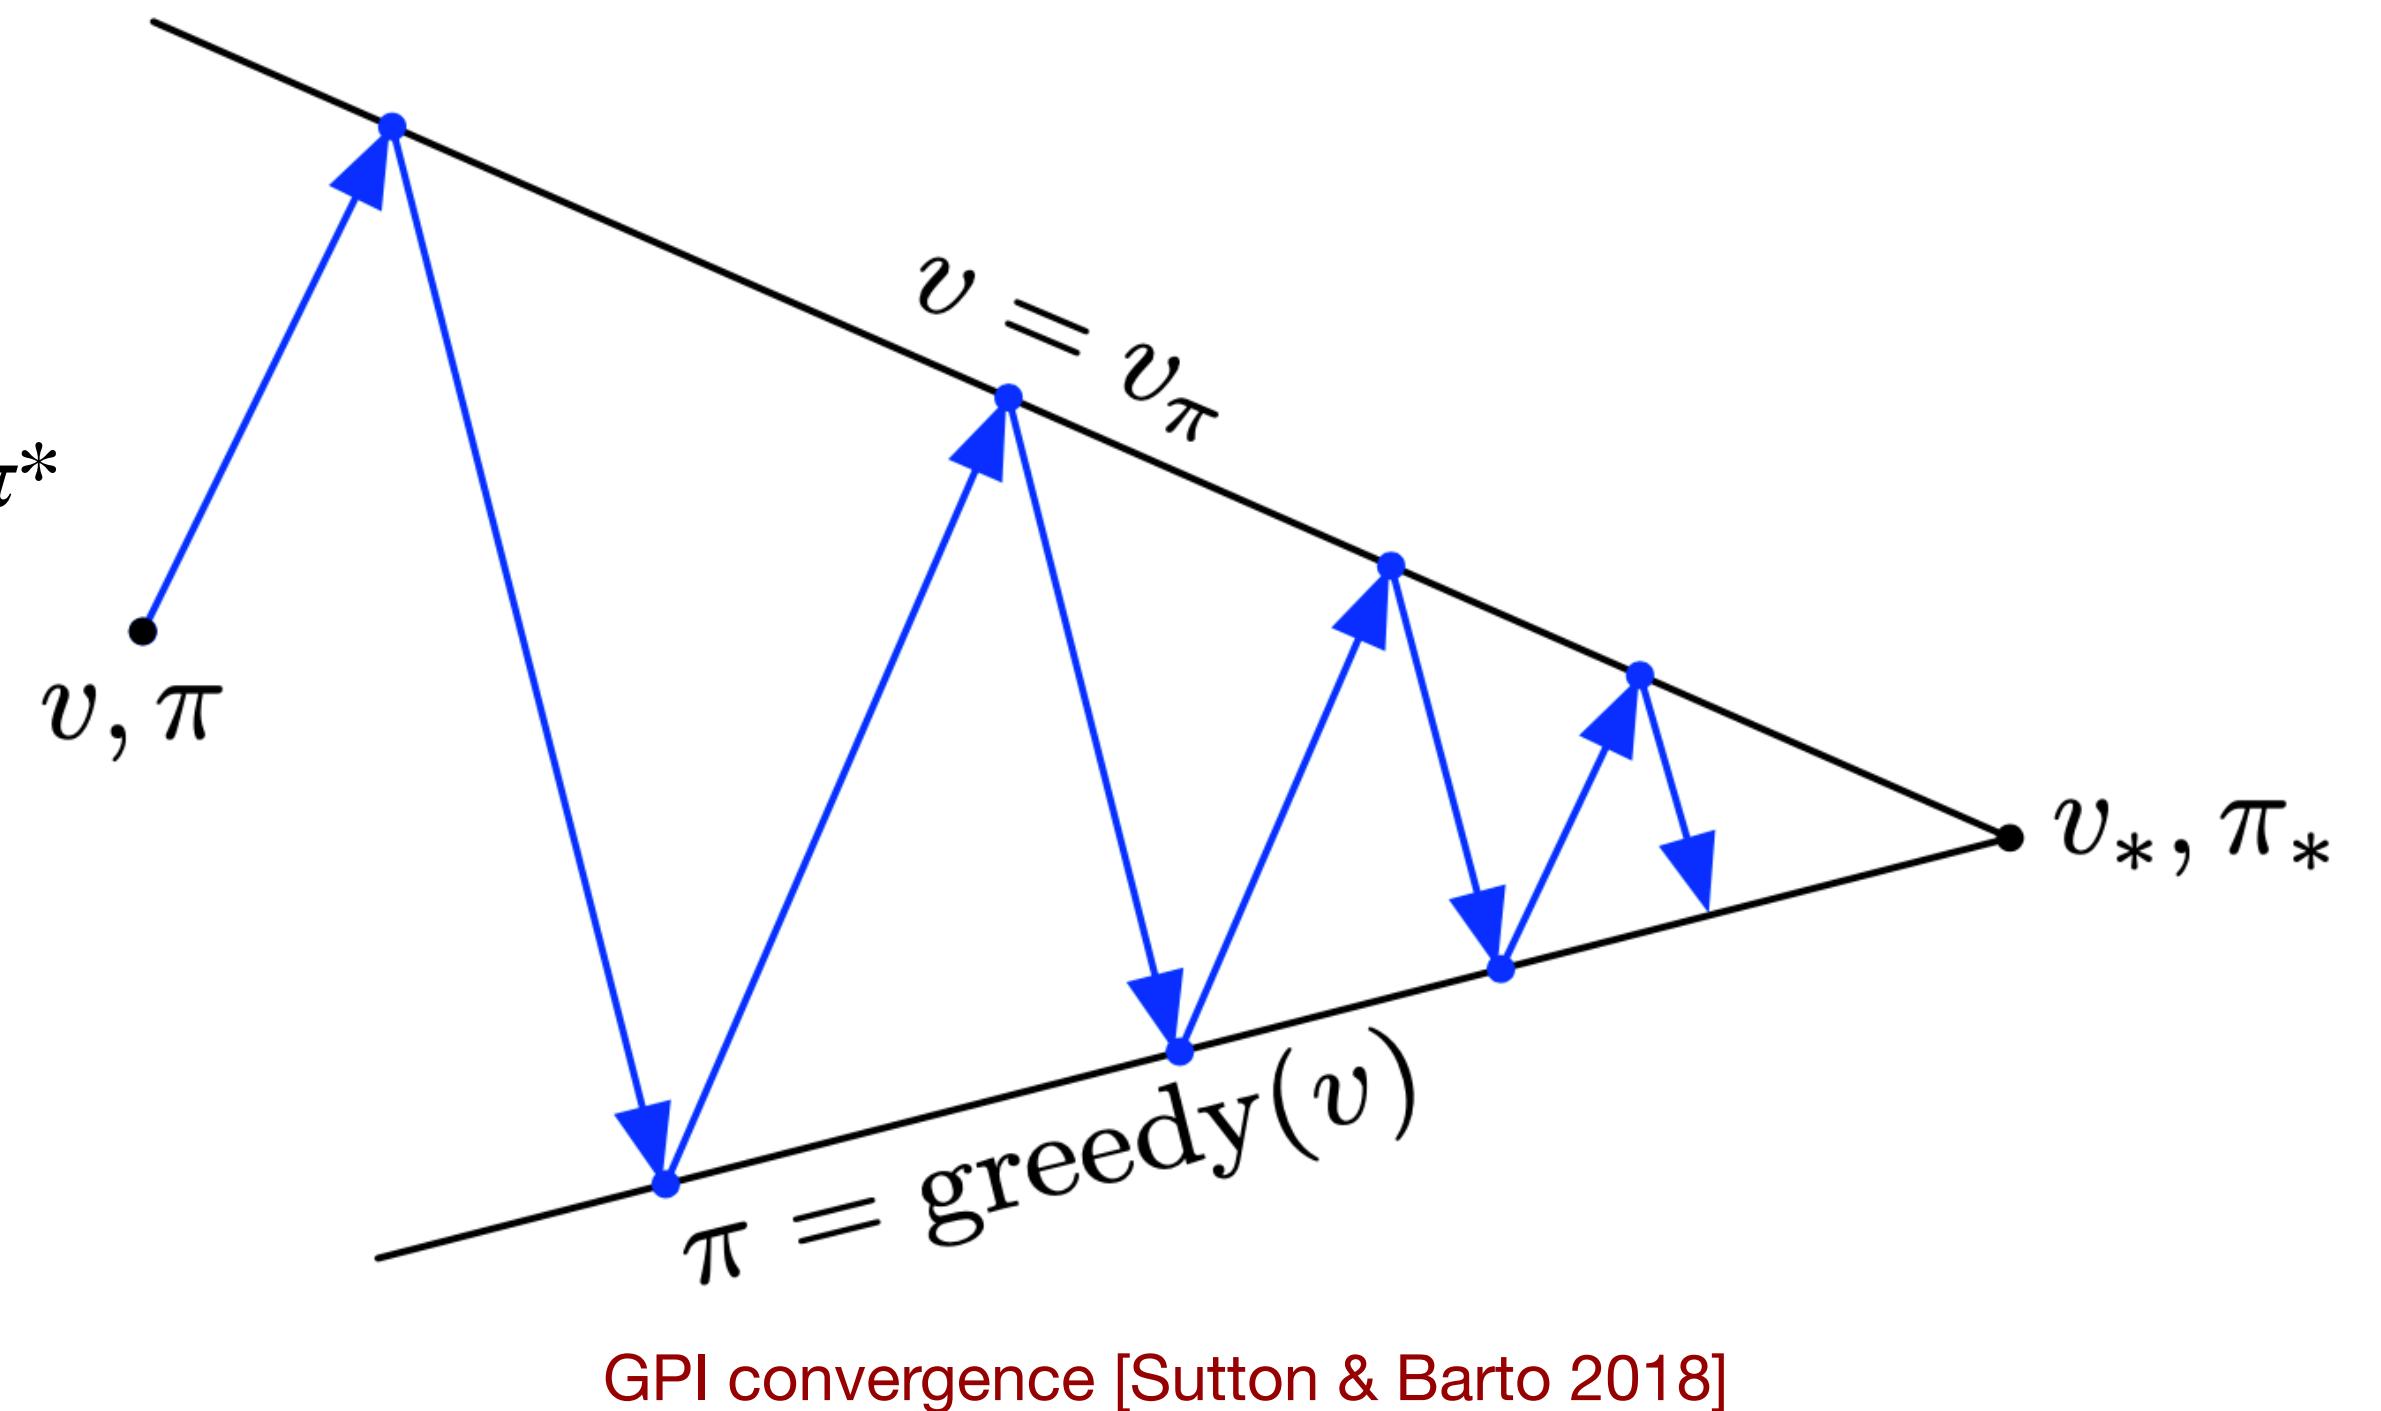
RL Formalisation

Generalised Policy Iteration (GPI)

- As a direct consequence of the **policy improvement theorem**:

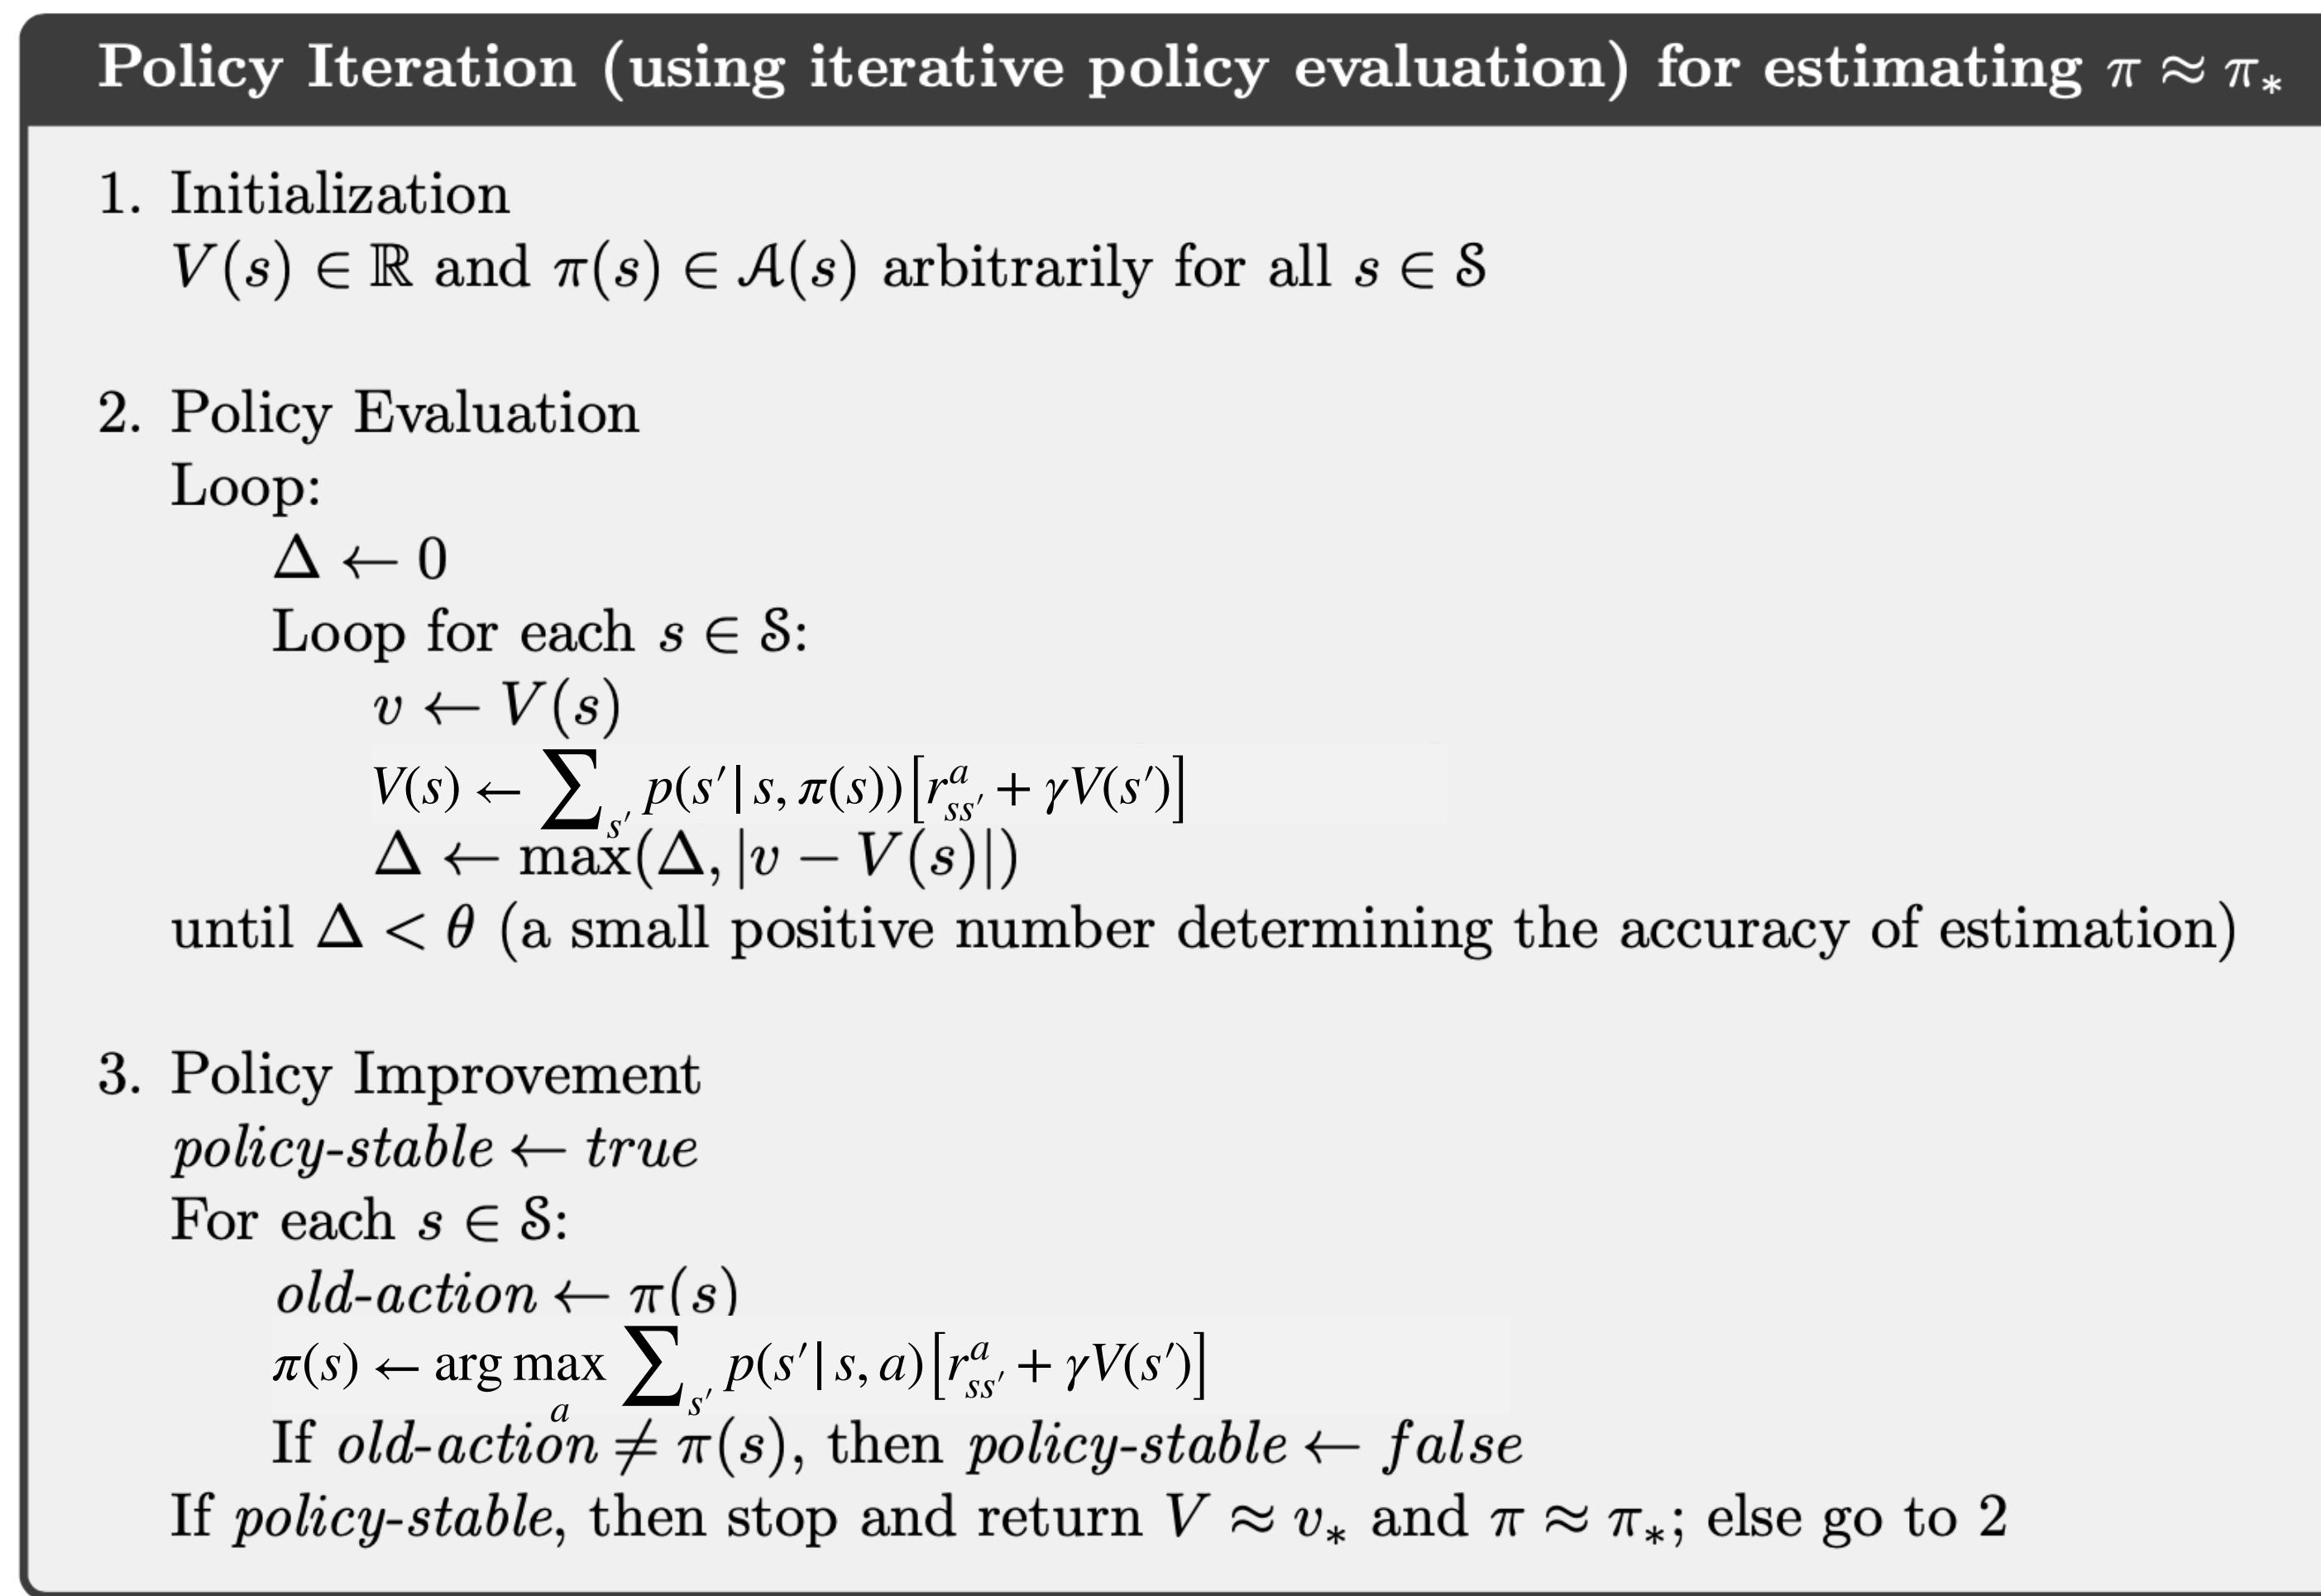
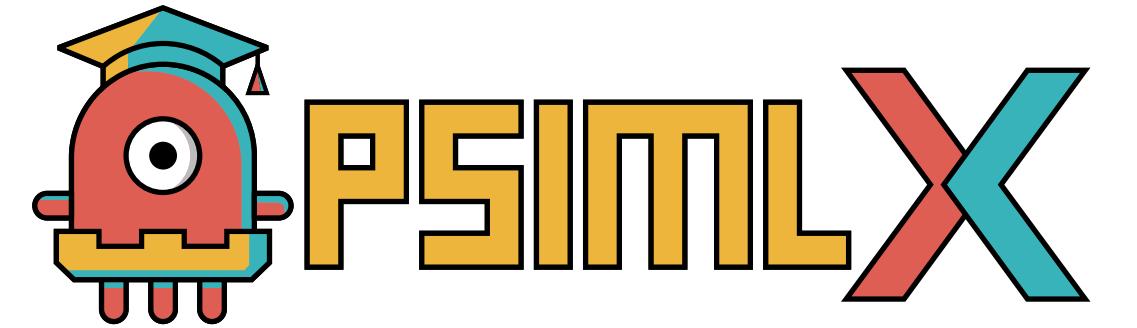
$$\pi_0 \xrightarrow{\text{Eval}} v_{\pi_0} \xrightarrow{\text{Impr}} \pi_1 \xrightarrow{\text{Eval}} v_{\pi_1} \xrightarrow{\text{Impr}} \dots \pi^* \xrightarrow{\text{Eval}} v_{\pi^*}$$

- Policy evaluation:** Estimate the true value $V \approx v_\pi$ iteratively
- Policy improvement:** Use estimated $V \approx v_\pi$ to select a better policy $\pi' \geq \pi$, $\pi' = \text{greedy}(V)$

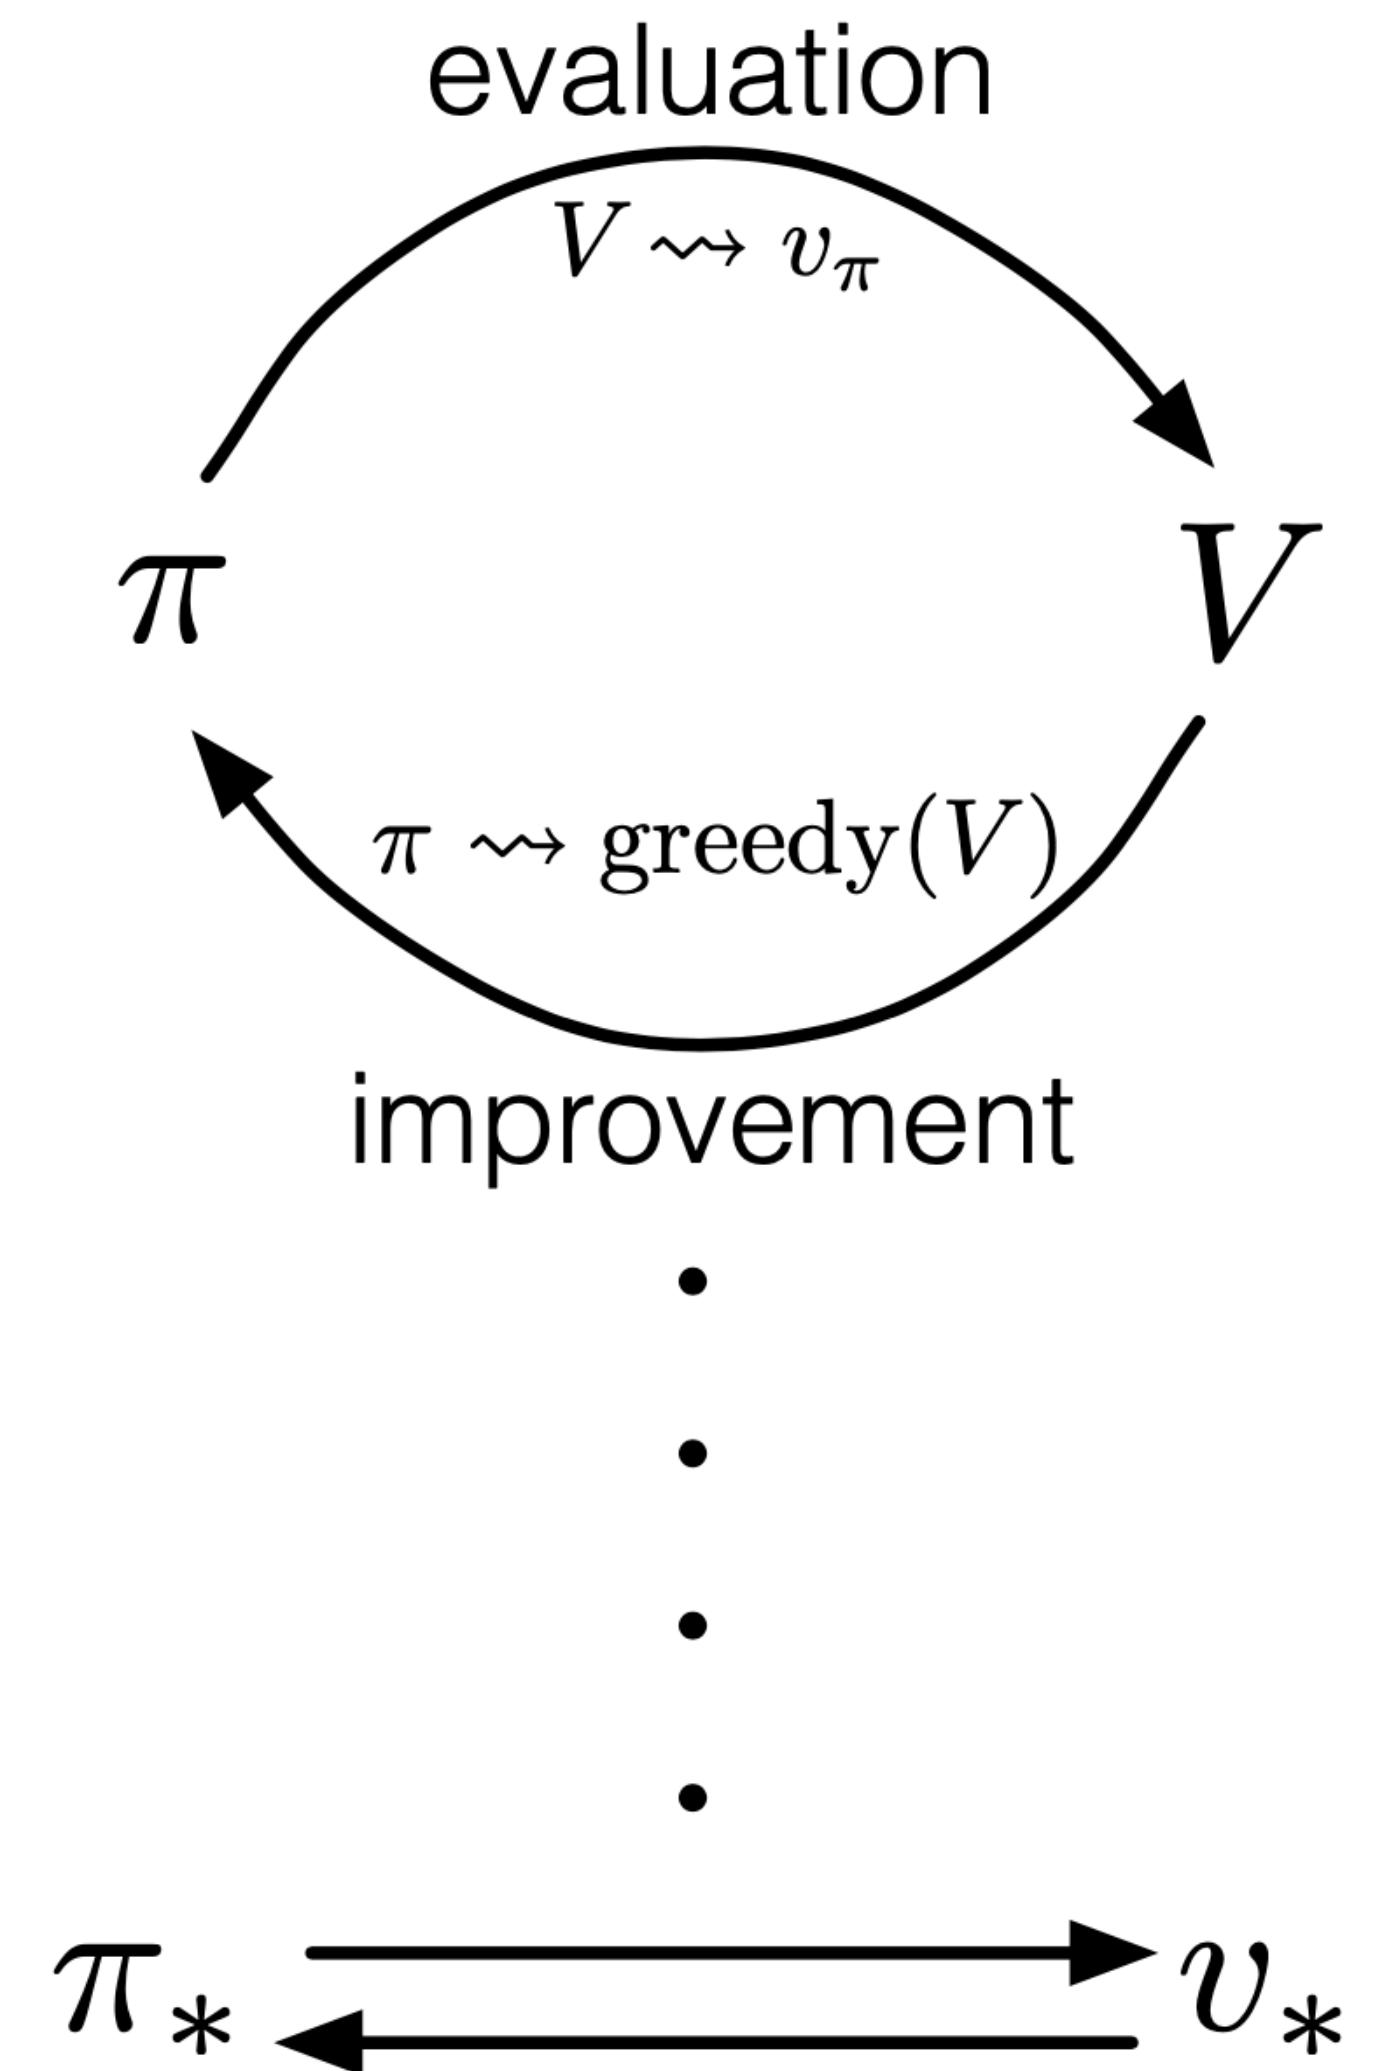


RL Formalisation

Generalised Policy Iteration (GPI)



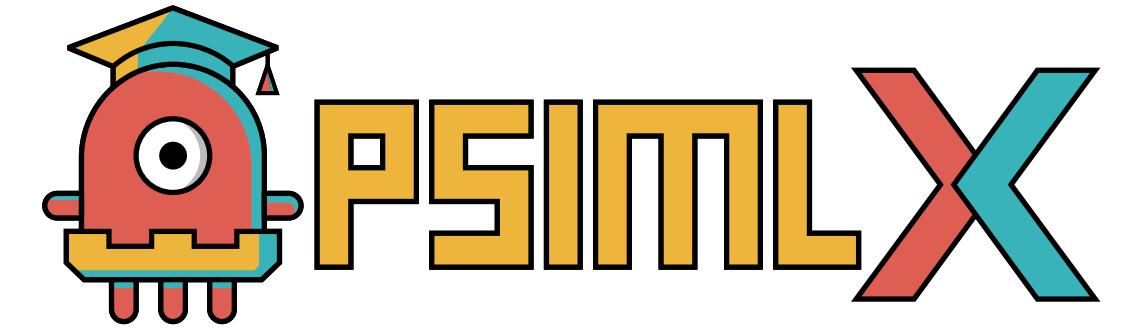
Policy Iteration algorithm [Sutton & Barto 2018]



GPI steps [Sutton & Barto 2018]

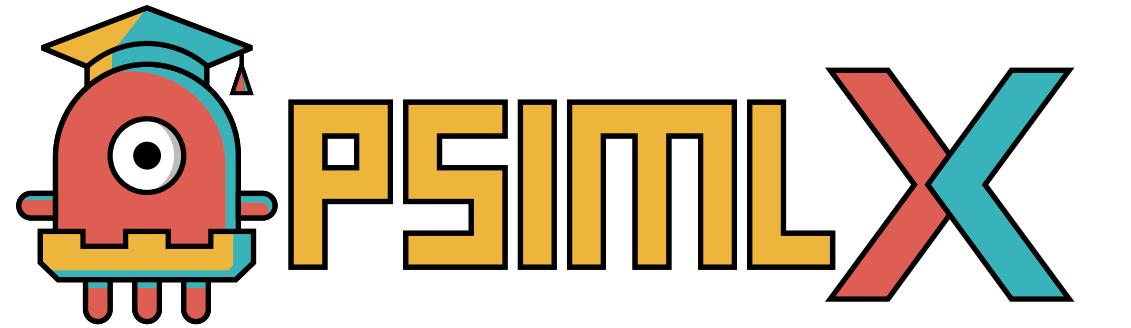
RL Formalisation

Exploration-Exploitation Trade-Off



- *Exploration*: Find more about the environment
- *Exploitation*: Utilise gained knowledge to garner higher returns
- The case of the agent tasked with garnering the highest return in a *continuing environment* throughout its entire lifetime:
 - If it commits to early-found schema for obtaining rewards, it may not find out possibly better schemas
 - If it overly explores, its return will suffer

Outline

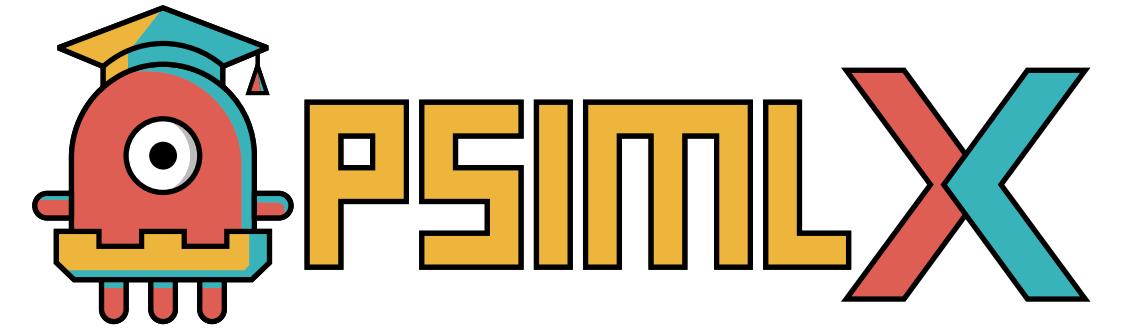


- Introduction
- Reinforcement Learning Formalisation
- **Model-Free Reinforcement Learning**
 - **Motivation**
 - **Monte Carlo and Temporal Difference Methods**
 - **MC and TD: Bias vs Variance Trade-Off**
 - **MC and TD: Future vs Previous Data**
 - **MC and TD: Summary**
 - **On-Policy vs Off-policy Learning**
 - **Example 1: First-visit MC**
 - **Example 2: Q-Learning**
 - **Conclusions**
 - **Beyond Tabular Methods**
- ...

Model free RL

Motivation

- MDP transition model $p(s' | s, a)$ is usually unknown, or using it is impractical
- Without it equations for v_π, q_π, v^*, q^* incomputable
- Two options:
 - Learn the model $p(s' | s, a)$, or use it to some degree if known – *model based RL (MBRL)*
 - Estimate v_π, q_π, v^*, q^* directly without learning the model – *model free RL (MFRL)*



Model free RL

Motivation

- MDP transition model $p(s' | s, a)$ is usually unknown, or using it is impractical
- Without it equations for v_π, q_π, v^*, q^* incomputable
- Two options:
 - Learn the model $p(s' | s, a)$, or use it to some degree if known – *model based RL (MBRL)*
 - Estimate v_π, q_π, v^*, q^* directly without learning the model – *model free RL (MFRL)*



Alpha Go, Silver et al. 2016



OpenAI Five, Barner et al. 2019

Model free RL

Motivation

- MDP transition model $p(s' | s, a)$ is usually unknown, or using it is impractical
- Without it equations for v_π, q_π, v^*, q^* incomputable
- Two options:
 - Learn the model $p(s' | s, a)$, or use it to some degree if known – *model based RL (MBRL)*
 - **Estimate v_π, q_π, v^*, q^* directly without learning the model – *model free RL (MFRL*)**
 - $V_t(s), Q_t(s, a)$ are (imperfect) estimates to v_π, q_π at computation time step t



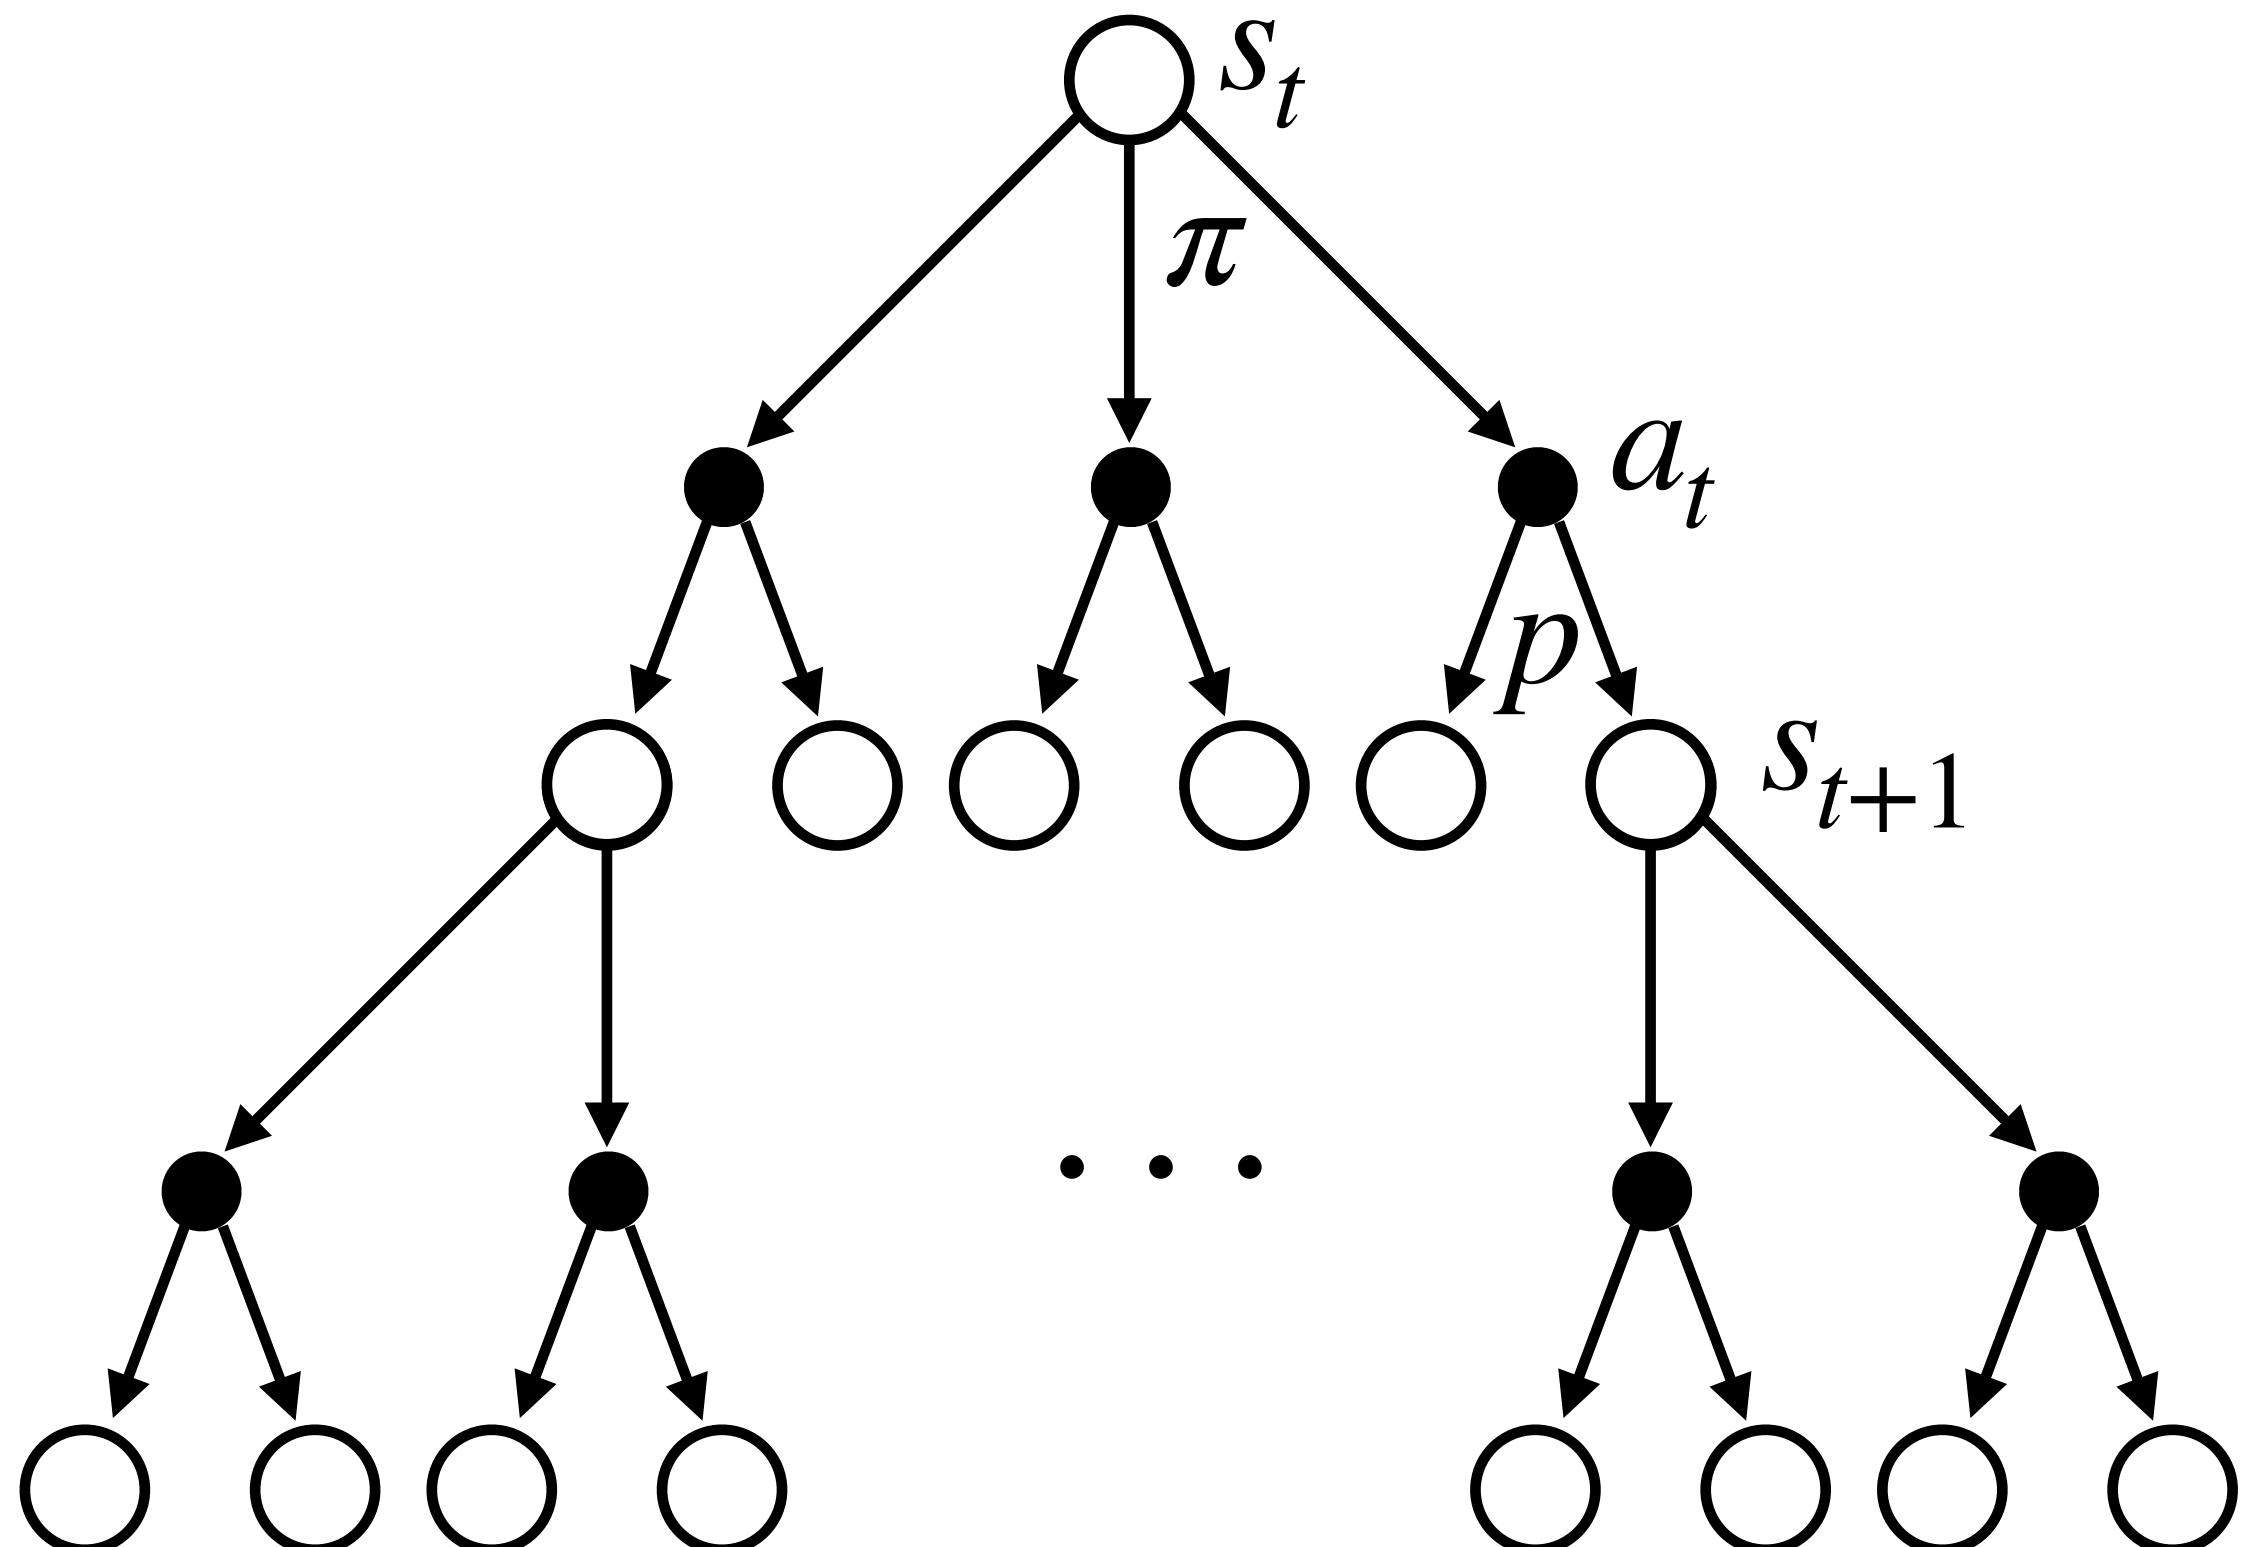
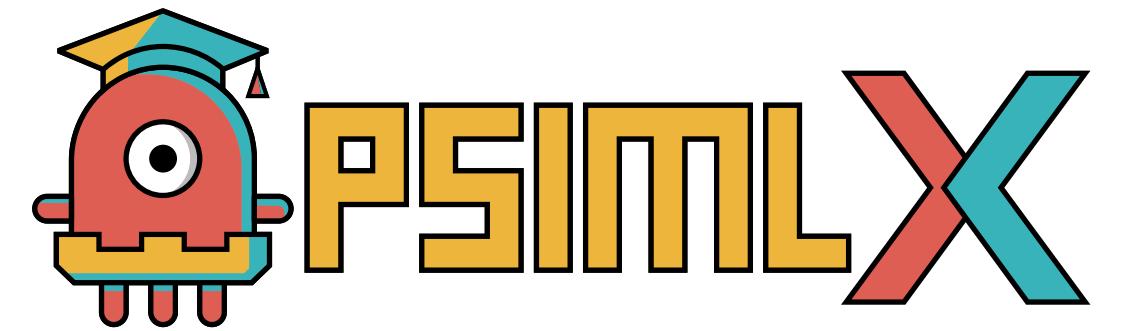
OpenAI Five, Barner et al. 2019

Model free RL

Motivation

$$\begin{aligned} v_\pi(s) &\doteq \mathbb{E}_\pi[G_t | S_t = s] \\ &= \mathbb{E}_\pi[r_{t+1} + \gamma G_{t+1} | S_t = s] \end{aligned}$$

Between MDPs and SMDPs [Sutton et al. 1999]

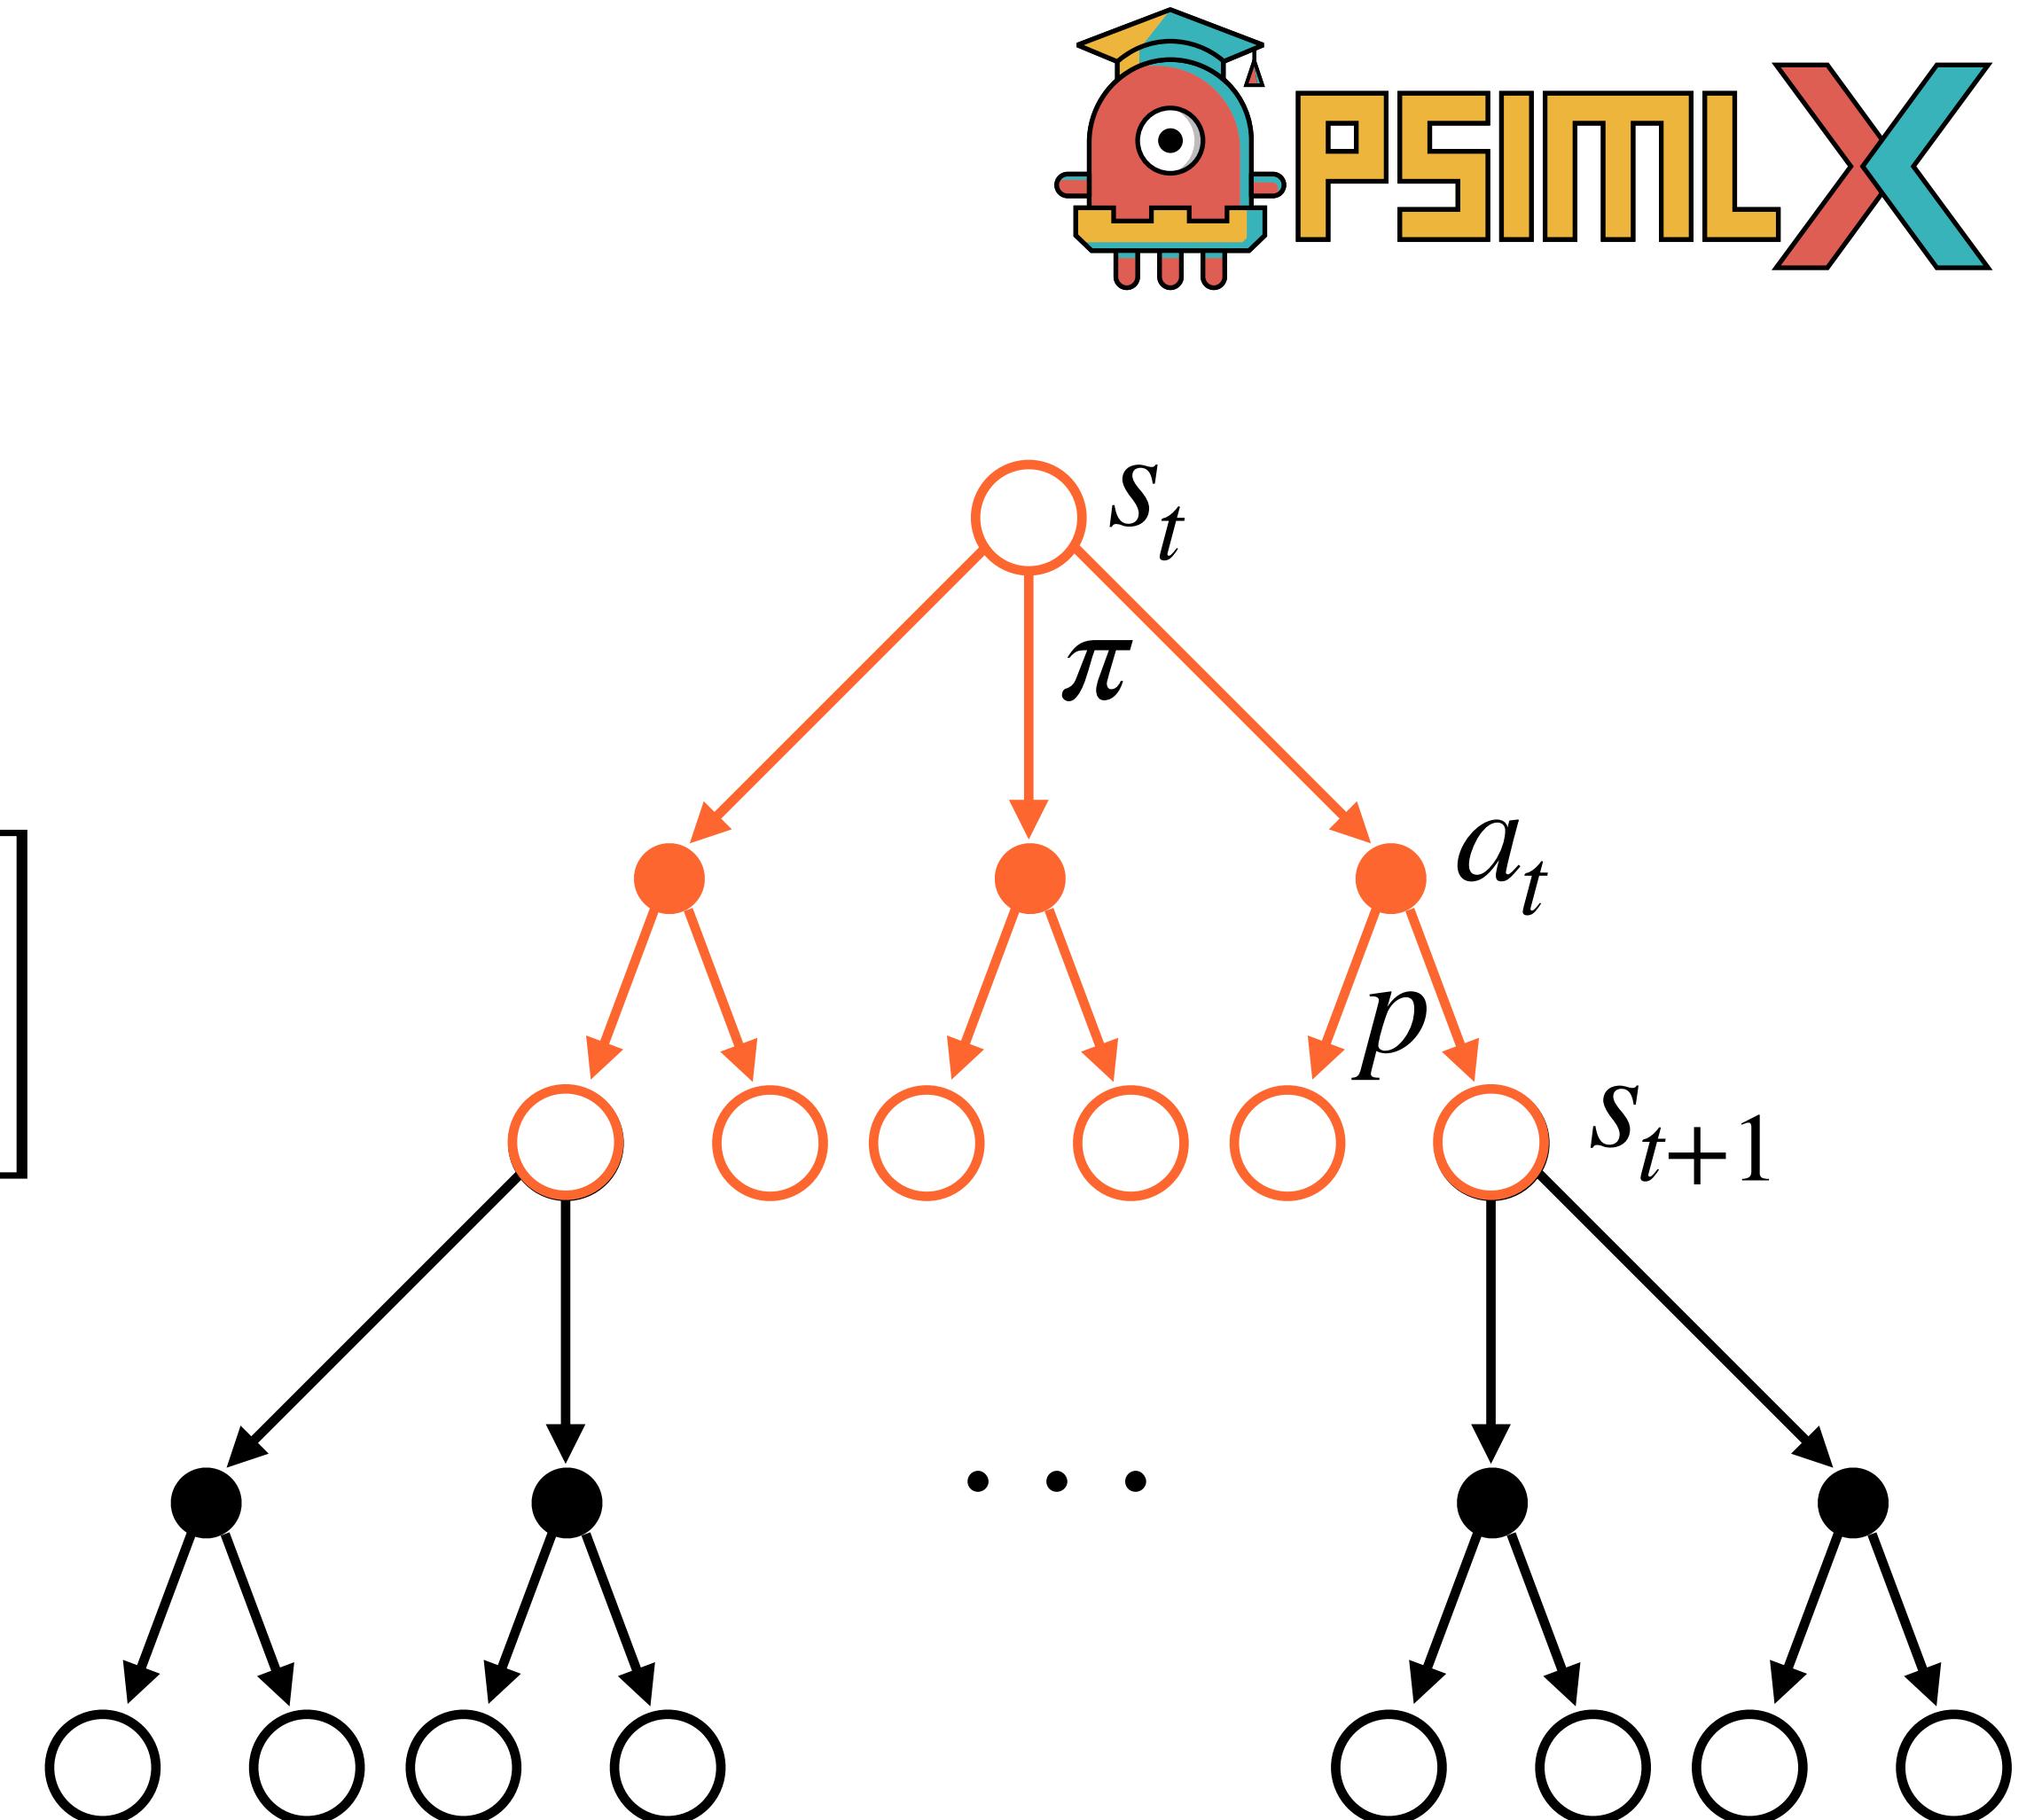


Model free RL

Motivation

$$\begin{aligned}
 v_\pi(s) &\doteq \mathbb{E}_\pi[G_t | S_t = s] \\
 &= \mathbb{E}_\pi[r_{t+1} + \gamma G_{t+1} | S_t = s] \\
 &= \sum_a \pi(a | s) \left[r_s^a + \gamma \sum_{s'} p(s' | s, a) v_\pi(s') \right]
 \end{aligned}$$

Between MDPs and SMDPs [Sutton et al. 1999]

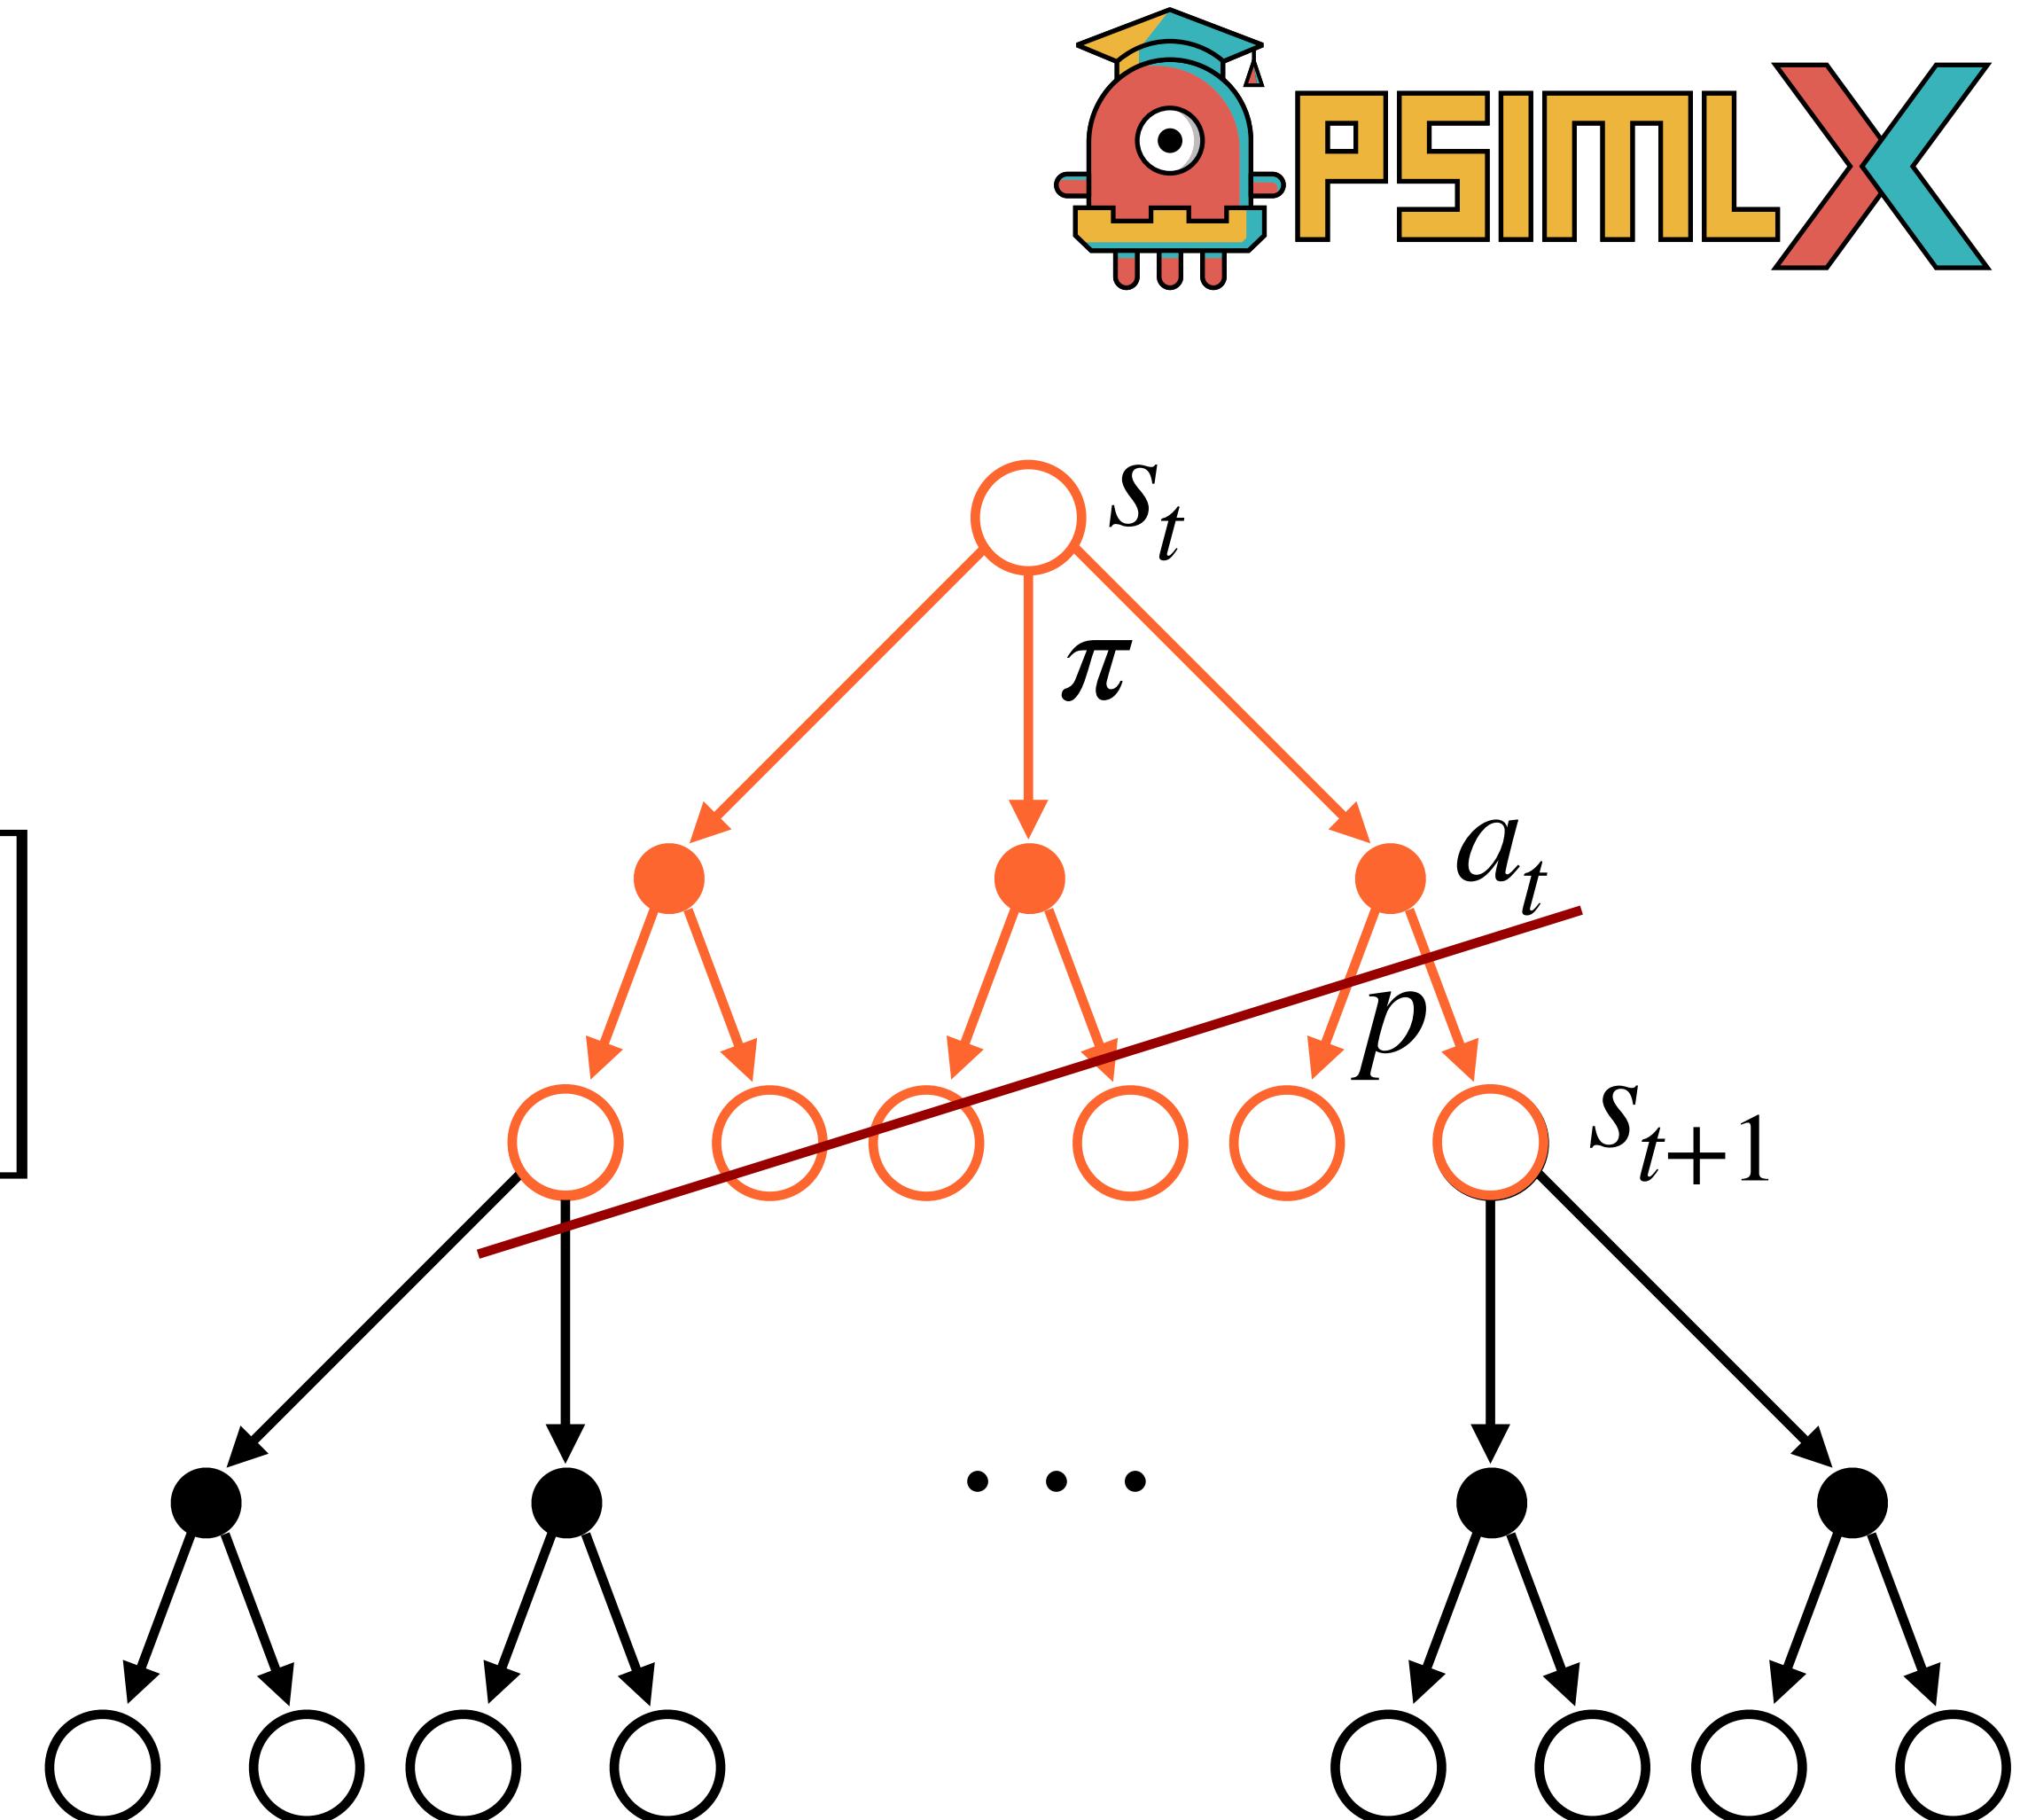


Model free RL

Motivation

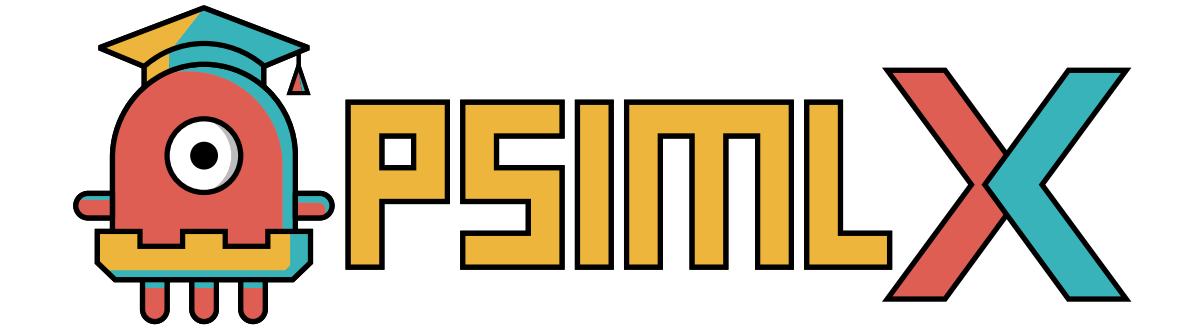
$$\begin{aligned}
 v_\pi(s) &\doteq \mathbb{E}_\pi[G_t | S_t = s] \\
 &= \mathbb{E}_\pi[r_{t+1} + \gamma G_{t+1} | S_t = s] \\
 &= \sum_a \pi(a | s) \left[r_s^a + \gamma \sum_{s'} p(s' | s, a) v_\pi(s') \right]
 \end{aligned}$$

Between MDPs and SMDPs [Sutton et al. 1999]



Model free RL

Monte Carlo (MC) and Temporal Difference (TD)



Monte Carlo (MC) methods

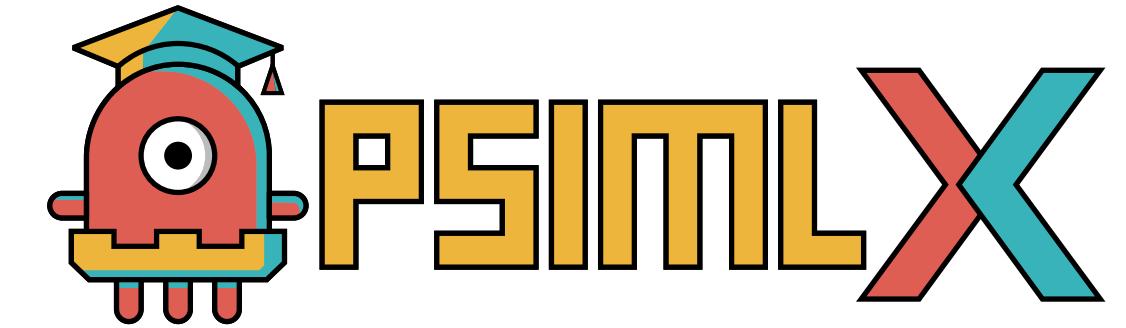
- **Core idea:**
 - Play entire episodes using a fixed policy π and estimate state values v_π as empirical means of returns
 - Can only be applied to *episodic tasks*

Temporal Difference (TD) methods

- **Core idea:**
 - Utilise the recursive nature of the Bellman equation to update state values v_π based on states agent transitions to
 - Learn from incomplete episodes, can be applied to *continuing tasks*
 - They *bootstrap* – instead of measuring the true return G_t they use $V(S_t)$ as its estimate, which in turn is also an estimate of the true $v_\pi(S_t)$

Model free RL

Monte Carlo (MC) and Temporal Difference (TD)



Criteria	Monte Carlo Methods	Temporal Difference Methods
Bias vs Variance		
Online		✓
Bootstrapping		✓ because v_t is based off of v_{t+1}
Estimation		
On-Policy		
Off-Policy		
Past vs Future Future data		

Model free RL

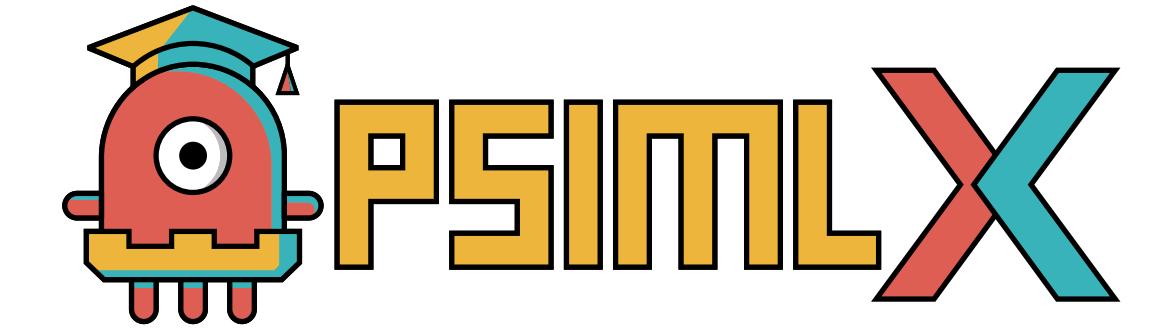
Monte Carlo (MC) and Temporal Difference (TD)

$$\begin{aligned}
 V(s)_{n+1} &= \frac{1}{n} \sum_{i=1}^n G_{s_i} \\
 &= \frac{1}{n} \left(G_{s_n} + \sum_{i=1}^{(n-1)} G_{s_i} \right) \\
 &= \frac{1}{n} \left(G_{s_n} + (n-1) \cdot \frac{1}{n-1} \sum_{i=1}^{(n-1)} G_{s_i} \right) \\
 &= \frac{1}{n} \left(G_{s_n} - (n-1) \cdot V(s)_n \right) \\
 &= V(s)_n + \frac{1}{n} \left(G_{s_n} - V(s)_n \right)
 \end{aligned}$$

Monte Carlo (MC) Update

$$V(s) \leftarrow V(s) + \frac{1}{n} [G_t - V(s)]$$

$$\begin{aligned}
 G_t &= R + \gamma G_{t+1} \\
 &\approx R + \gamma V_{t+1}
 \end{aligned}$$

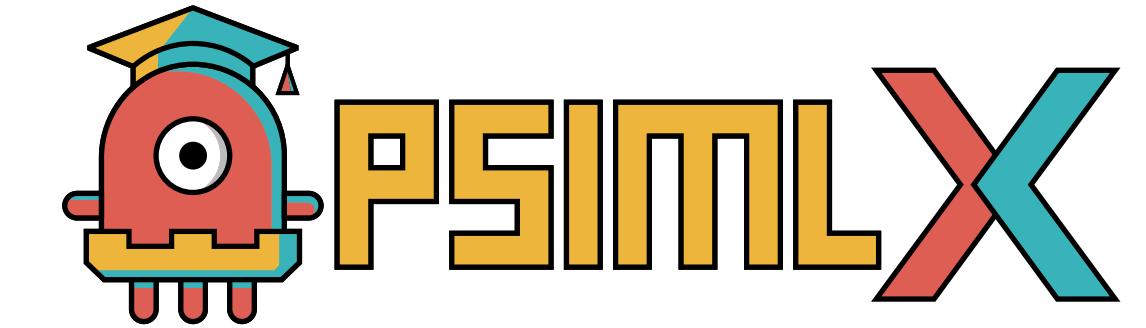


Temporal Difference (TD) Update

$$V(s) \leftarrow V(s) + \alpha [R + \gamma V(s') - V(s)]$$

Model free RL

Monte Carlo (MC) and Temporal Difference (TD)



$$\begin{aligned}
 V(s)_{n+1} &= \frac{1}{n} \sum_{i=1}^n G_{s_i} \\
 &= \frac{1}{n} \left(G_{s_n} + \sum_{i=1}^{(n-1)} G_{s_i} \right) \\
 &= \frac{1}{n} \left(G_{s_n} + (n - 1)V(s)_n \right) \\
 &= V(s)_n + \frac{1}{n} \left(G_{s_n} - V(s)_n \right)
 \end{aligned}$$

$$\begin{aligned}
 G_t &= R + \gamma G_{t+1} \\
 &\approx R + \gamma V_{t+1}
 \end{aligned}$$

Monte Carlo (MC) Update

$$V(s) \leftarrow V(s) + \frac{1}{n} [G_t - V(s)]$$

step size
target
error

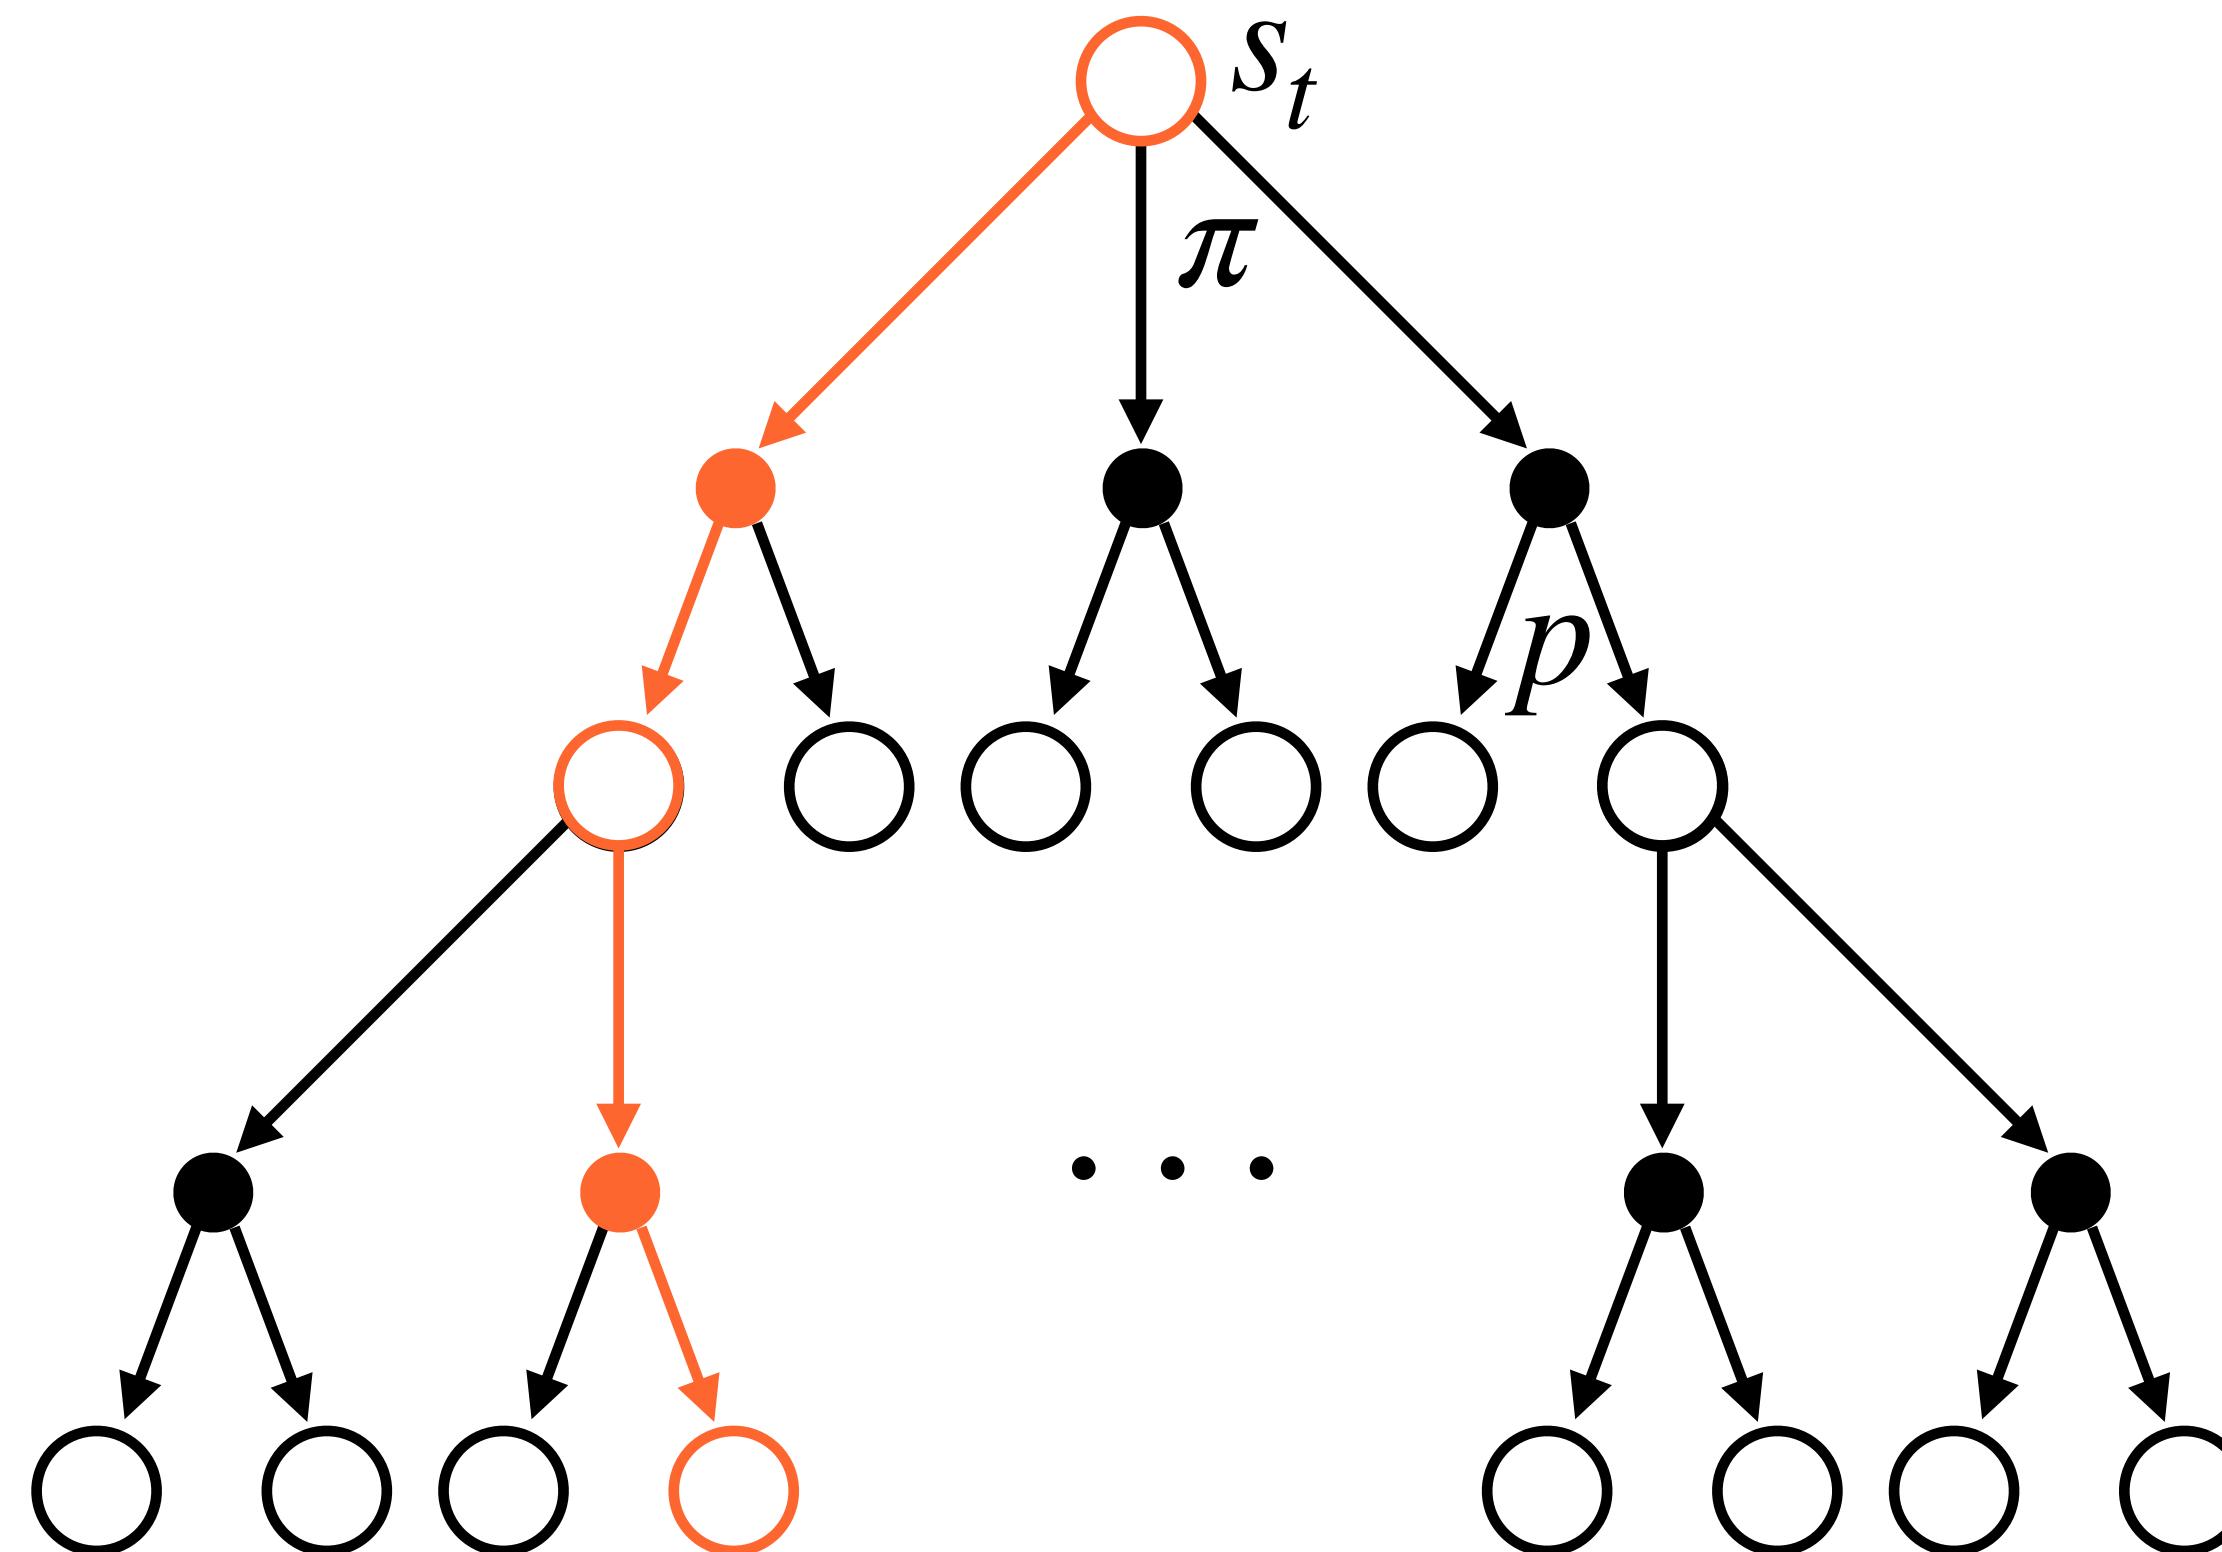
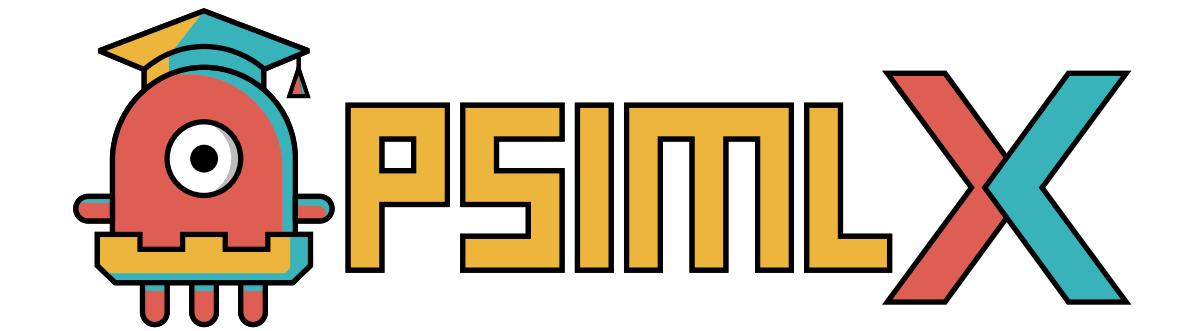
Temporal Difference (TD) Update

$$V(s) \leftarrow V(s) + \alpha [R + \gamma V(s') - V(s)]$$

δ_t

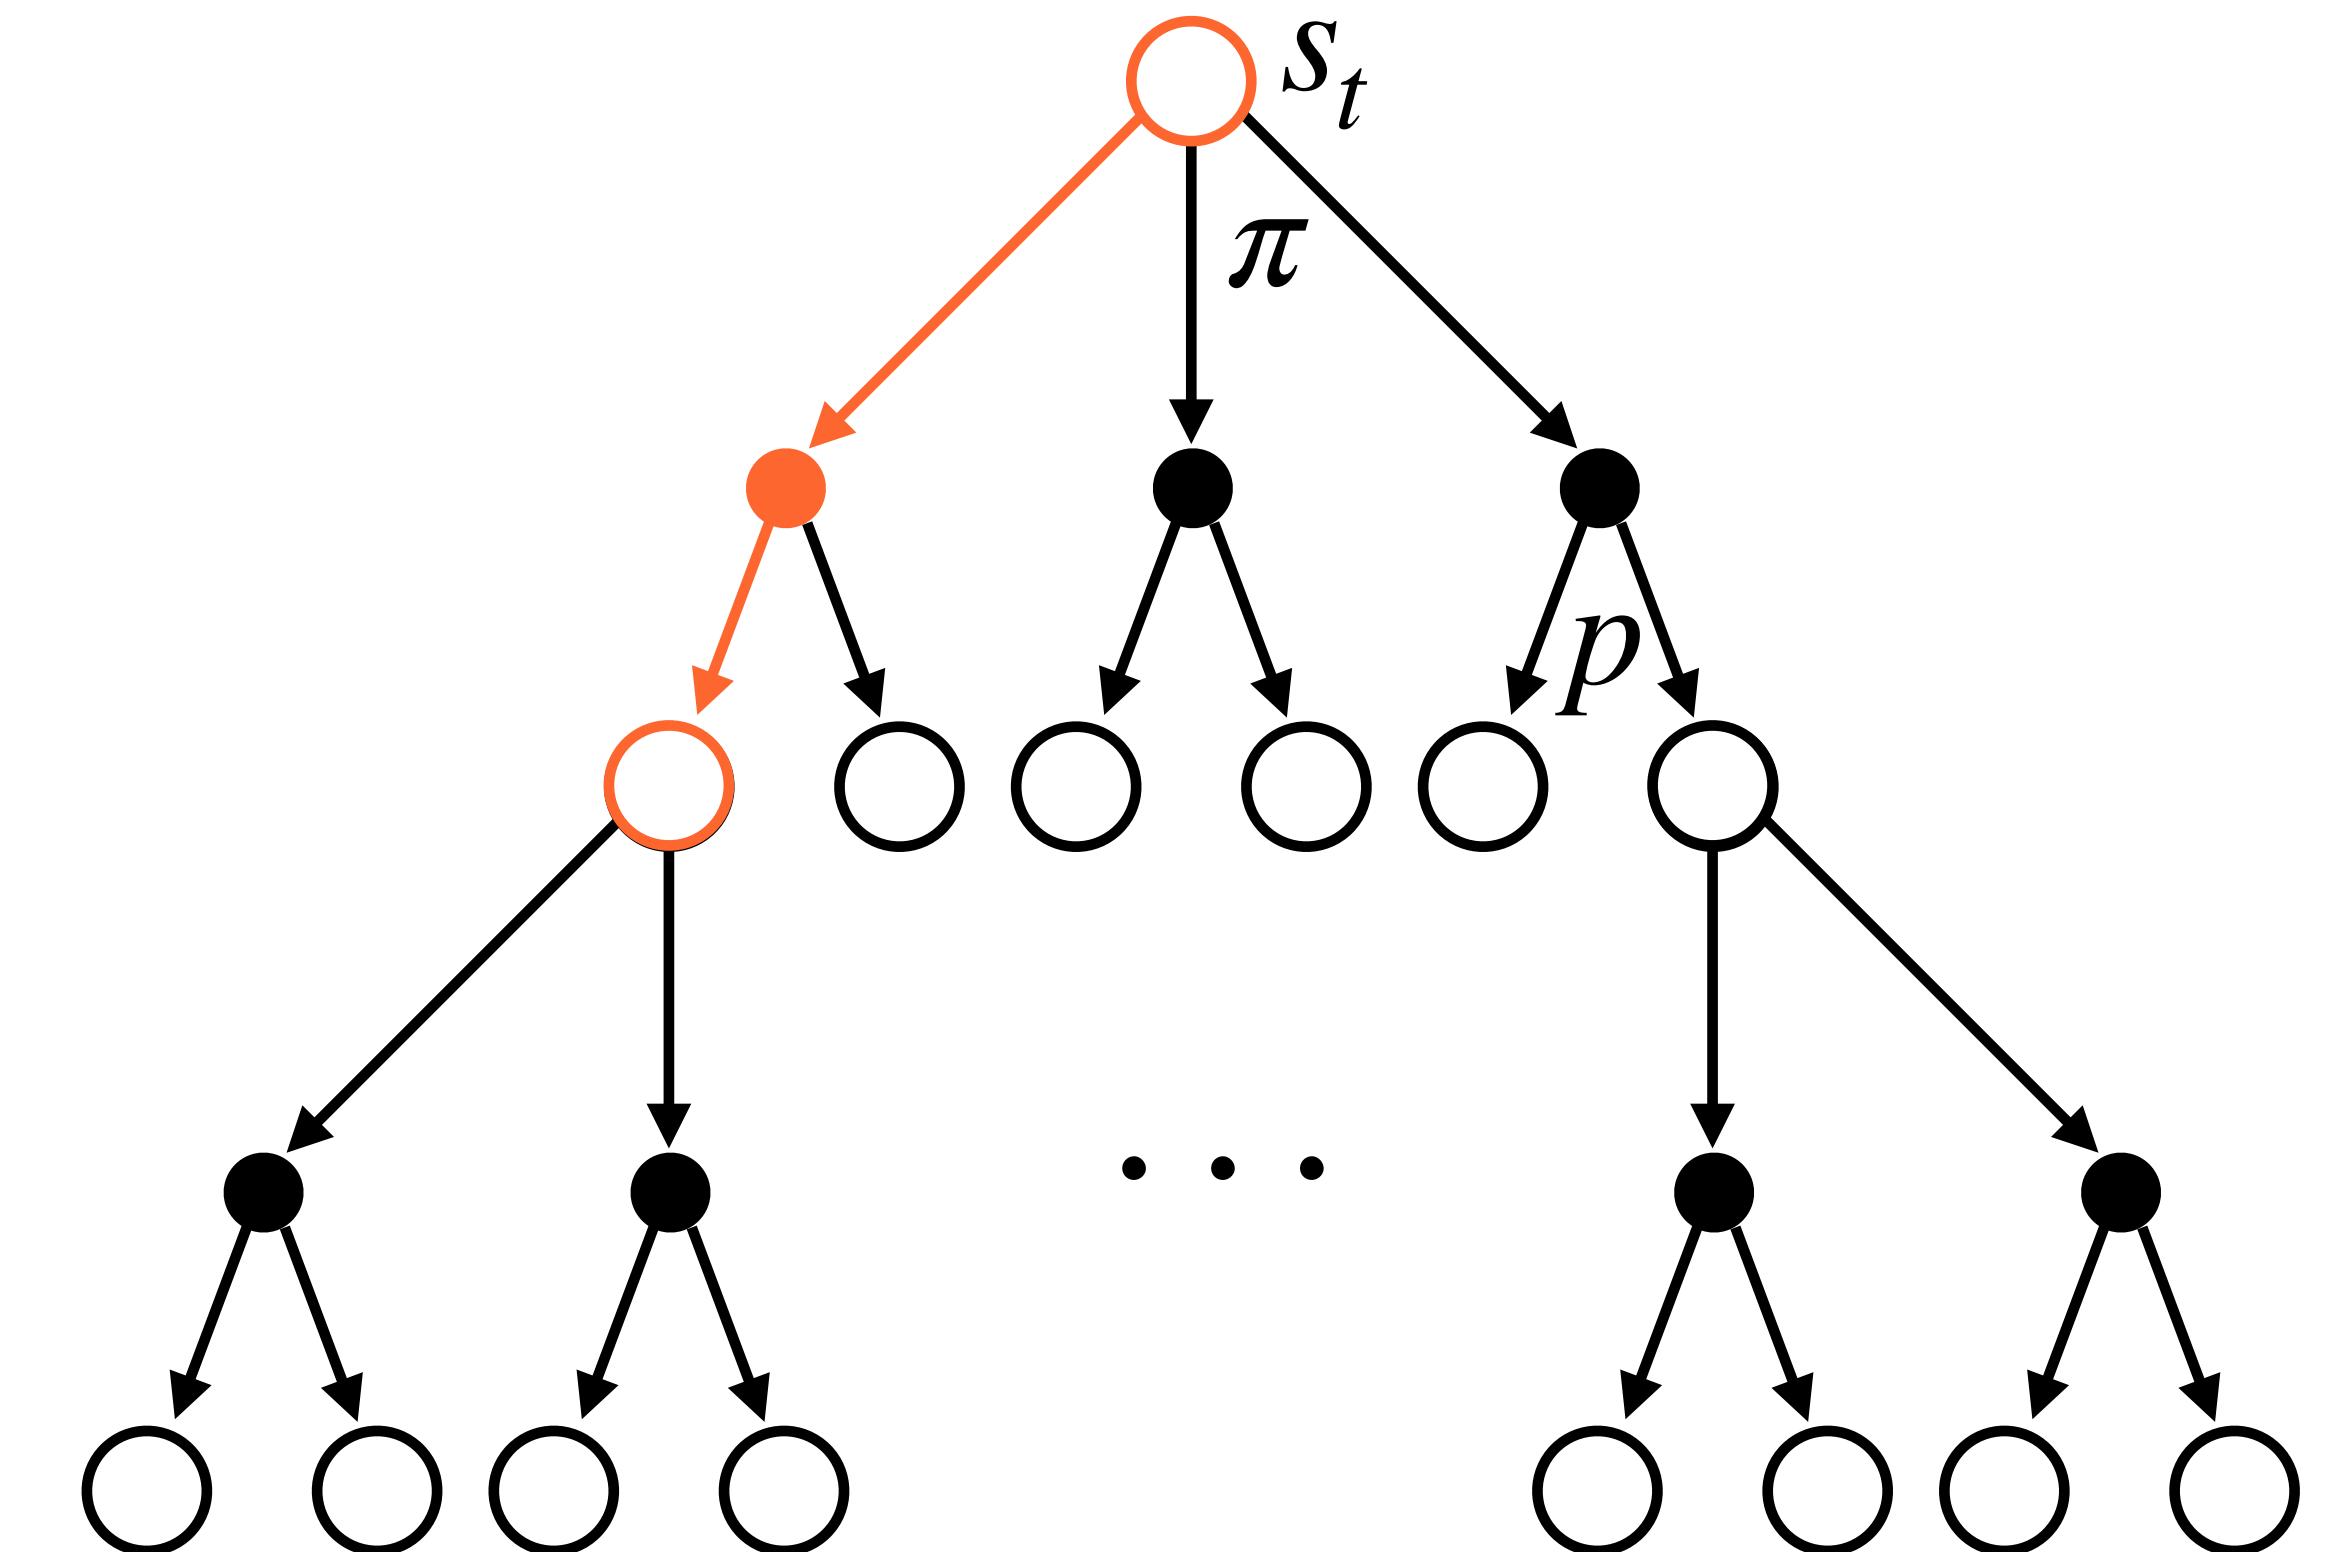
Model free RL

Monte Carlo (MC) and Temporal Difference (TD)



Monte Carlo (MC) Update

$$V(s) \leftarrow V(s) + \frac{1}{n} [G_t - V(s)]$$

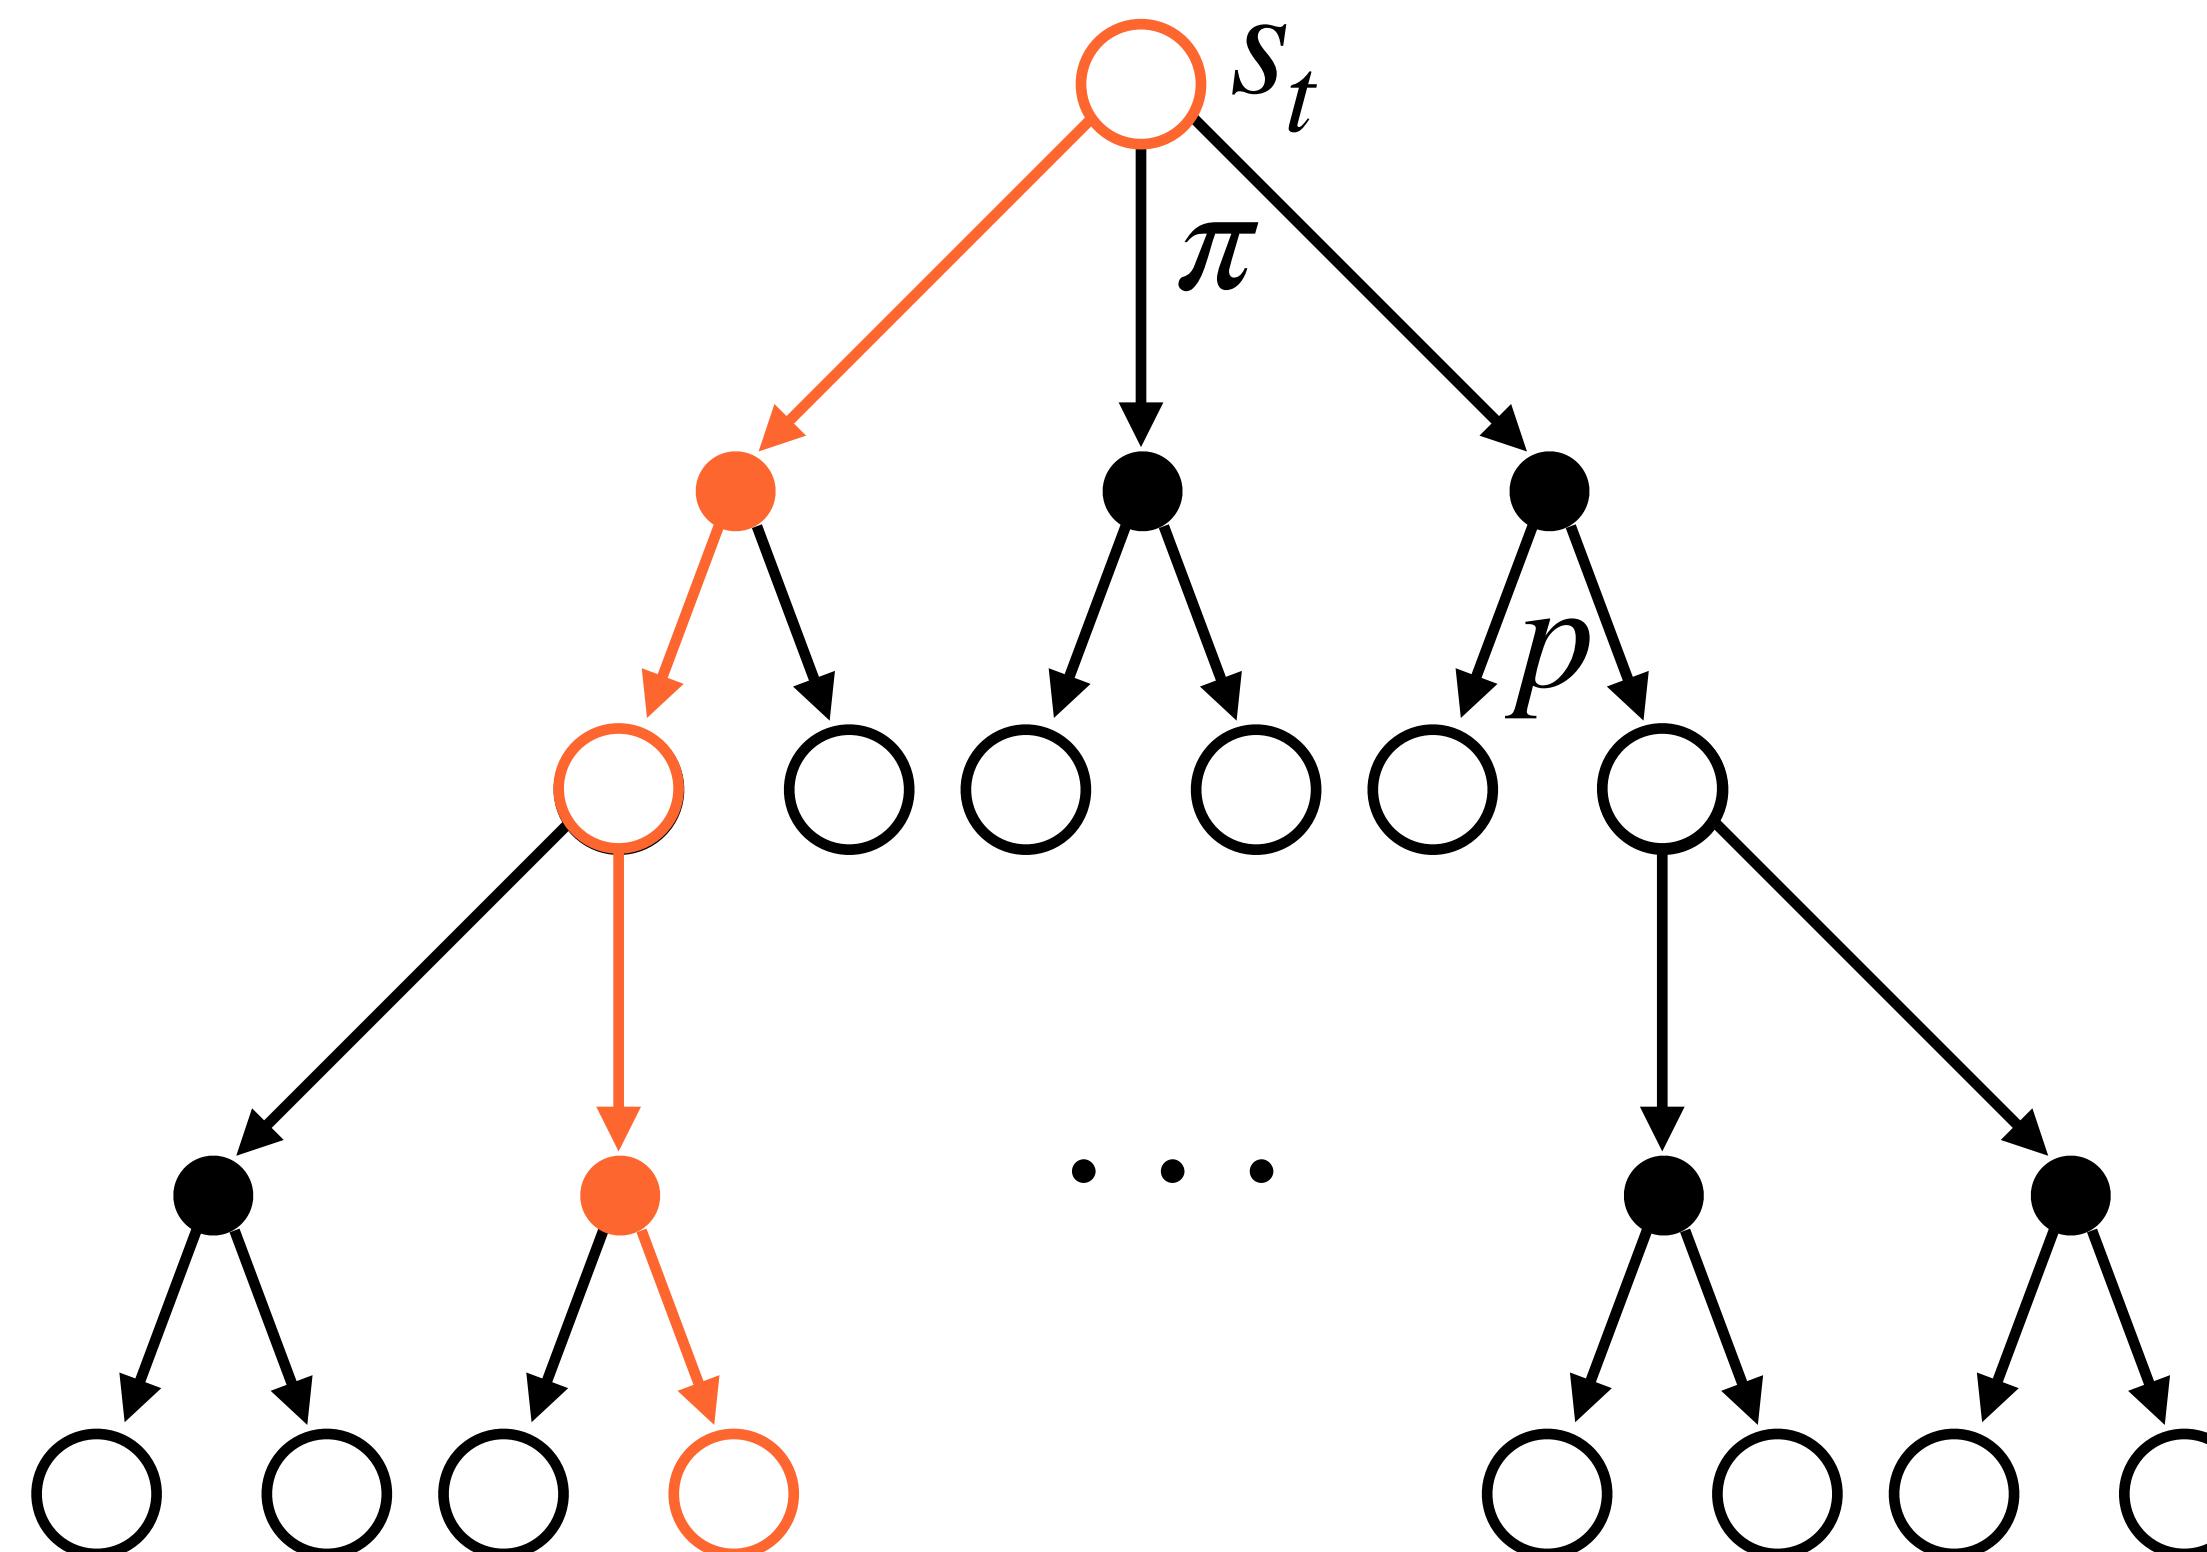
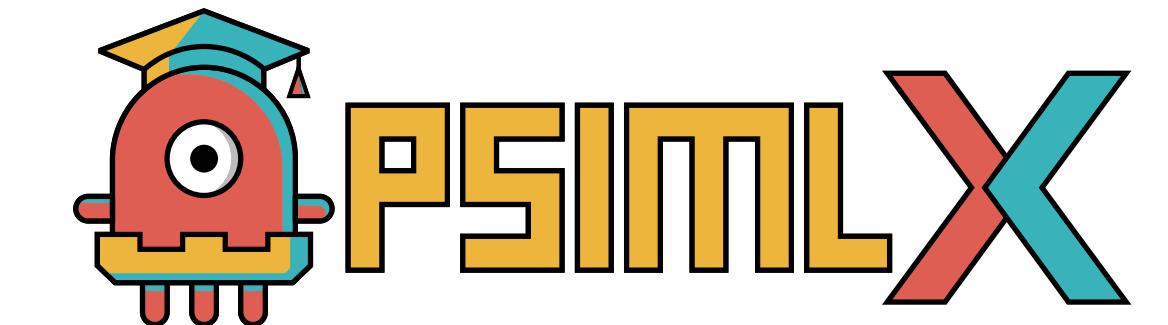


Temporal Difference (TD) Update

$$V(s) \leftarrow V(s) + \alpha [R + \gamma V(s') - V(s)]$$

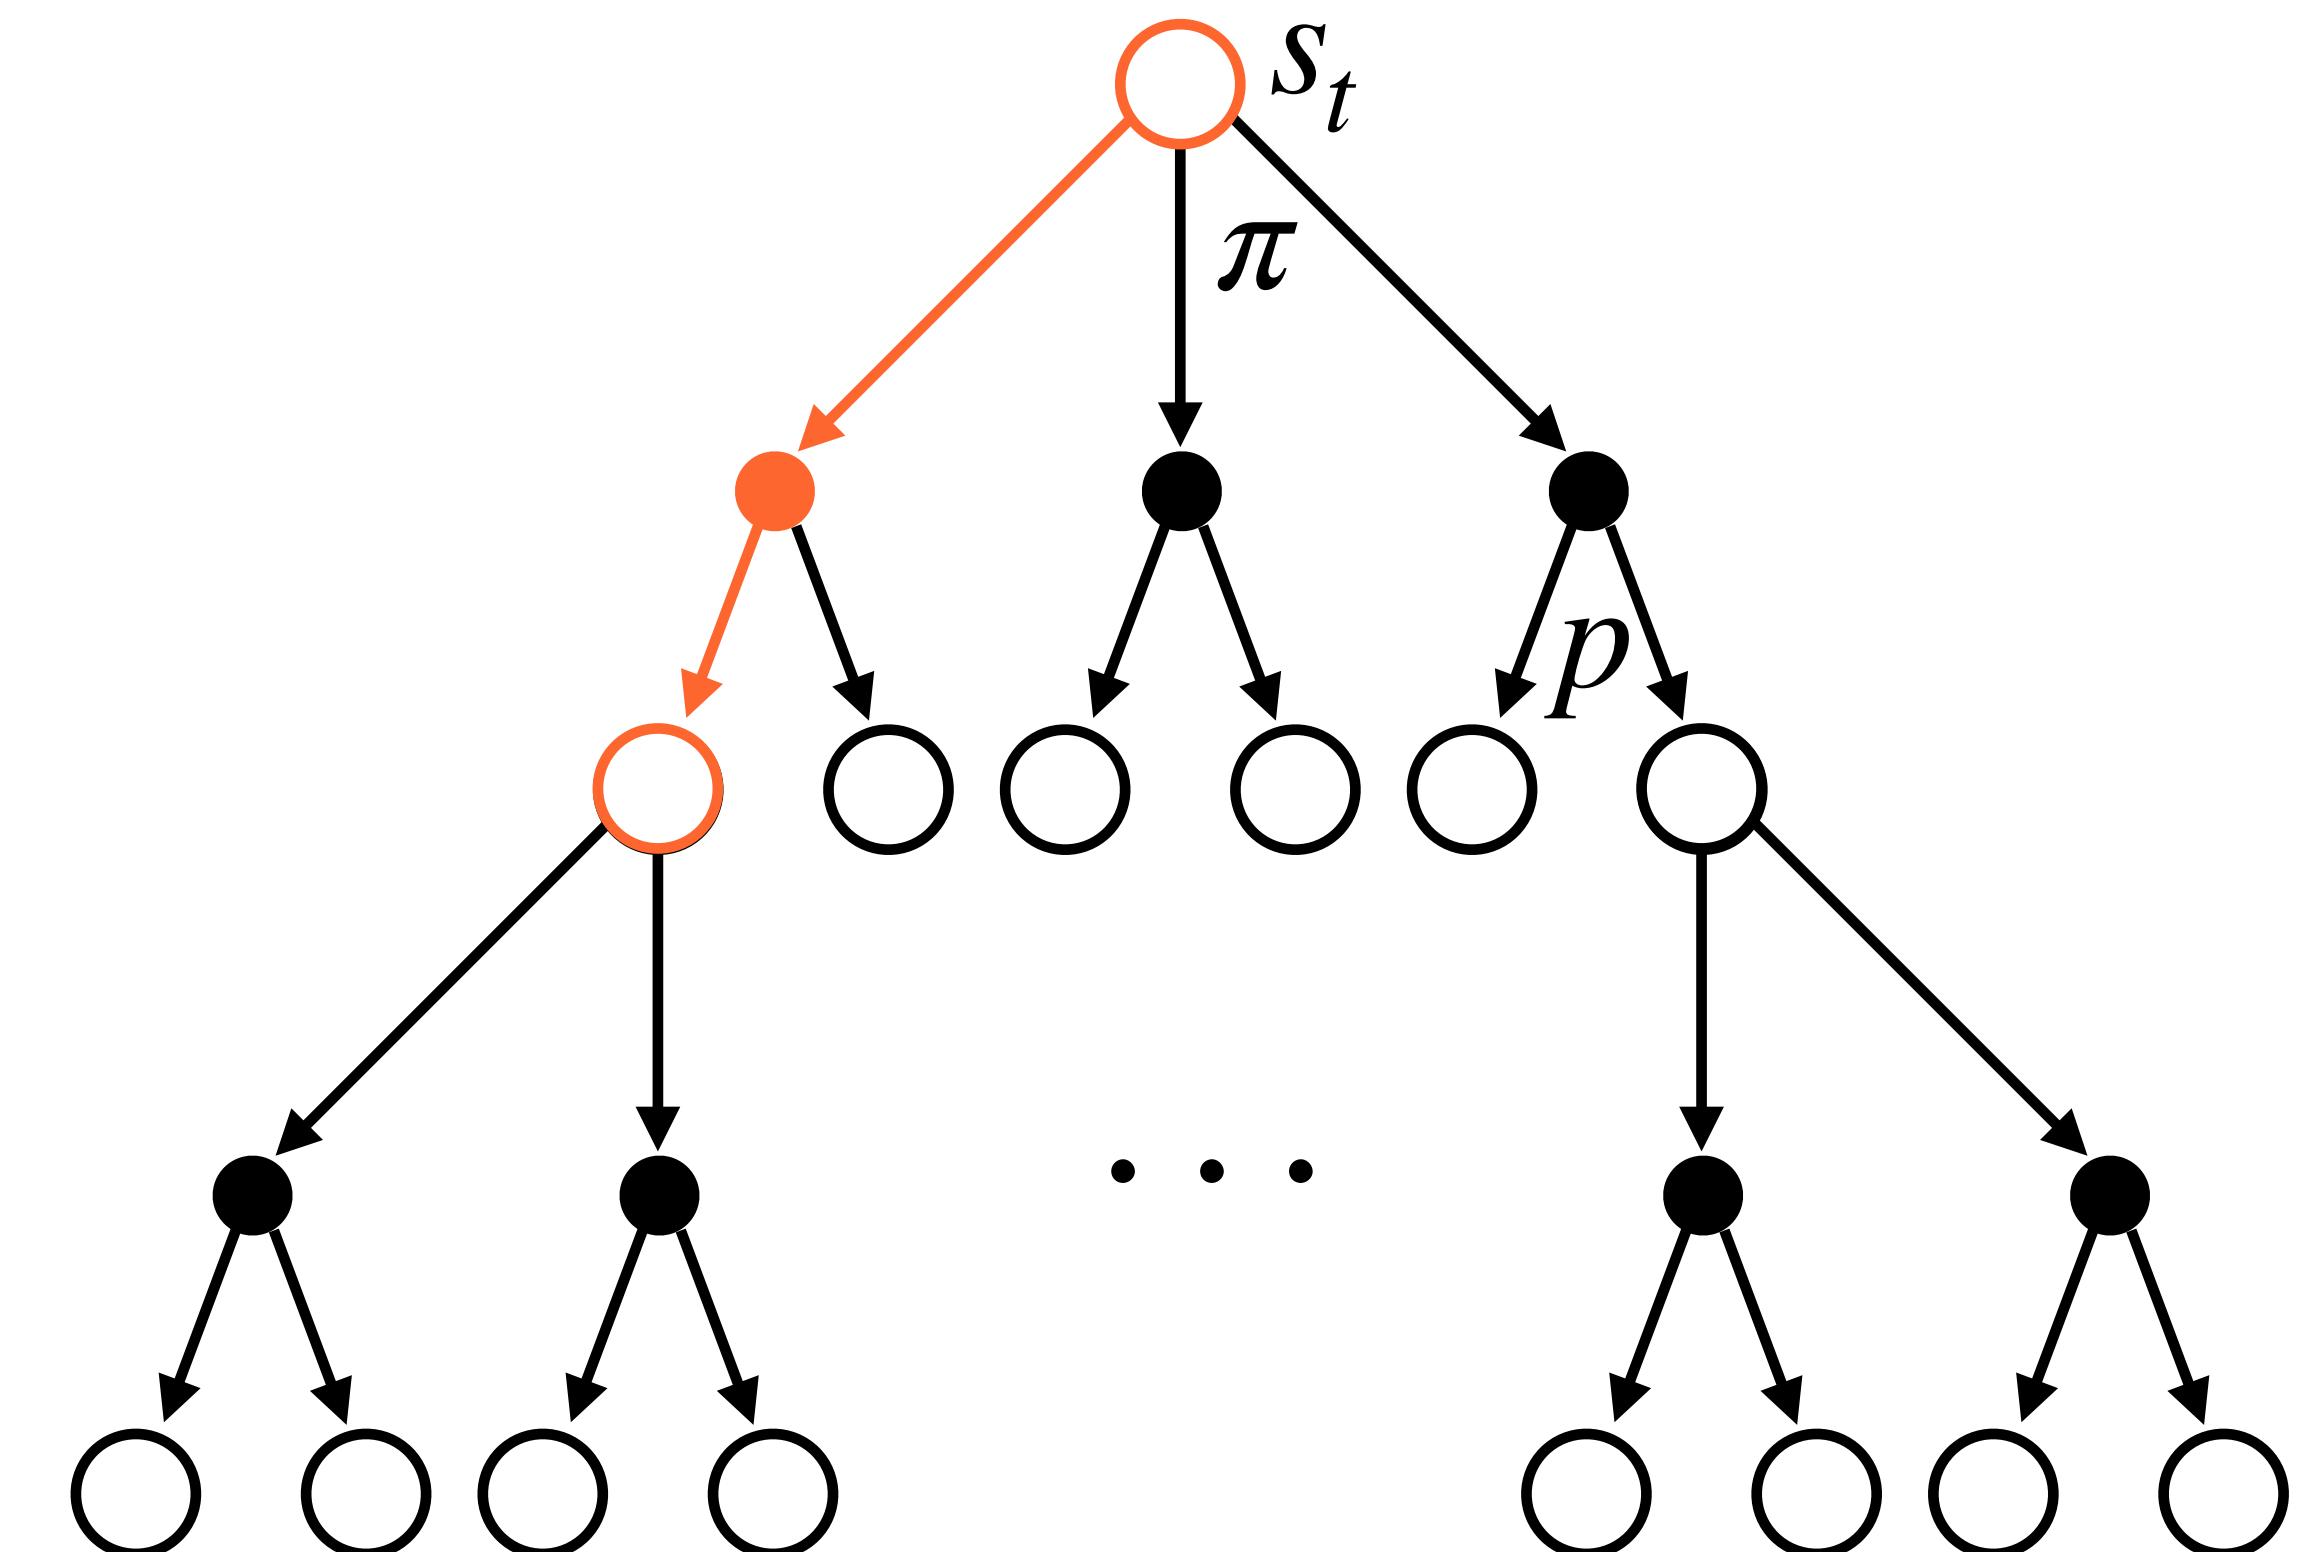
Model free RL

Monte Carlo (MC) and Temporal Difference (TD)



Monte Carlo (MC) Update

- Is an estimate because expectation over return is not known and is estimated by sample mean

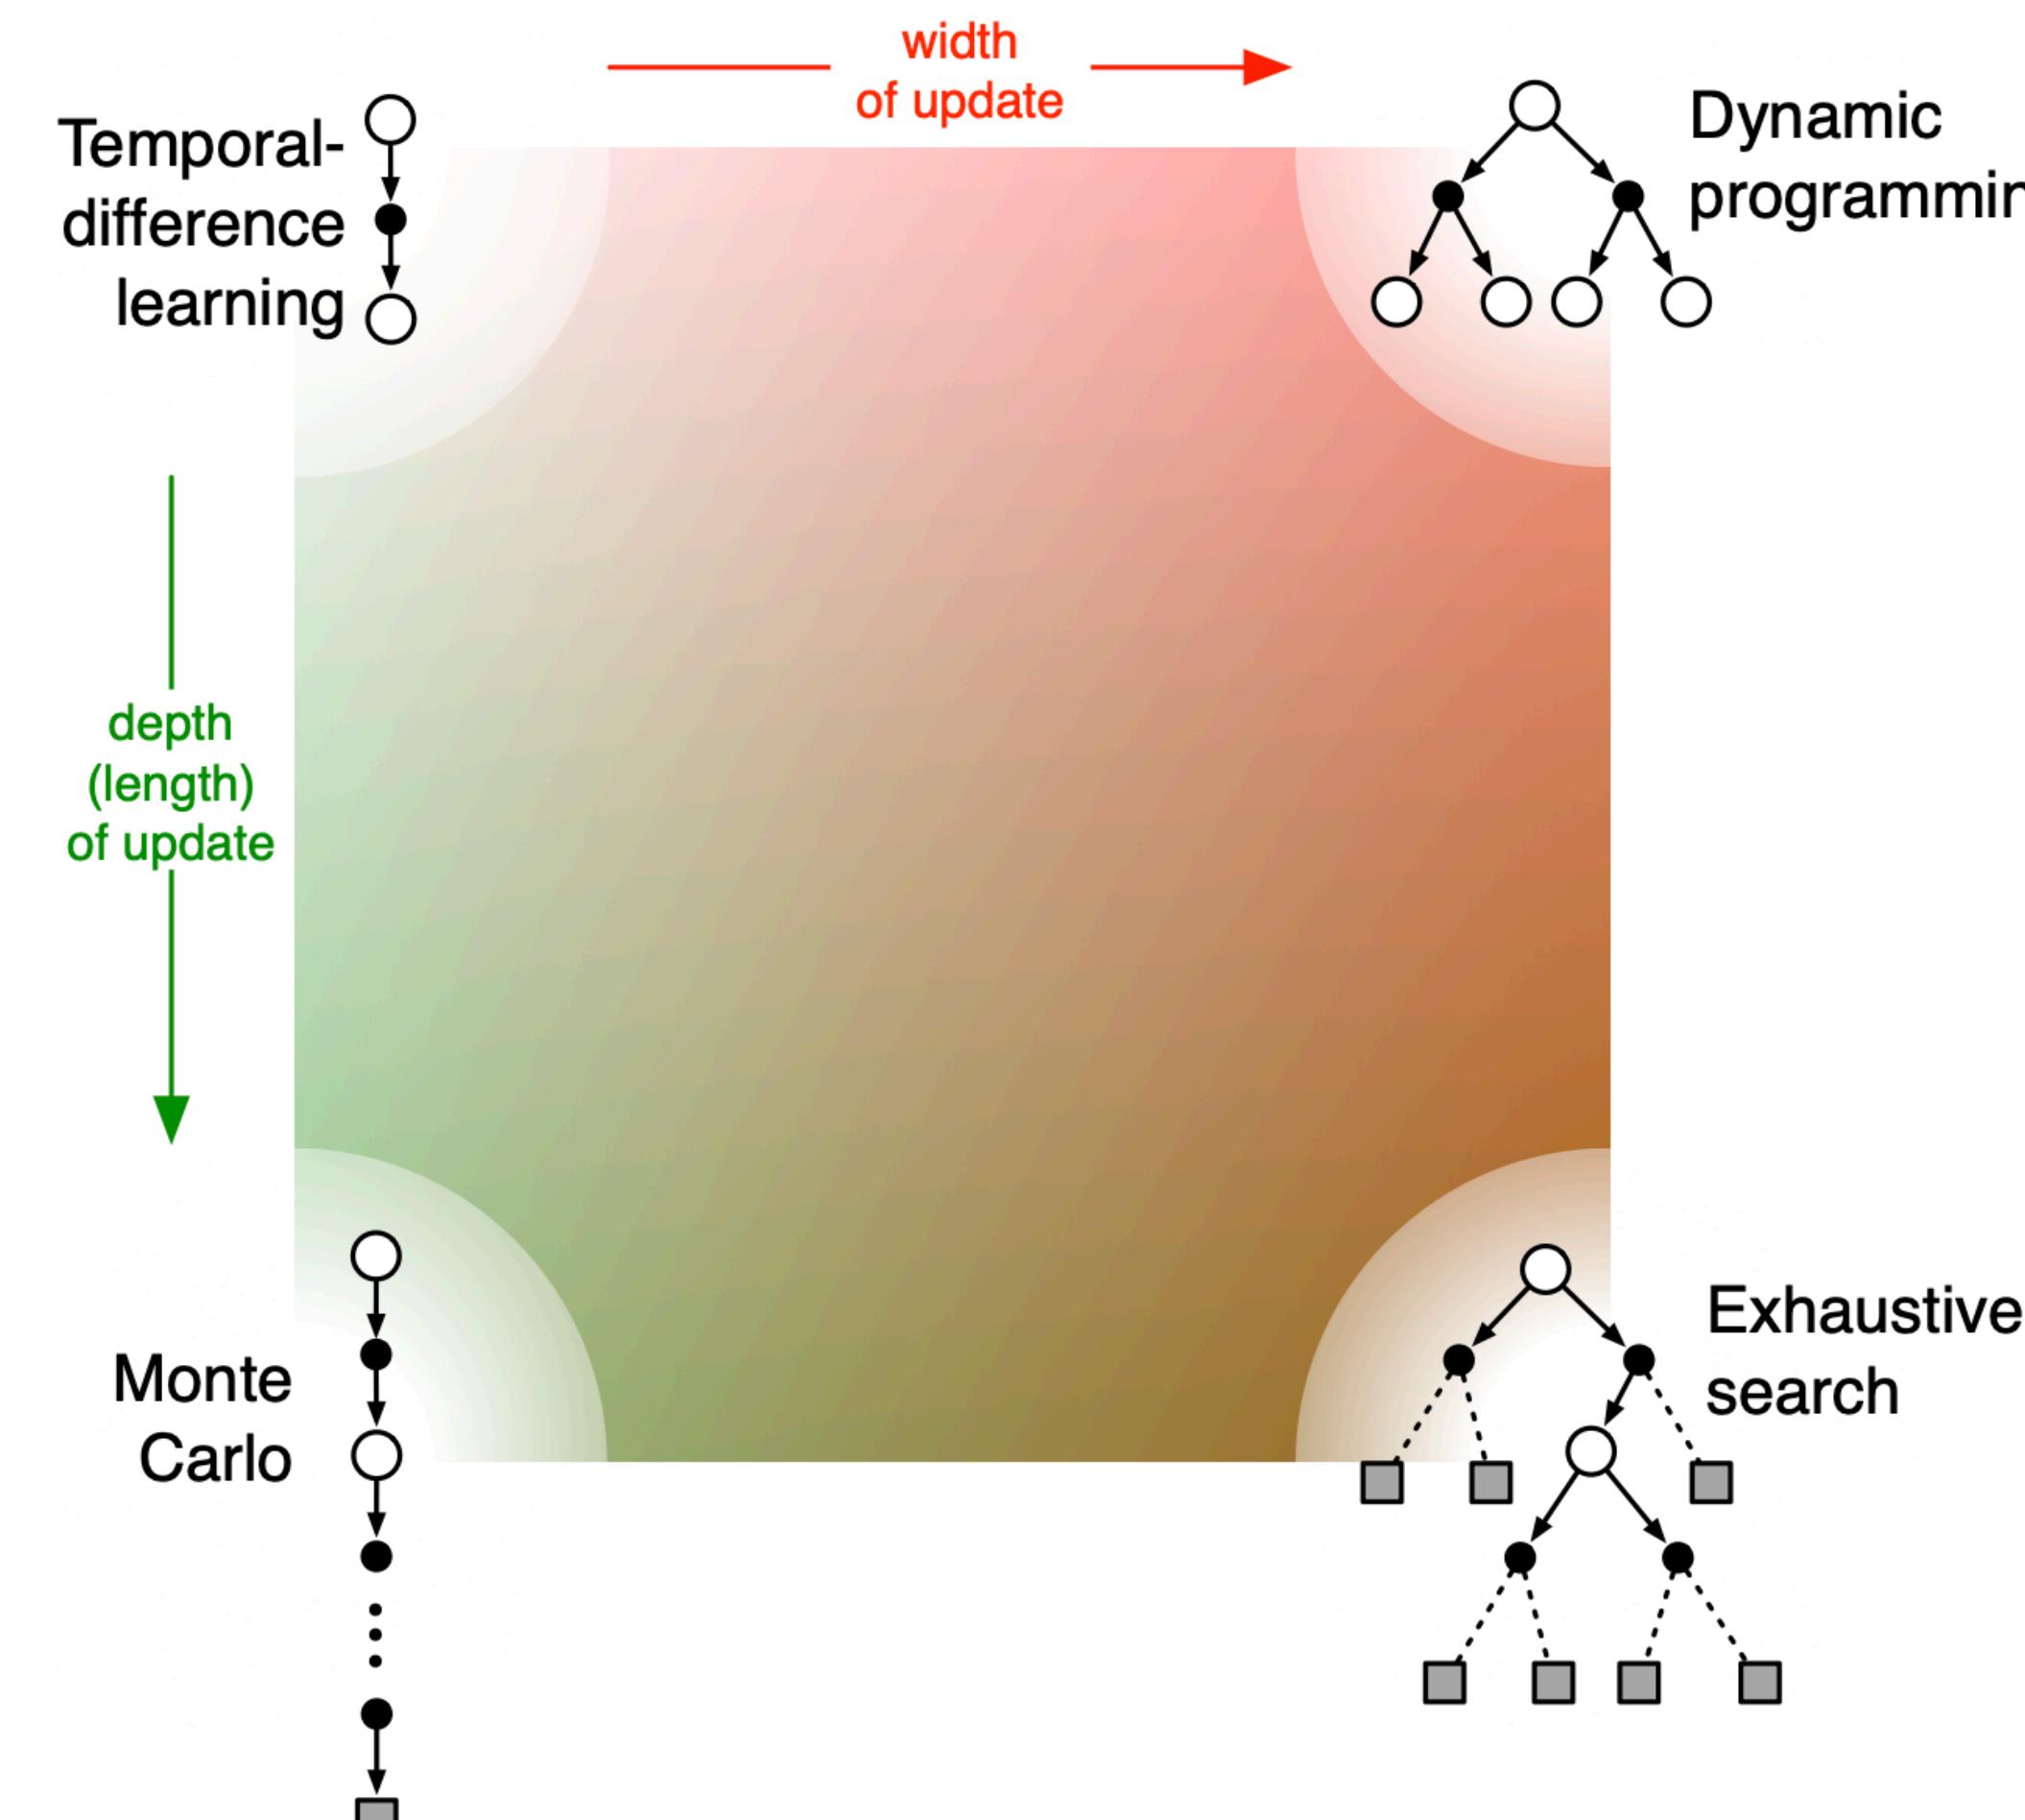
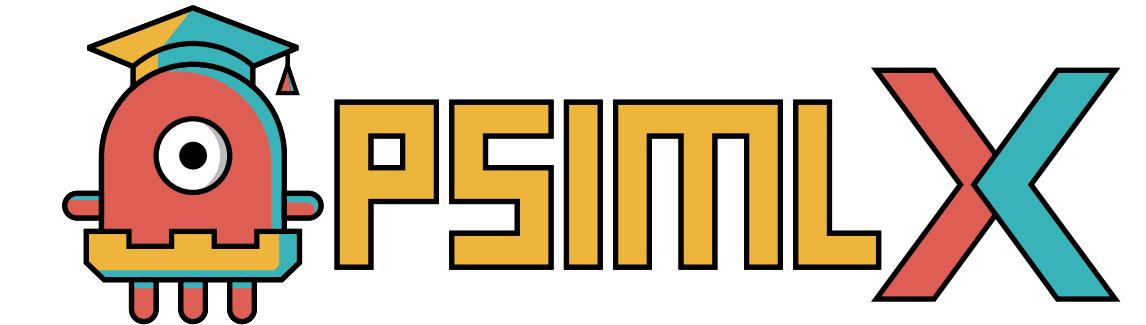


Temporal Difference (TD) Update

- Is an estimate both because the expectation is unknown, and true $v_\pi(S_{t+1})$ is approximated by current estimate $V(S_{t+1})$

Model free RL

Monte Carlo (MC) and Temporal Difference (TD)

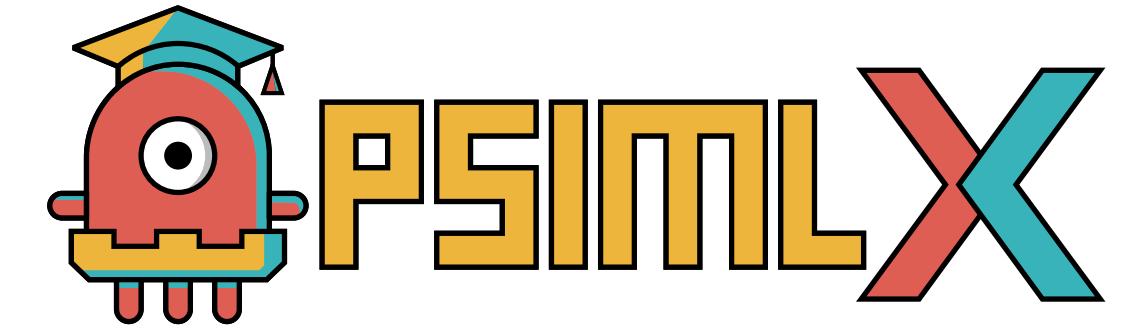


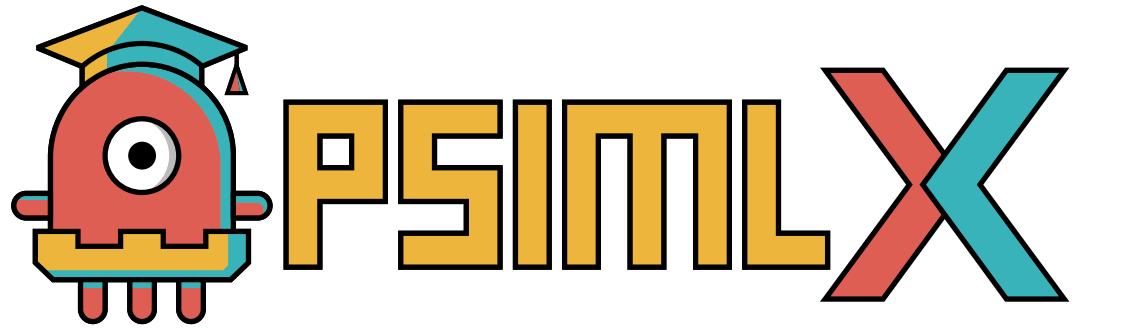
Comparison of RL methods [Sutton & Barto 2018]

Model free RL

MC and TD: Bias vs Variance Trade-Off

- MC Target [return G_t] is an *unbiased estimate* of $v_\pi(S_t)$
- TD Target [$R_{t+1} + \gamma V(S_{t+1})$] is *biased estimate* of $v_\pi(S_t)$
- On the other hand, TD target has much lower variance compared to MC target:
 - Return depends on many steps during the episode, each potentially affected by the environment stochasticity
 - TD target depends only on a single step
- MC methods do not bootstrap, while TD methods bootstrap because estimating $V(S_t)$ uses another estimate $V(S_{t+1})$



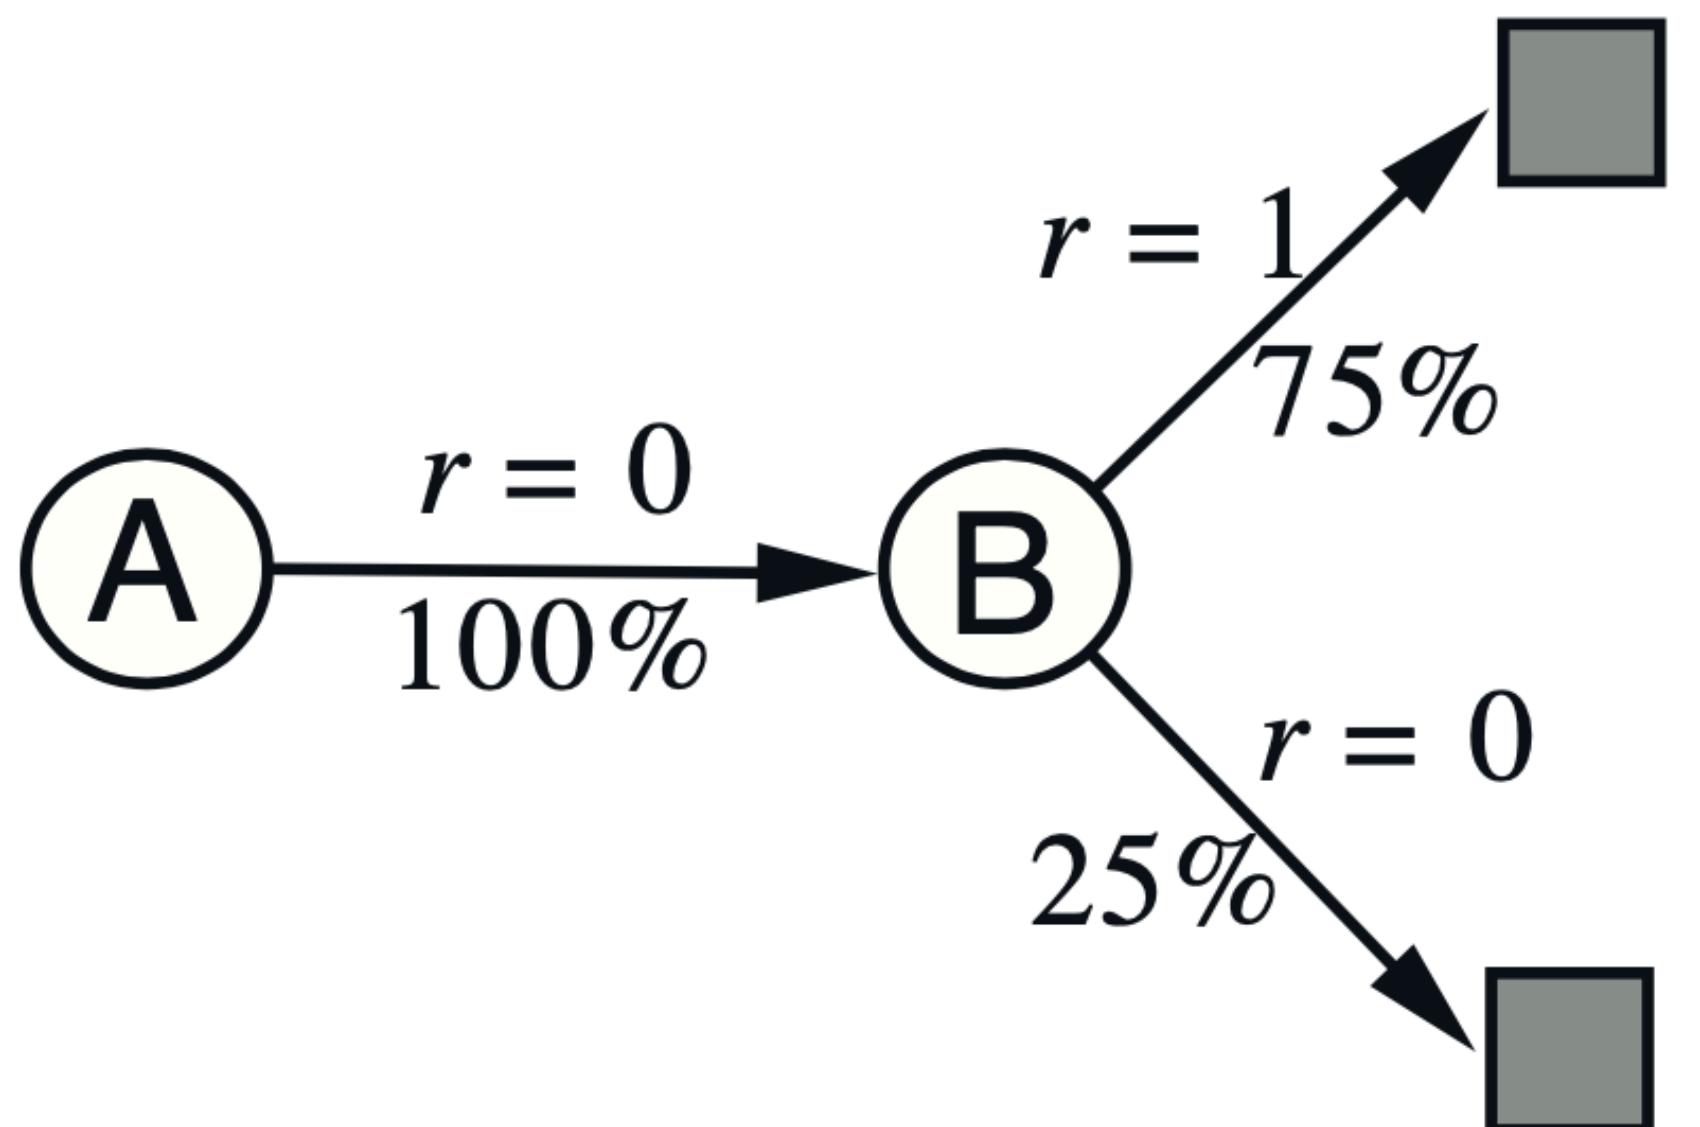


Model free RL

MC and TD: Future vs Previous Data

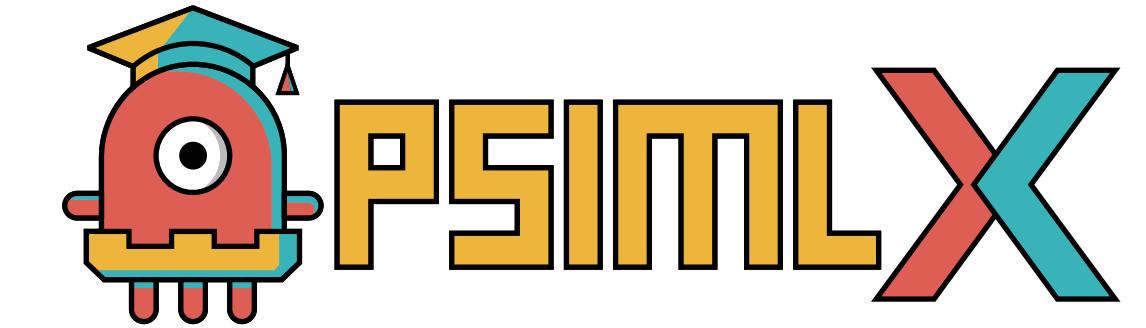
Example

- A, 0, B, 0
- B, 1
- B, 0
- $V(A) = ?; V(B) = ?$



Model free RL

MC and TD: Bias vs Variance Trade-Off



Example

- A, 0, B, 0
- B, 1
- B, 0
- $V(A) = ?; V(B) = ?$

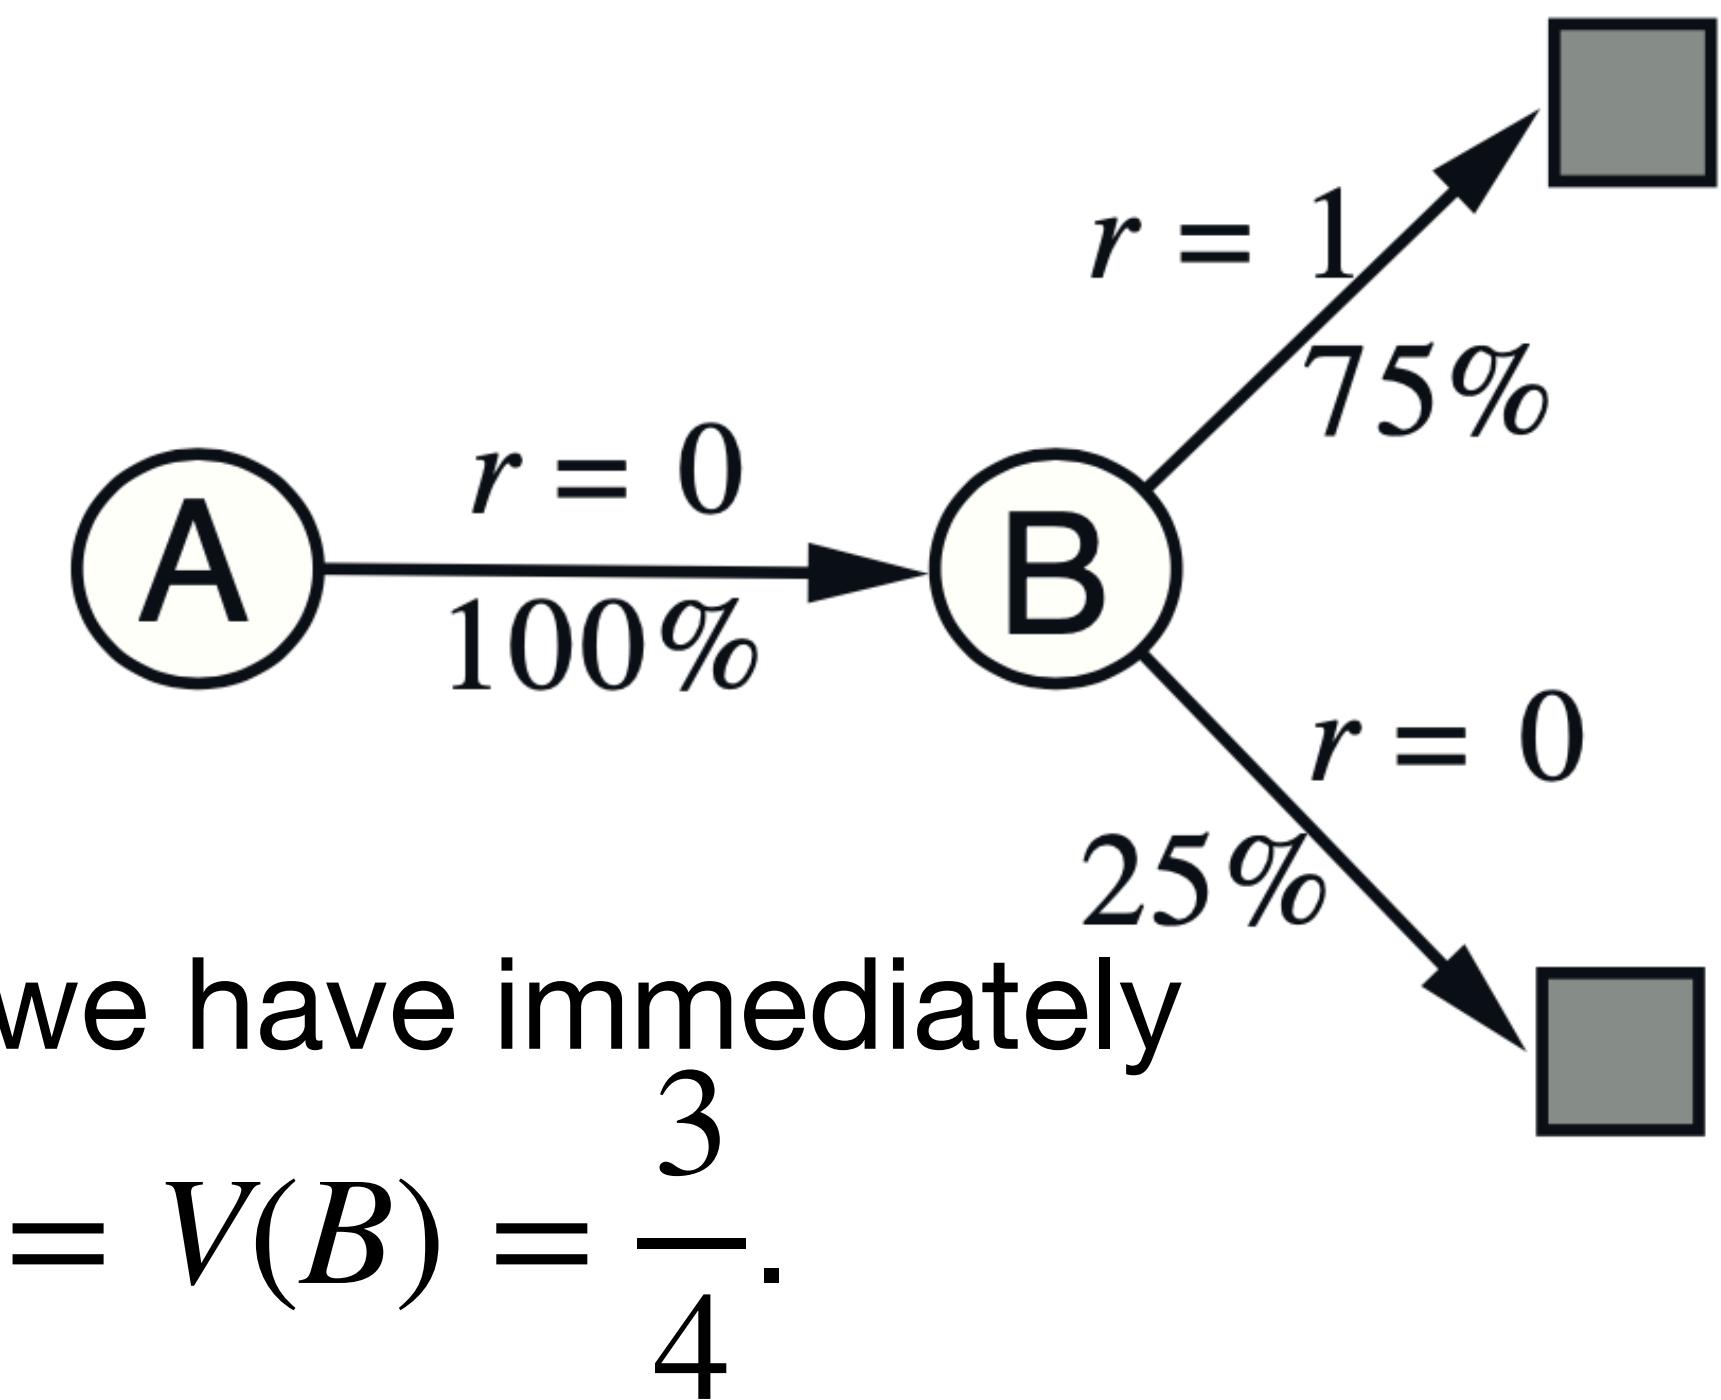
Solution

V(B): In 6/8 examples G_t for B was 1, in 2/8 examples it was 0. Both TD and MC would

agree $V(B) = \frac{3}{4}$.

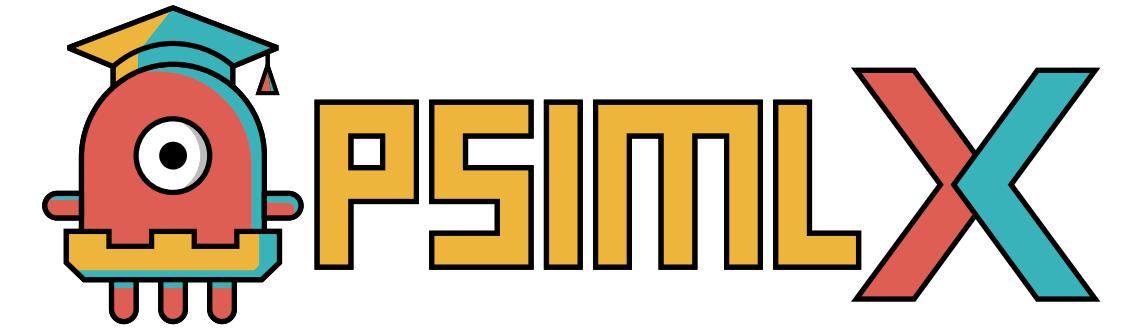
V(A): Two lines of reasoning:

- TD: Every time A was seen, we have immediately transitioned into B , so $V(A) = V(B) = \frac{3}{4}$.
- MC: A was seen only once, and the return was 0, therefore $V(A) = 0 \setminus (\vee) \setminus \lceil$



Model free RL

MC and TD: Summary

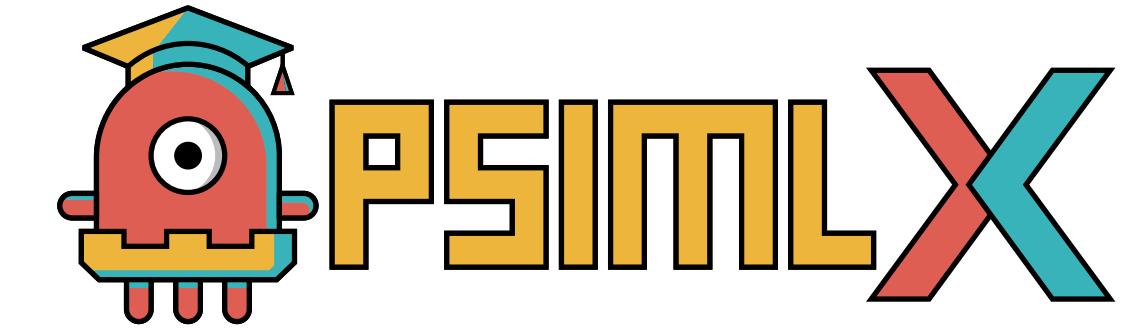


- MC methods minimise the error on the training set
- TD methods try to estimate the underlying MDP and give its maximum likelihood estimate
- MC methods need to wait for episodes to finish in order to update value functions
- TD methods can update value functions at each time-step
- MC methods have low (no) bias but high variance
- TD methods have hight bias and low variance
- MC methods are estimates because the expectation $\mathbb{E}_\pi[G_t | S_t = s]$ is unknown
- TD methods are estimates both because the expectation $\mathbb{E}_\pi[R_{t+1} + G_{t+1} | S_t = s]$ is unknown, and because $V(S_{t+1})$ is used instead of $v_\pi(S_{t+1})$

Model free RL

On-Policy vs Off-Policy Learning

- All learning methods face a dilemma:
 - They seek to learn action values conditional on subsequent optimal behaviour
 - But they need to behave non-optimally in order to explore all actions
- On-policy learning:
 - Learn about policy π from experience sampled by π
 - π is neither fully greedy, nor fully exploratory — we need to keep sufficient exploration in order to converge to good behaviour
 - Conceptually very simple
- Off-policy learning:
 - Learn about *target* policy π from experience sampled by *behaviour policy* b
 - Allows for learning from past policy experiences or people
 - Alleviates exploration-exploitation tradeoff, but is more conceptually challenging



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

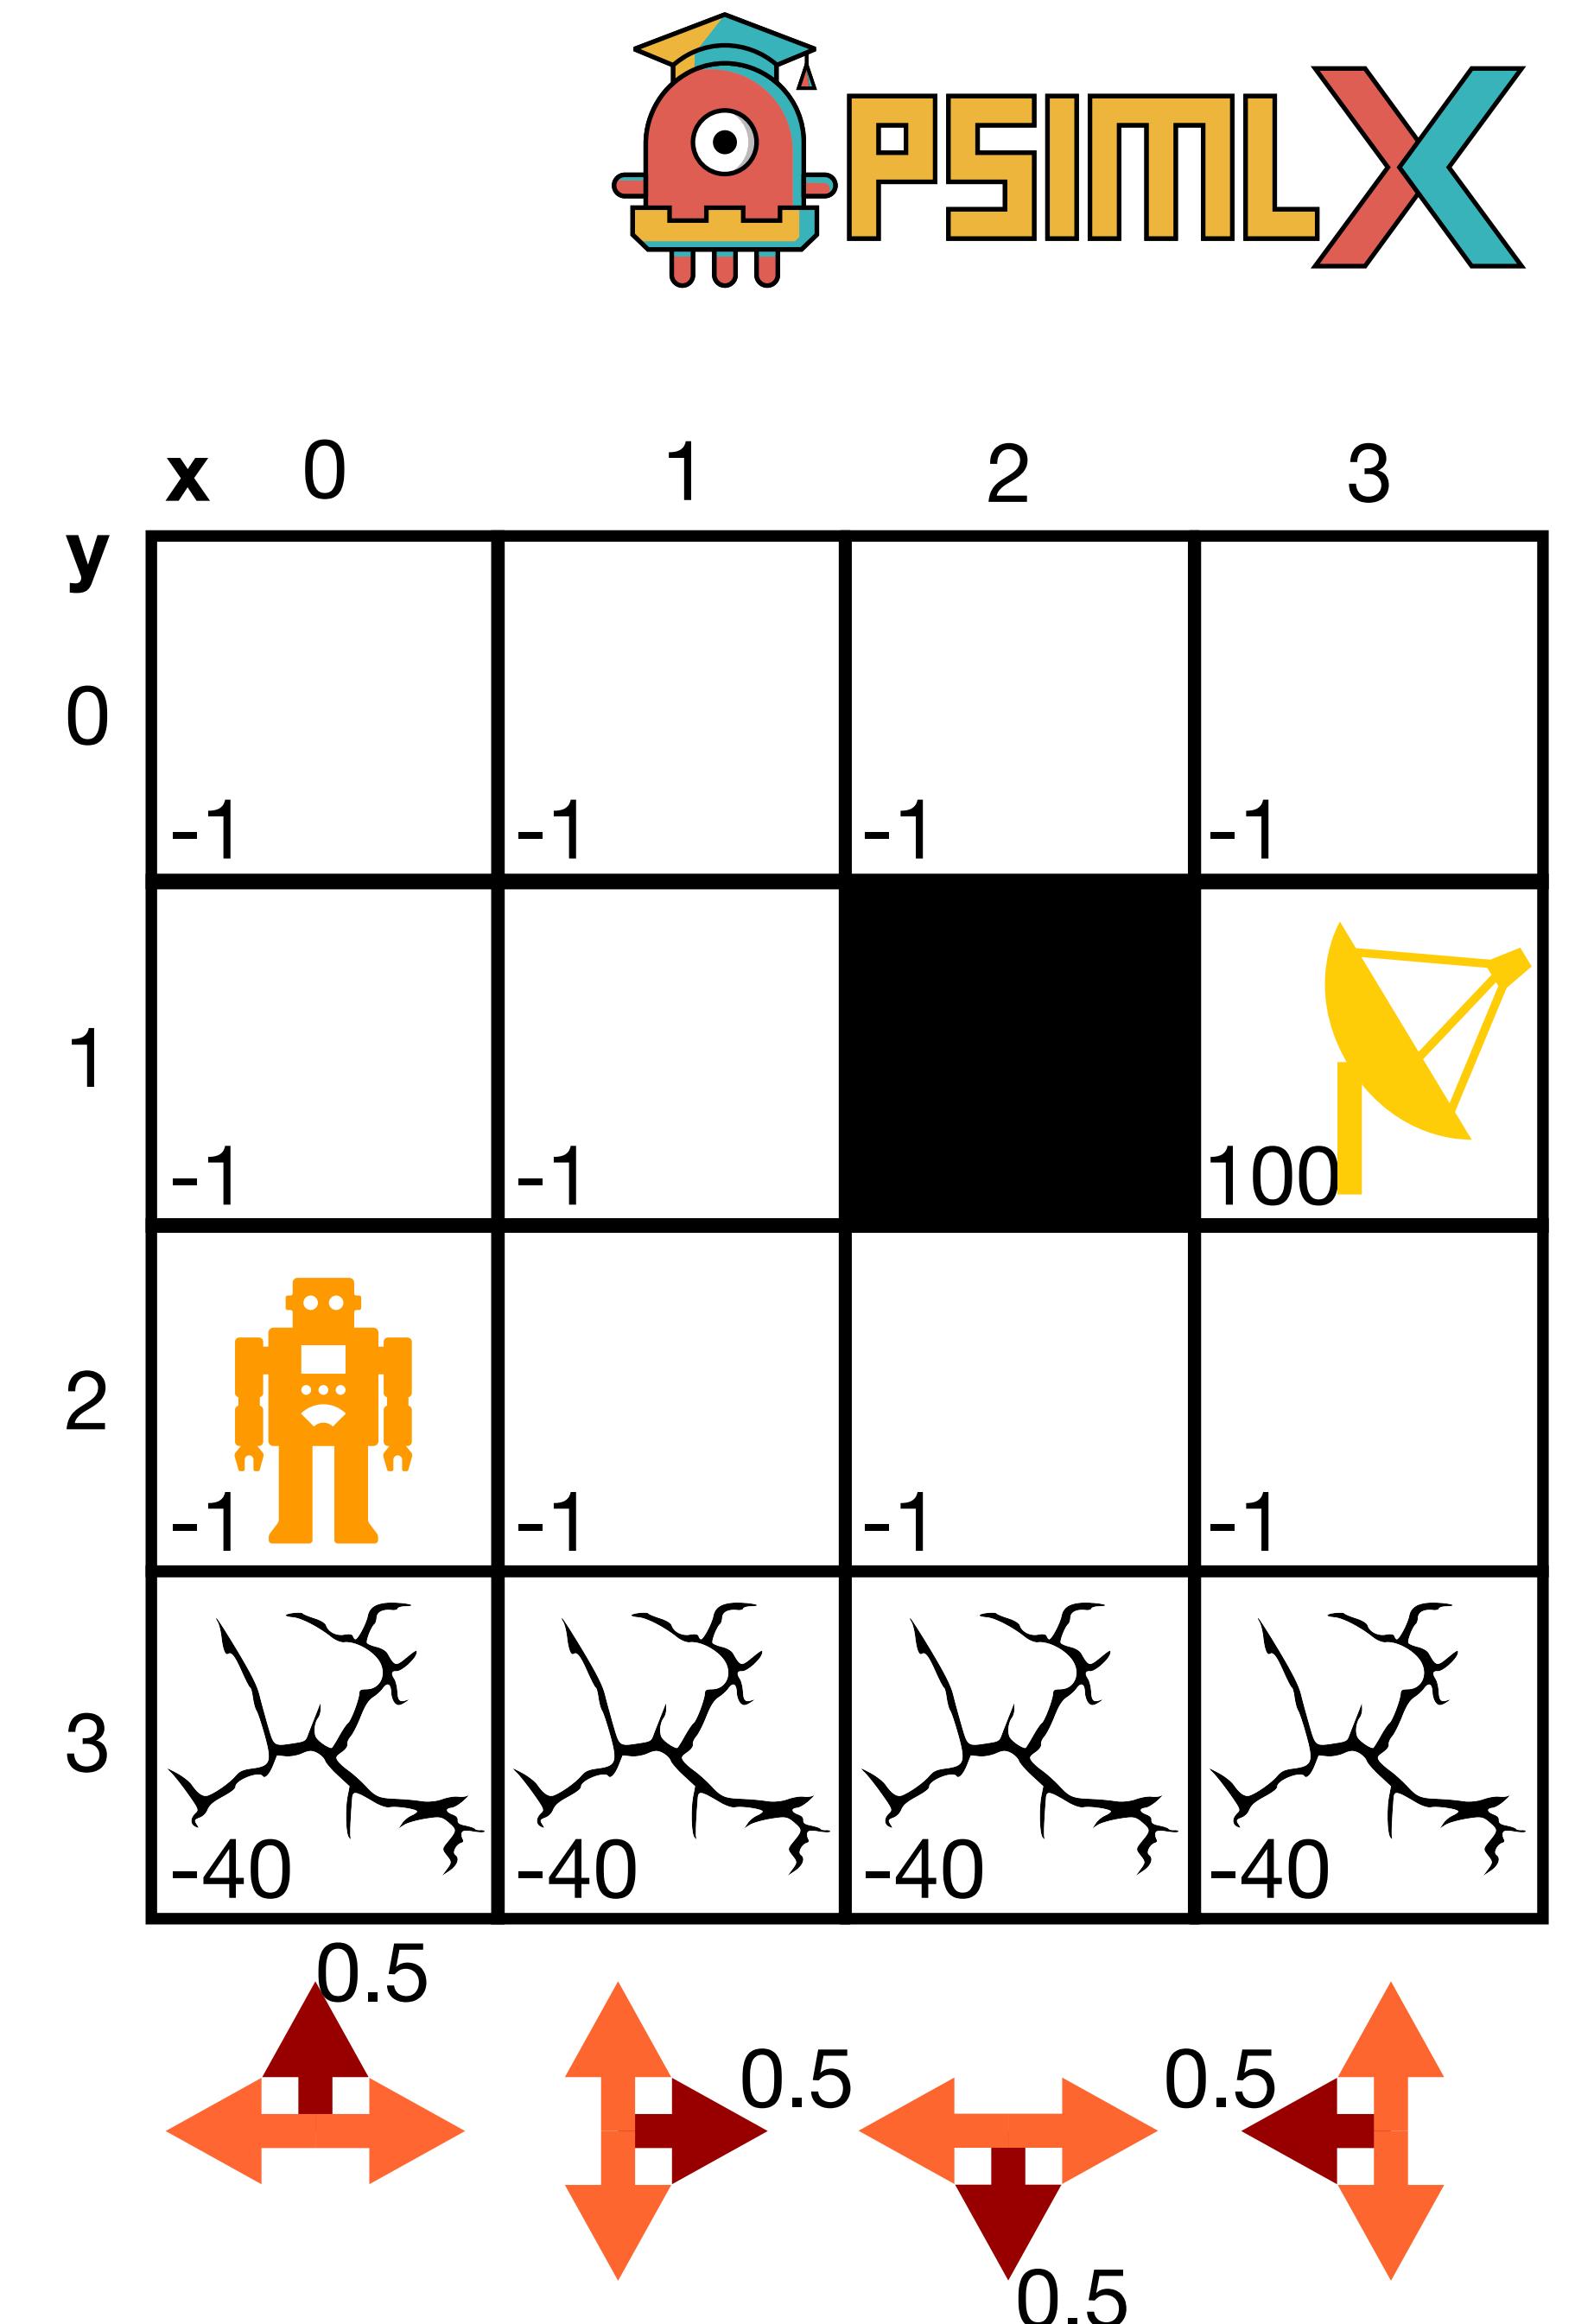
Algorithm parameter: small $\varepsilon > 0$

Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
- $G \leftarrow \gamma G + R_{t+1}$
- Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1 \dots, S_{t-1}, A_{t-1}$:
- Append G to $Returns(S_t, A_t)$
- $Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$
- $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
- For all $a \in \mathcal{A}(S_t)$:
- $$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

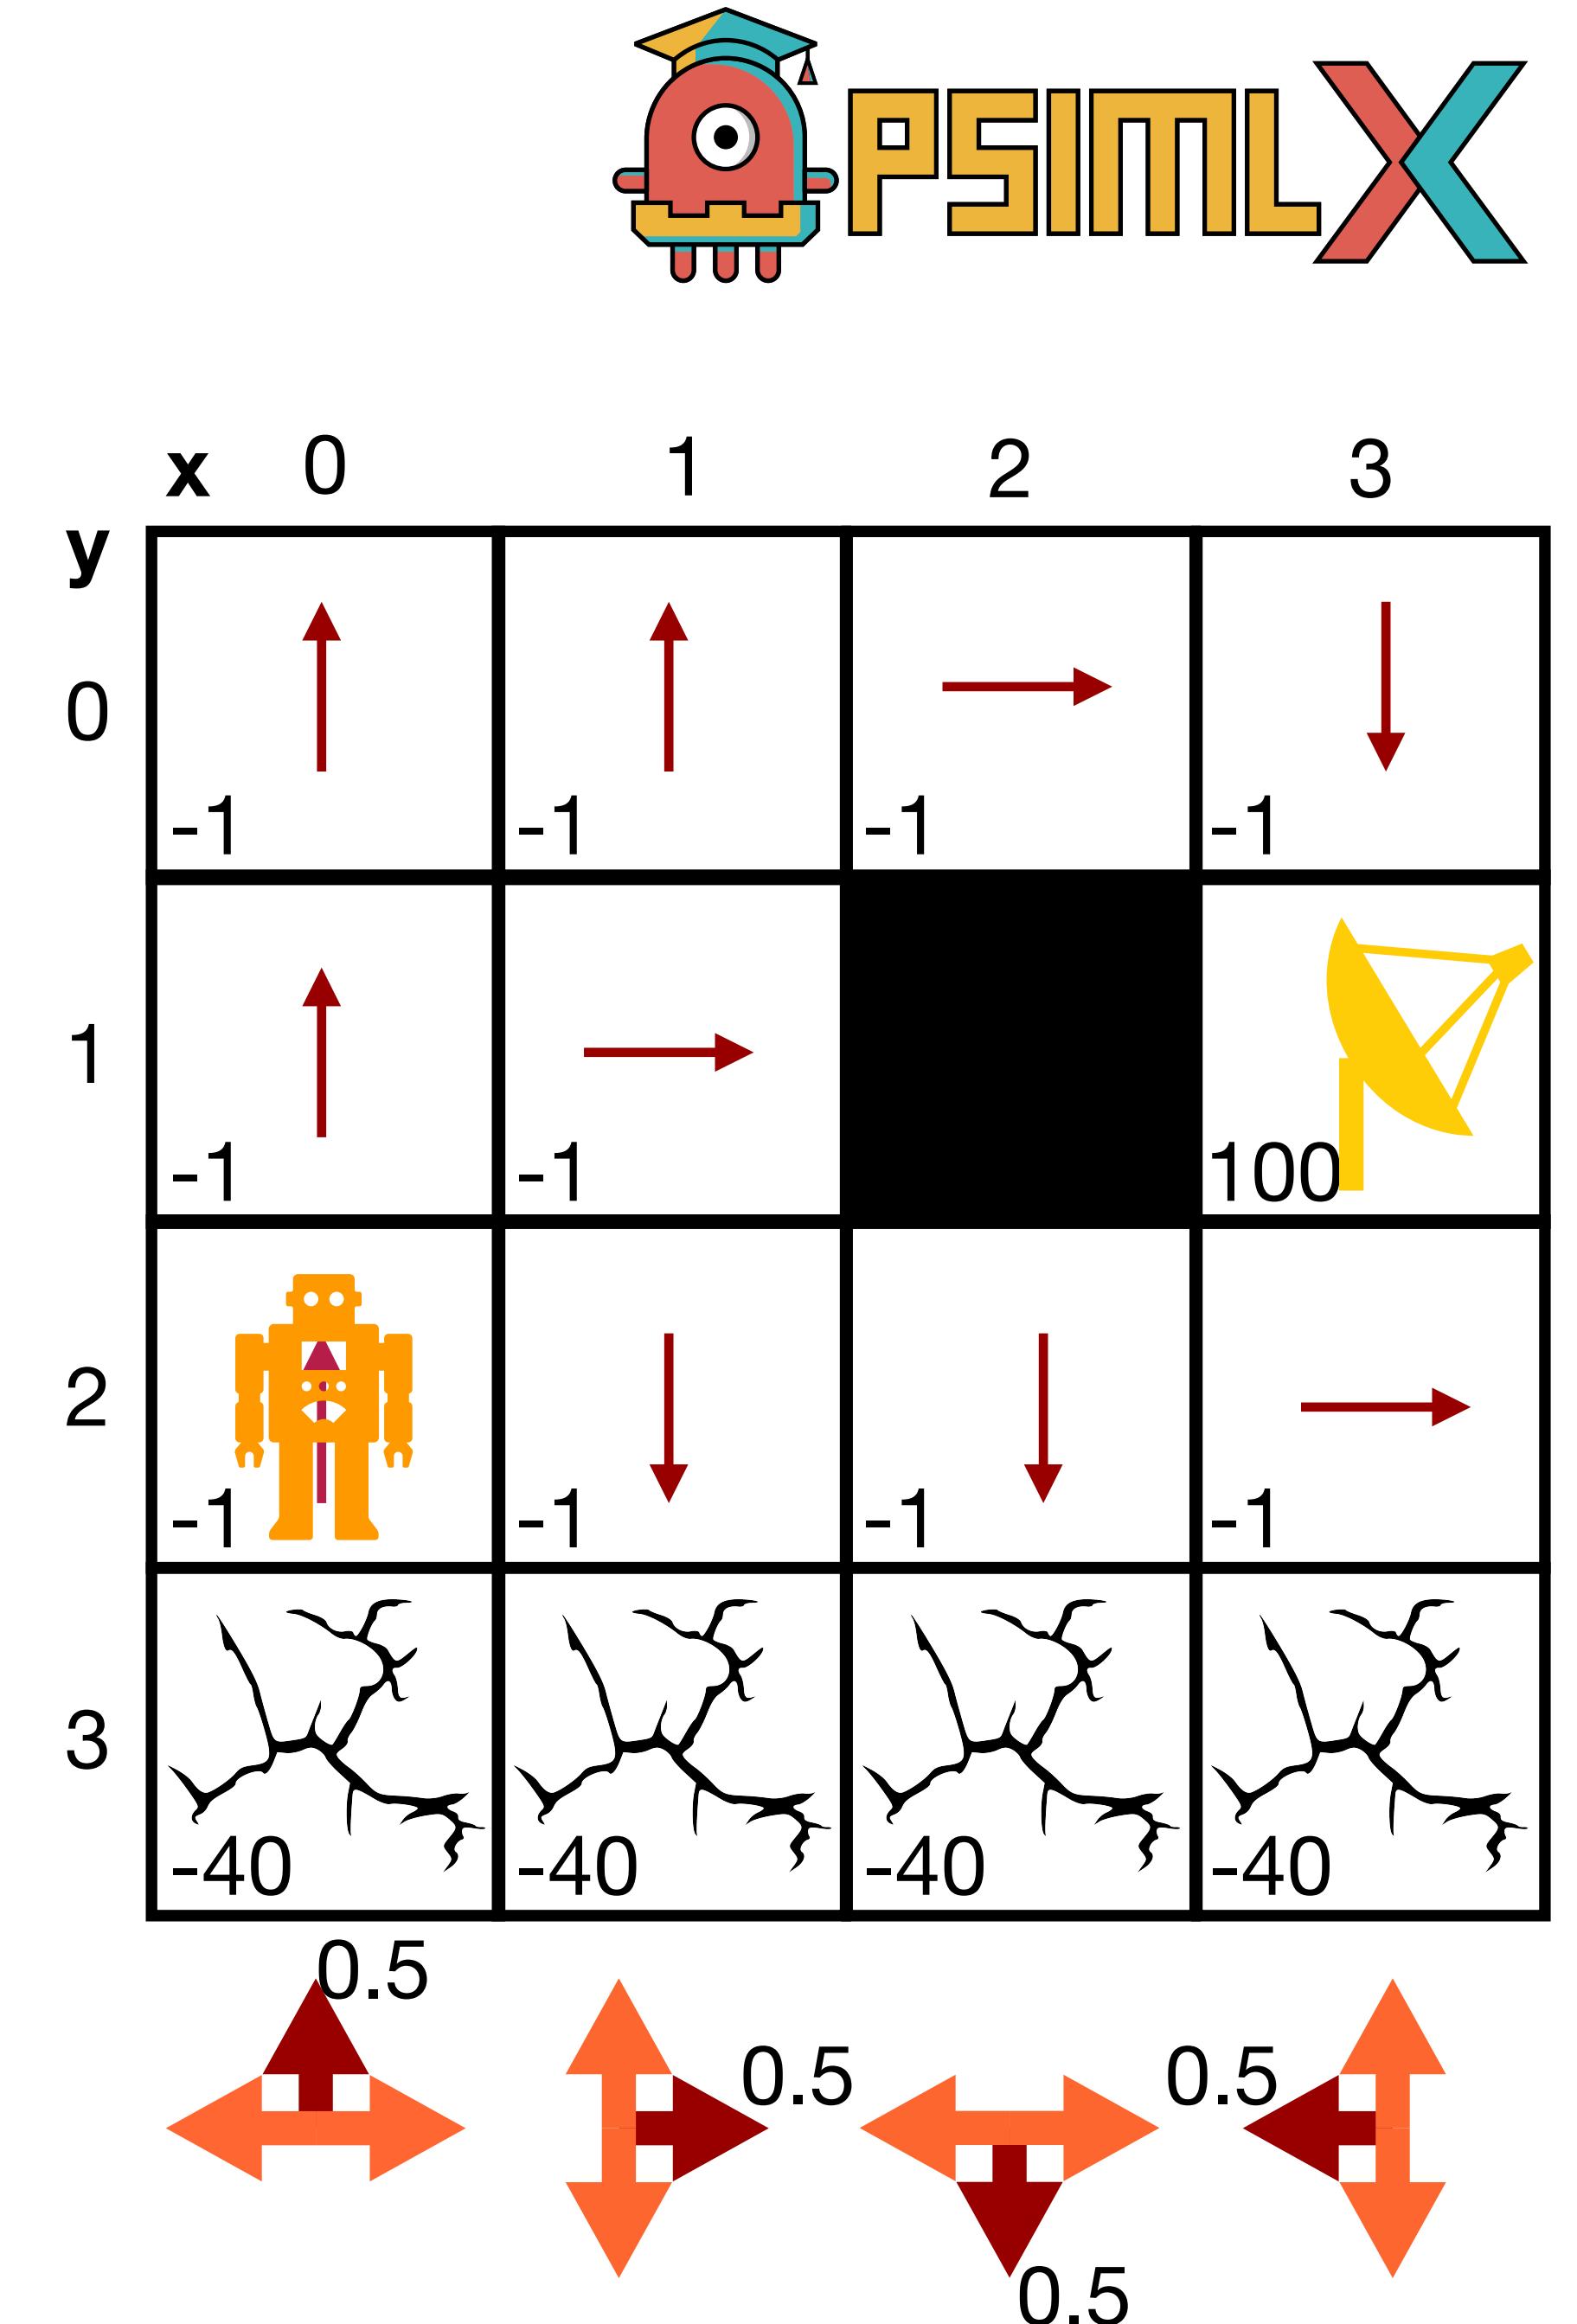
Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

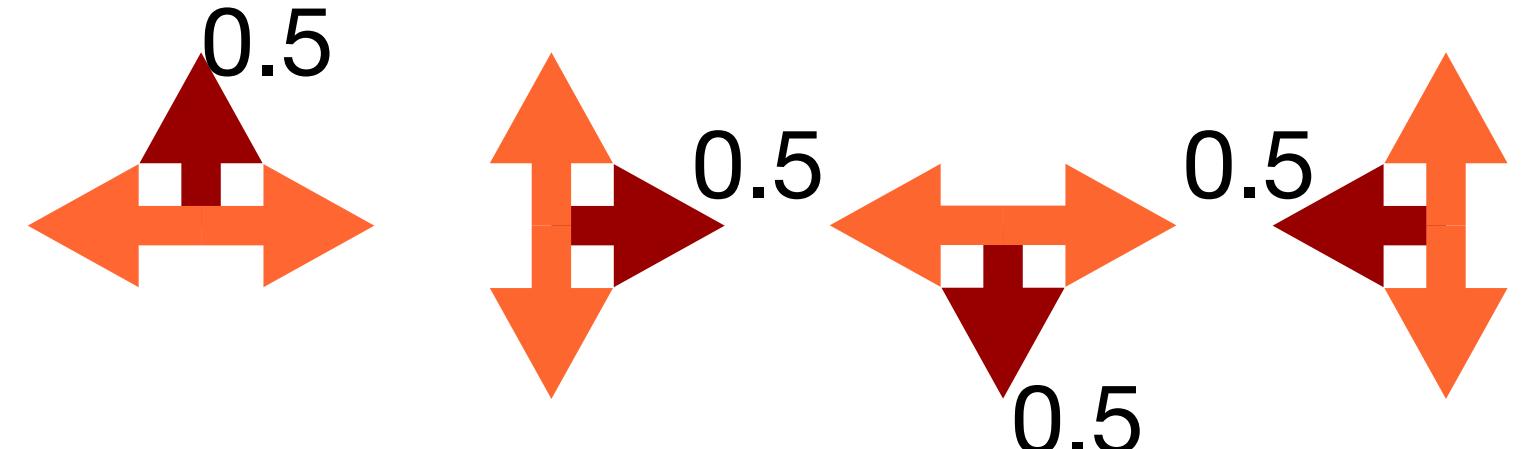
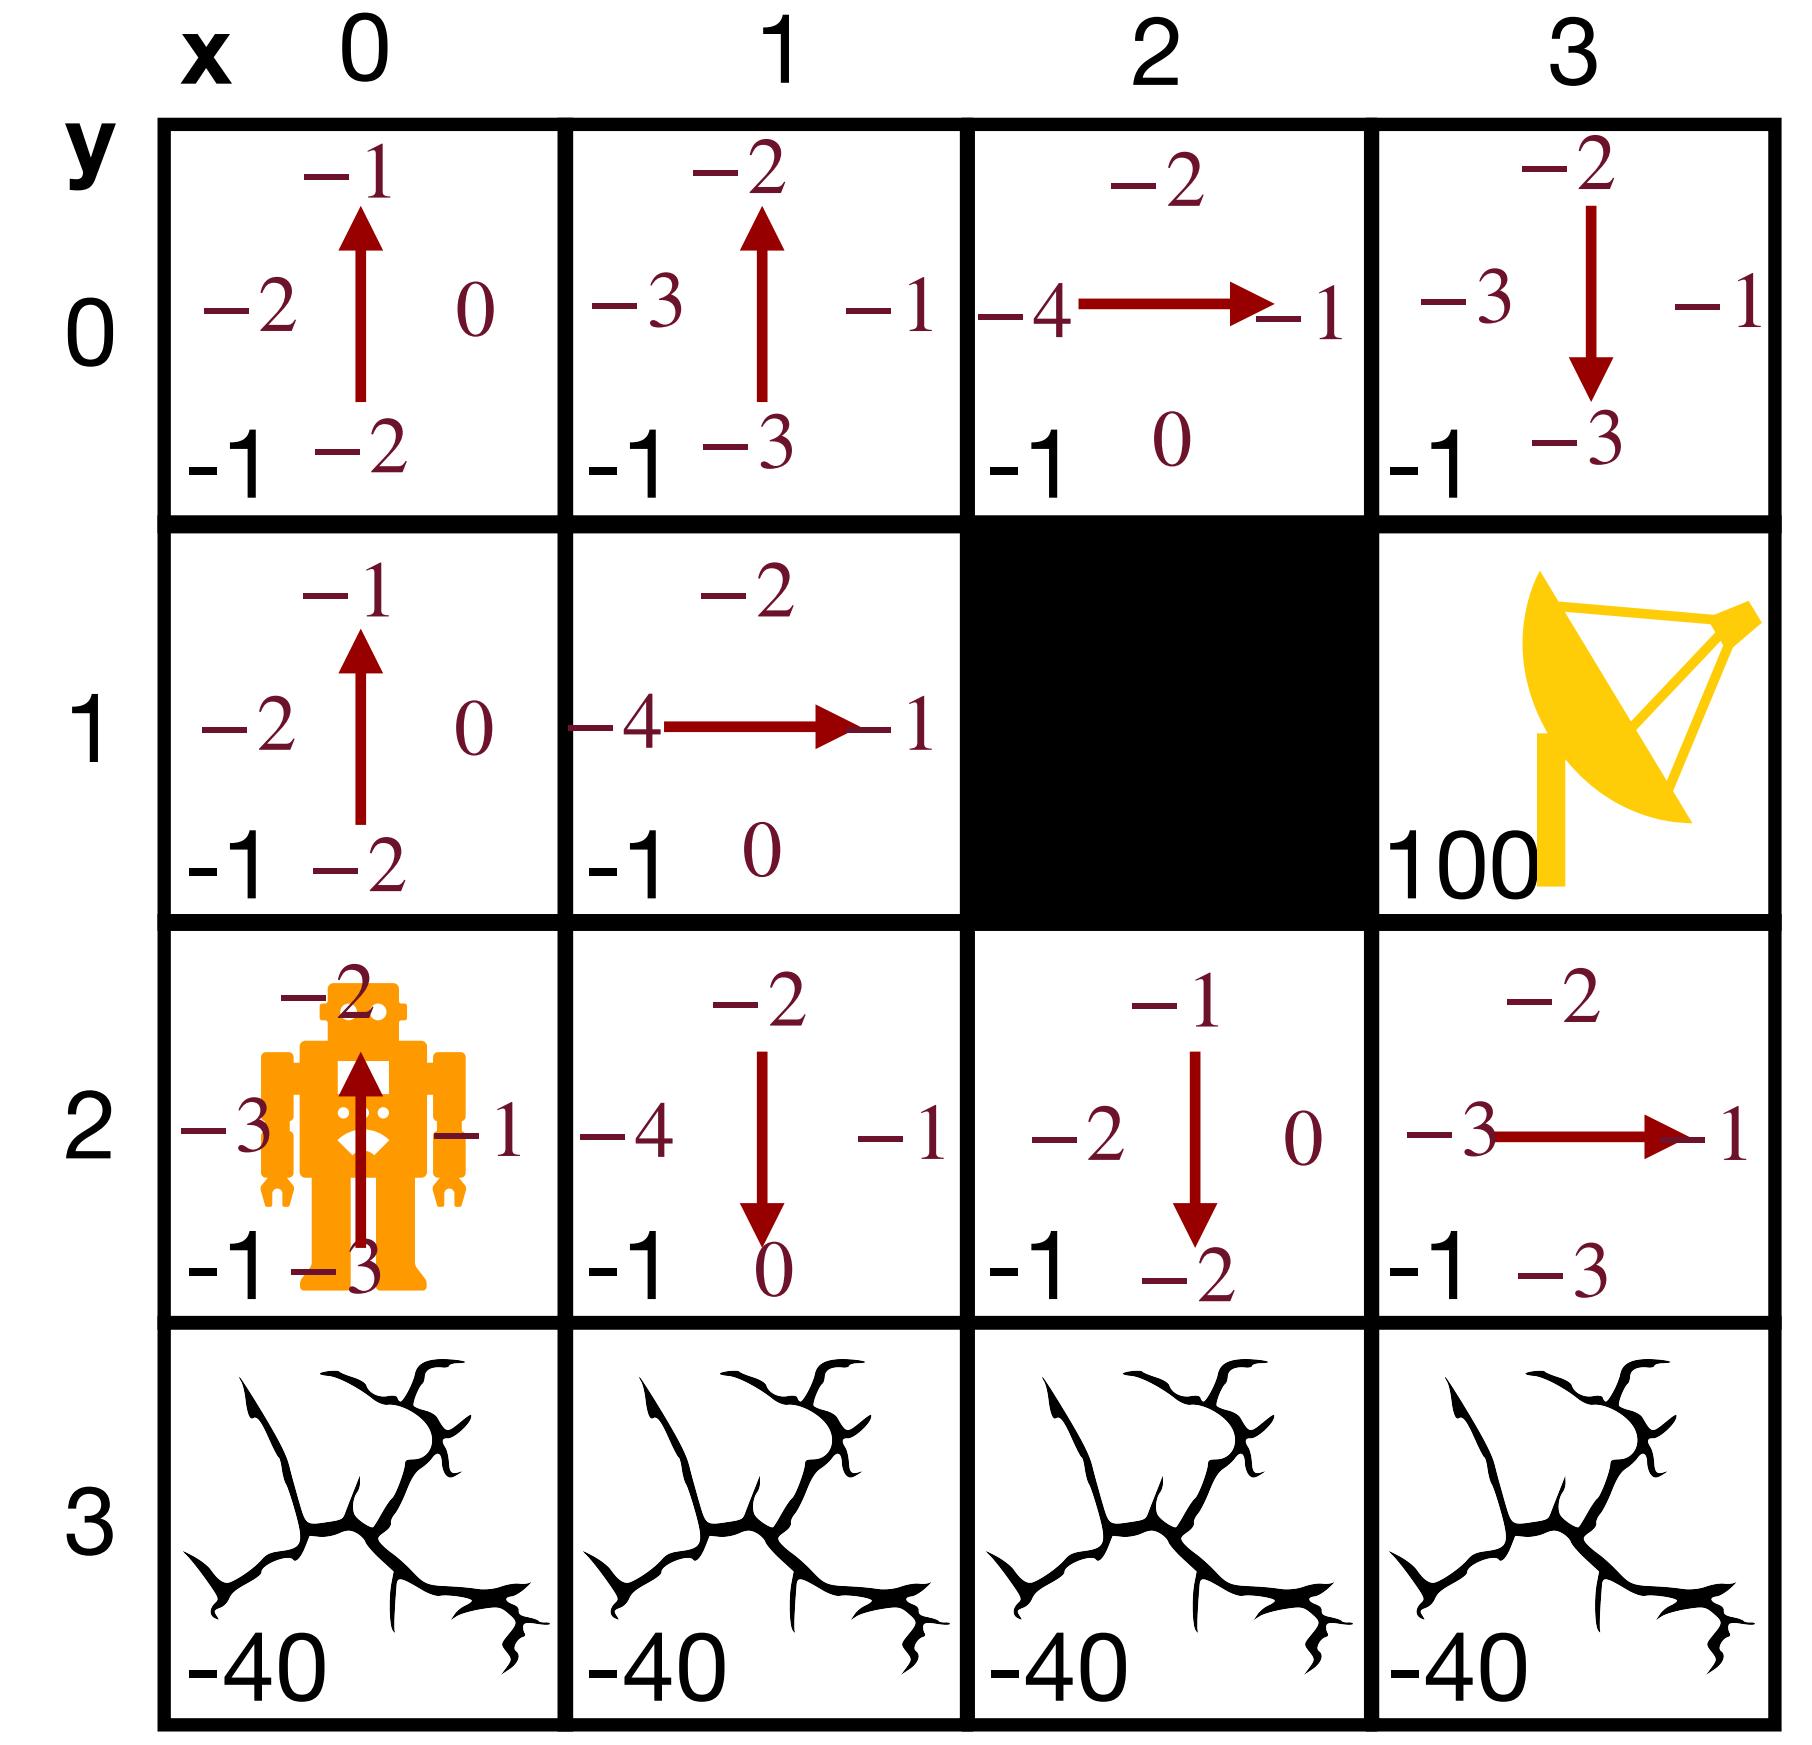
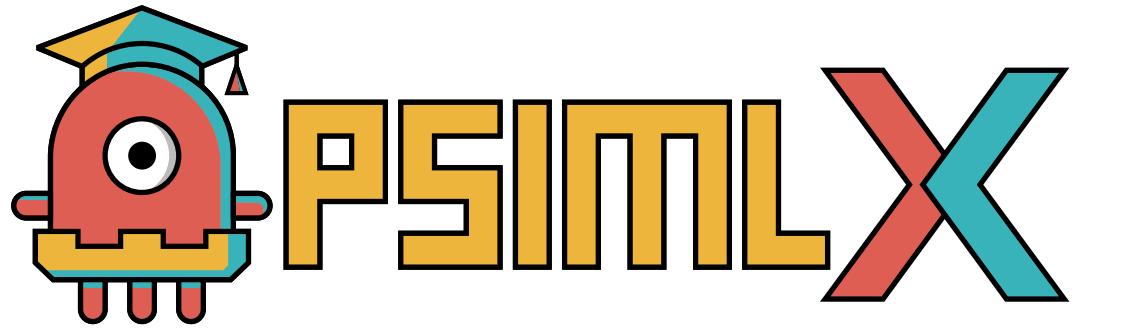
Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

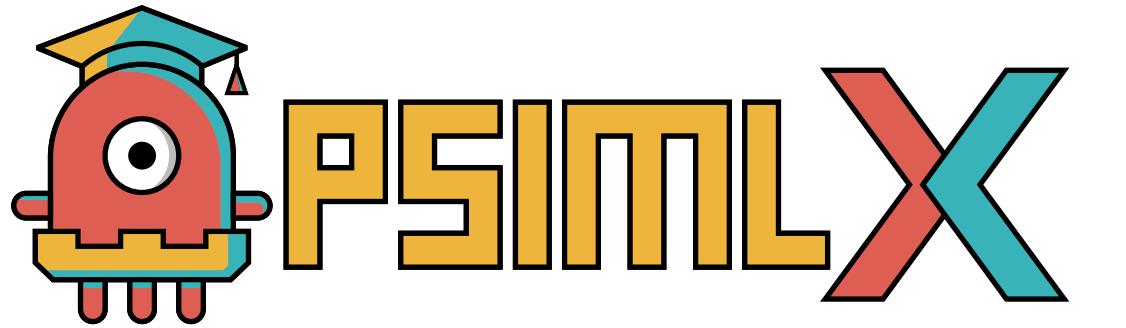
Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1 \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



x	0	1	2	3	4	5	6	7
y	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0

Returns(s, a)

Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

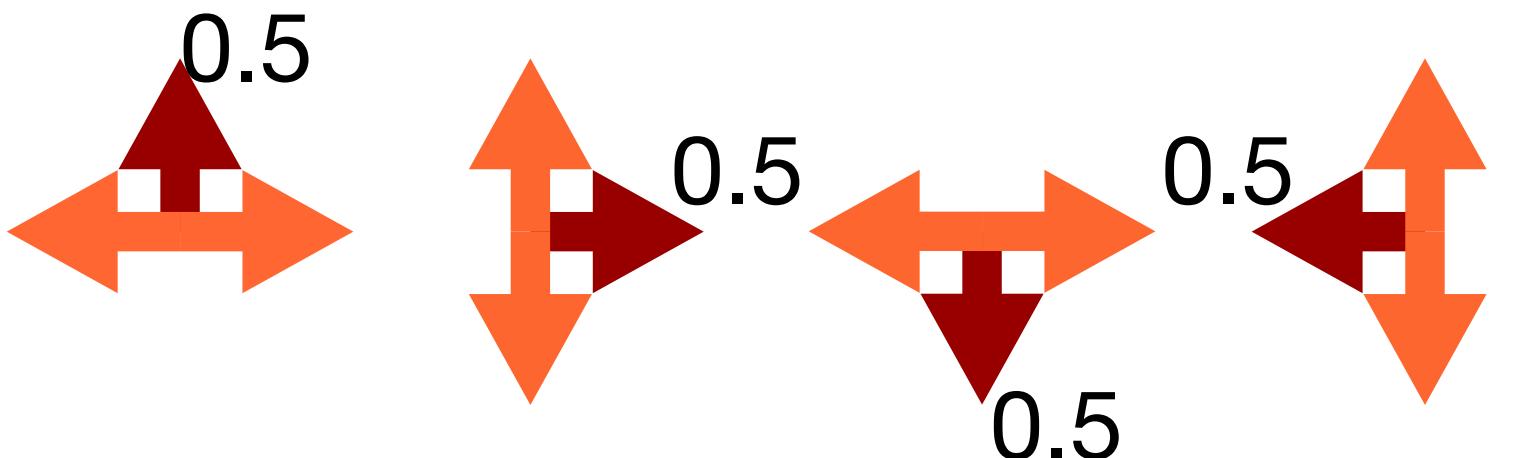
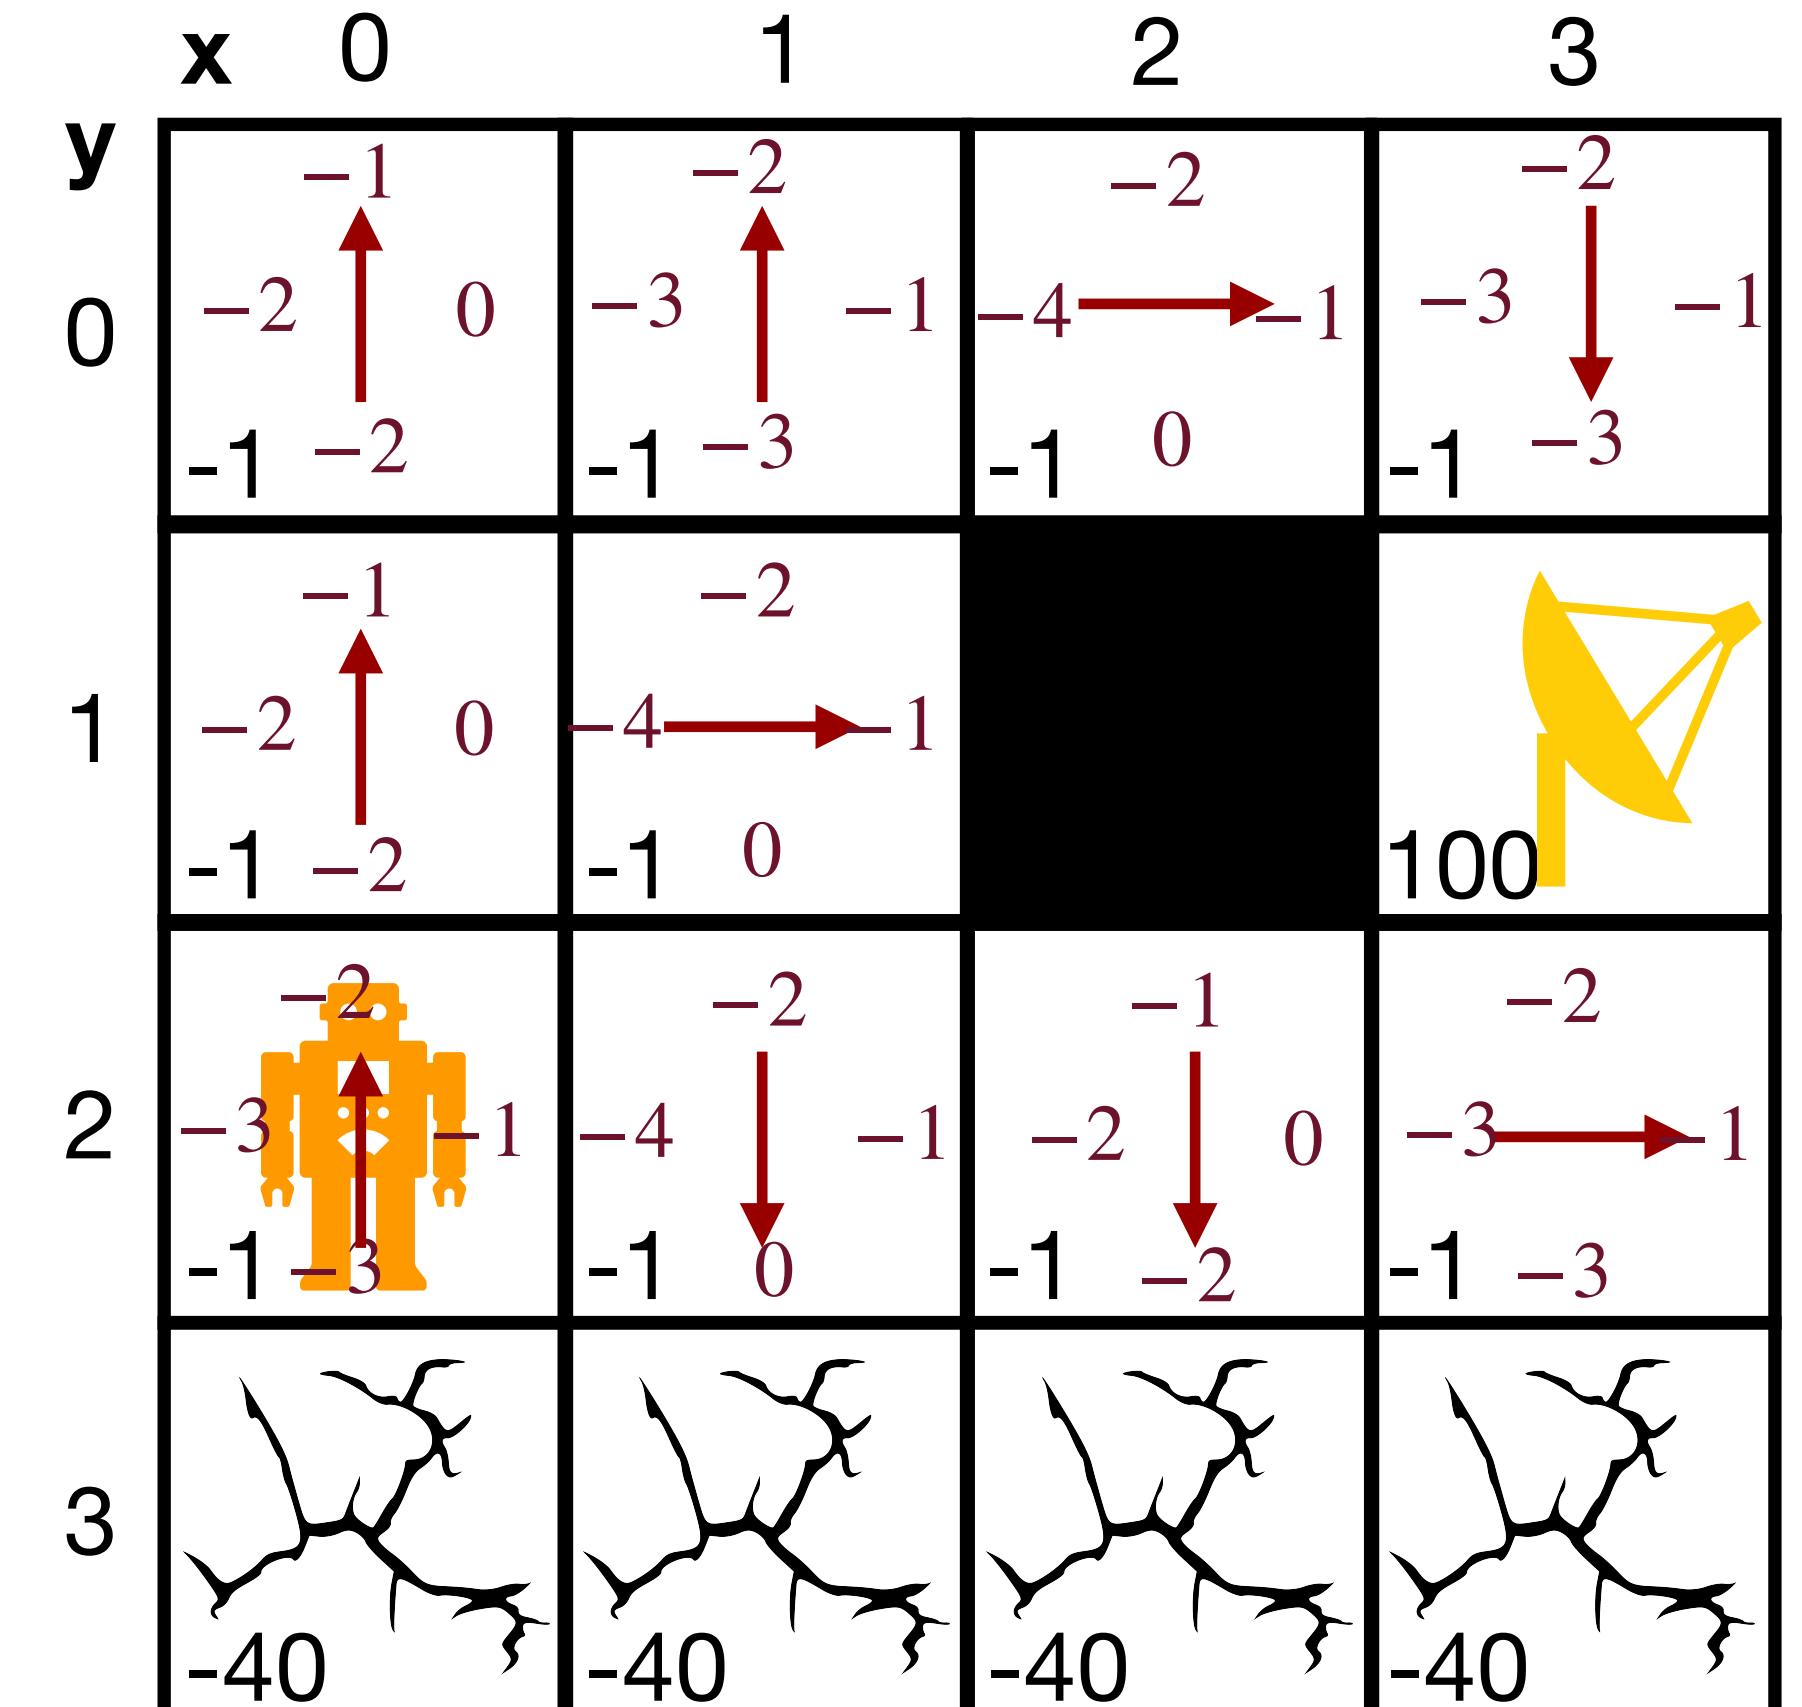
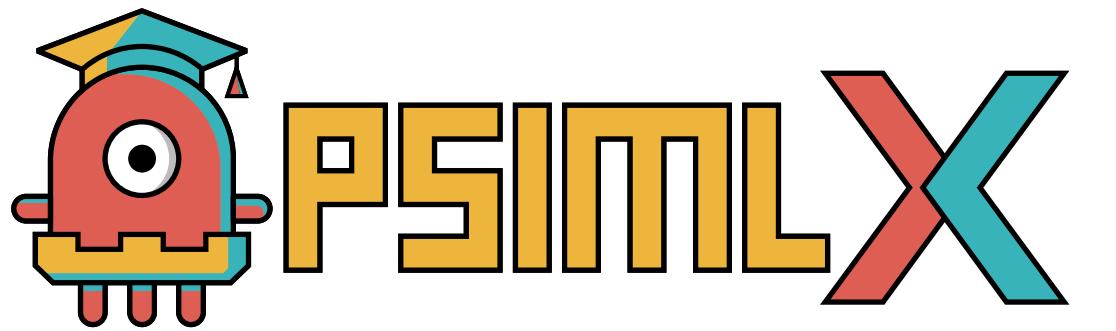
$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

- Append G to $Returns(S_t, A_t)$
- $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
- $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
- For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\tau_\pi = []$$



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

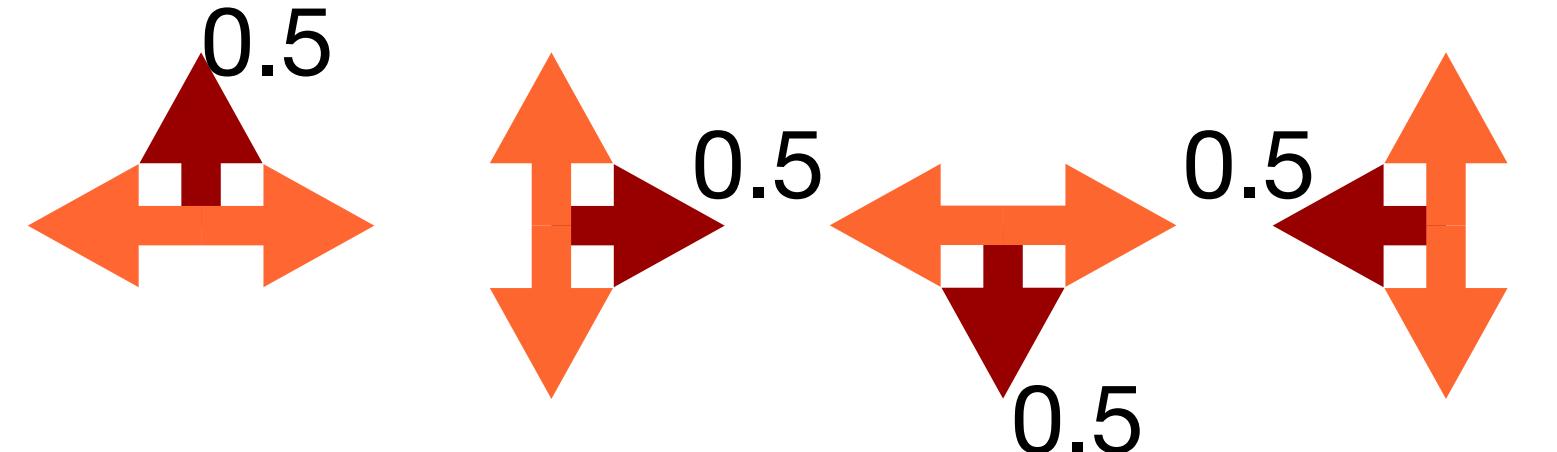
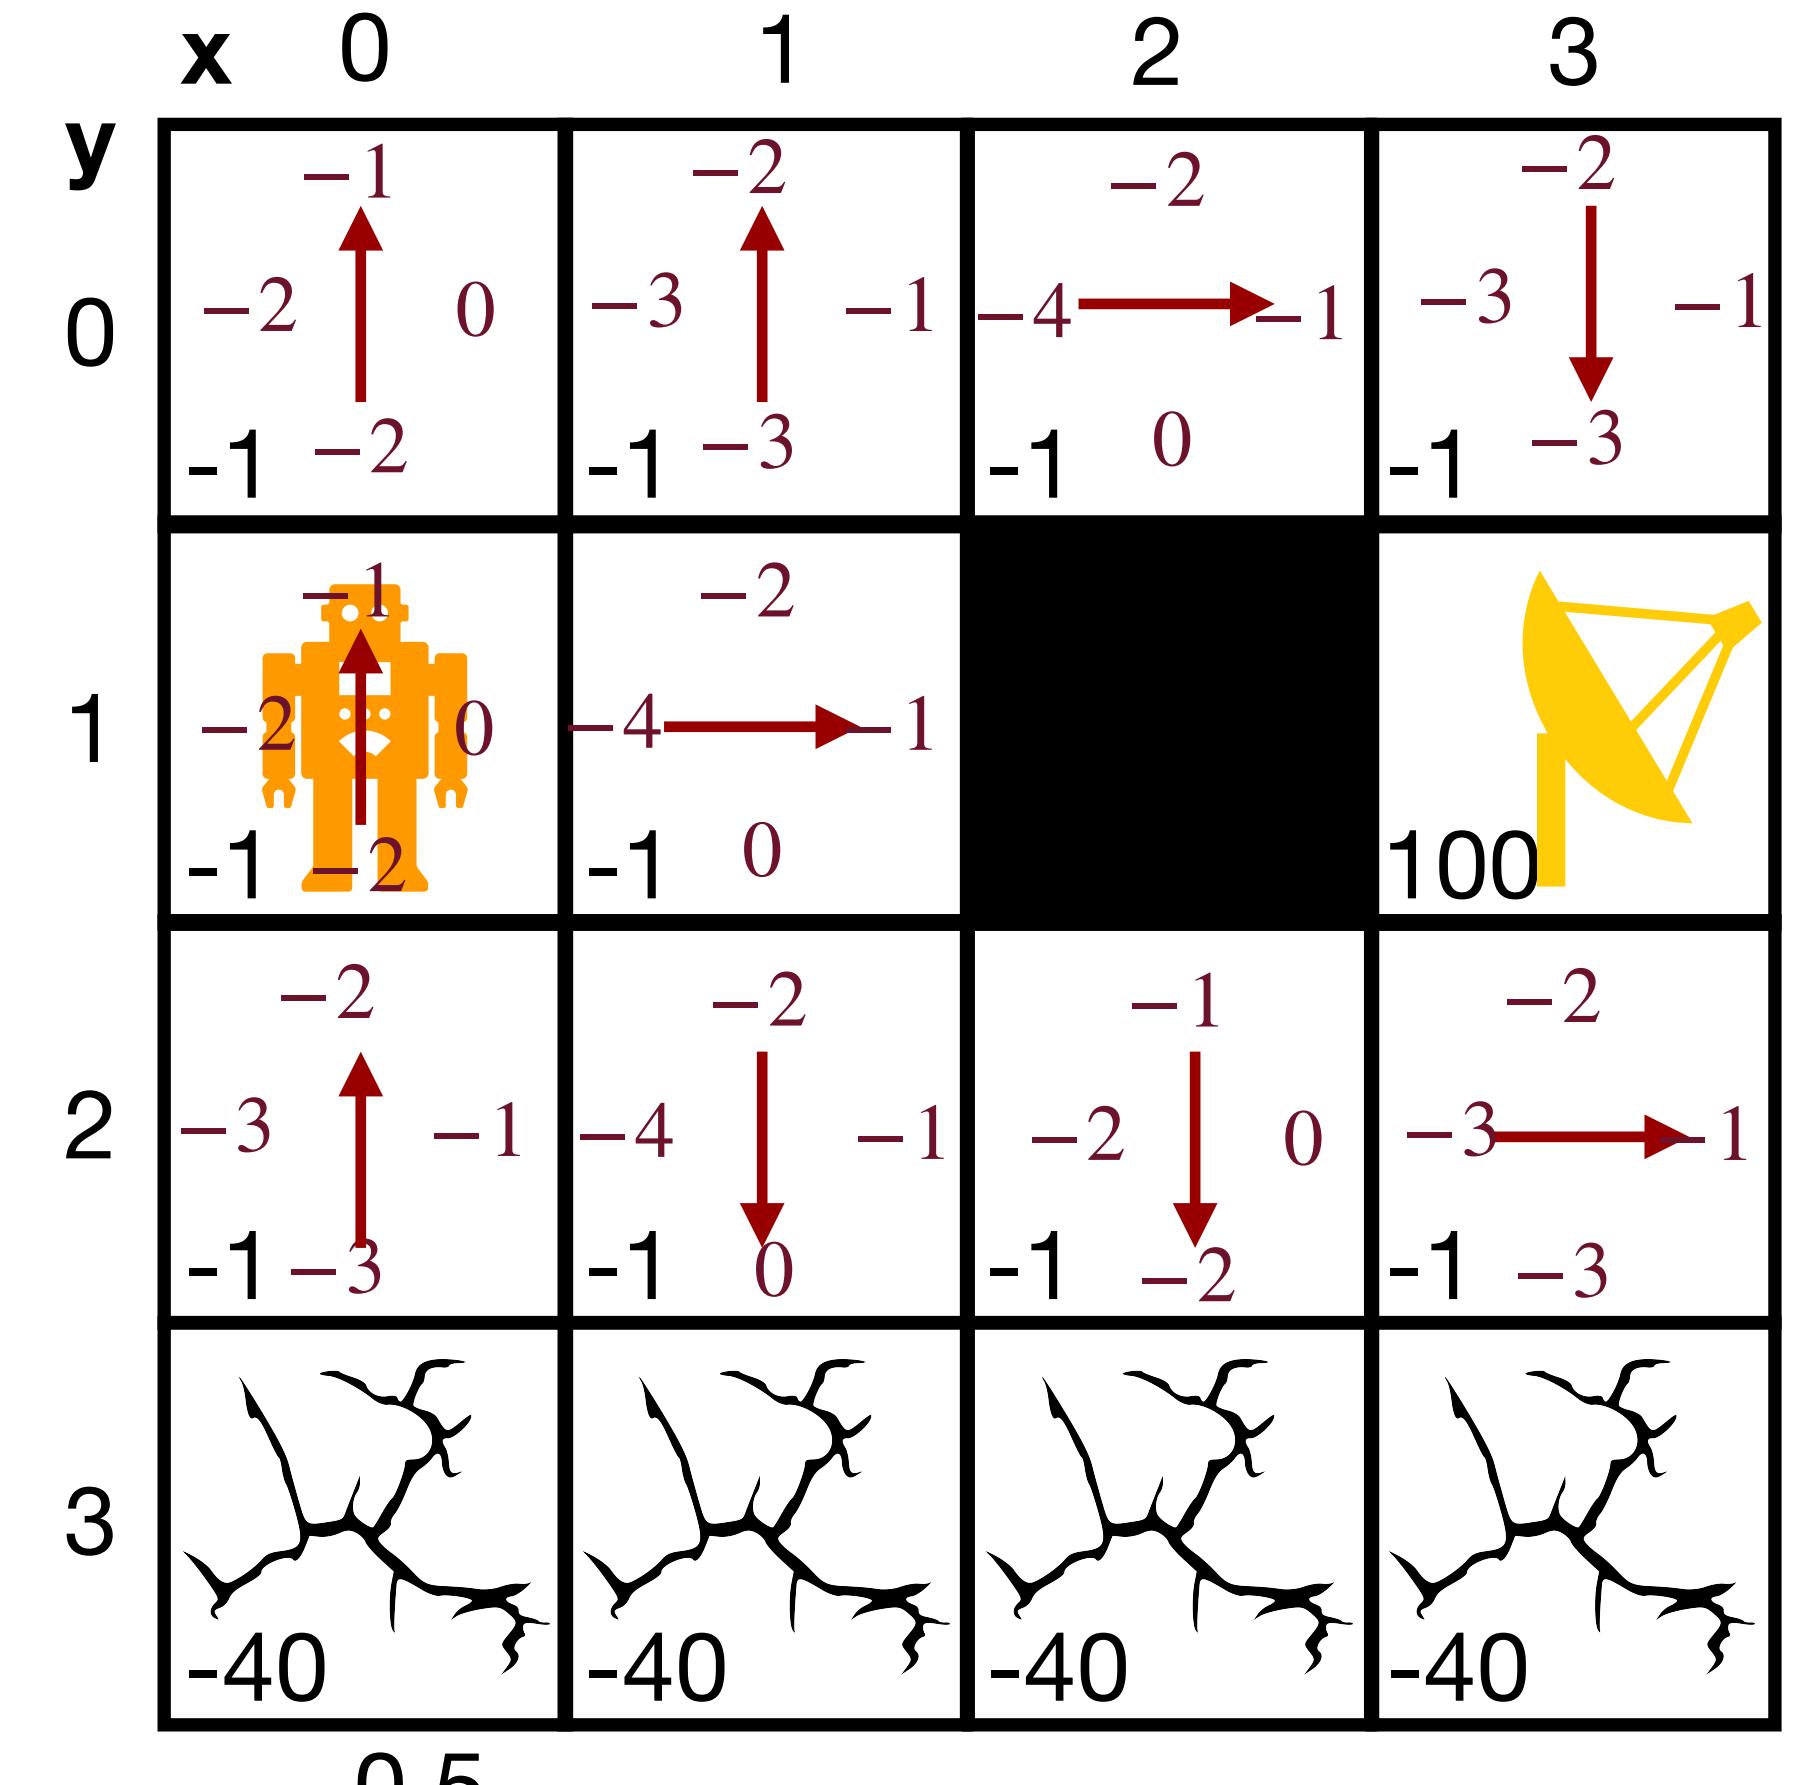
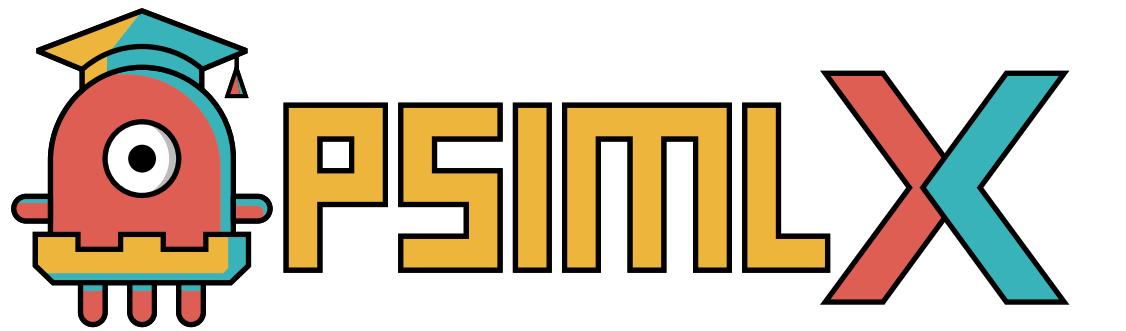
- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\tau_\pi = [(0,2), \text{UP}, -1]$$



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

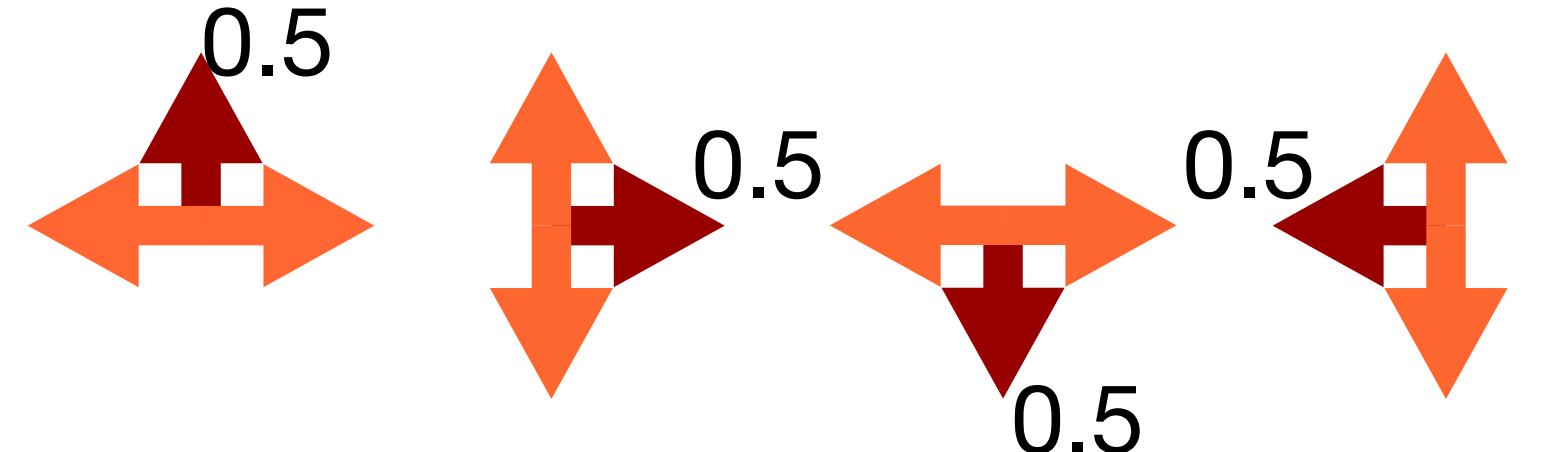
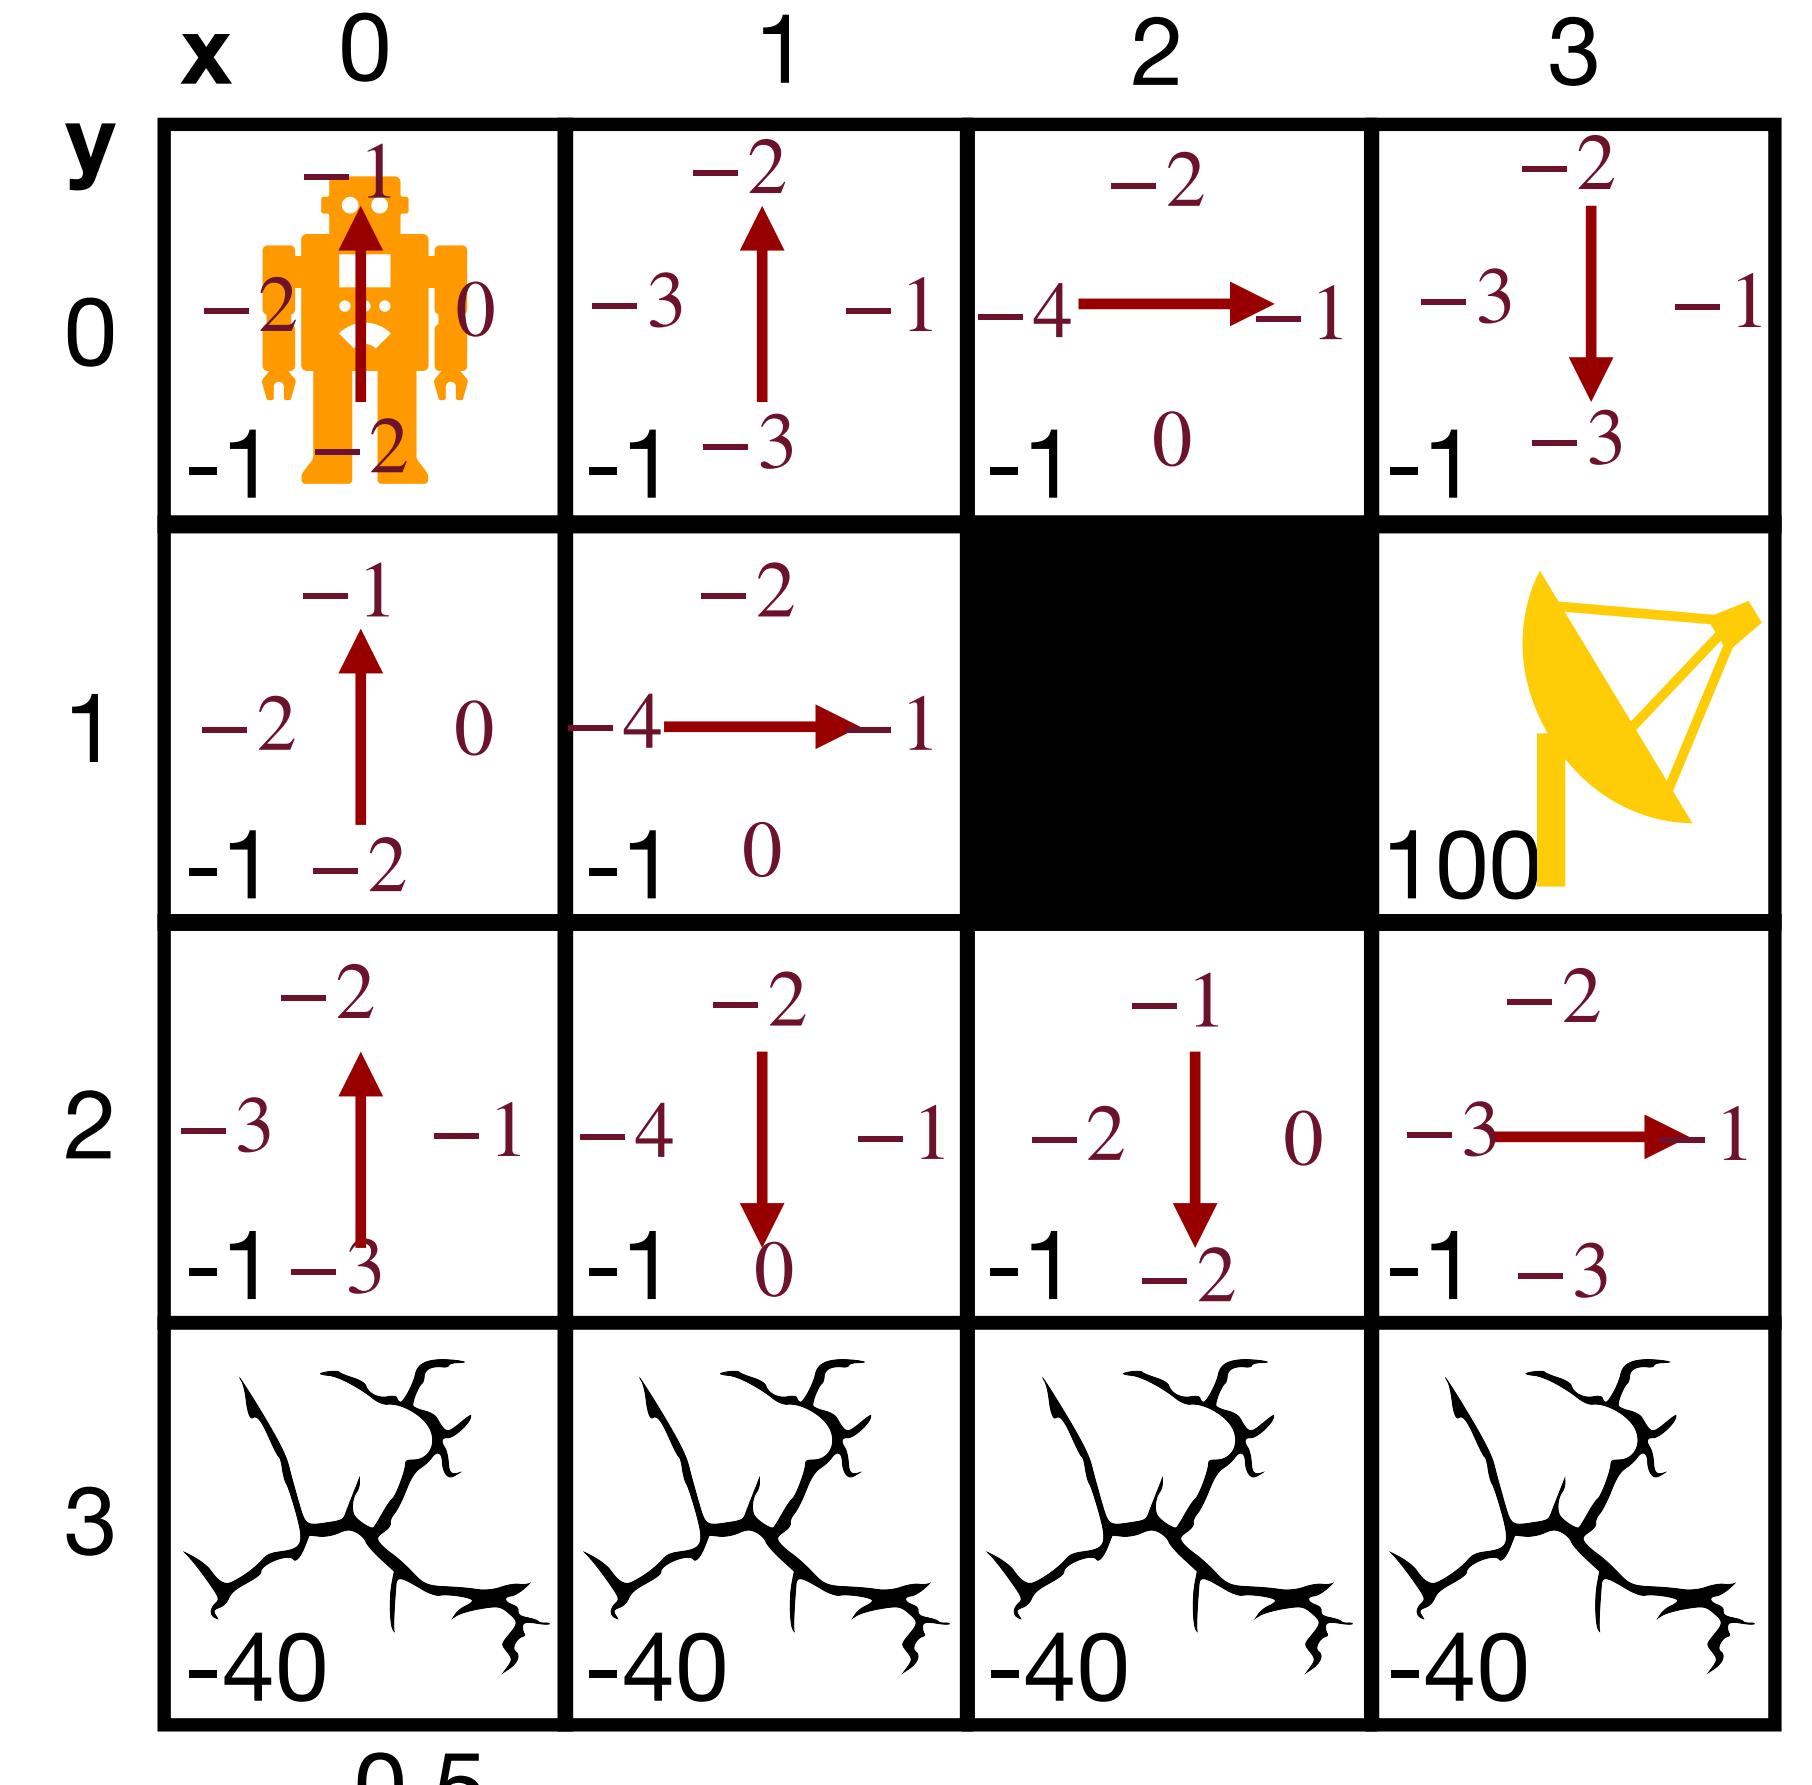
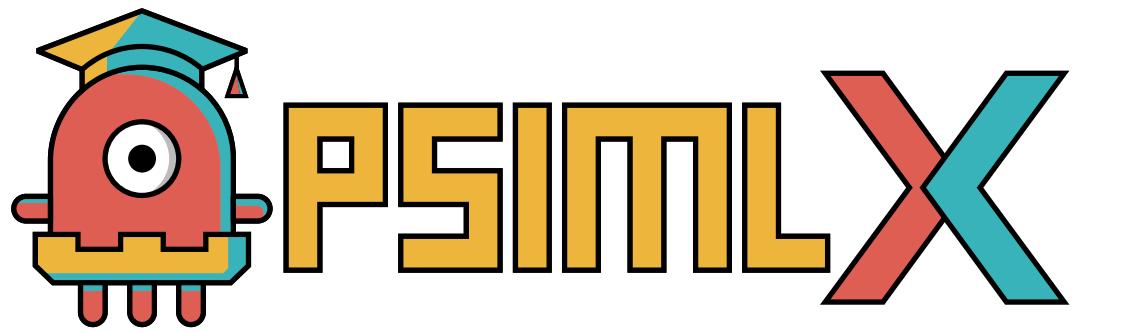
- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1]$$



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

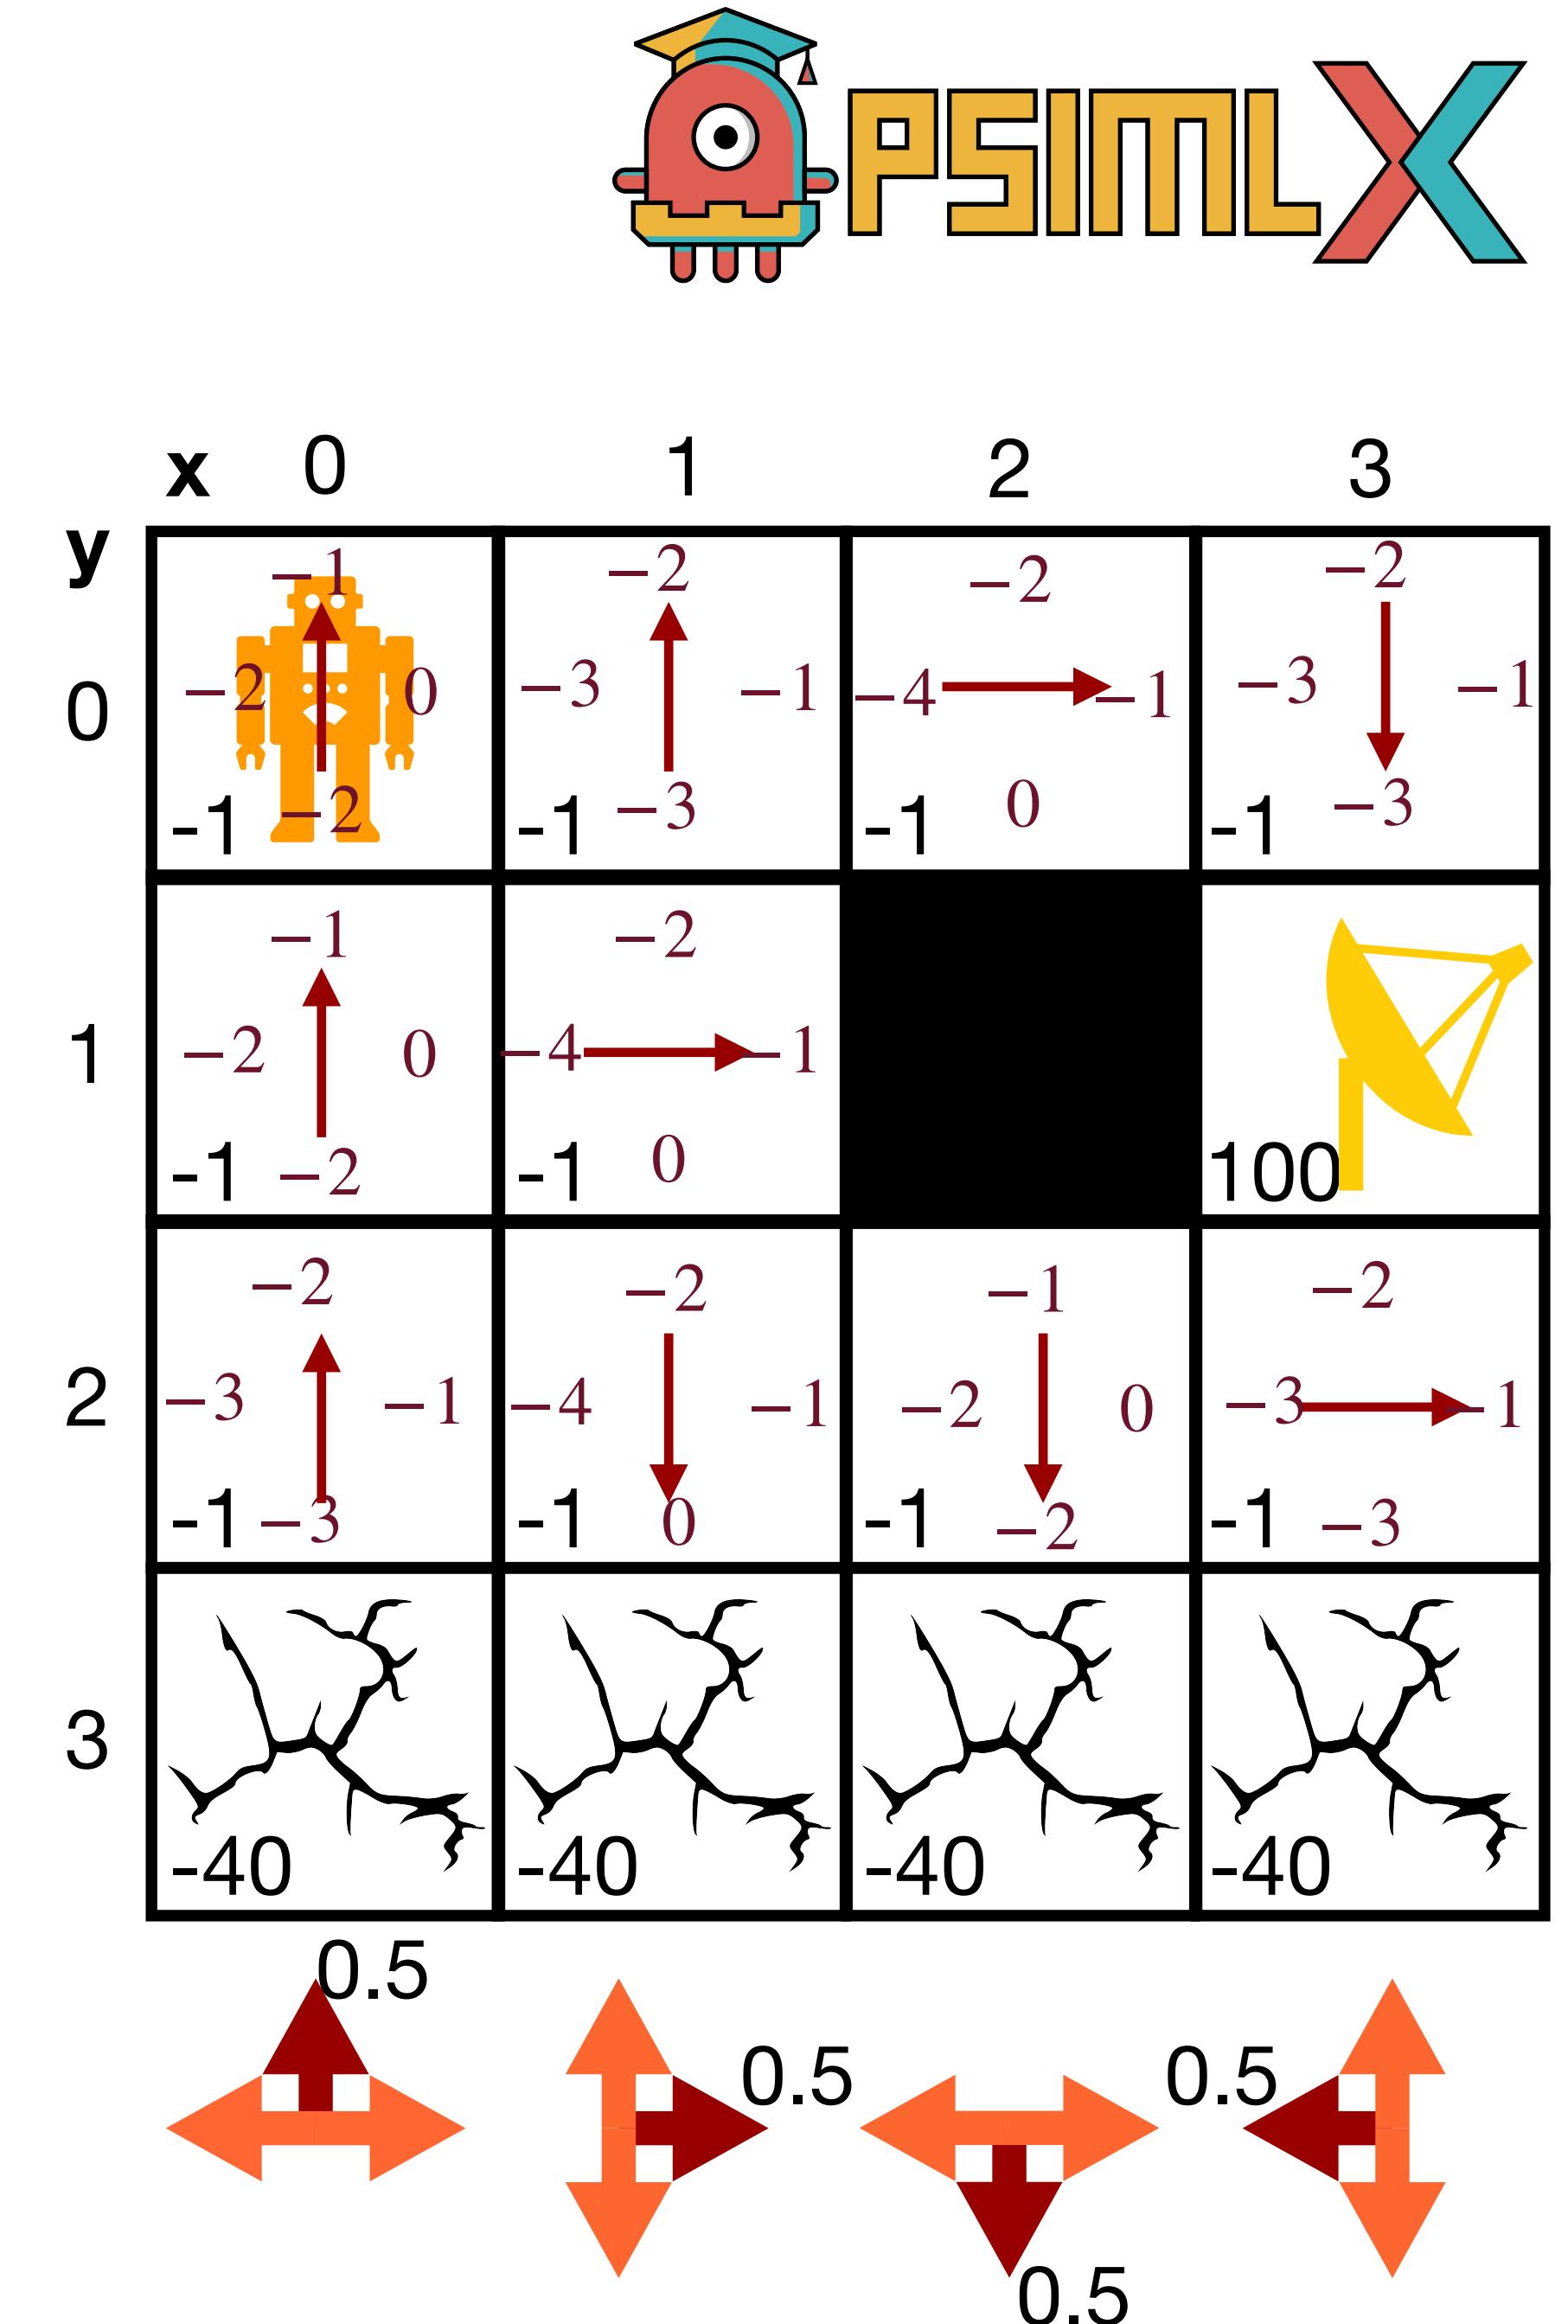
- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

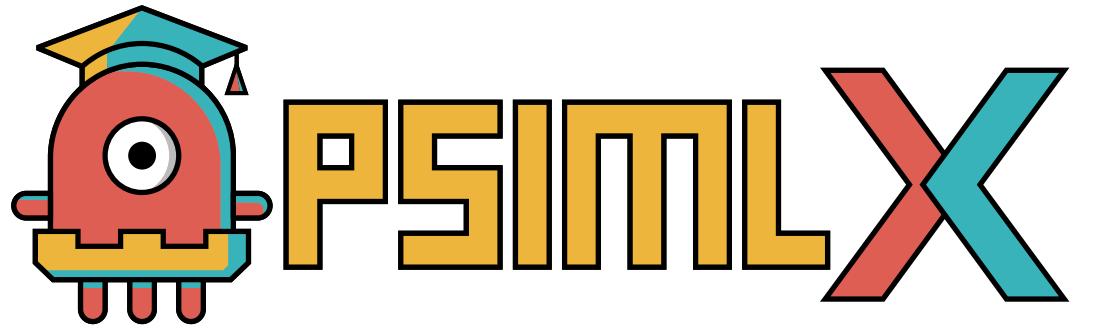
$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1]$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

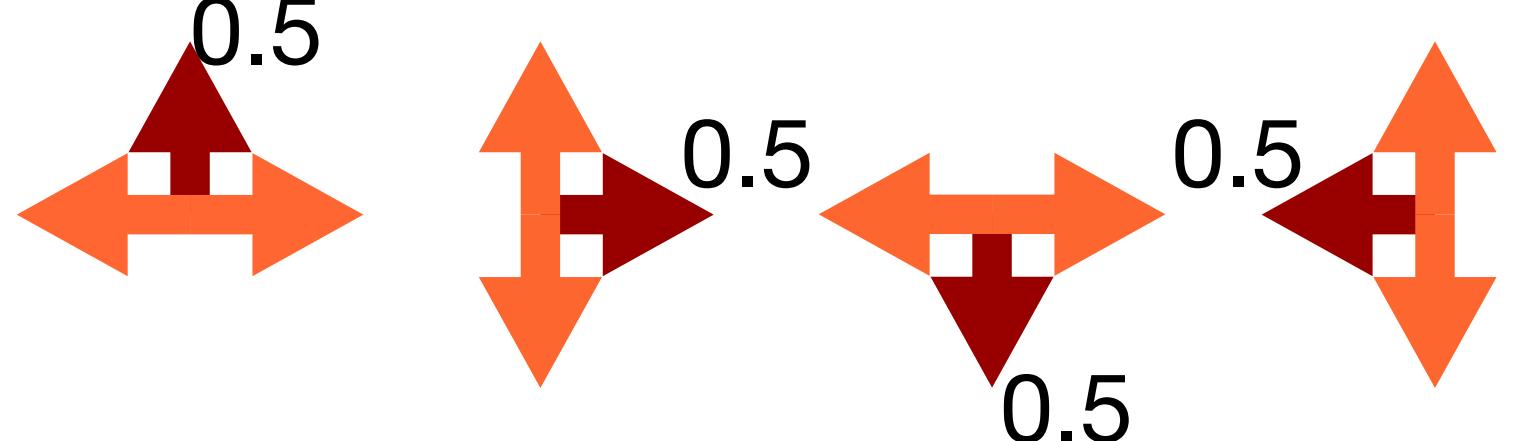
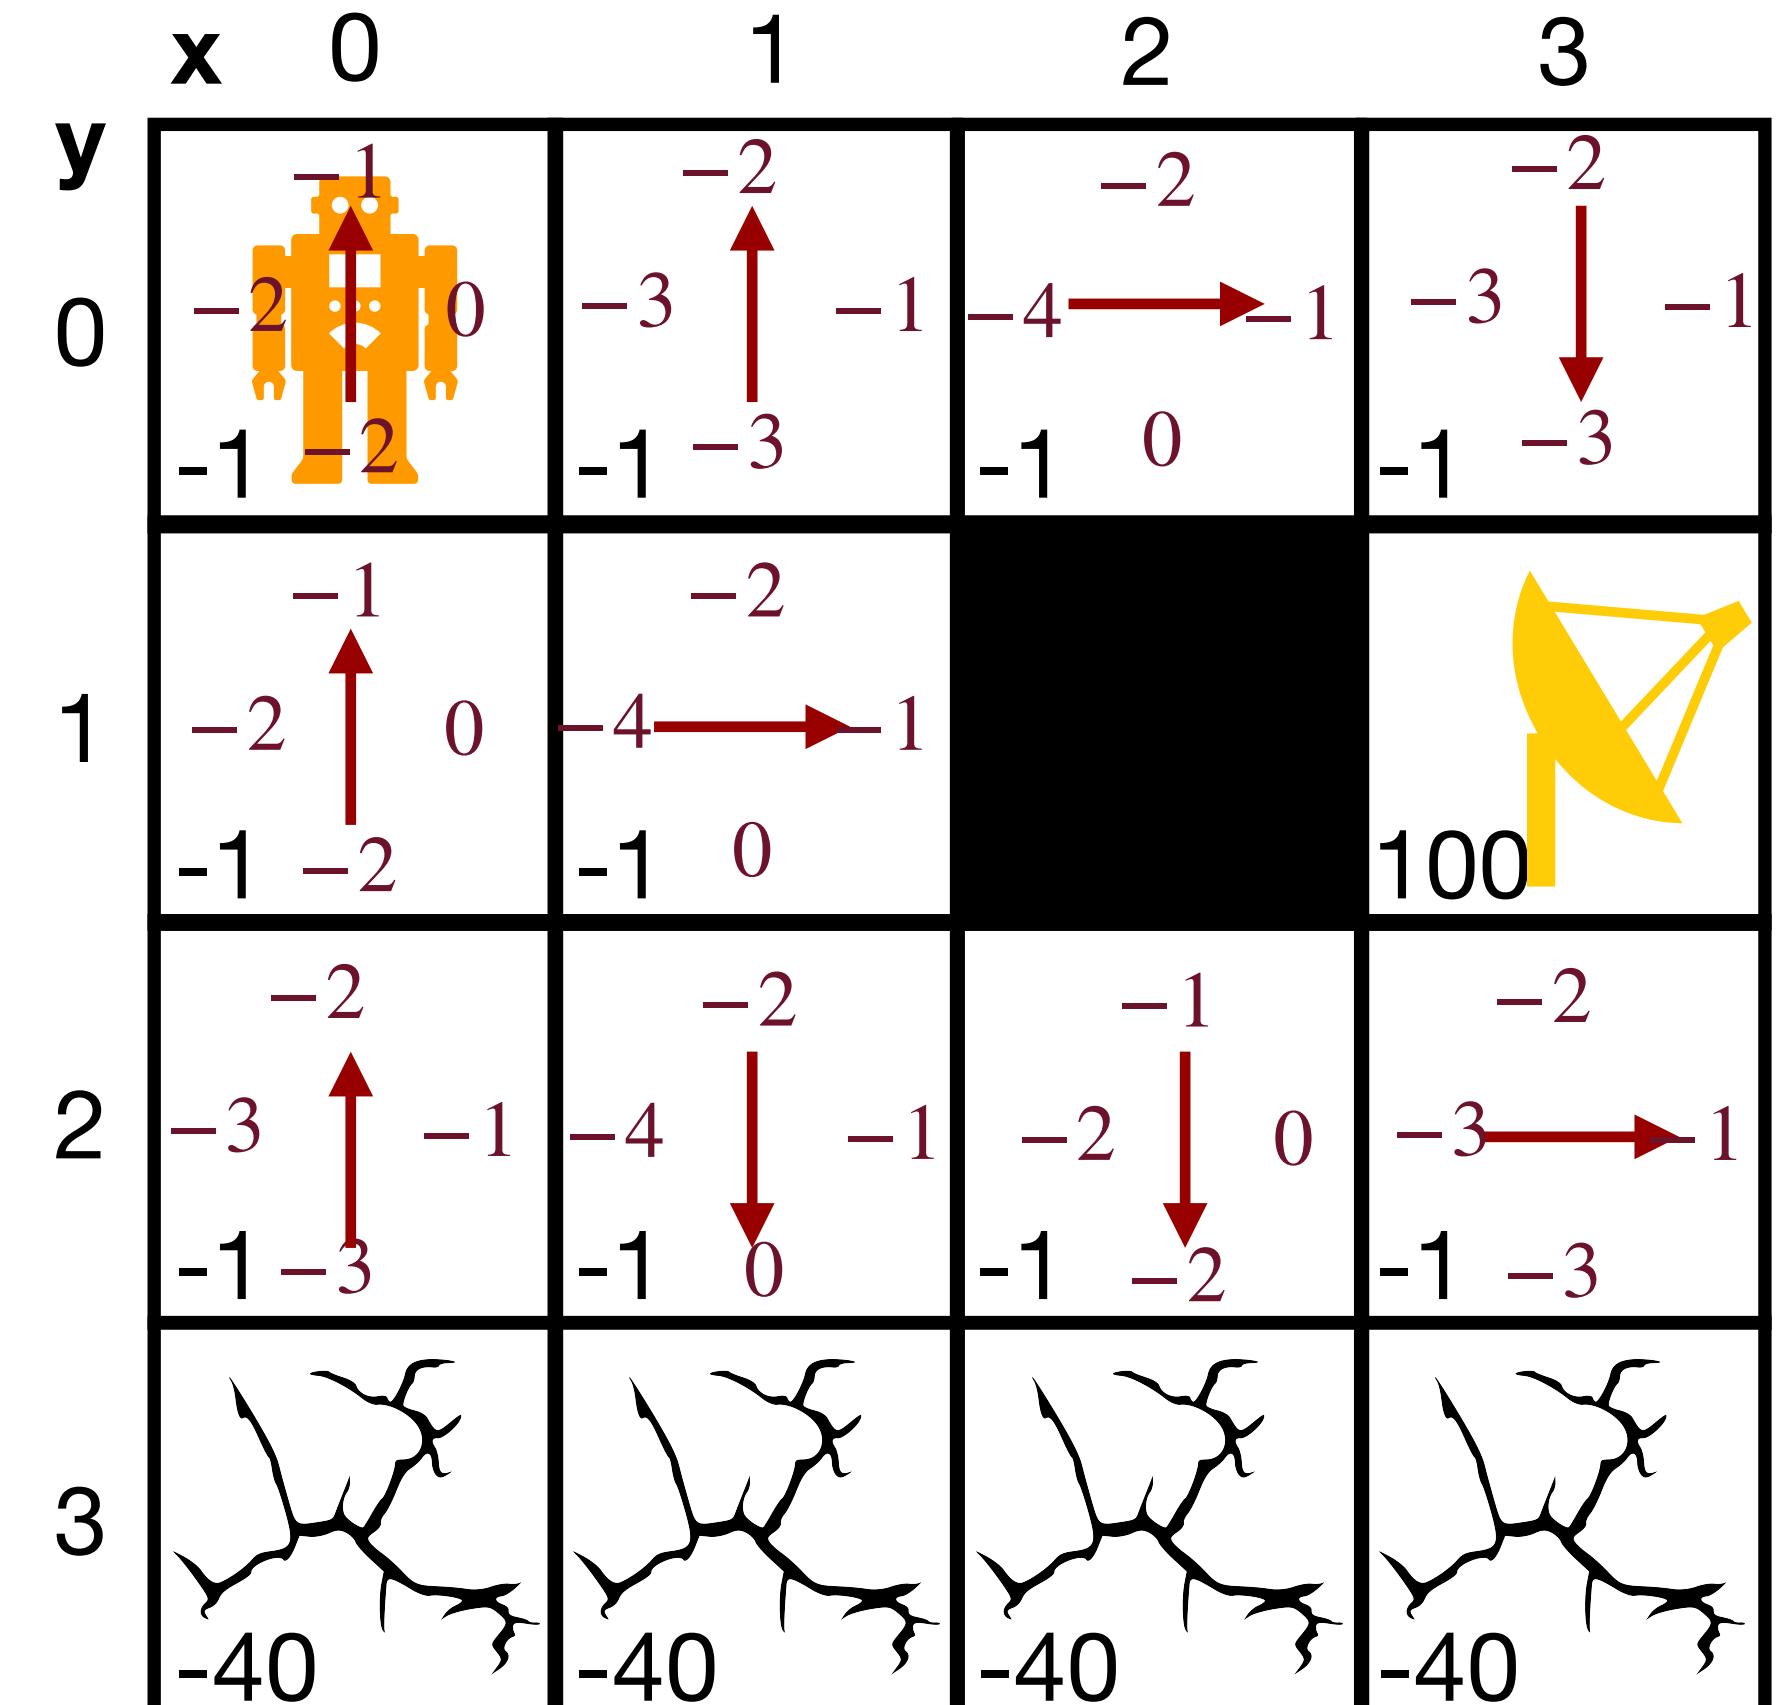
$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

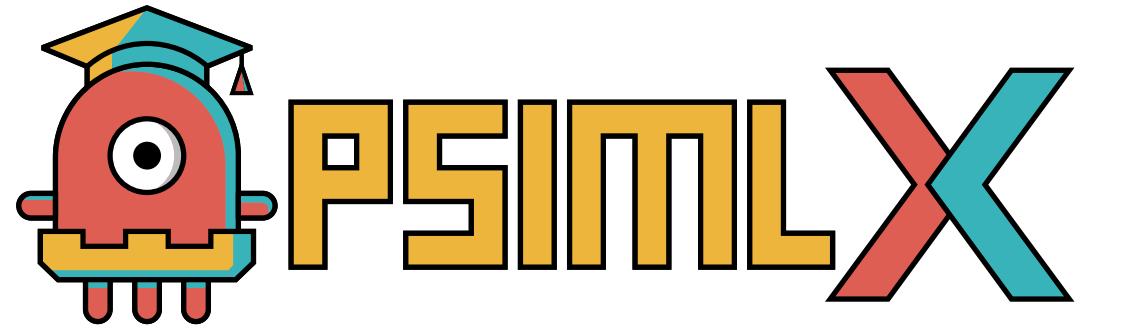
$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1]$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

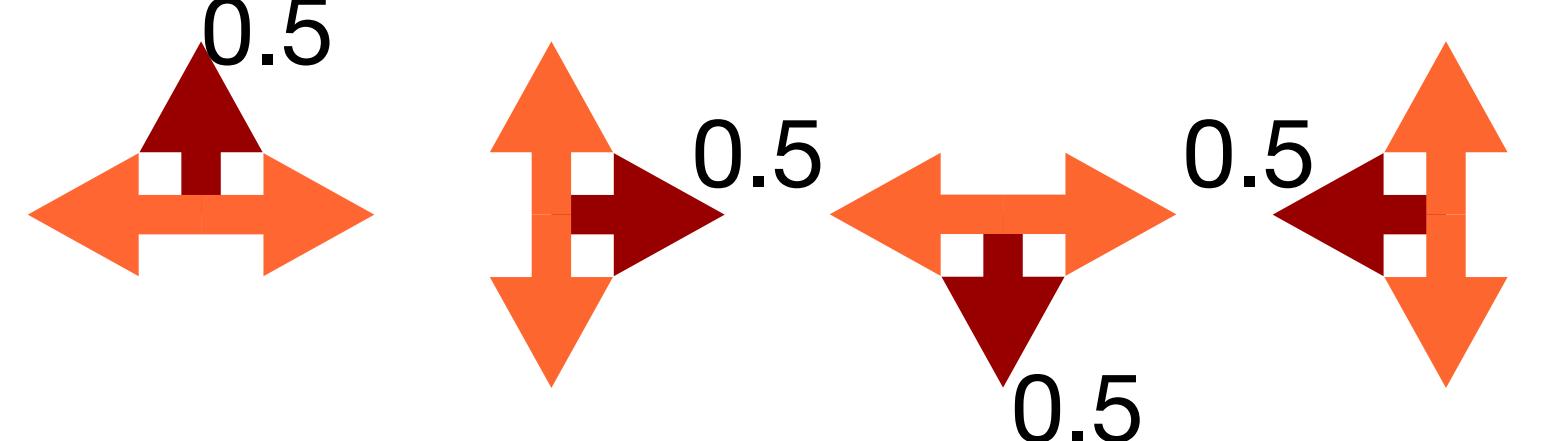
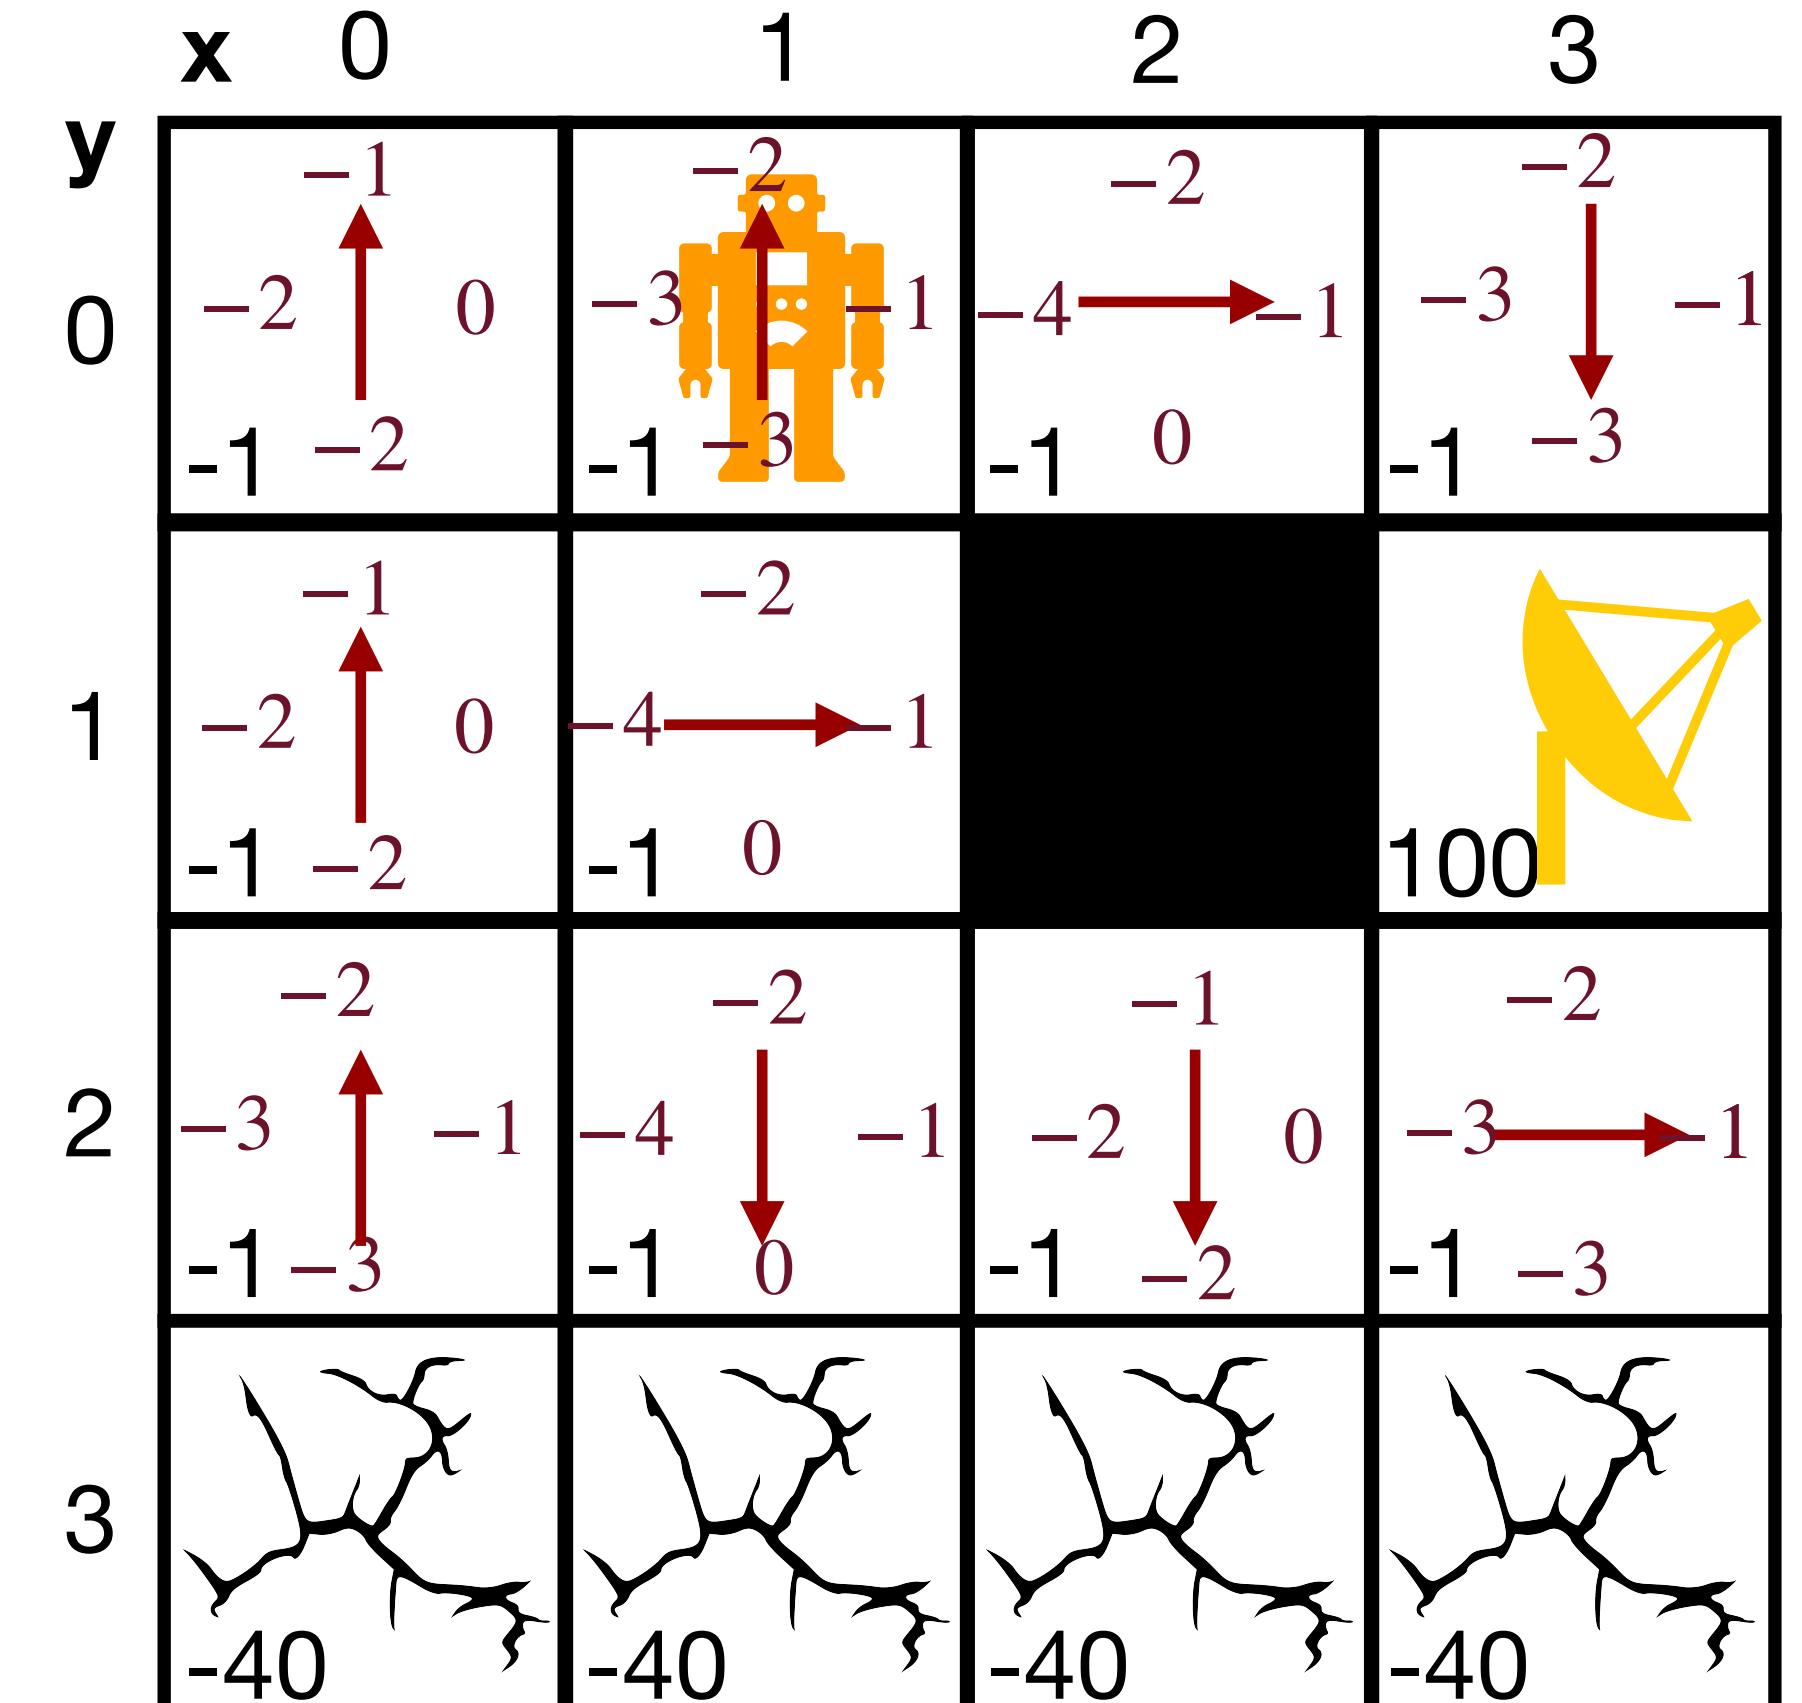
$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\begin{aligned} \tau_\pi = & [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ & [(0,0), \text{UP}, -1] \end{aligned}$$



Model free RL

Example 1: On-Policy MC Algorithm

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

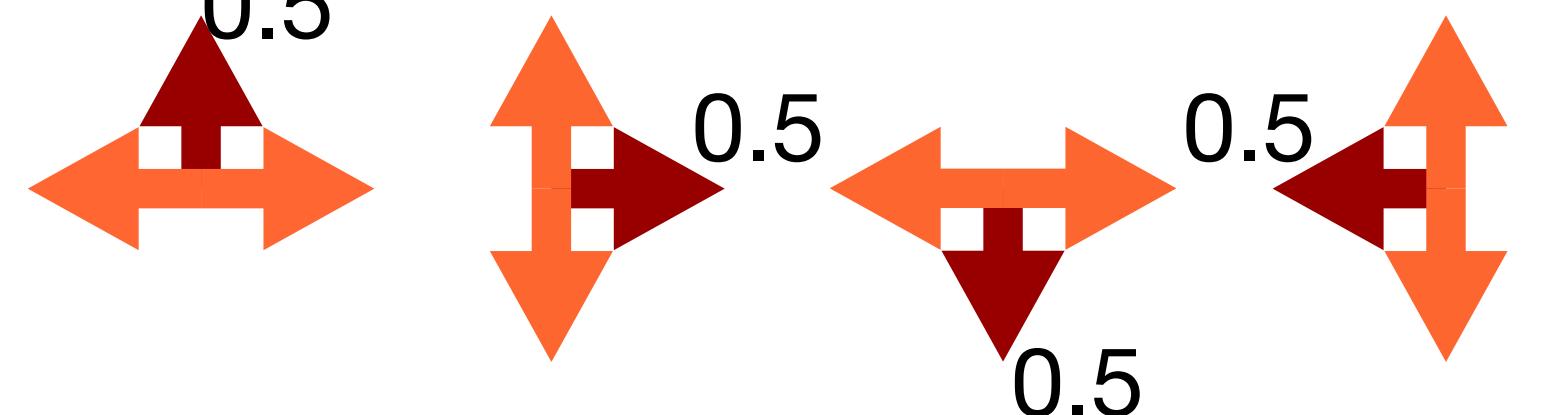
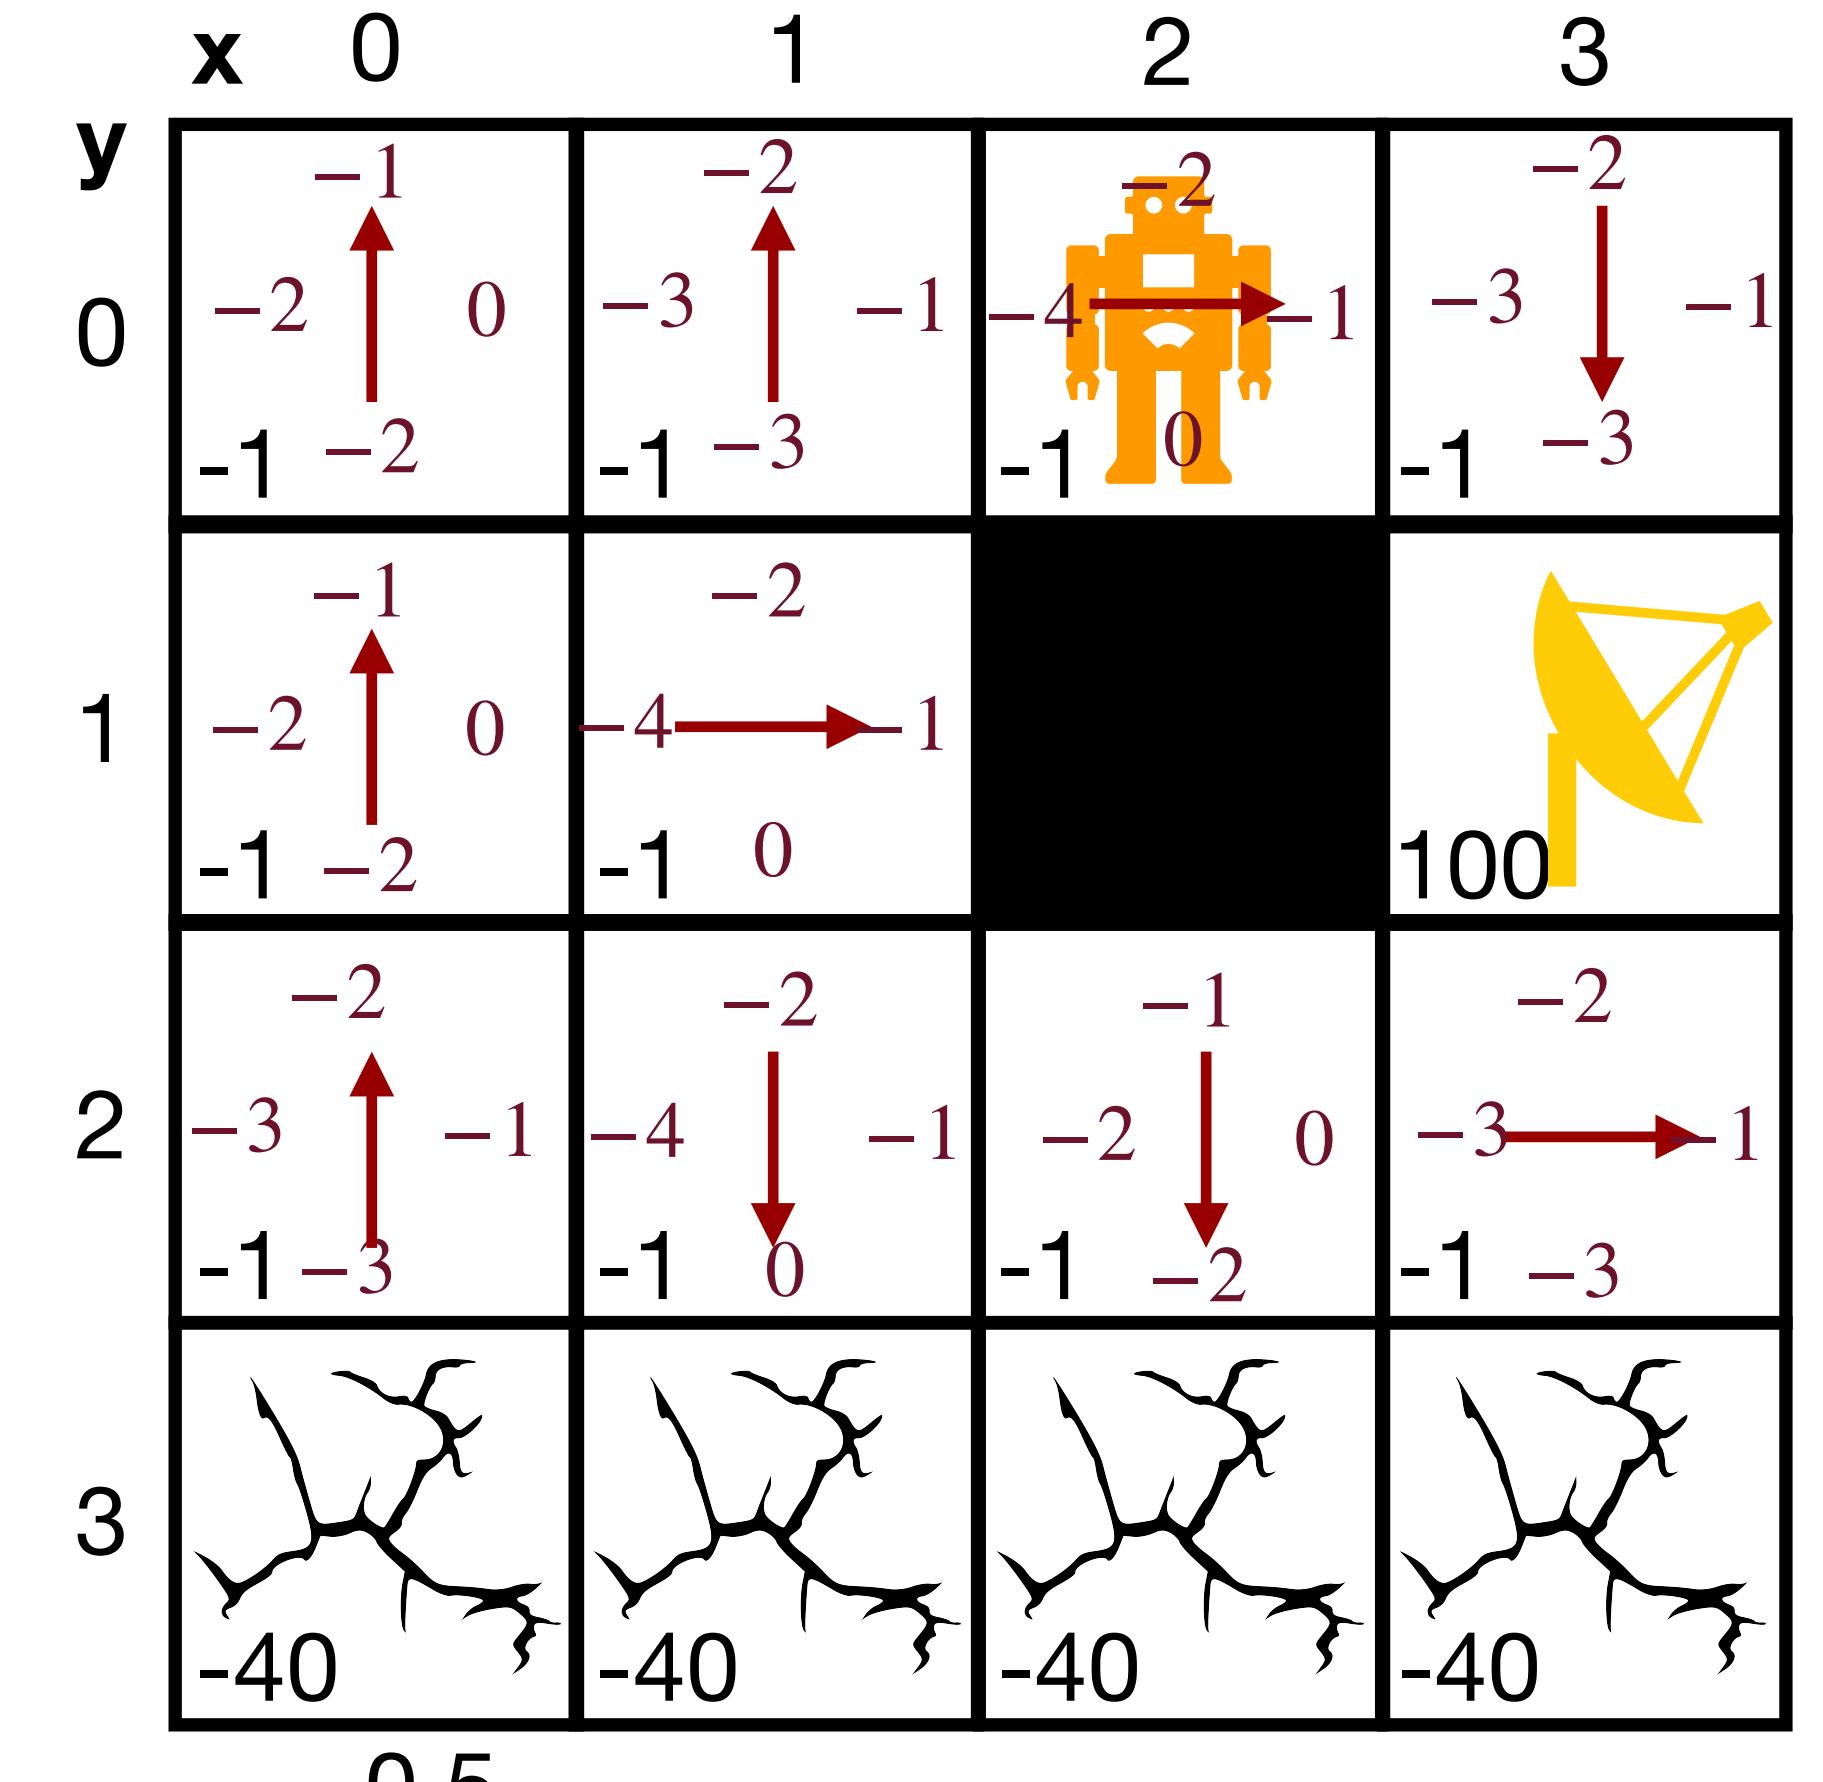
- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

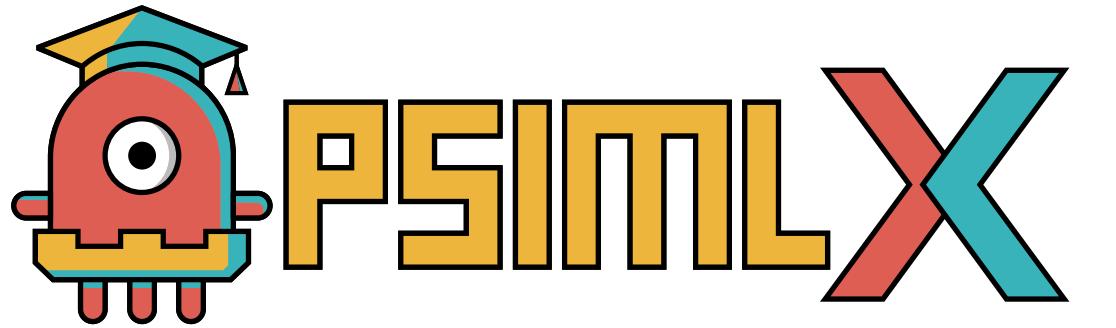
$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\begin{aligned} \tau_\pi = & [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ & [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1] \end{aligned}$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

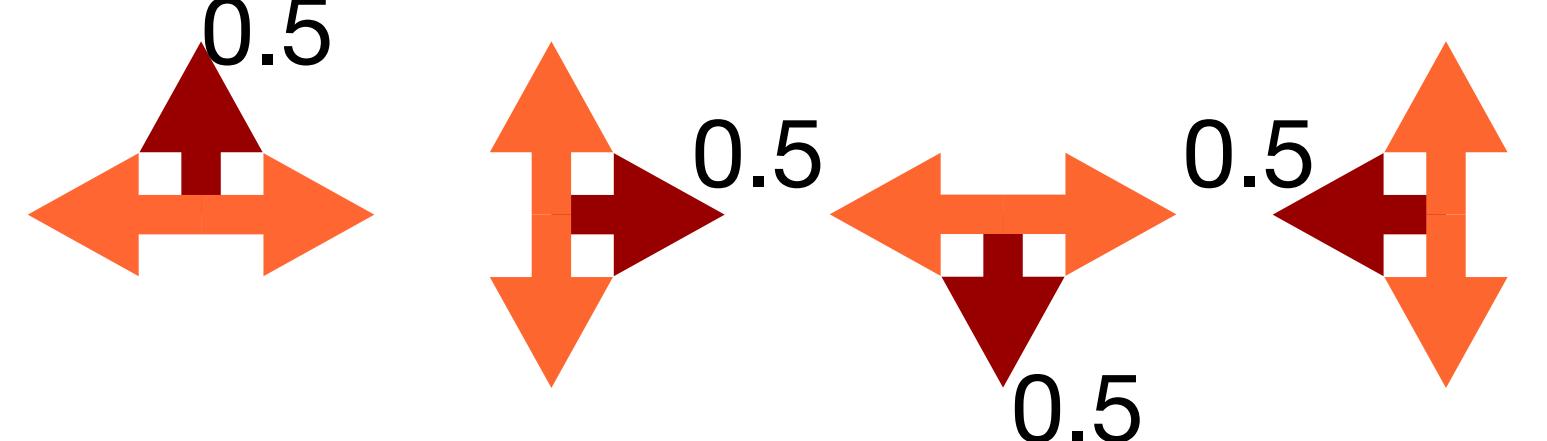
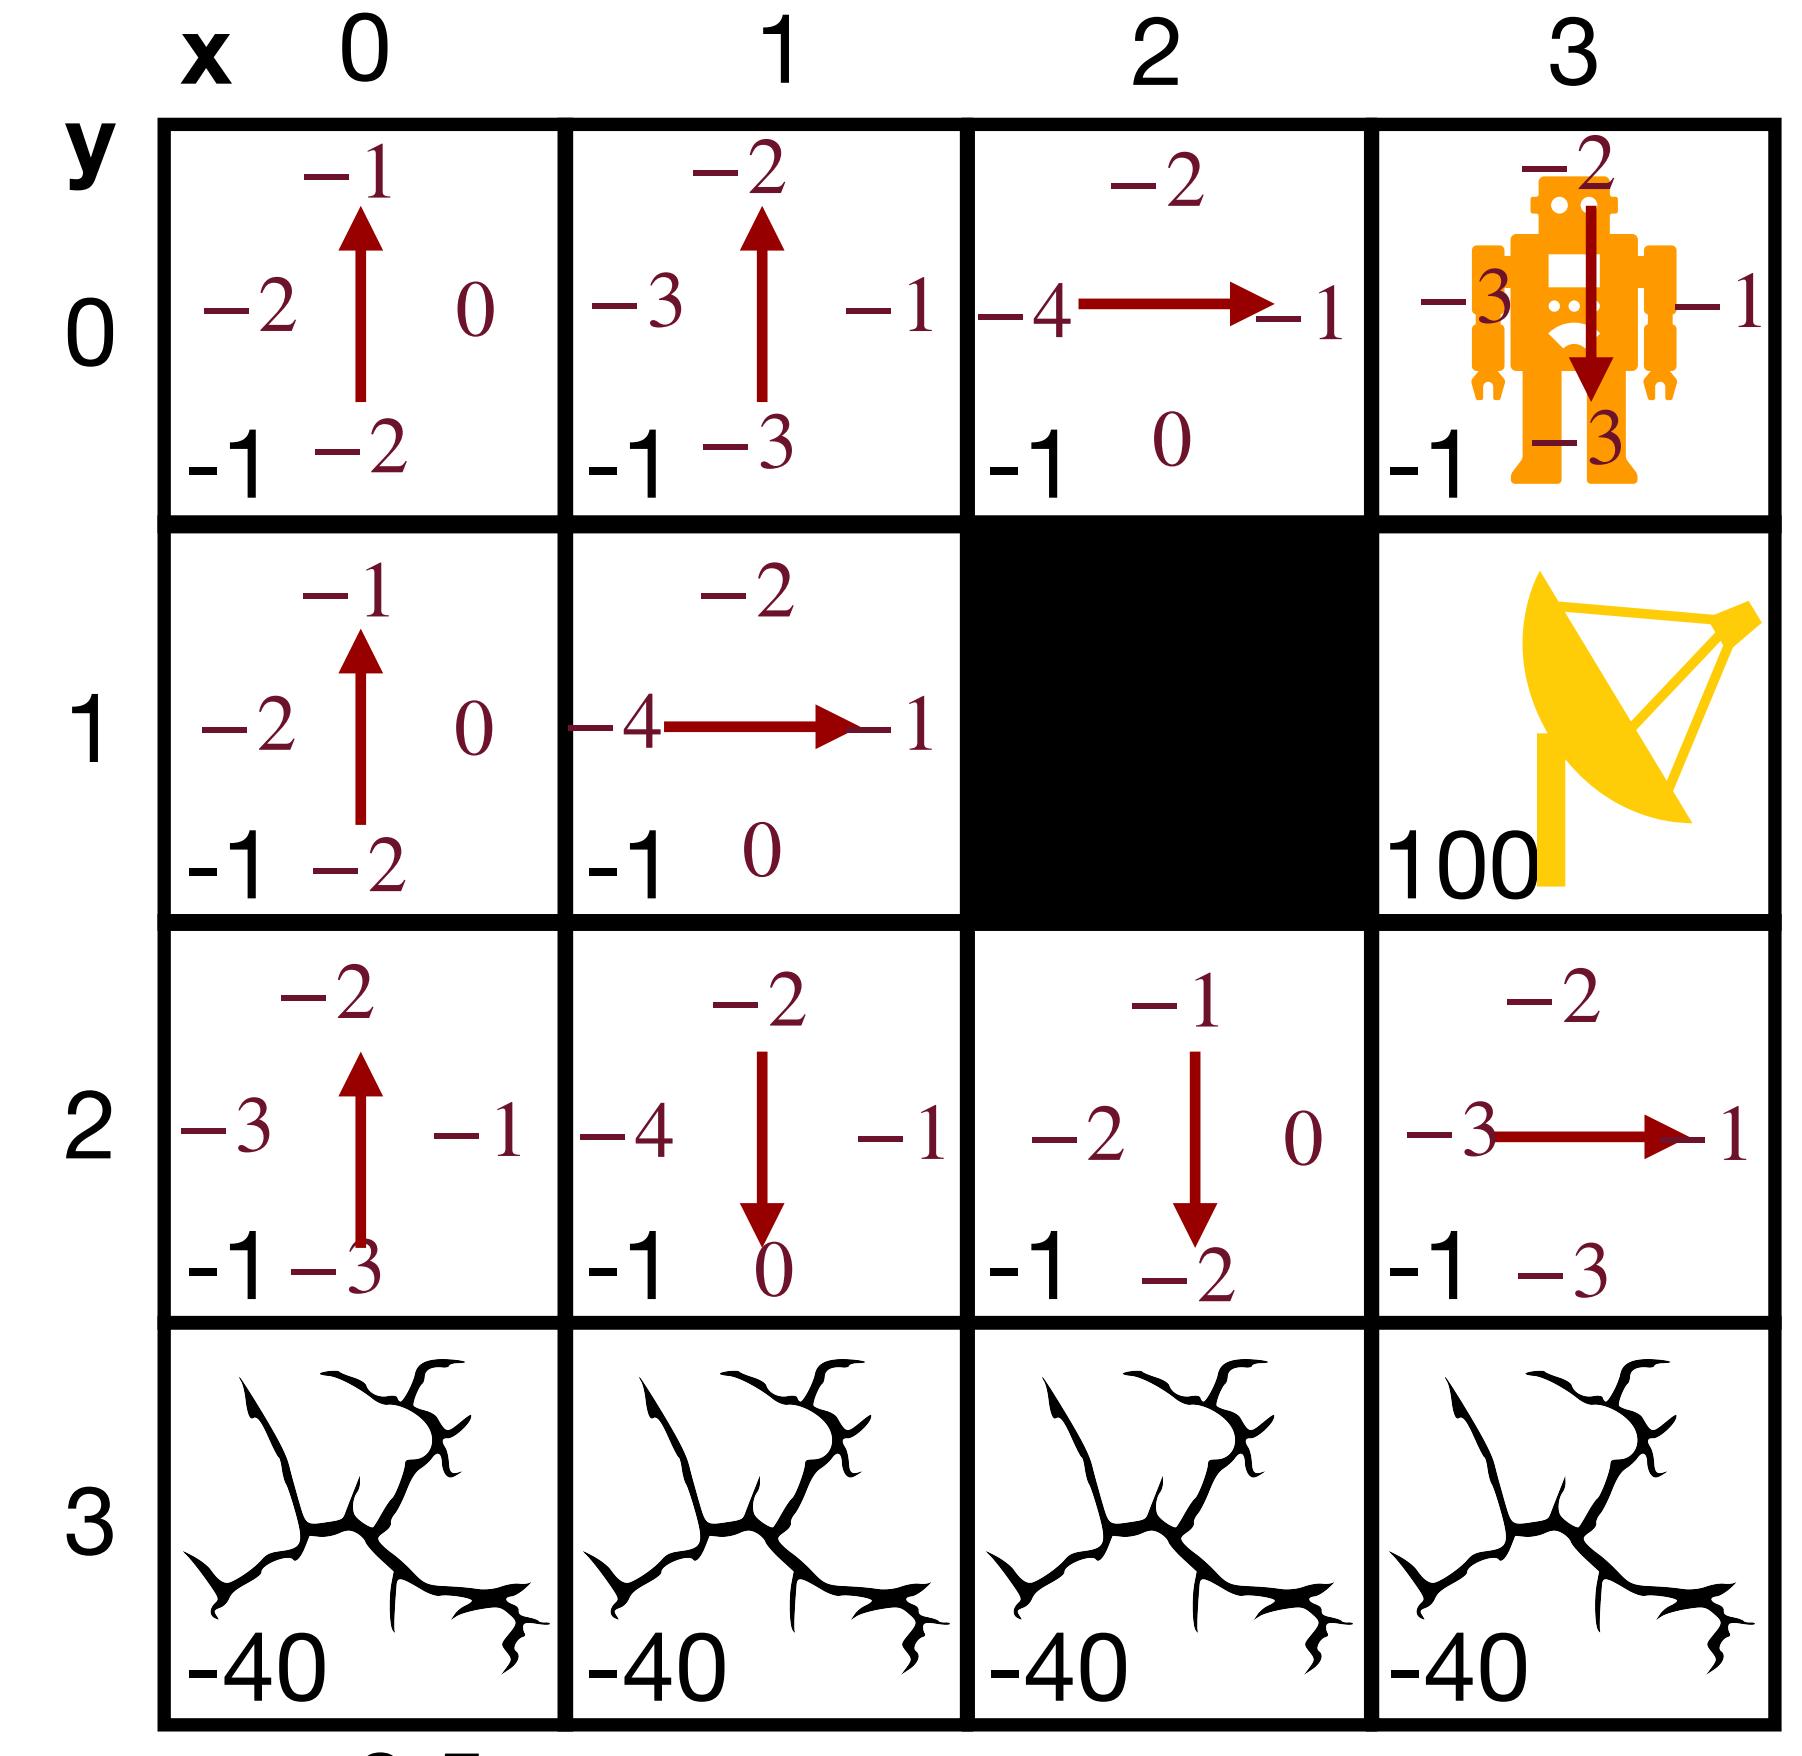
Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

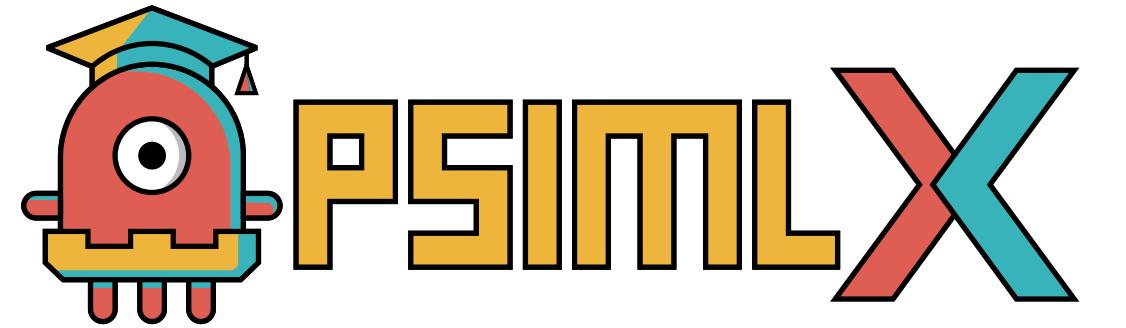
- Append G to $Returns(S_t, A_t)$
- $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
- $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
- For all $a \in \mathcal{A}(S_t)$:
- $\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$

$$\begin{aligned} \tau_\pi = & [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ & [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1] \end{aligned}$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

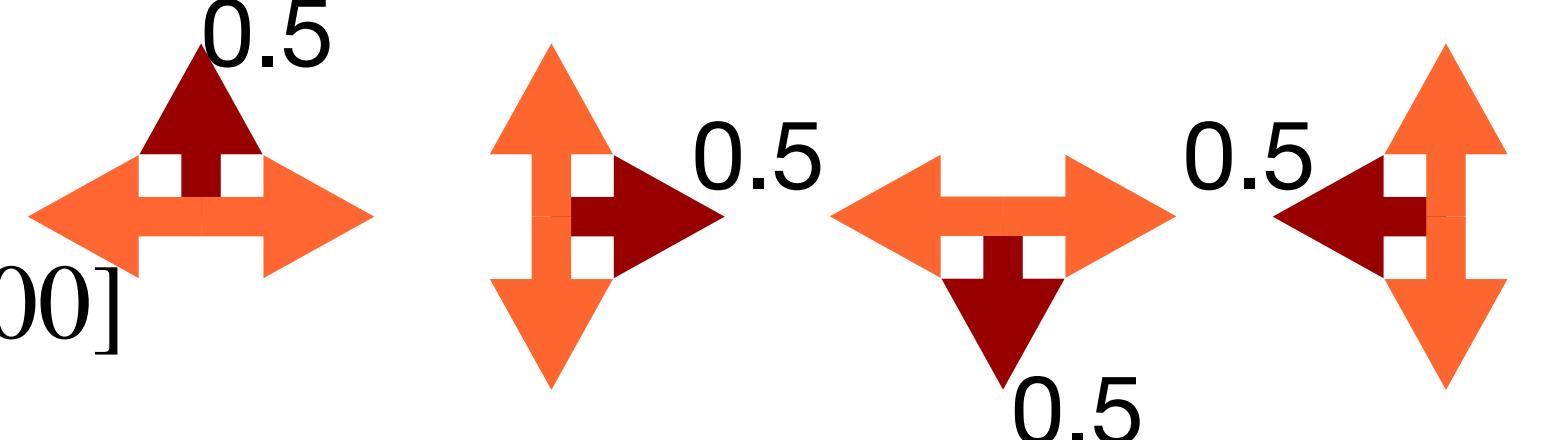
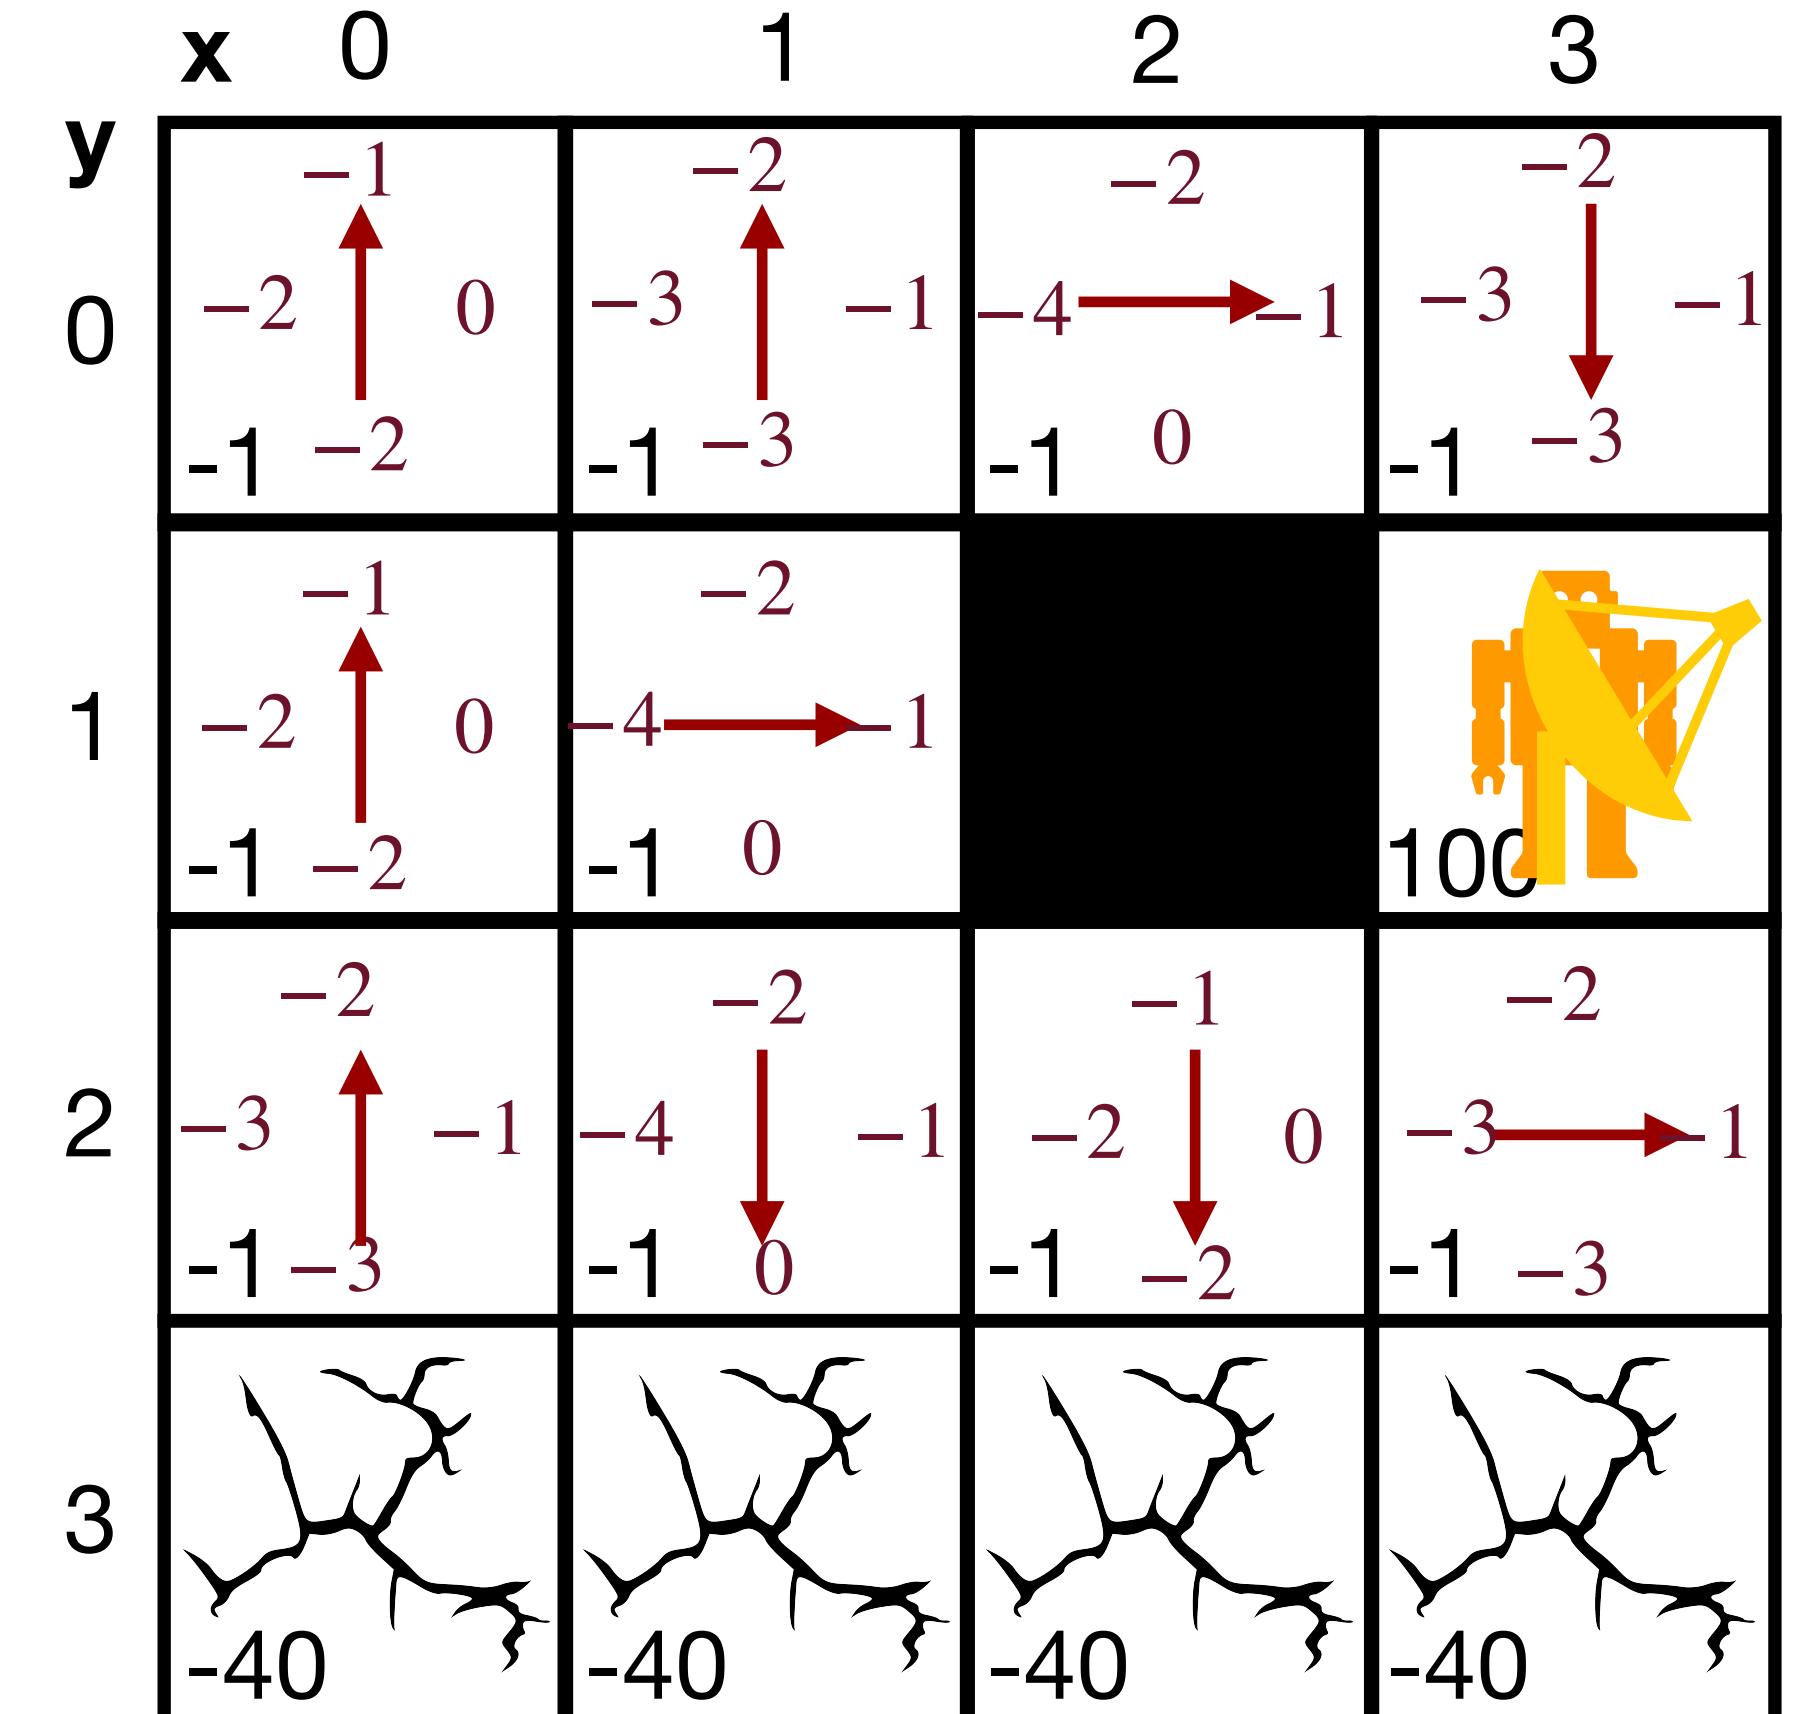
$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$$\begin{aligned} \tau_\pi = & [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ & [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100] \end{aligned}$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

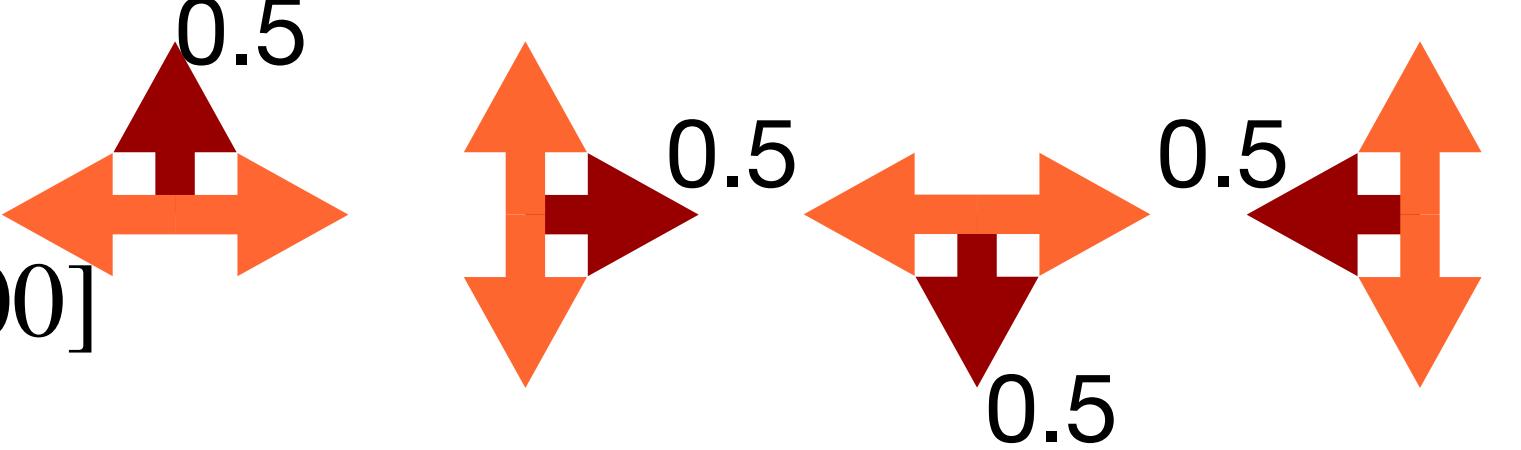
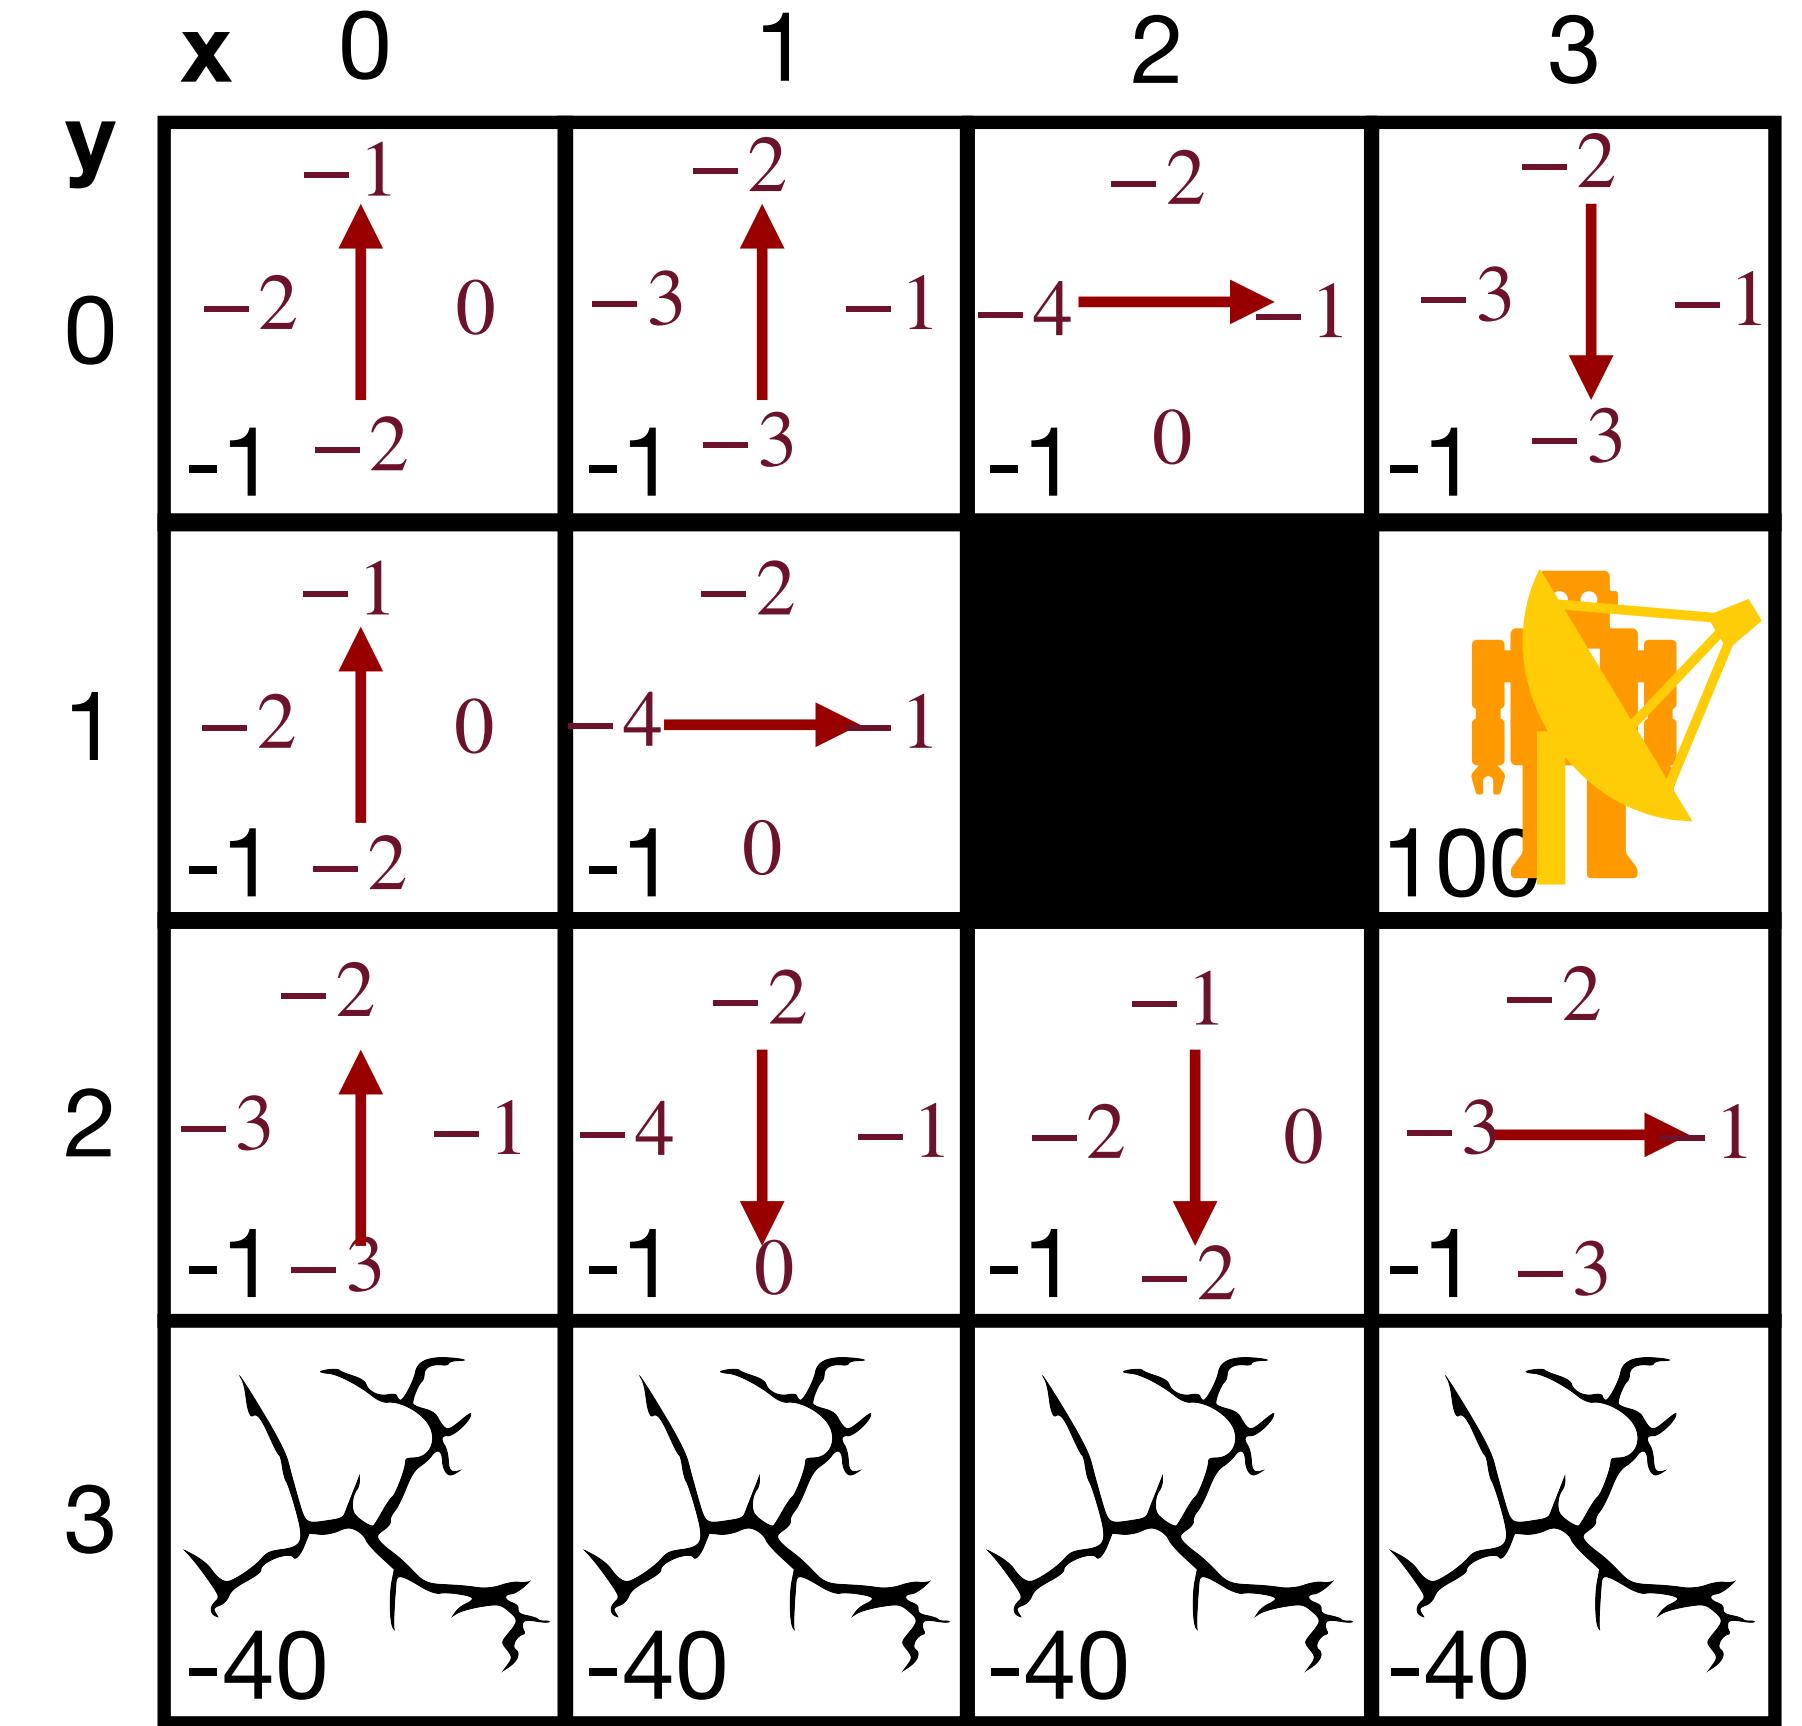
$A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

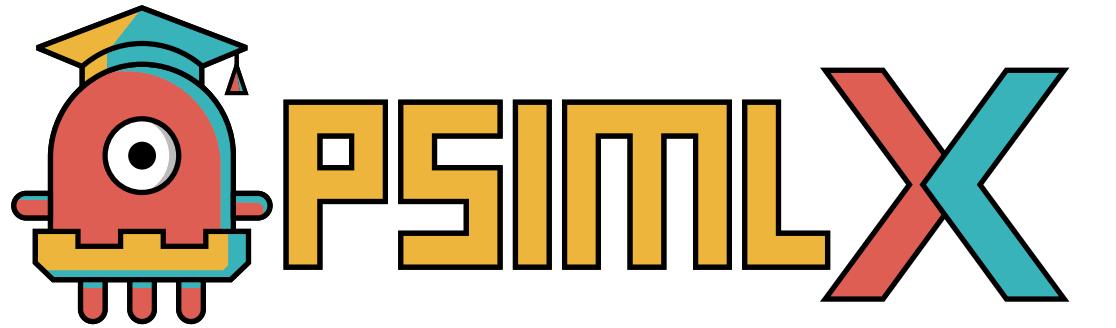
$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100]$$

$$G = 0$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

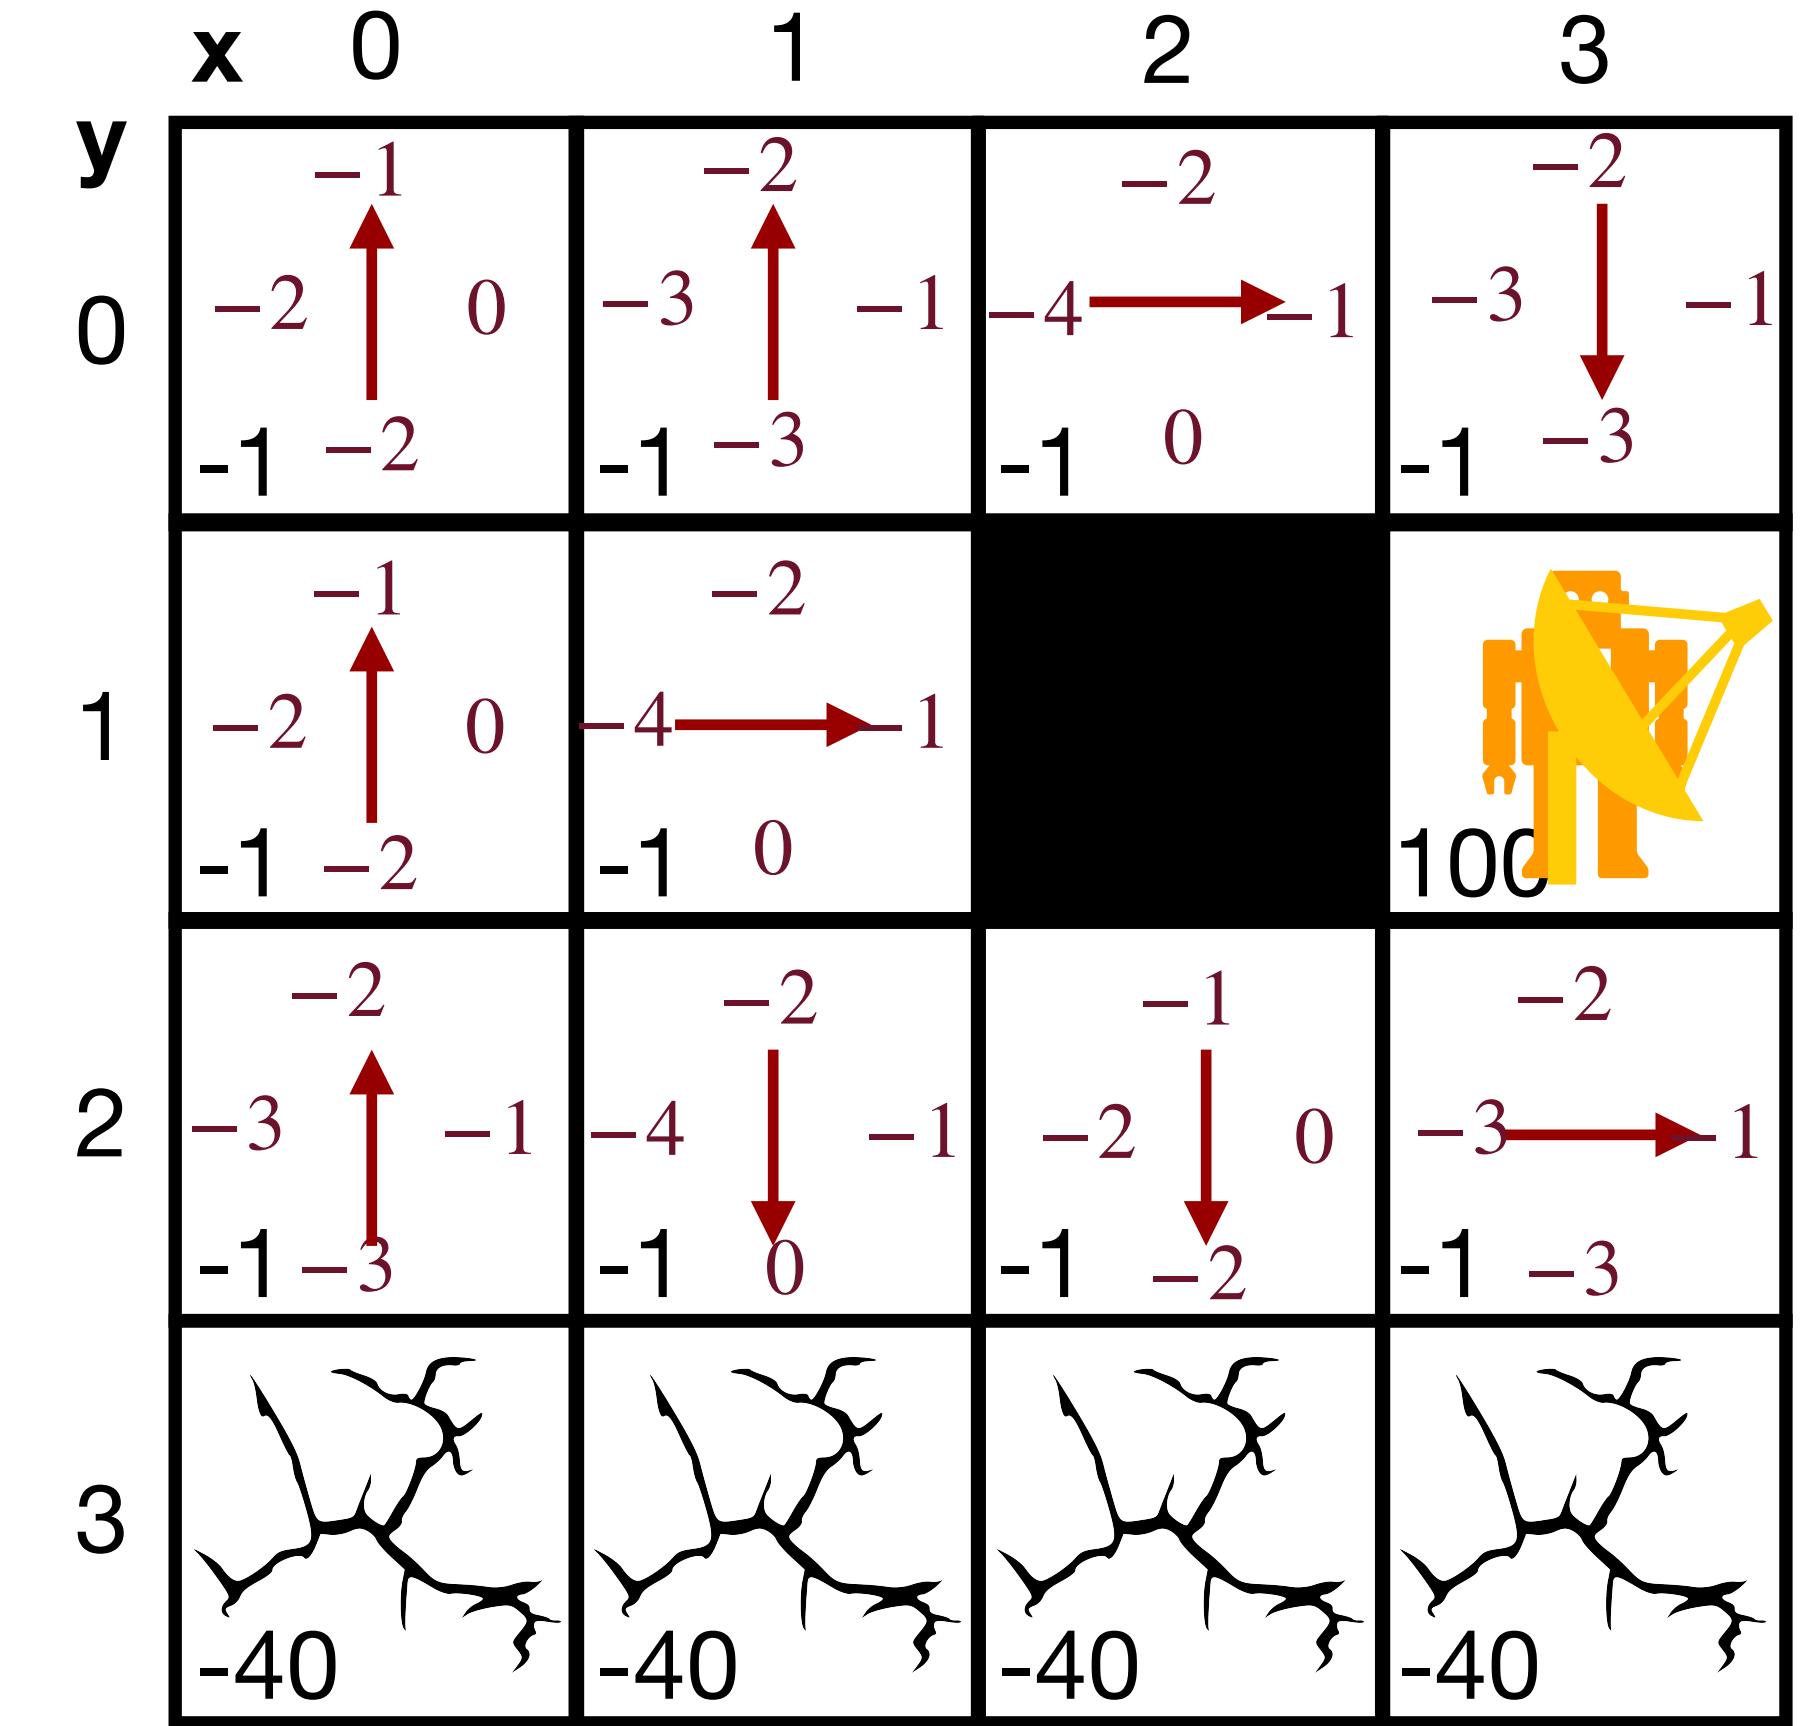
Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

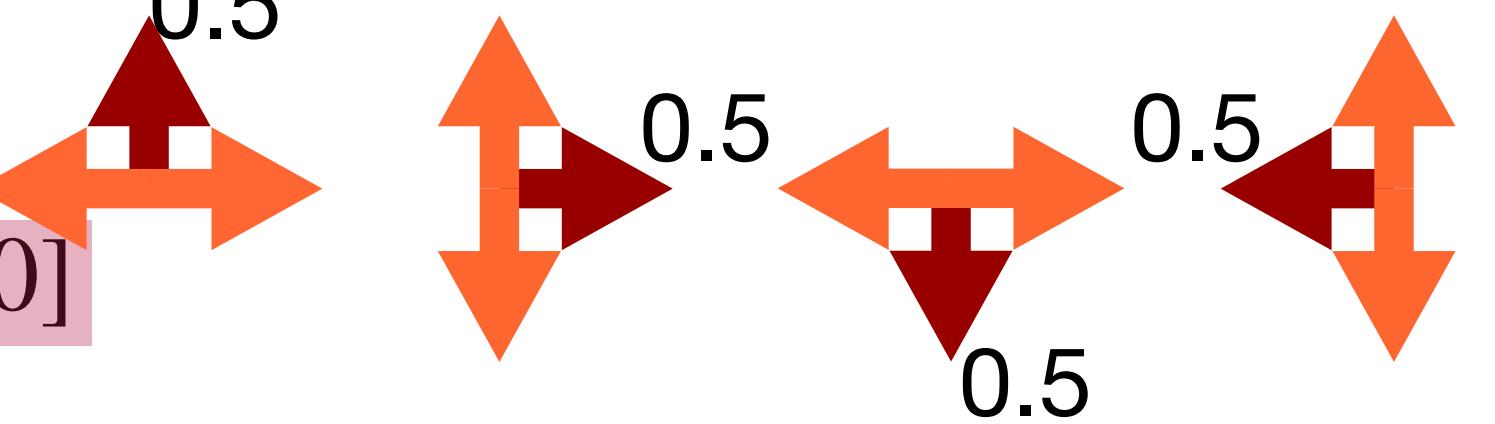
- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1 \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



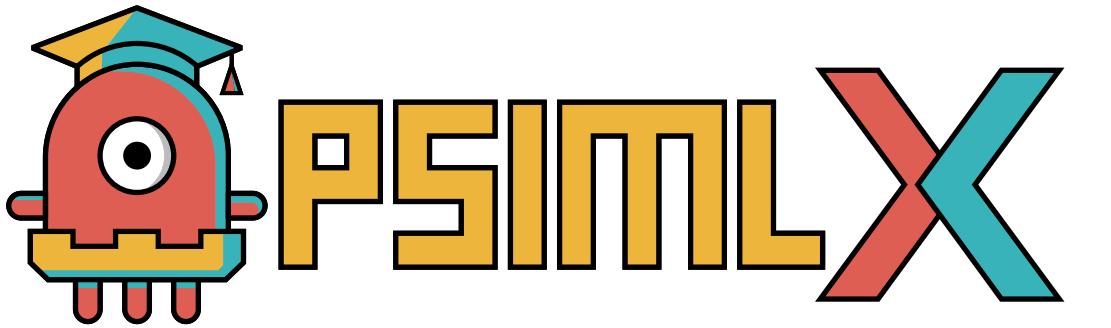
$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100]$$

$$G = 100$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)

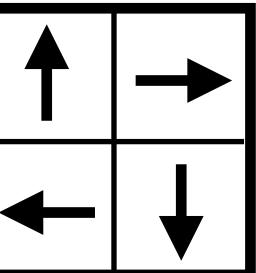
For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

x	0	1	2	3				
y	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	100
	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	

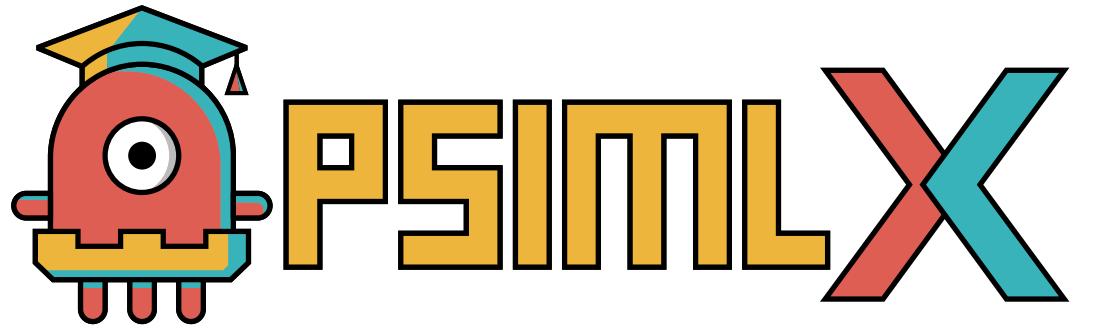
$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \text{ Returns}(s, a) \\ [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100]$$

$$G = 100$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

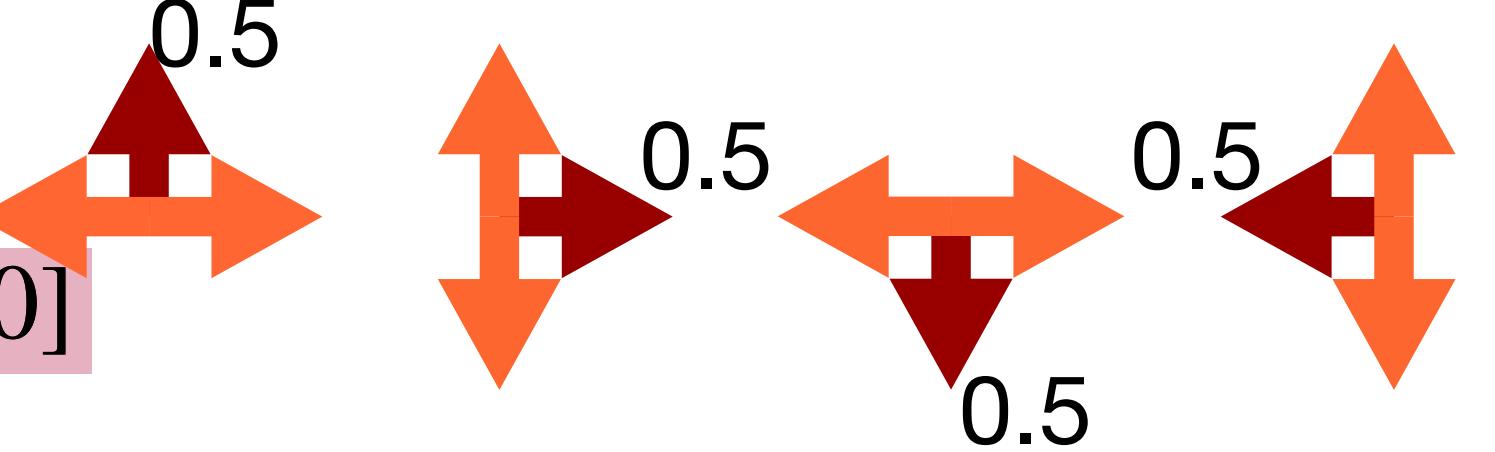
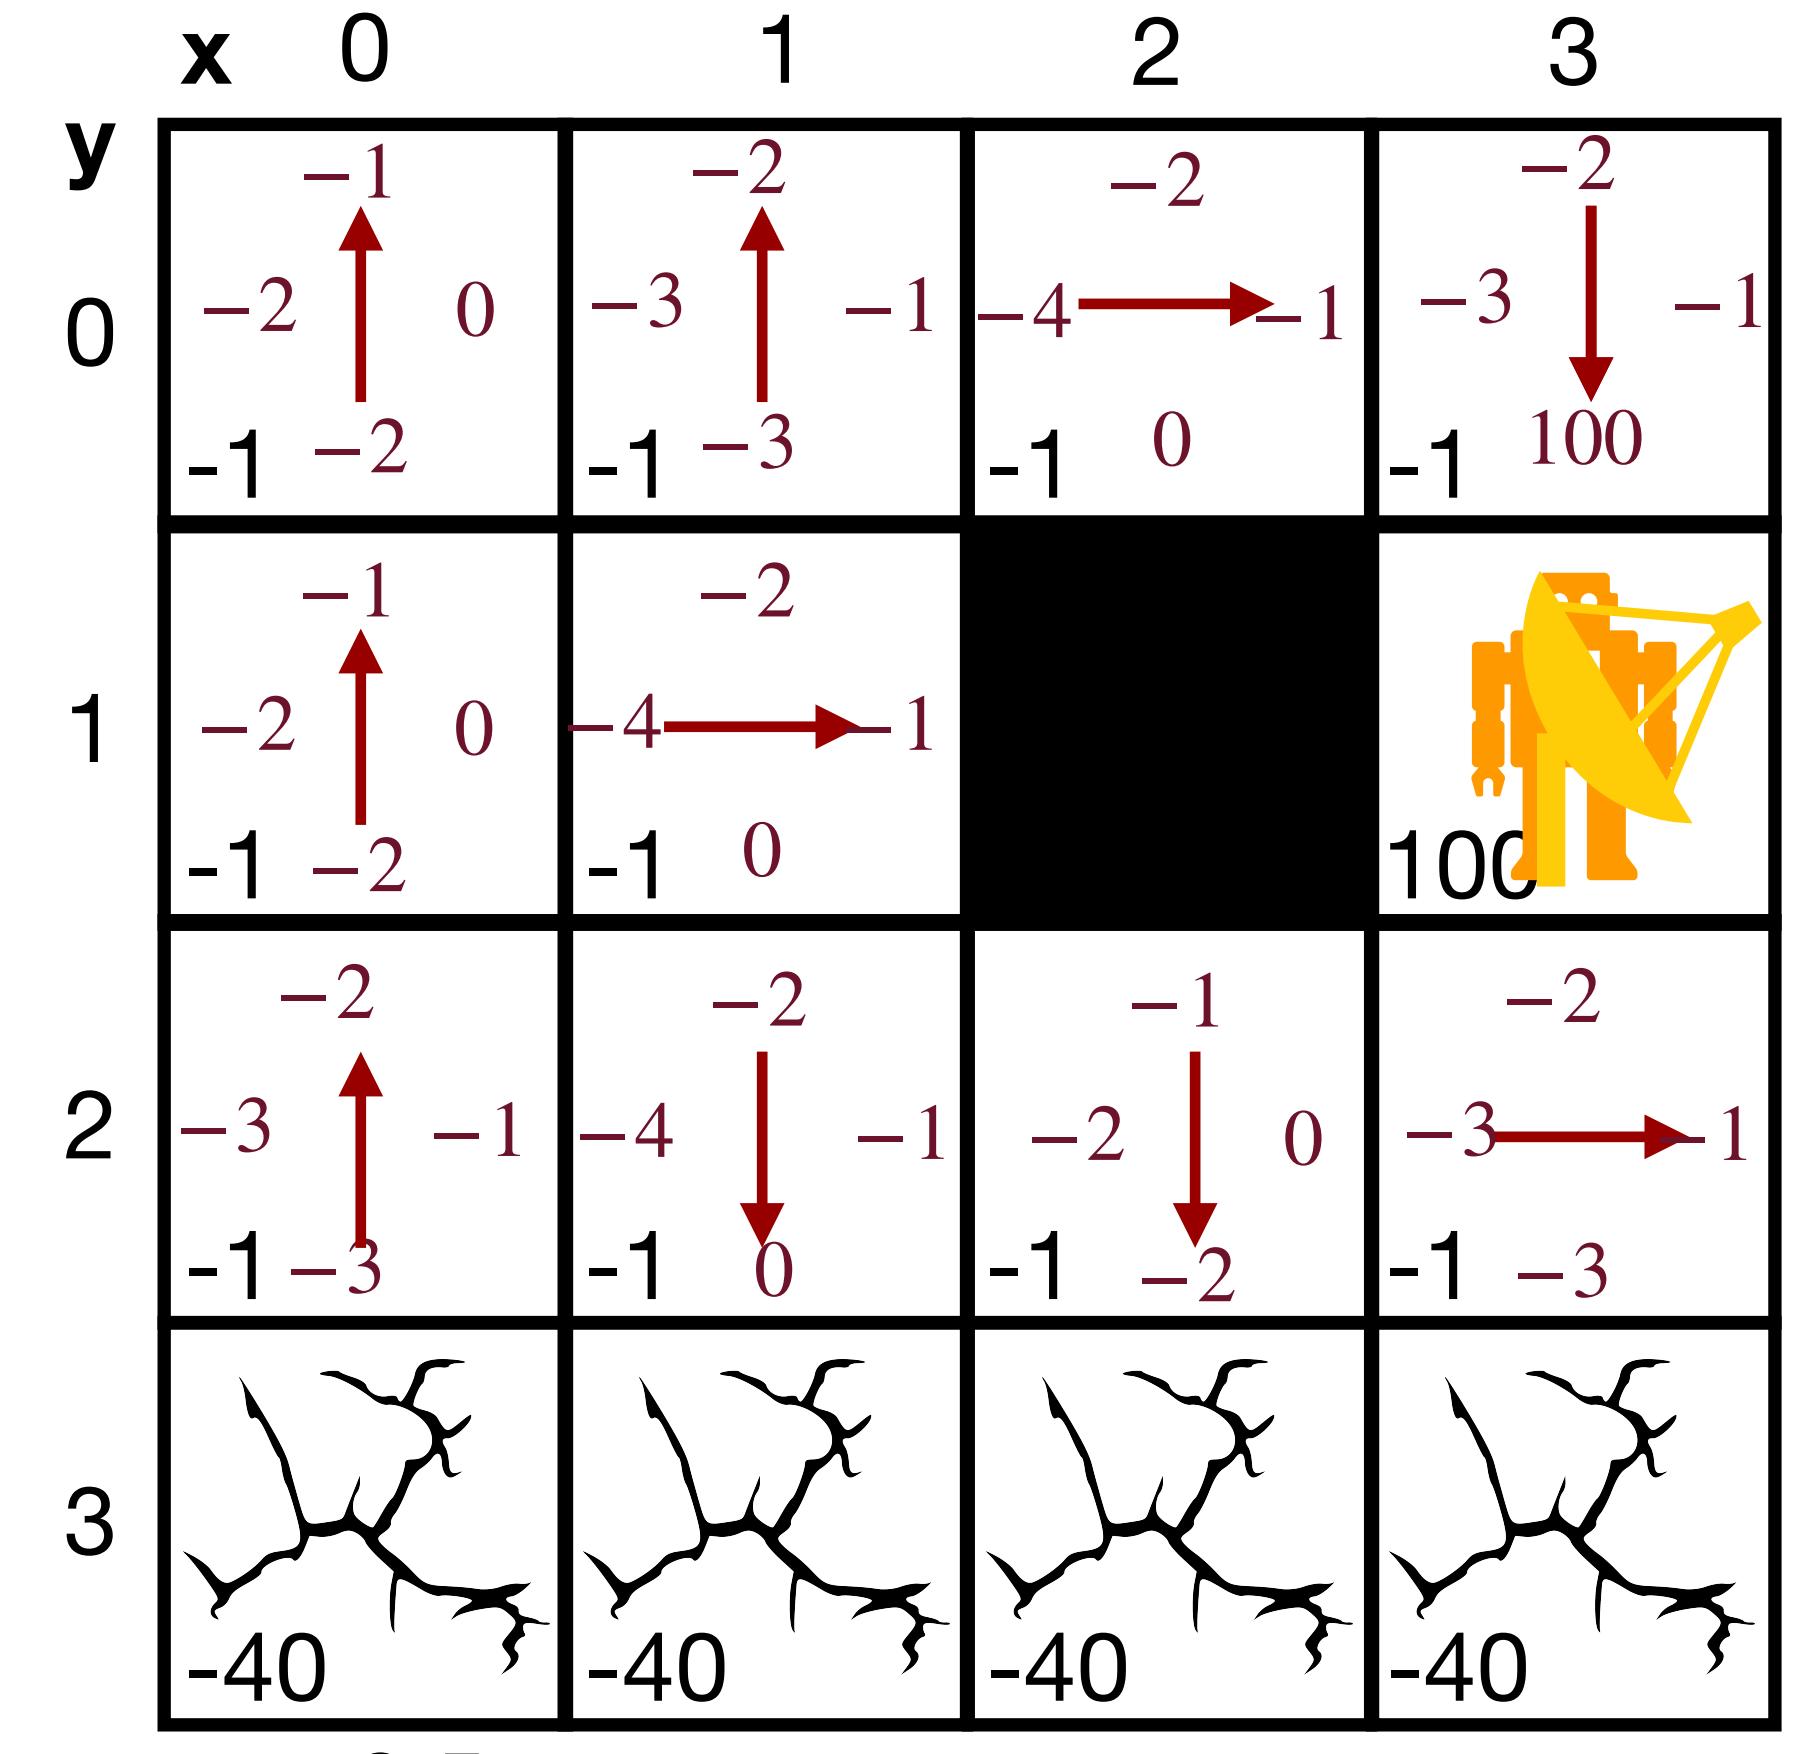
- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1 \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:
 - $\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$

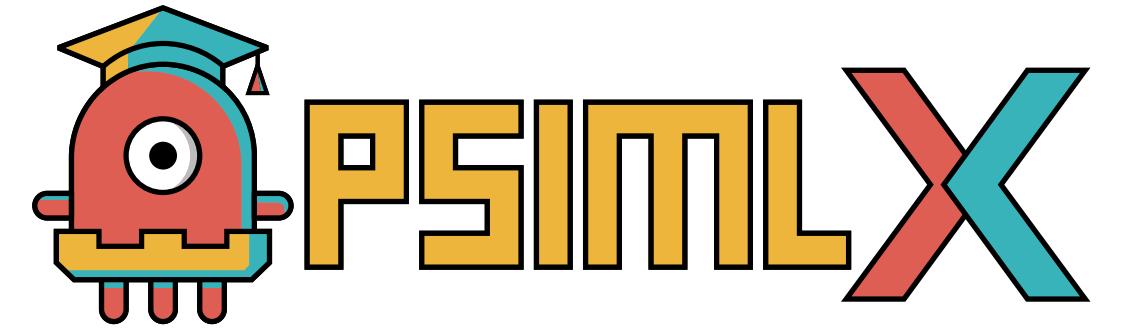
$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100]$$

$$G = 100$$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

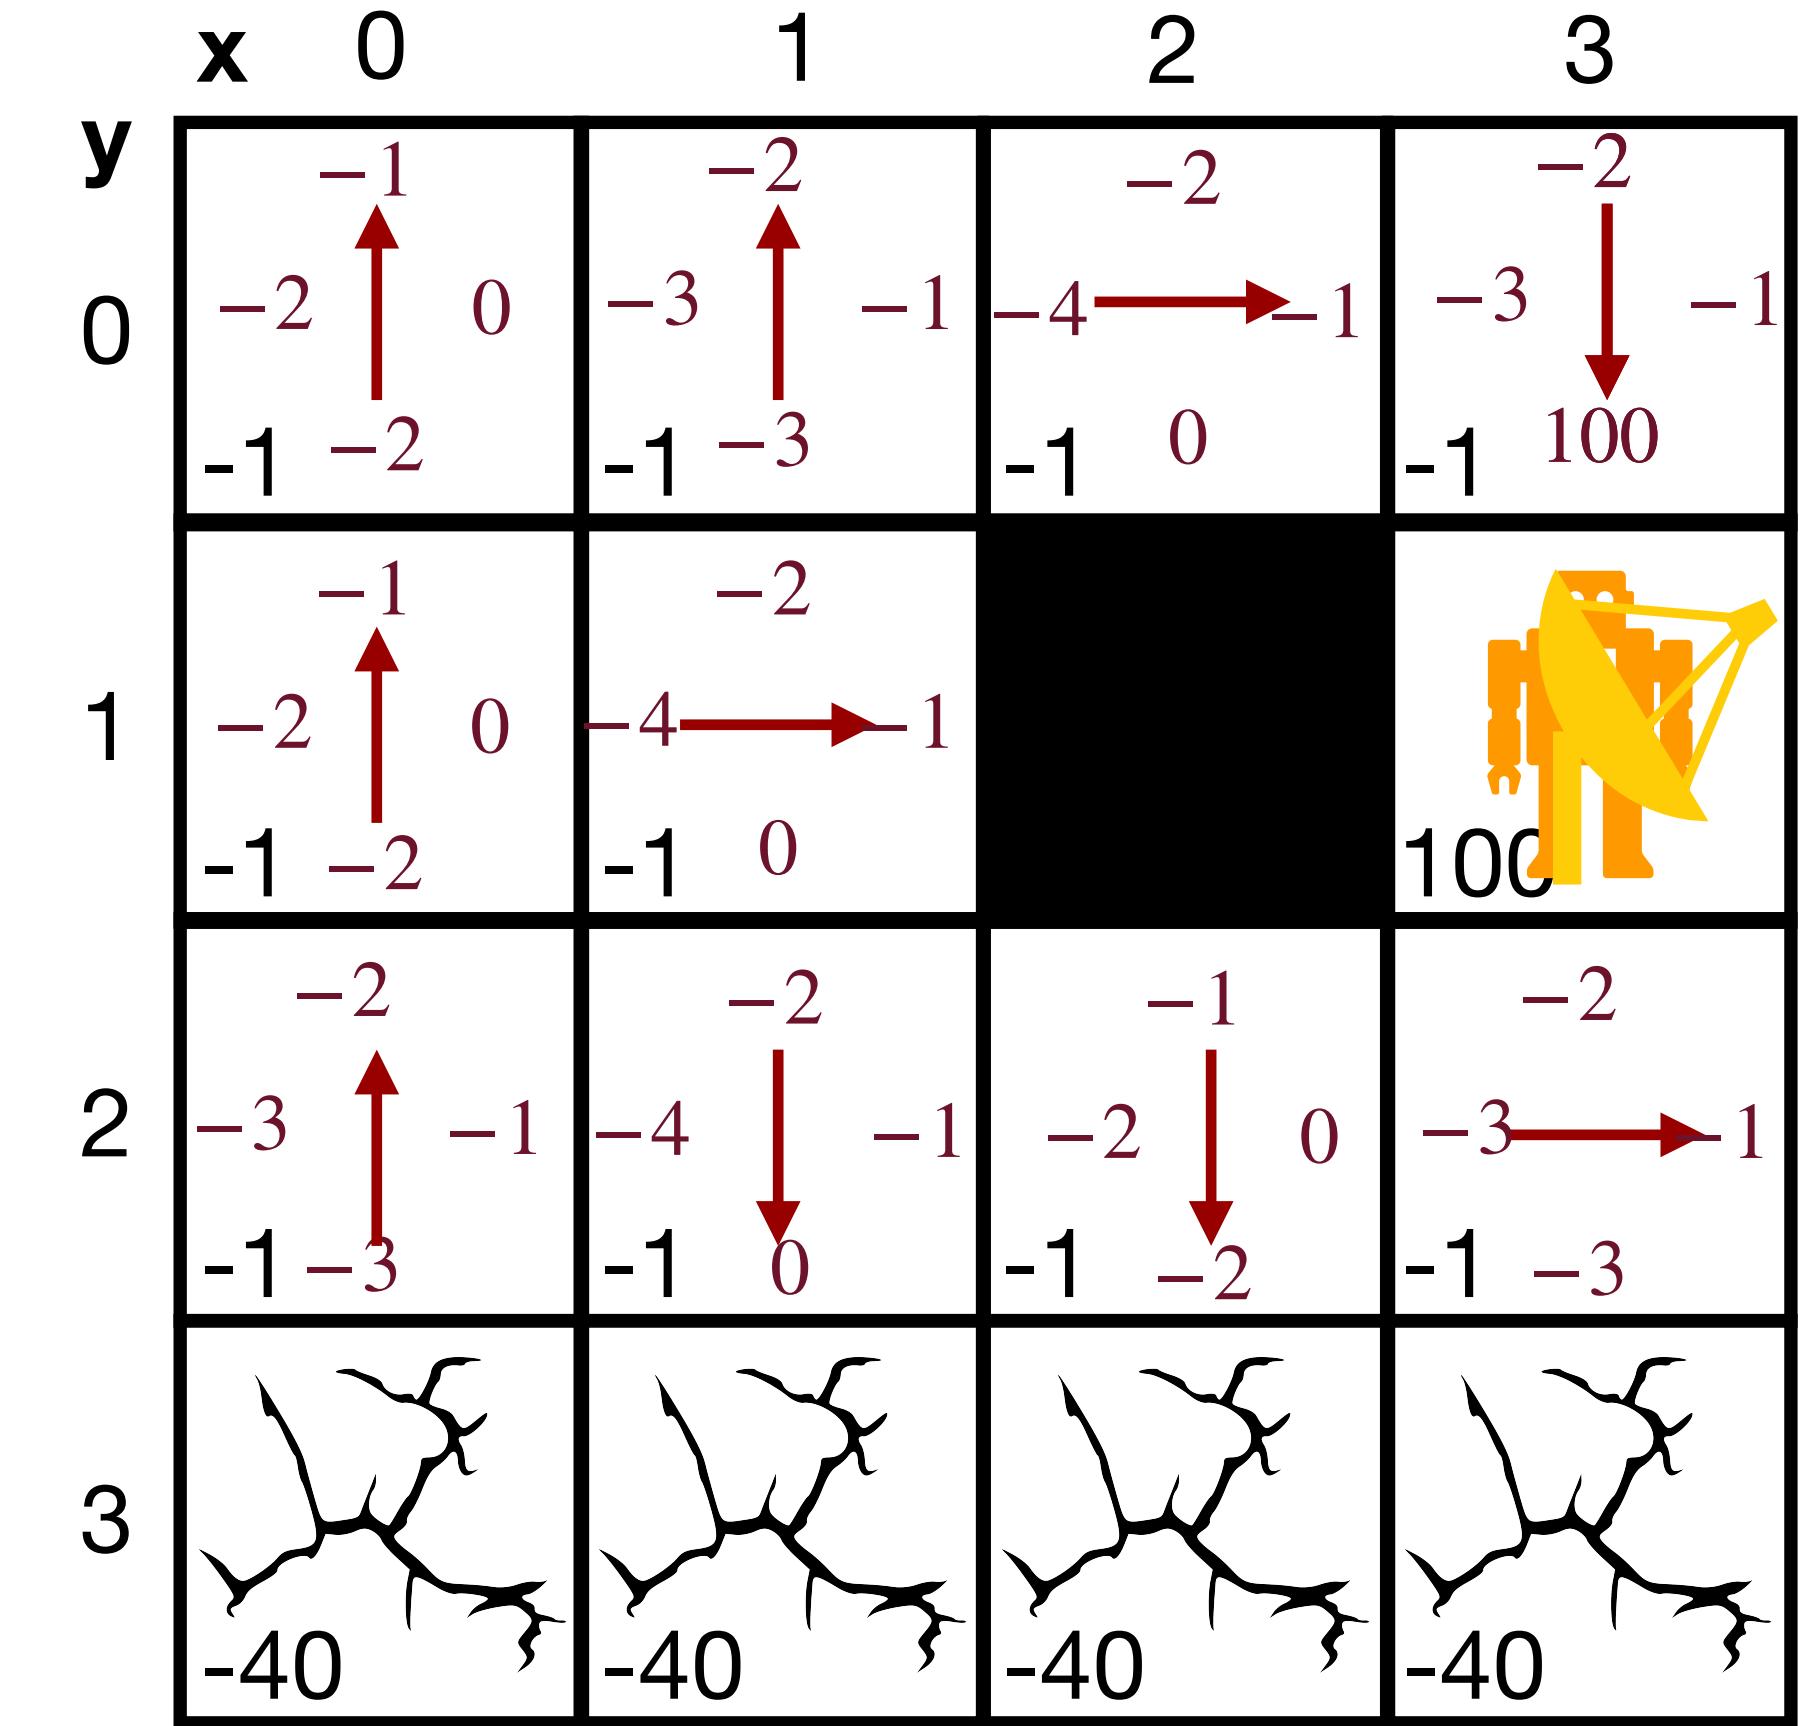
Initialize:

- $\pi \leftarrow$ an arbitrary ε -soft policy
- $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$
- $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

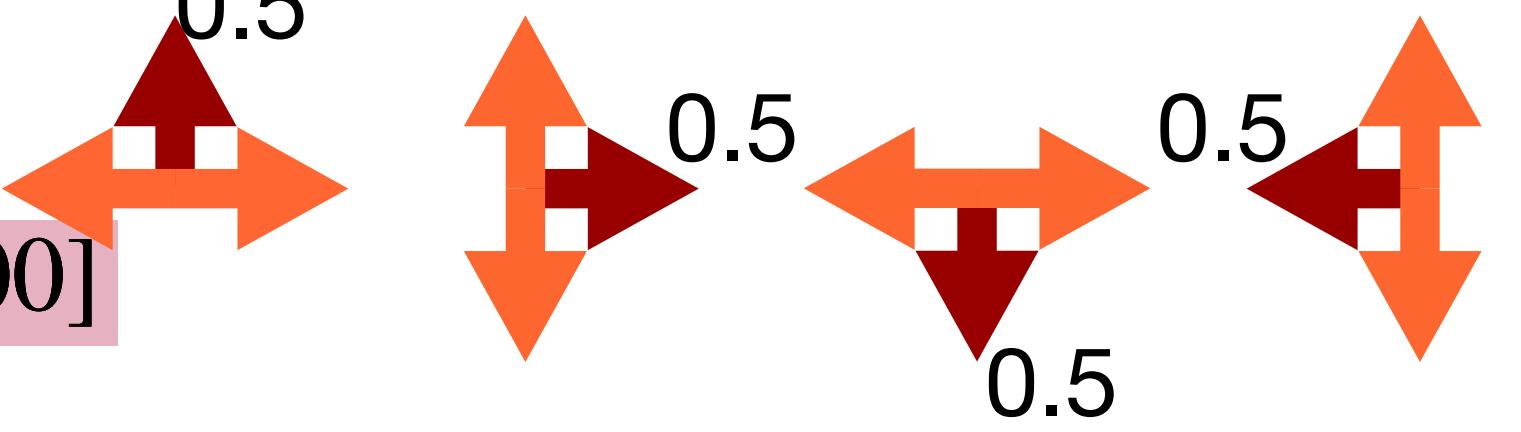
- Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
- $G \leftarrow 0$
- Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 - $G \leftarrow \gamma G + R_{t+1}$
 - Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1 \dots, S_{t-1}, A_{t-1}$:
 - Append G to $Returns(S_t, A_t)$
 - $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 - $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 - For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



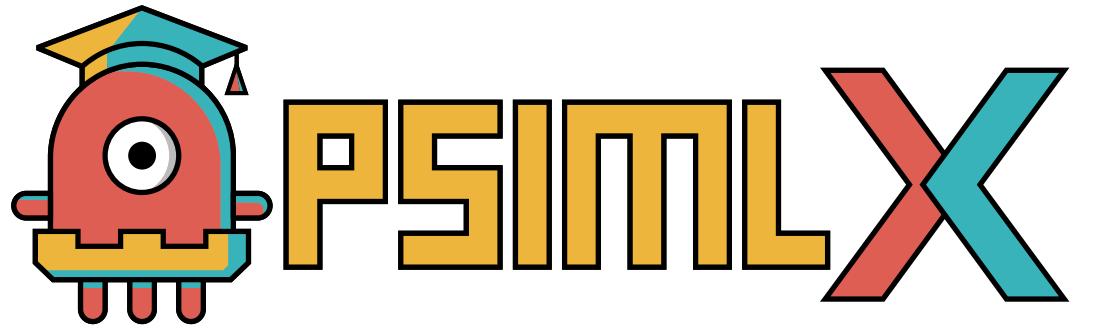
$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100]$$

$G = 100 \quad A^* = \text{DOWN}$



Model free RL

Example 1: On-Policy MC Algorithm



On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

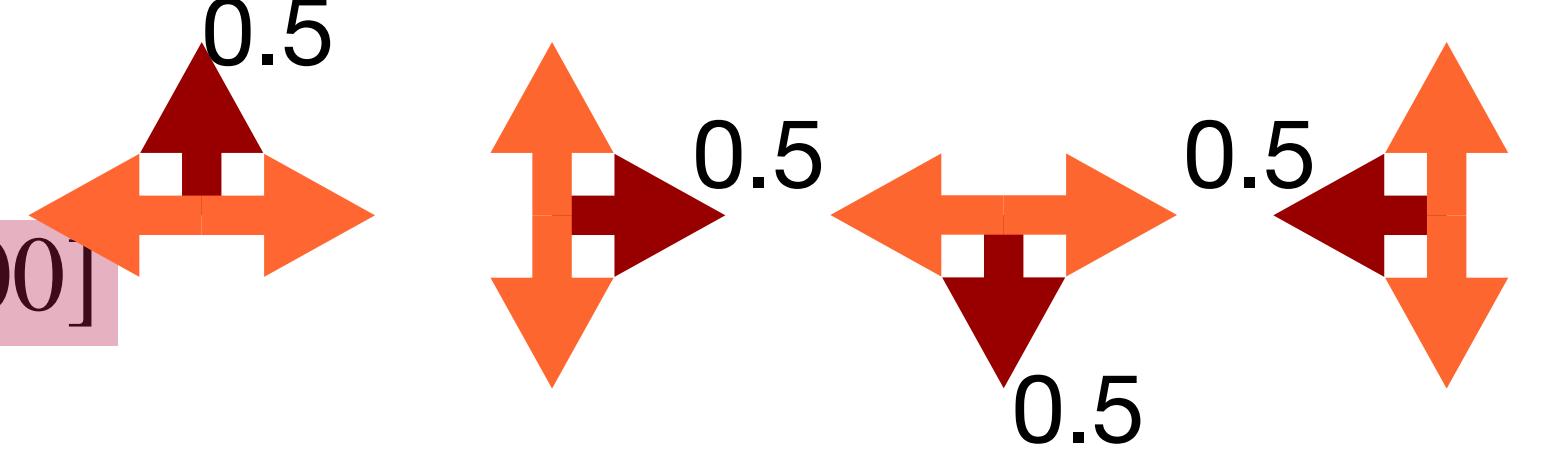
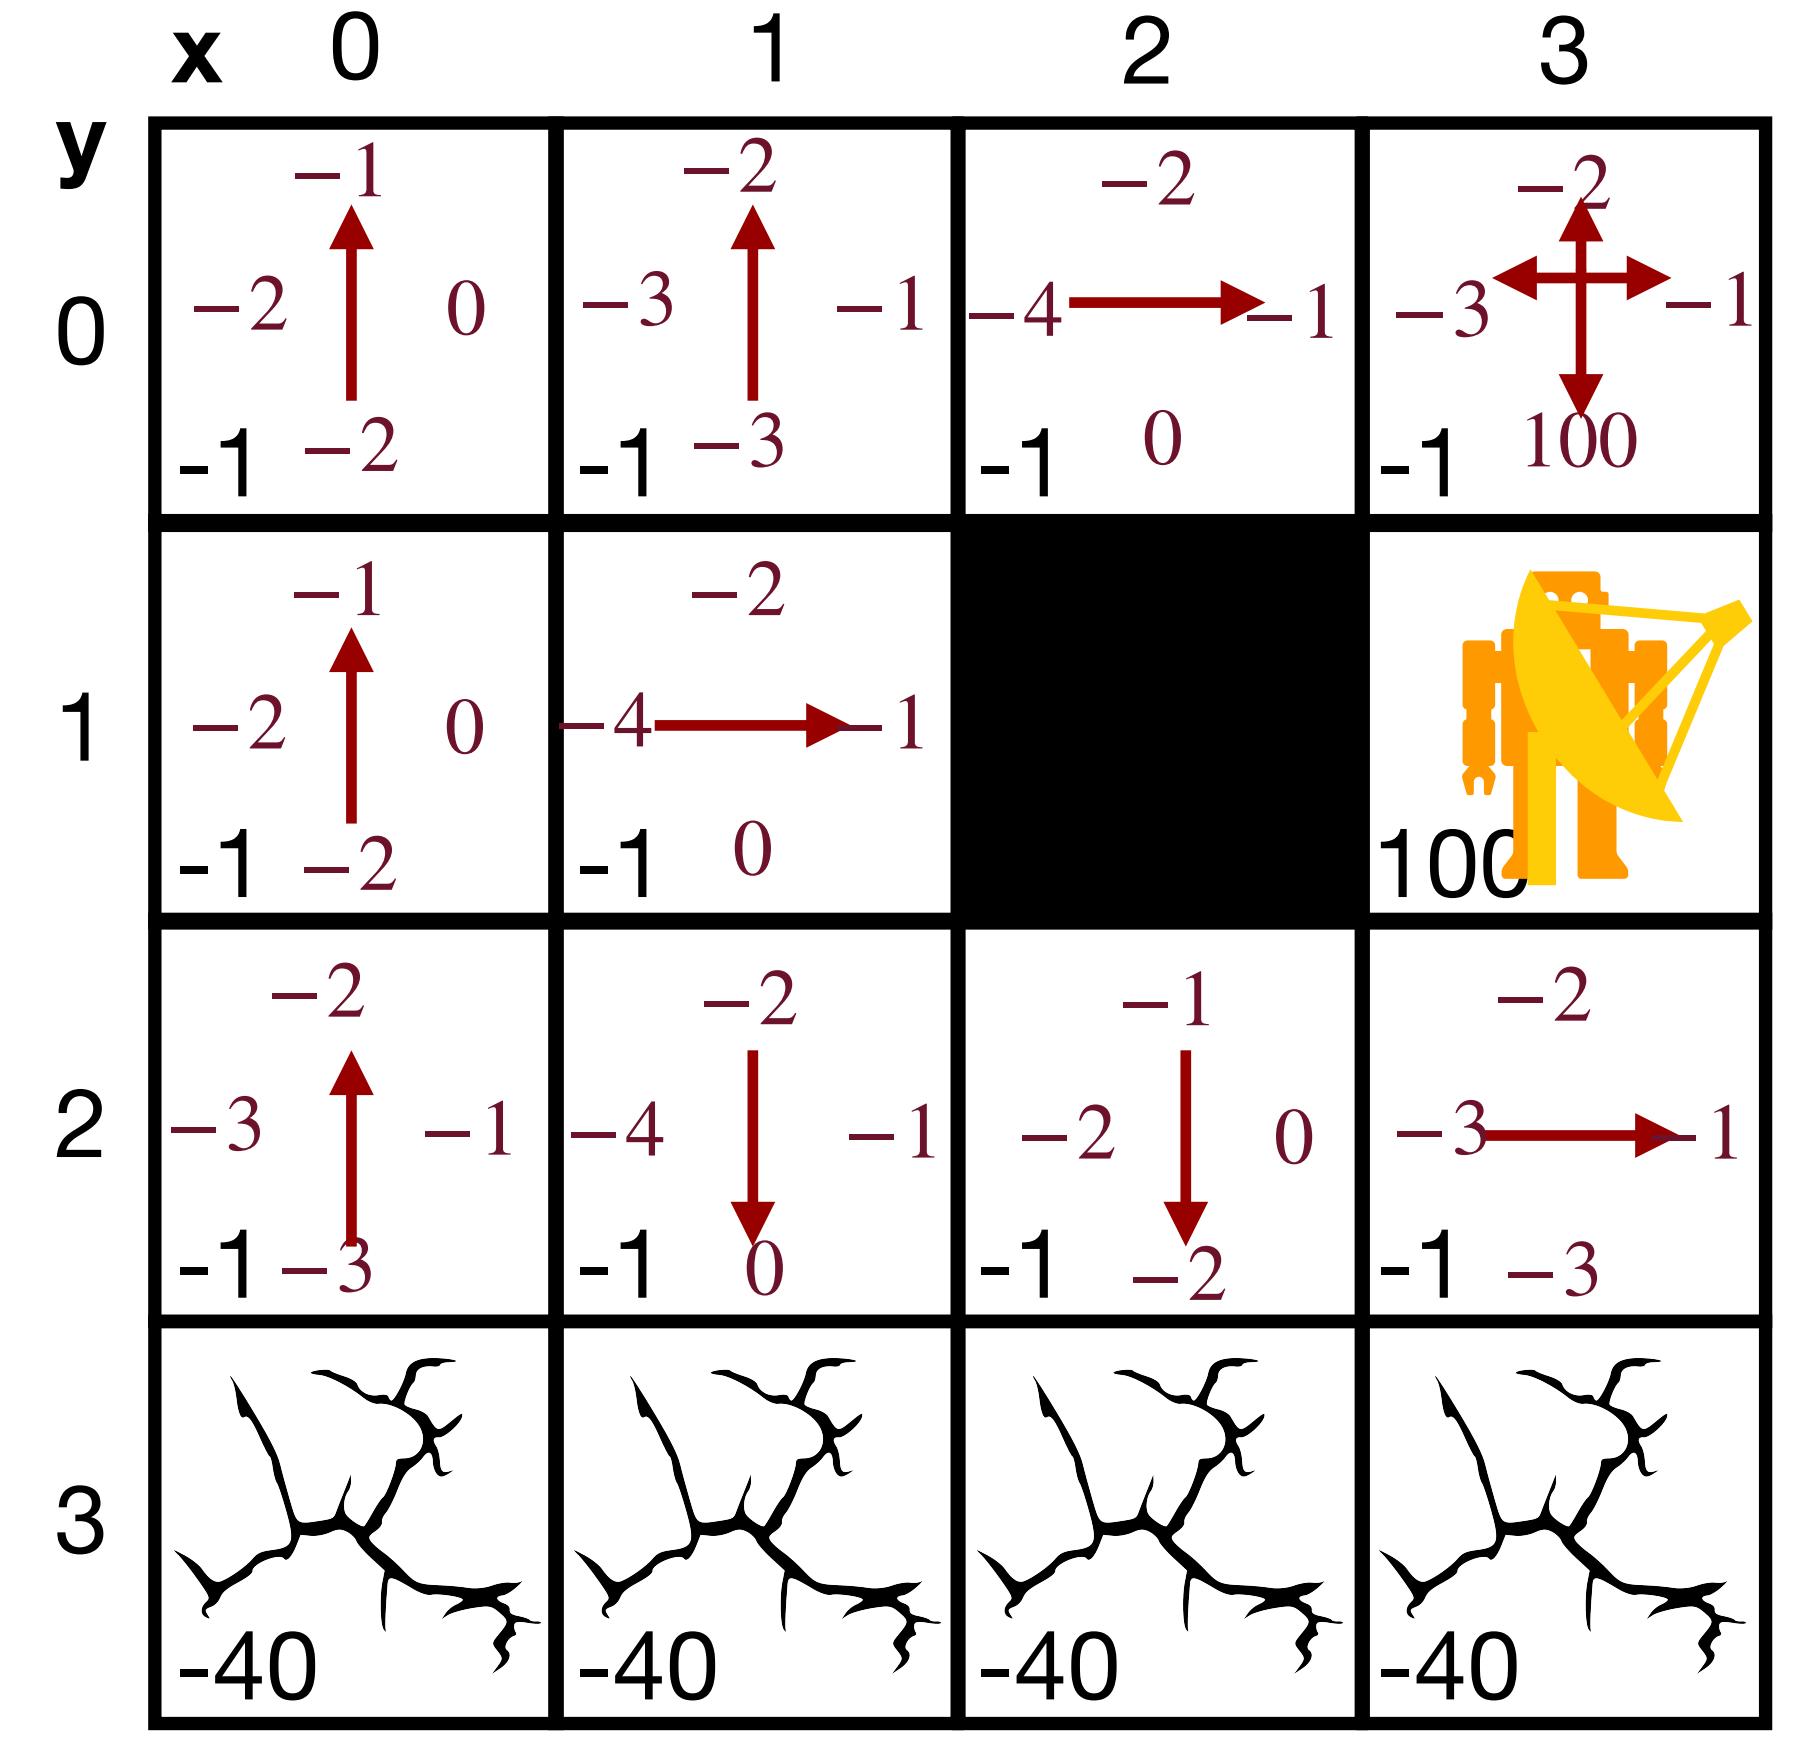
Algorithm parameter: small $\varepsilon > 0$
 Initialize:
 $\pi \leftarrow$ an arbitrary ε -soft policy
 $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
 $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):
 Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
 $G \leftarrow 0$
 Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 $G \leftarrow \gamma G + R_{t+1}$
 Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:
 Append G to $Returns(S_t, A_t)$
 $Q(S_t, A_t) \leftarrow$ average($Returns(S_t, A_t)$)
 $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

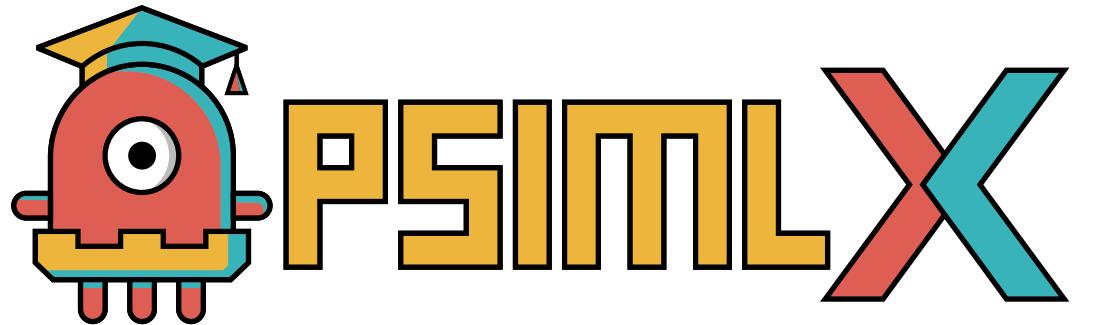
$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100]$$

$$G = 100 \quad A^* = \text{DOWN}$$



Model free RL

Example 1: On-Policy MC Algorithm

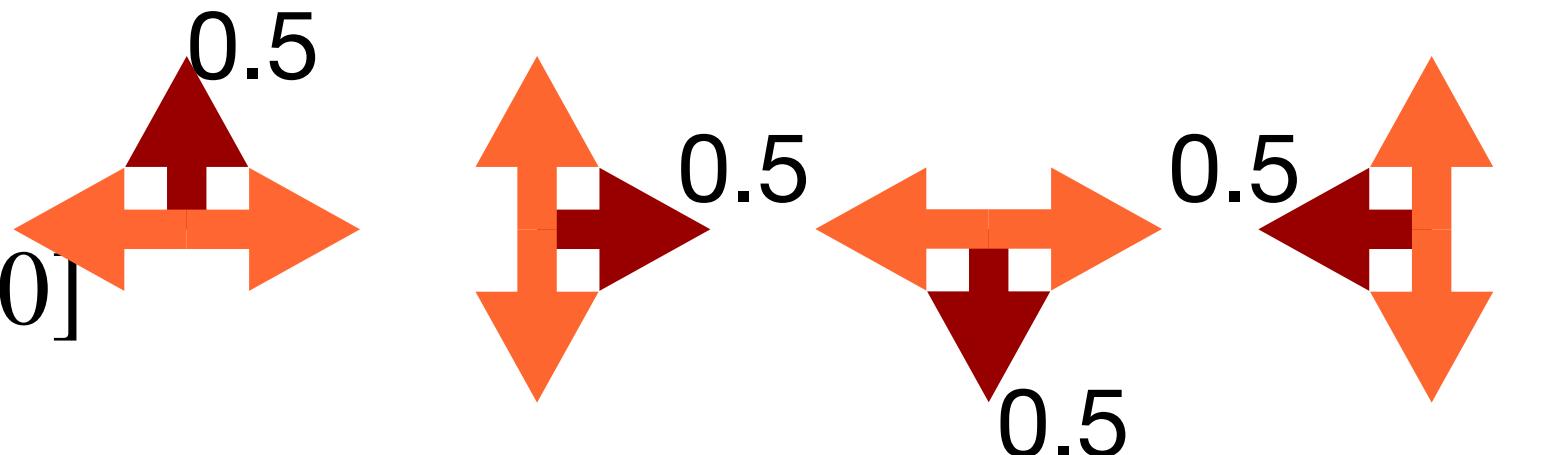
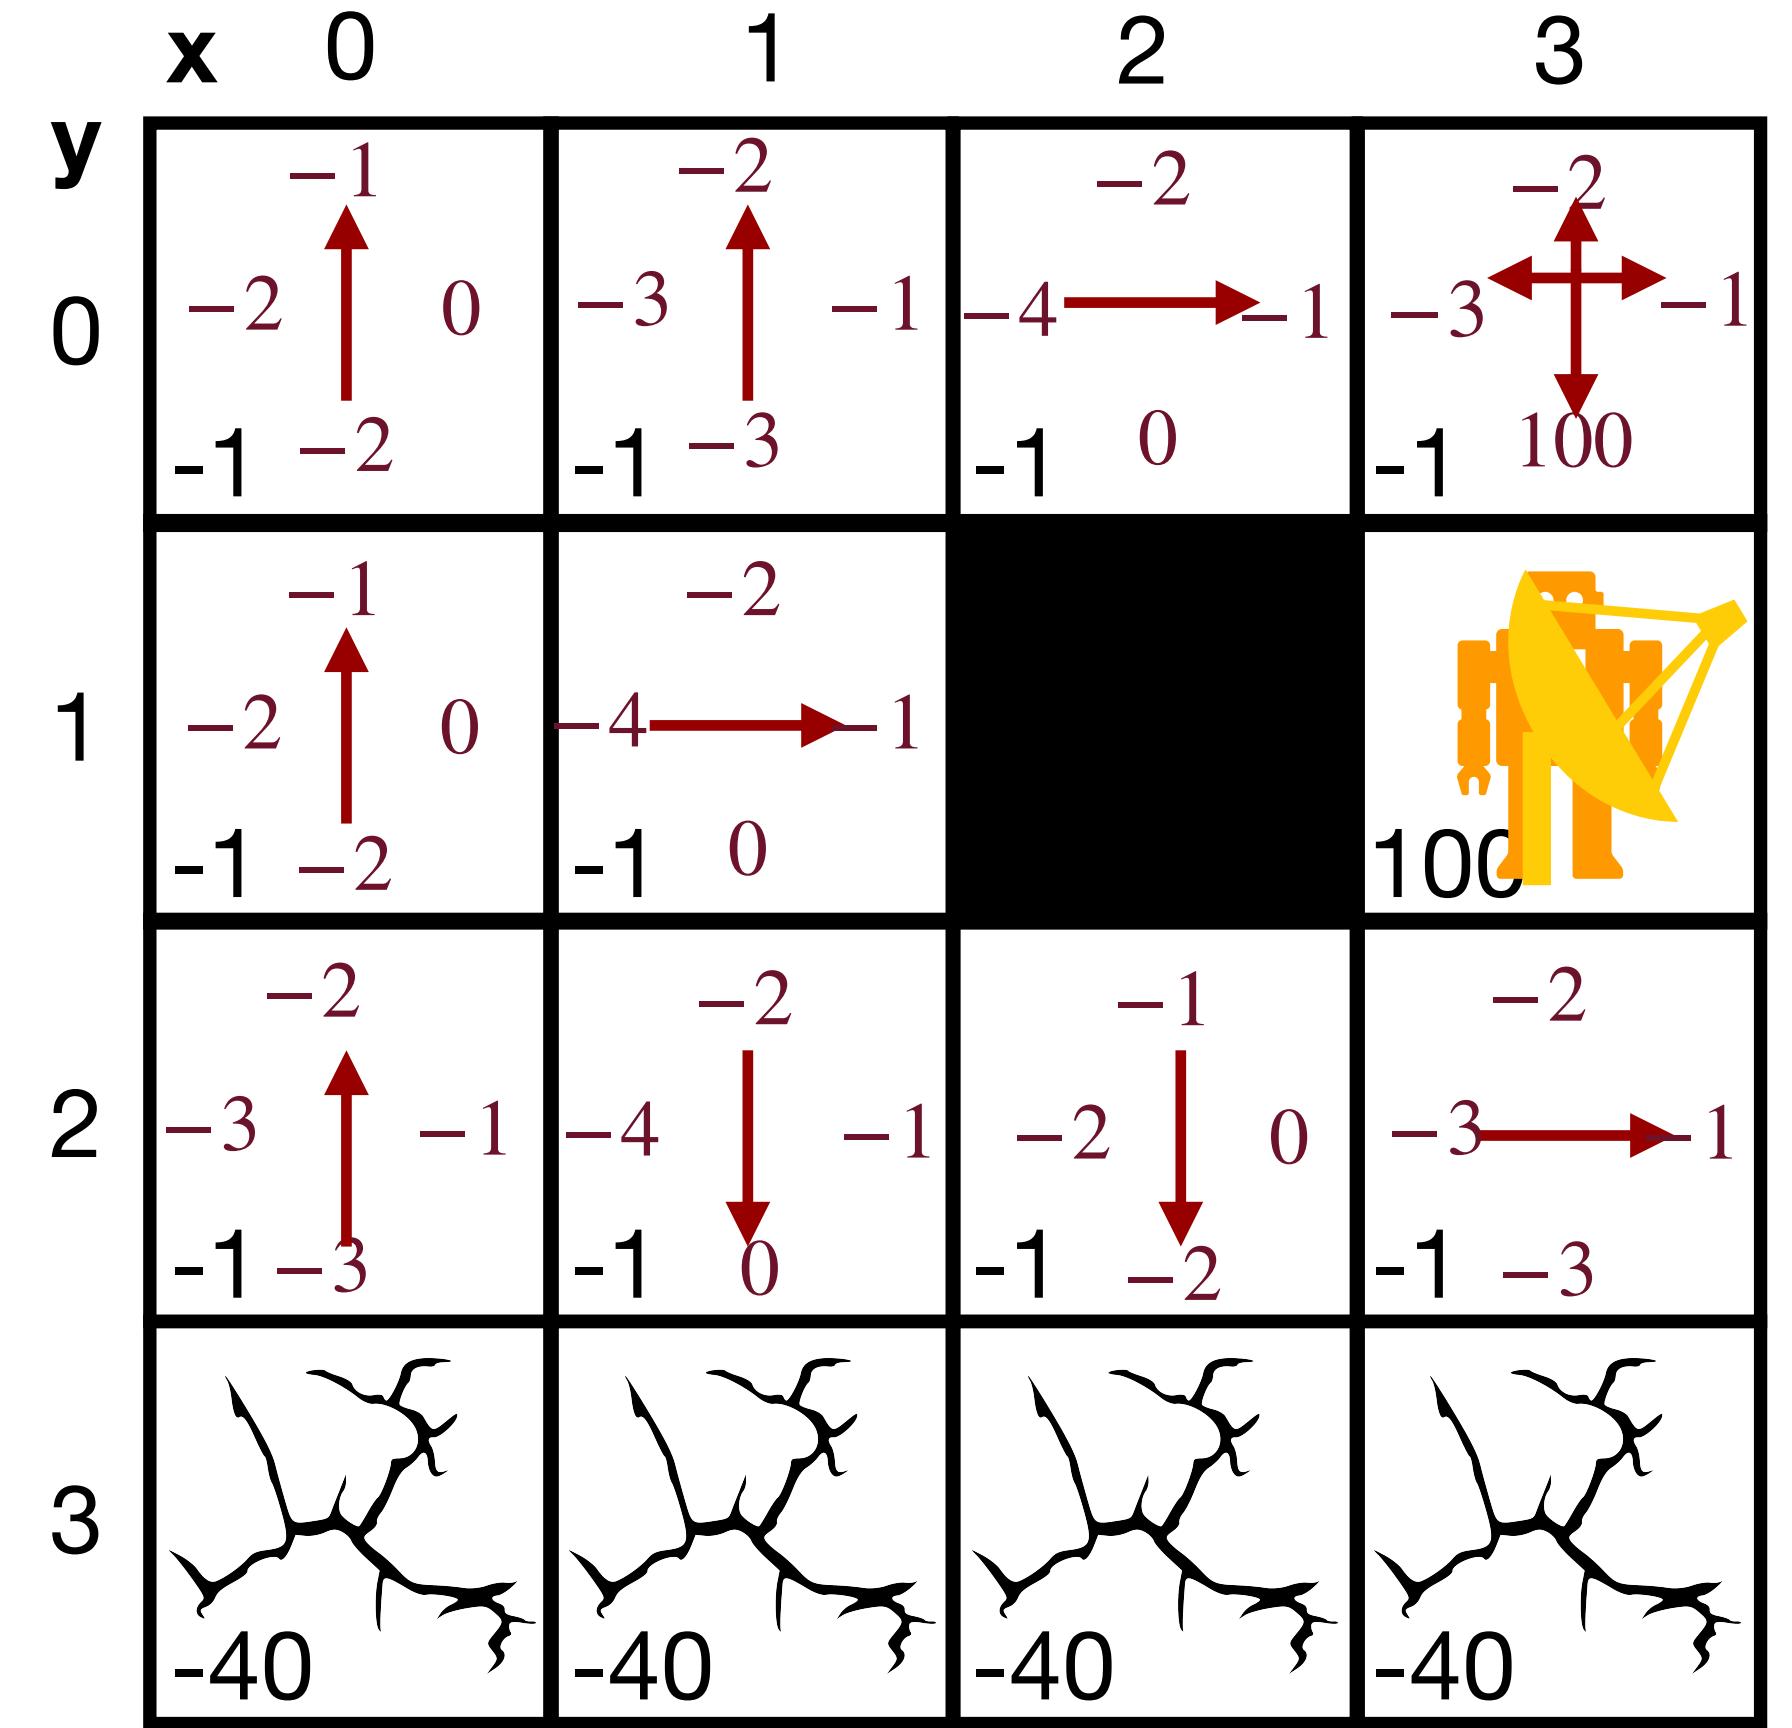


On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$
 Initialize:
 $\pi \leftarrow$ an arbitrary ε -soft policy
 $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
 $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

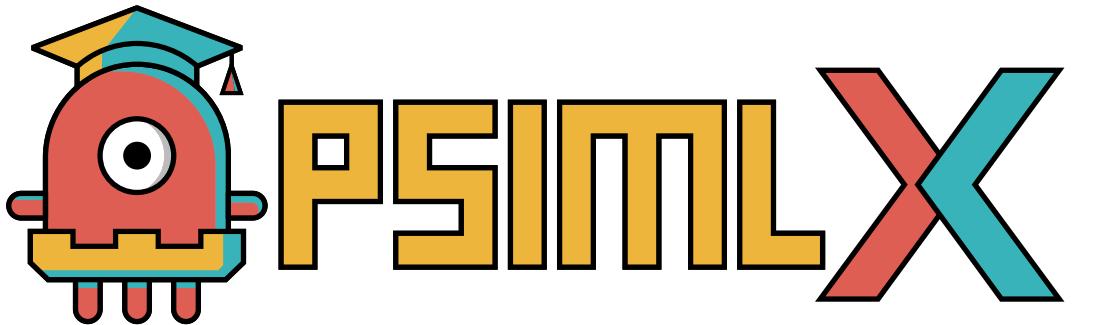
Repeat forever (for each episode):
 Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
 $G \leftarrow 0$
 Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 $G \leftarrow \gamma G + R_{t+1}$
 Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1 \dots, S_{t-1}, A_{t-1}$:
 Append G to $Returns(S_t, A_t)$
 $Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$
 $A^* \leftarrow \arg \max_a Q(S_t, a)$ (with ties broken arbitrarily)
 For all $a \in \mathcal{A}(S_t)$:
 $\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$

$$\tau_\pi = [(0,2), \text{UP}, -1], [(0,1), \text{UP}, -1], [(0,0), \text{UP}, -1], [(0,0), \text{UP}, -1] \\ [(0,0), \text{UP}, -1], [(1,0), \text{UP}, -1], [(2,0), \text{RIGHT}, -1], [(3,0), \text{DOWN}, 100]$$



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

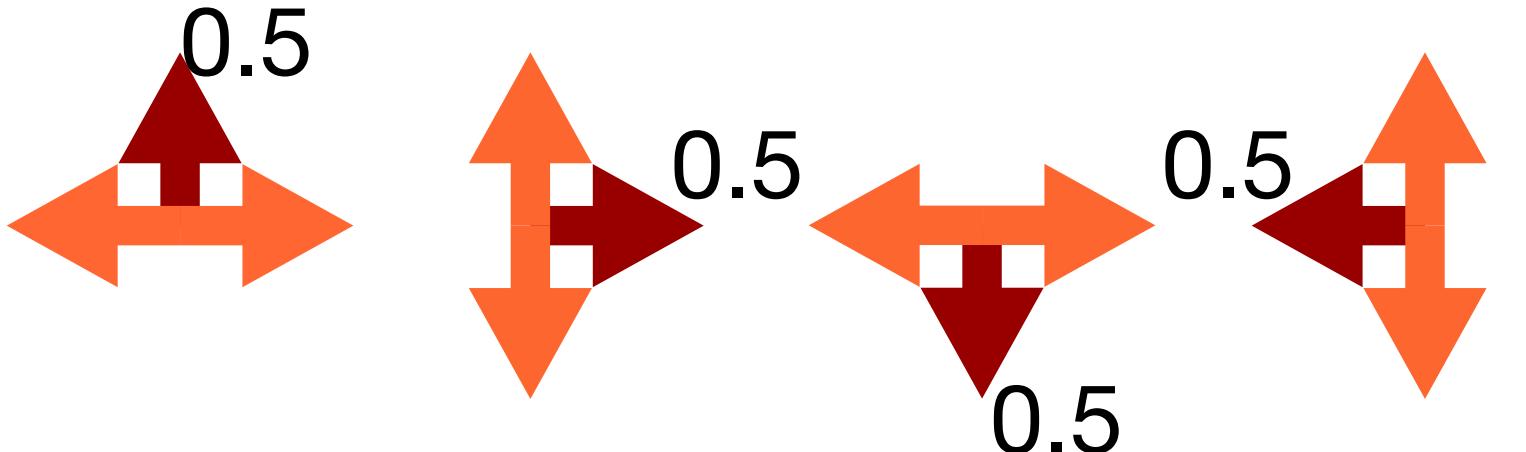
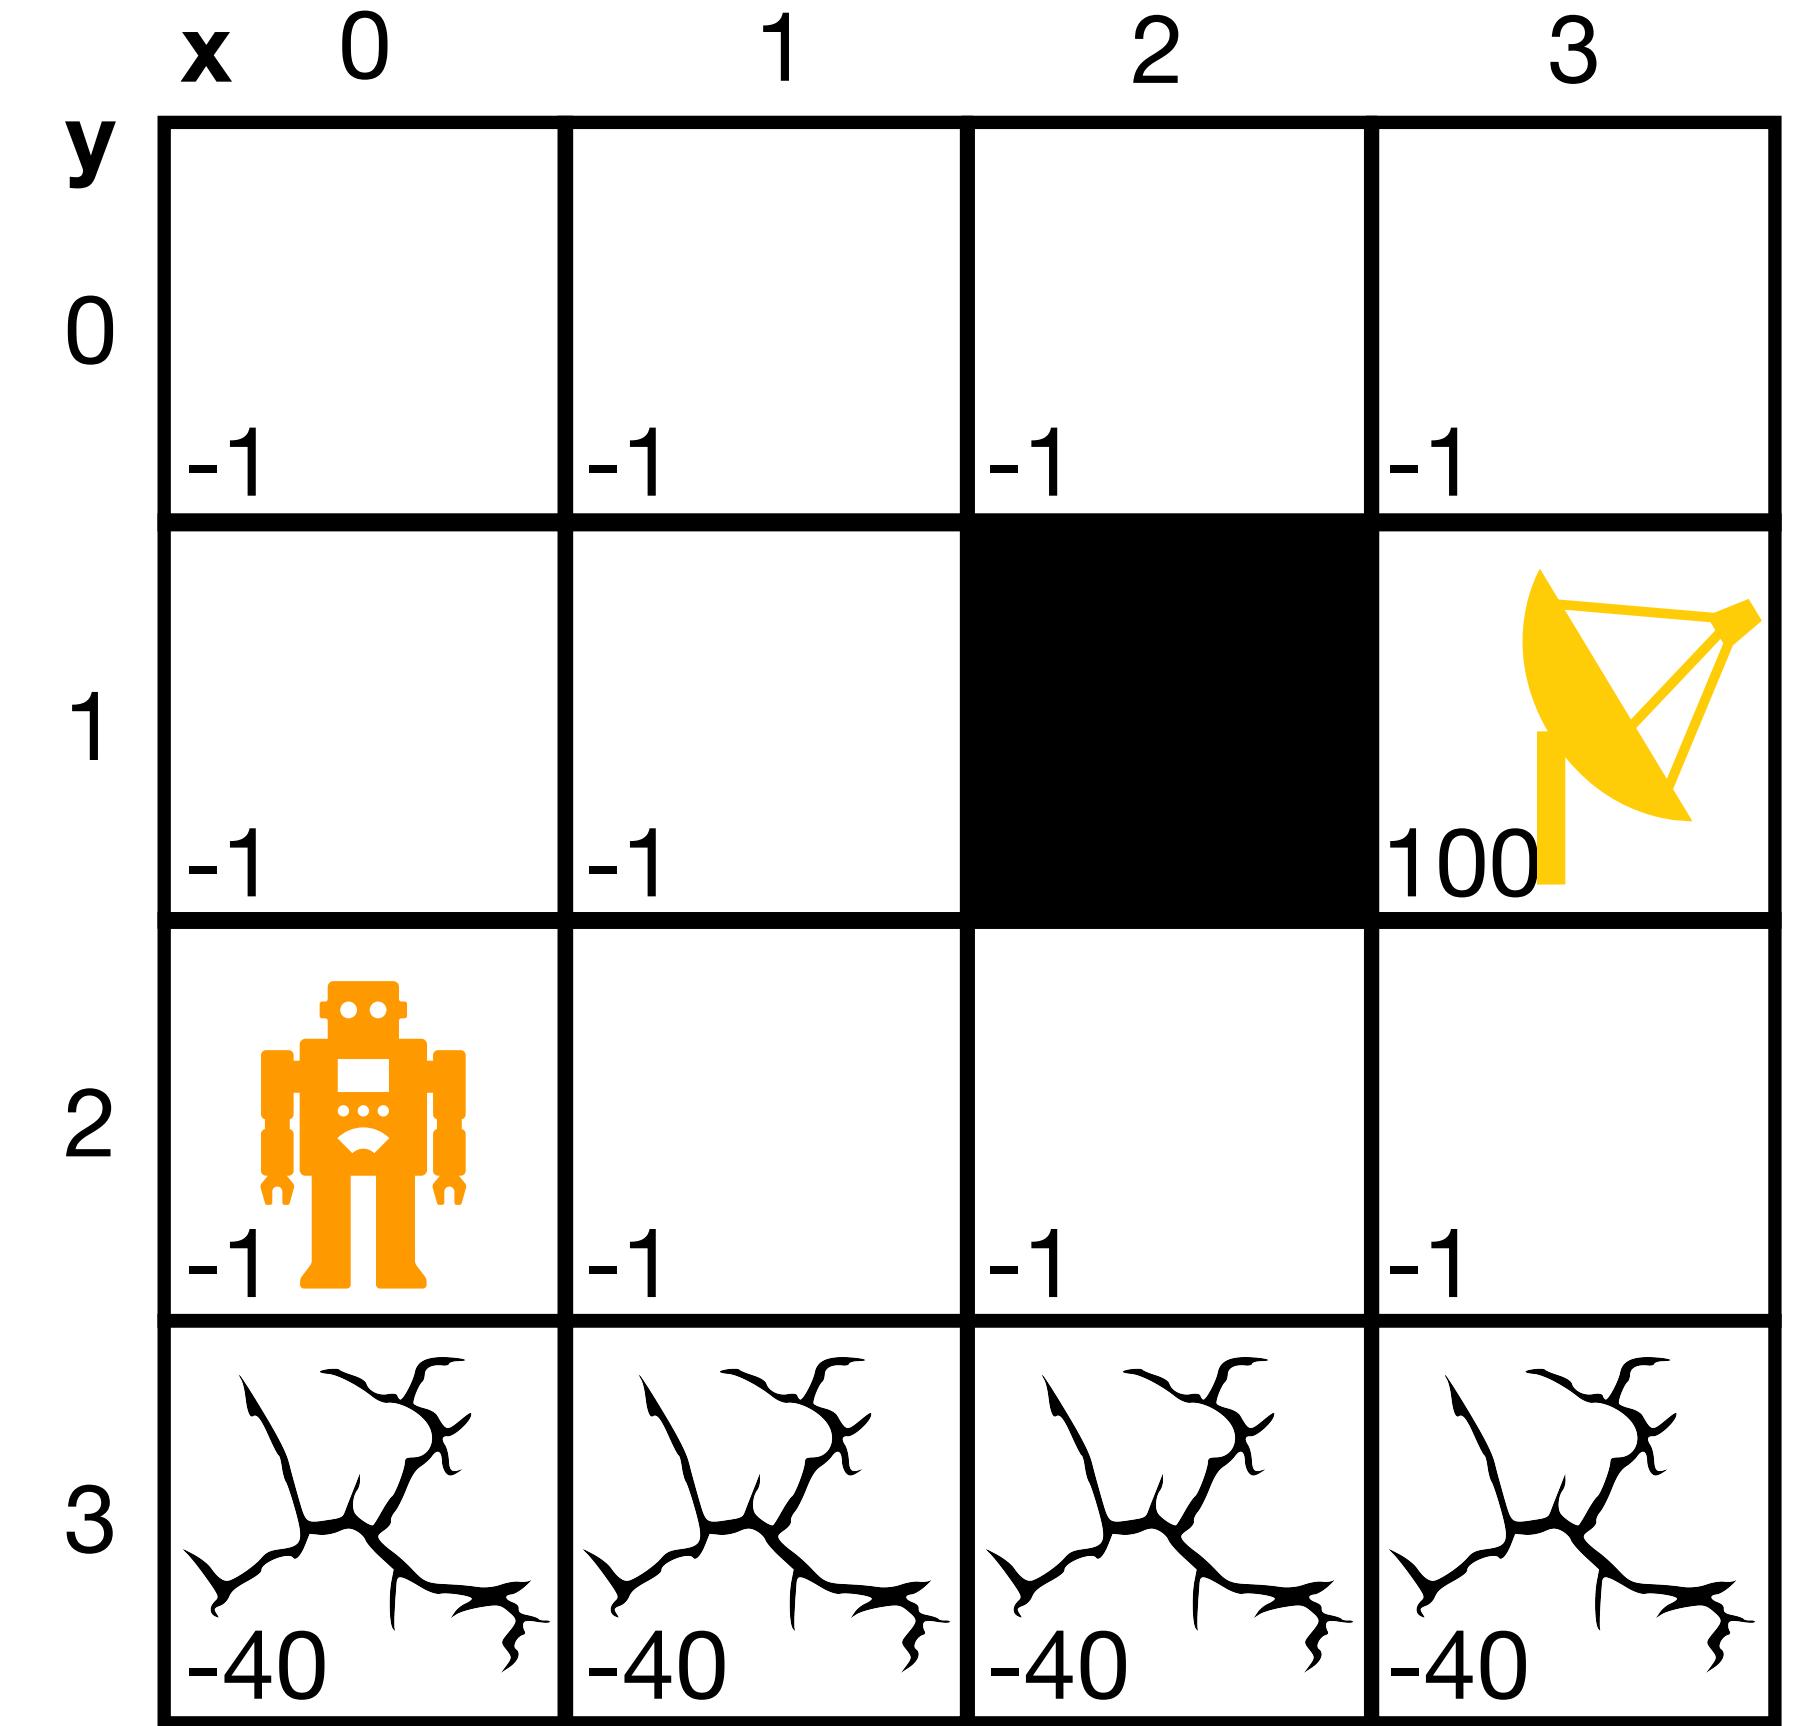
 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

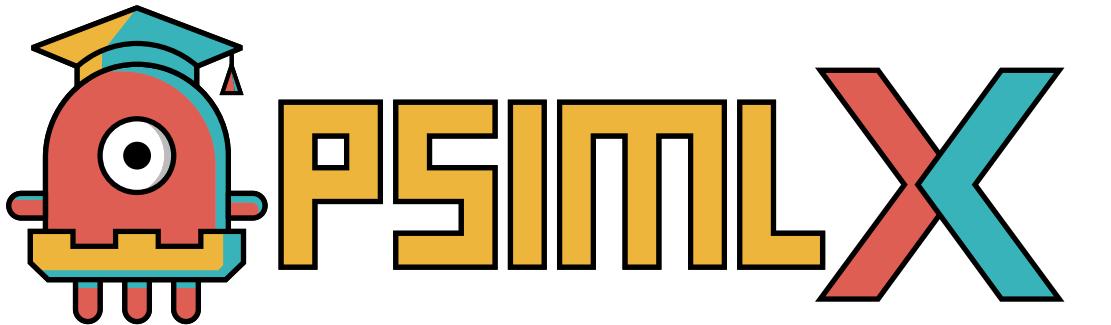
$S \leftarrow S'$

 until S is terminal



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

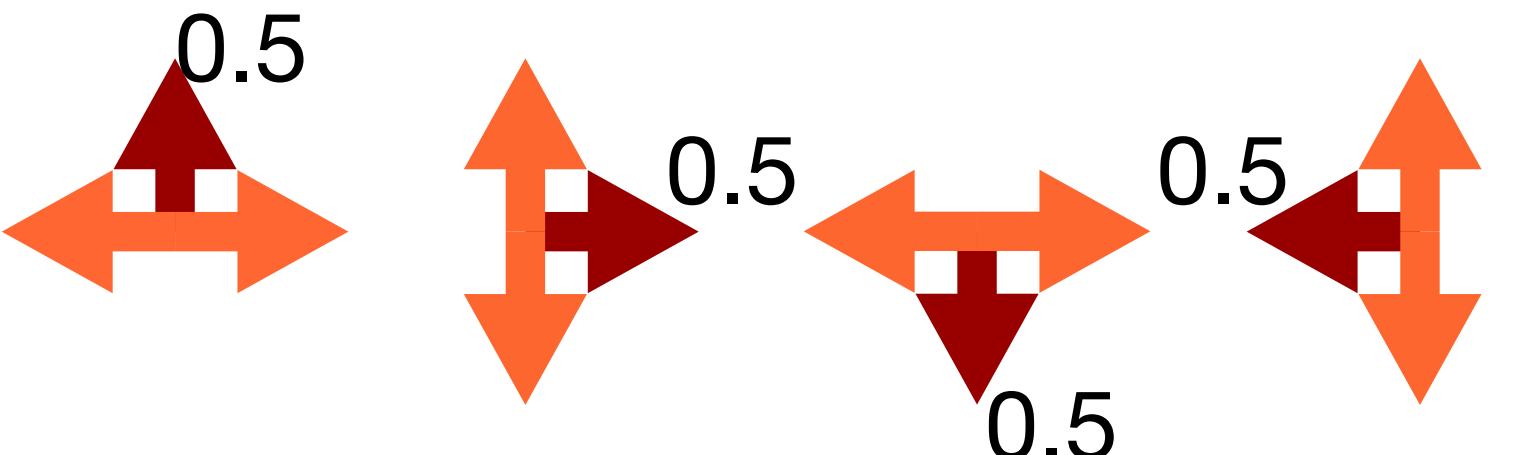
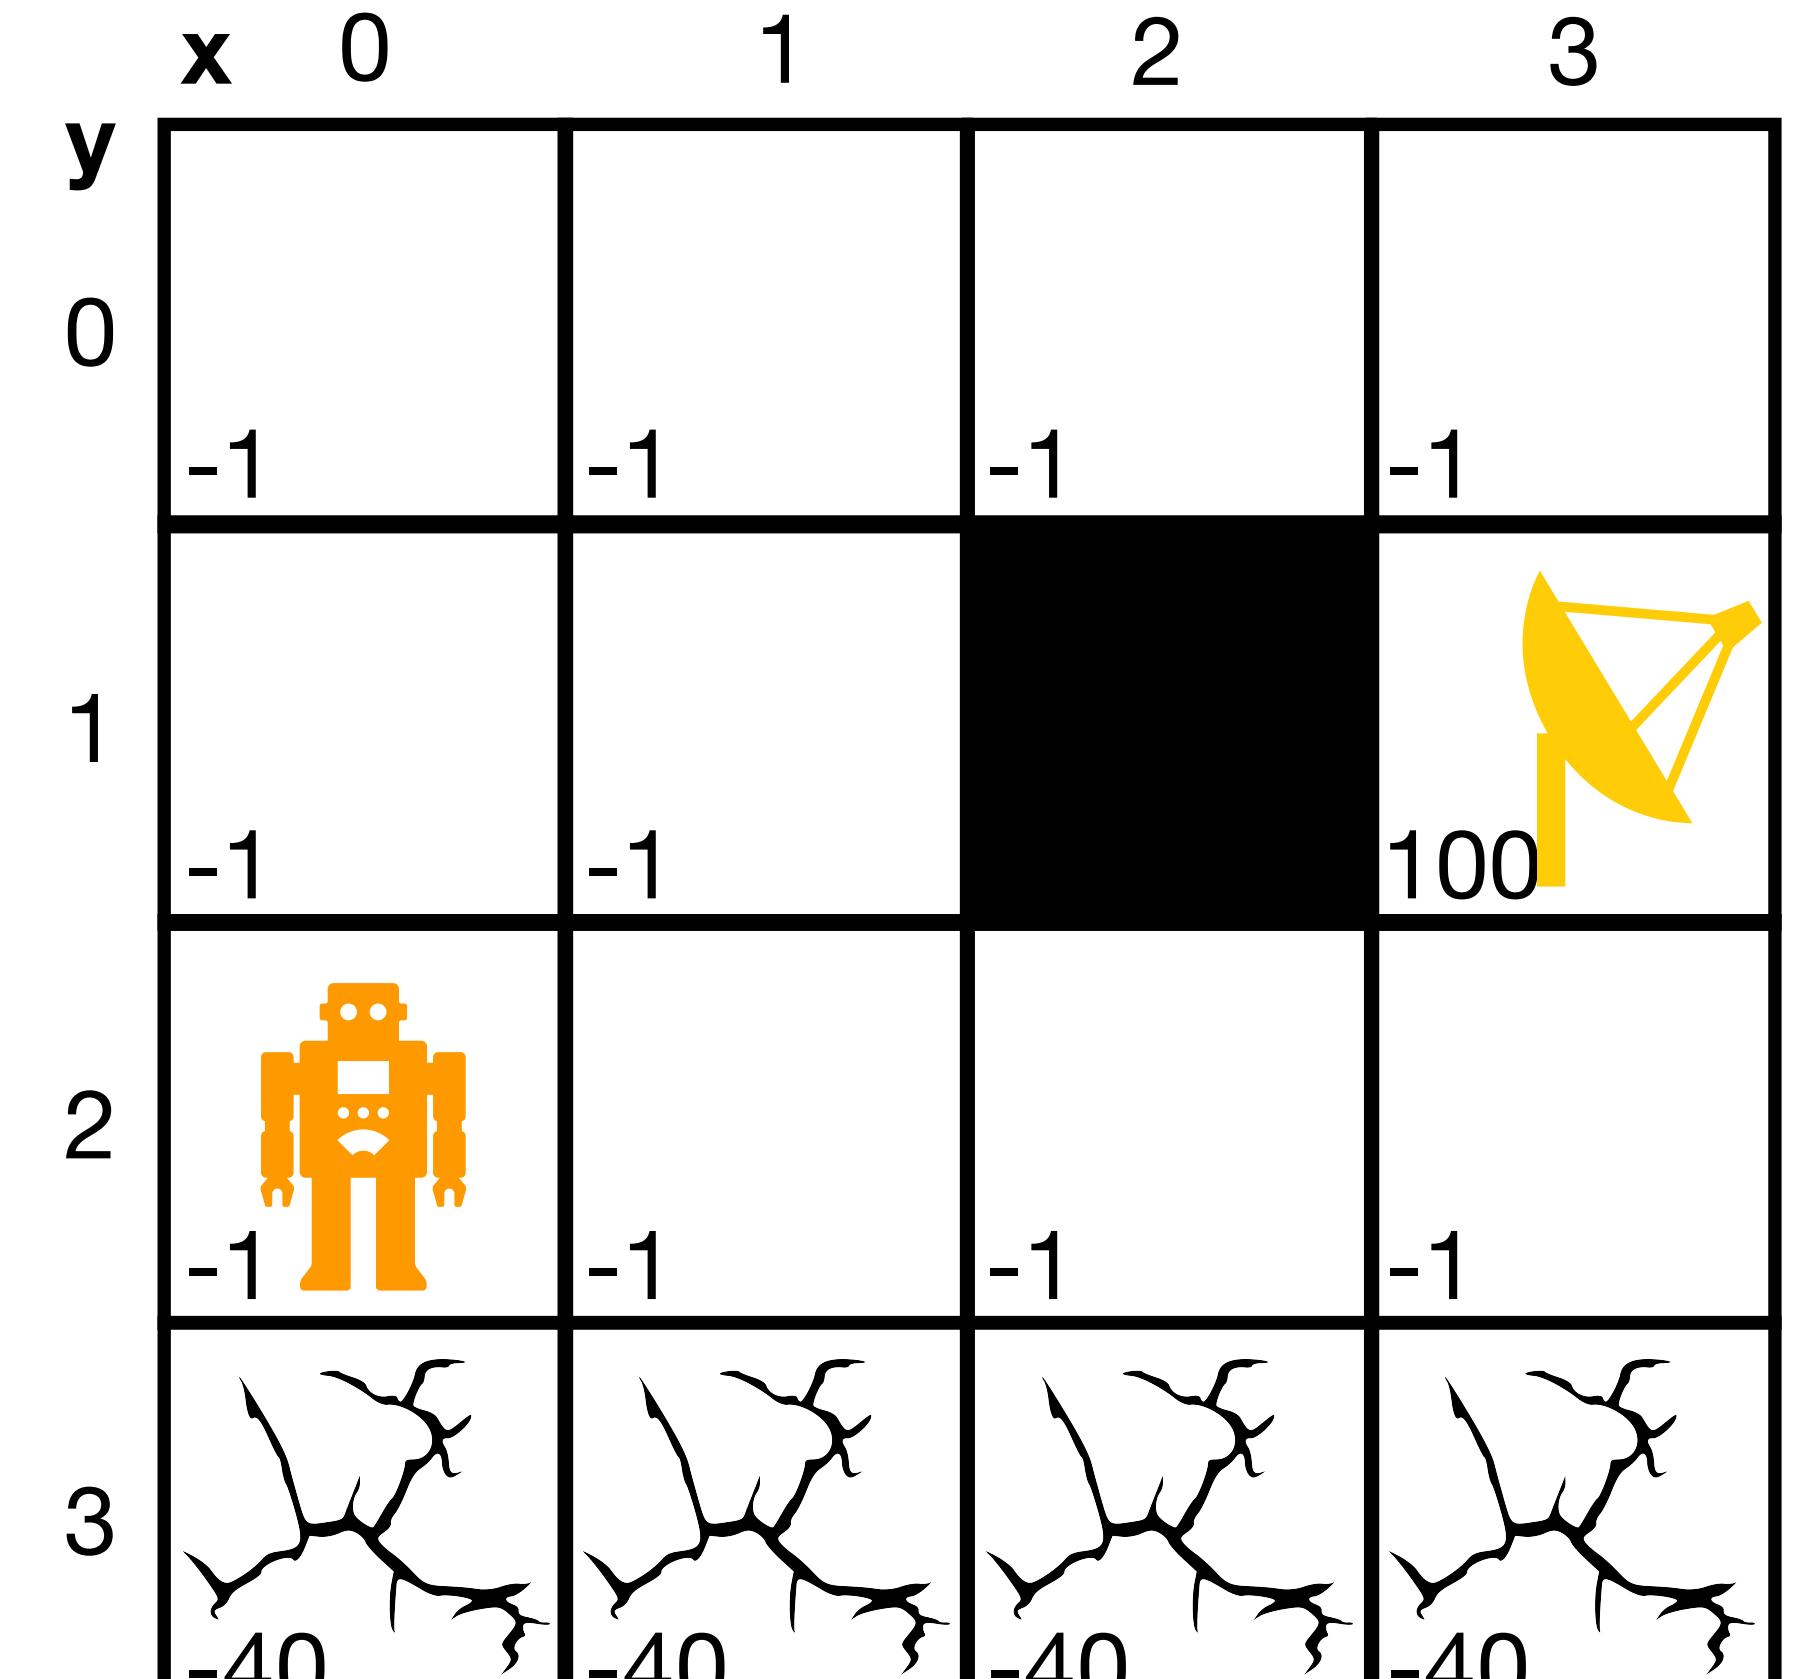
 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

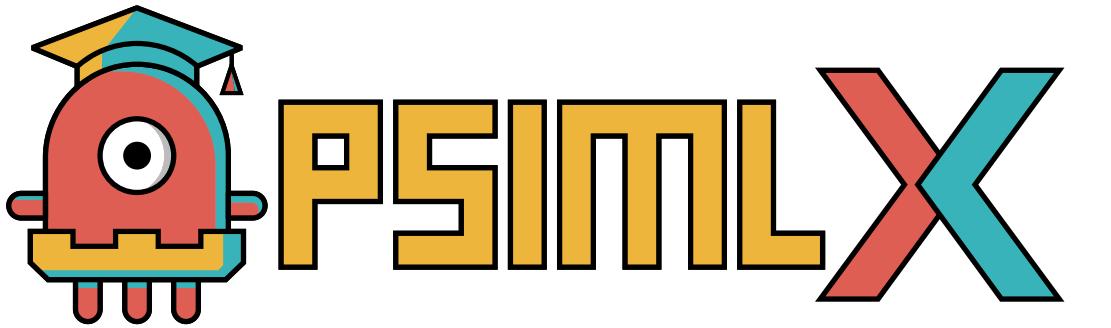
 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

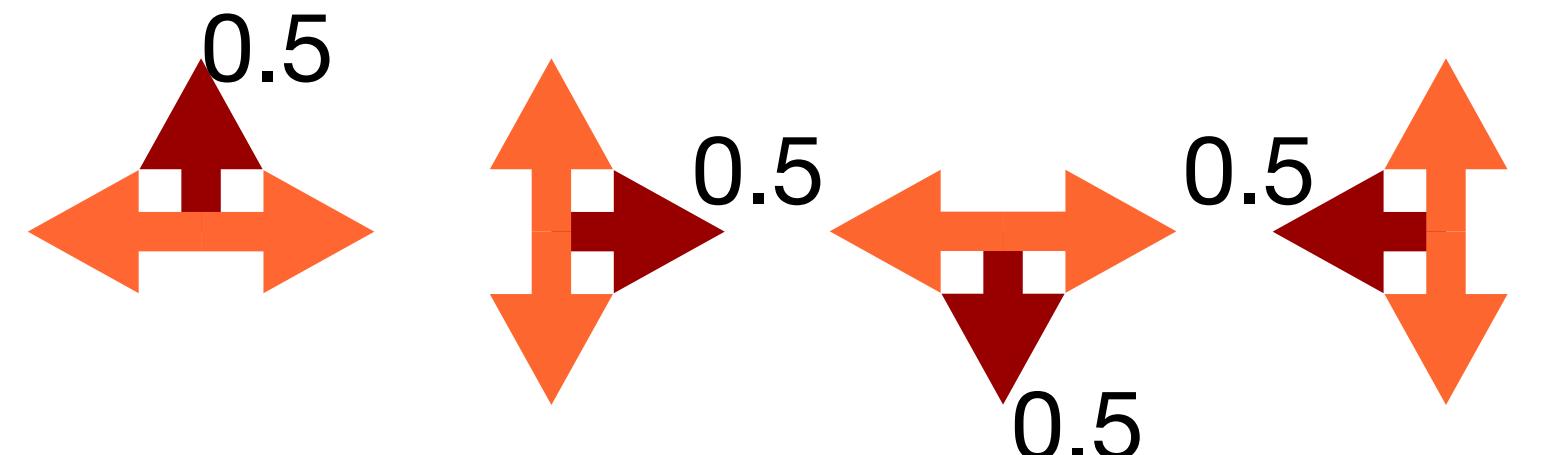
$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

 until S is terminal

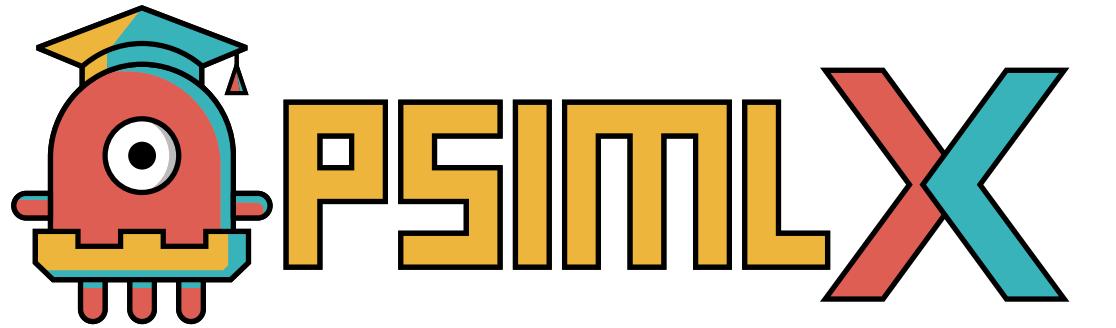
$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	-2	-1	-2
-3	-3	-1	-2	
3	0	0	0	0
-4	-40	0	-40	0
4	100			



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

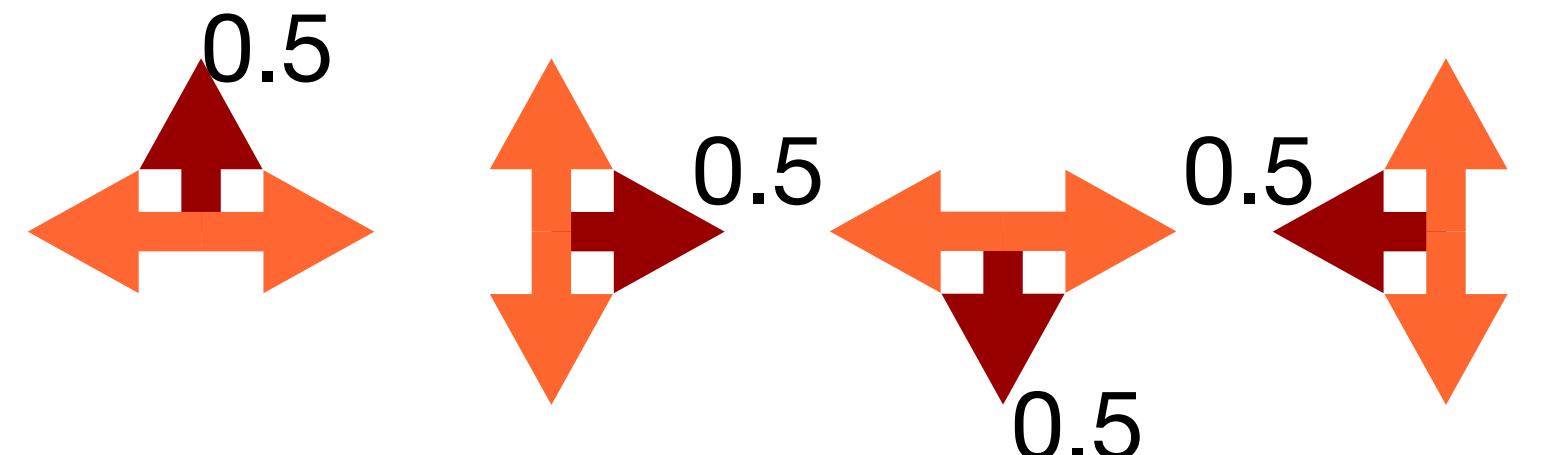
$S \leftarrow S'$

 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

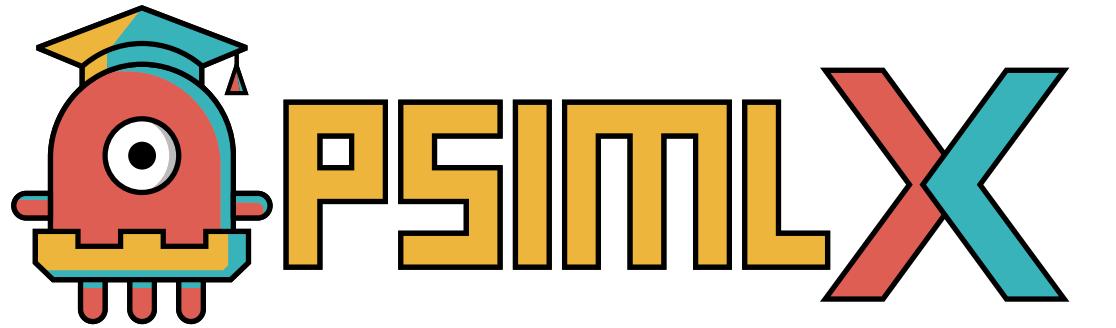
$$S = (0, 2)$$

x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	-2	-1	-2
-3	-3	-1	-2	
3	-1	-2	-2	-1
-4	-1	-1	-1	-3
4	-40	0	0	0



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

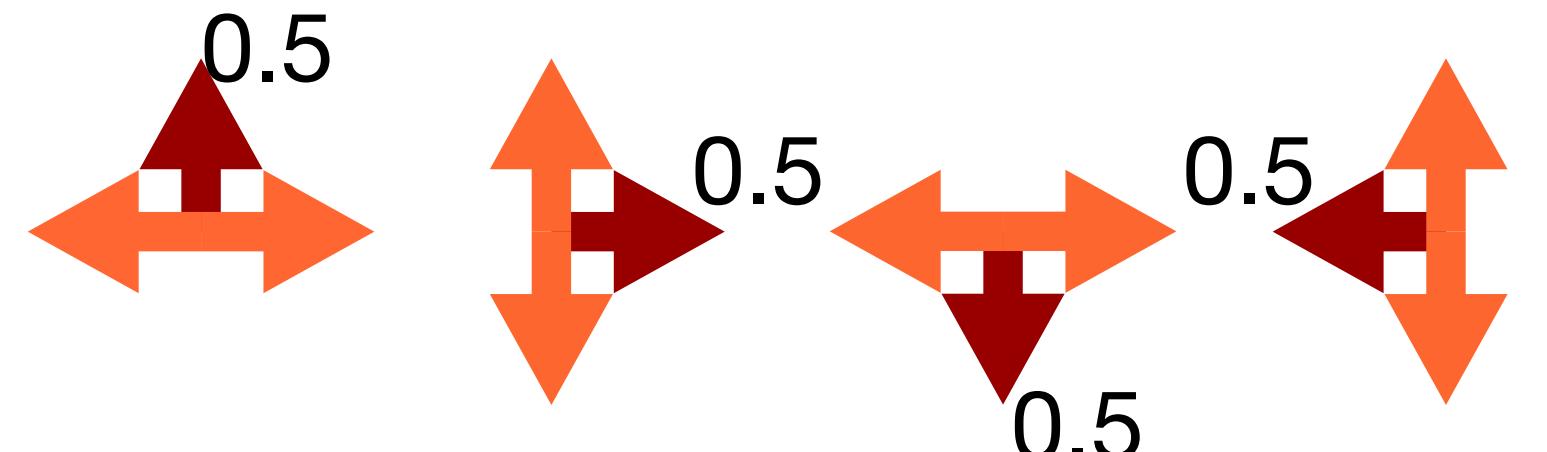
$S \leftarrow S'$

 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

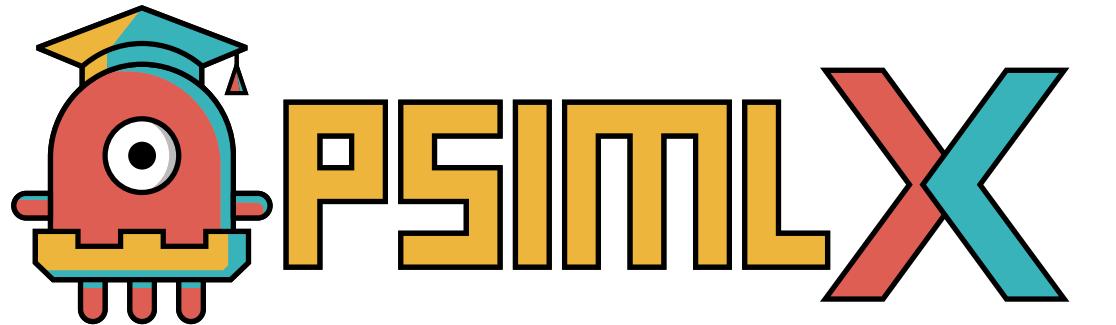
$$S = (0, 2) \quad c \sim U_{[0,1]} = 0.42 > \varepsilon \Rightarrow A = \text{RIGHT}$$

x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	-1	0	
-3	-2	-1	-1	-2
3	-1	-2	-2	-2
-4	-1	-2	-1	-1
0	0	0	0	0
-1	-40	0	-40	0
1	-40	0	-40	0
-2	-40	0	-40	0
2	-40	0	-40	0
-3	-40	0	-40	0
3	-40	0	-40	0



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

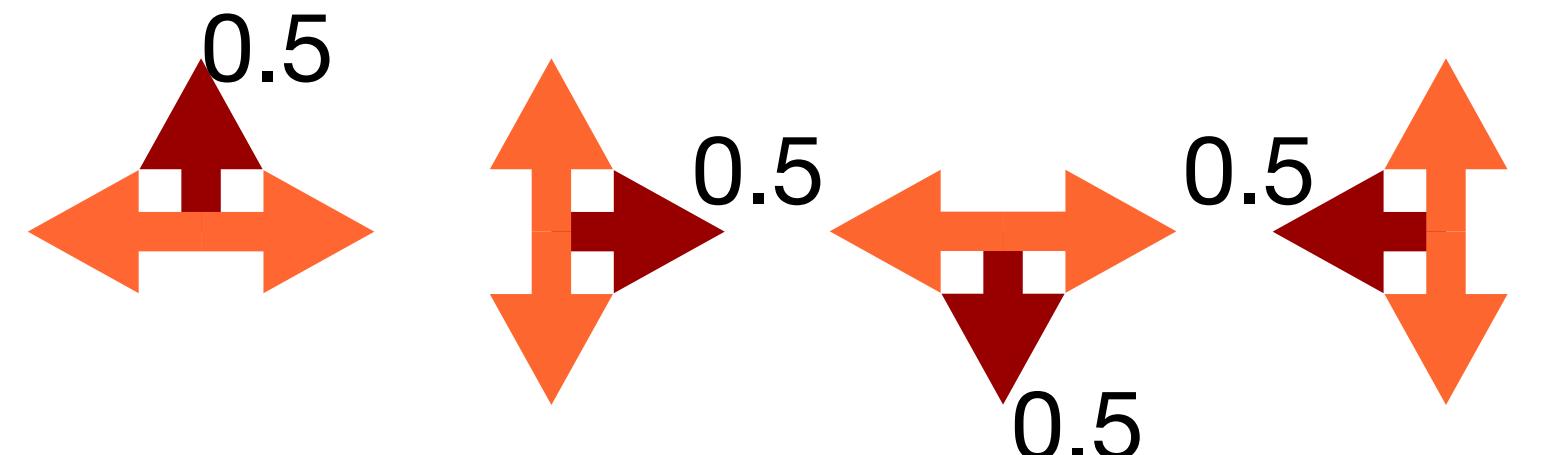
 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

$$S = (0, 2) \quad c \sim U_{[0,1]} = 0.42 > \varepsilon \Rightarrow A = \text{RIGHT} \quad R = -0.1$$

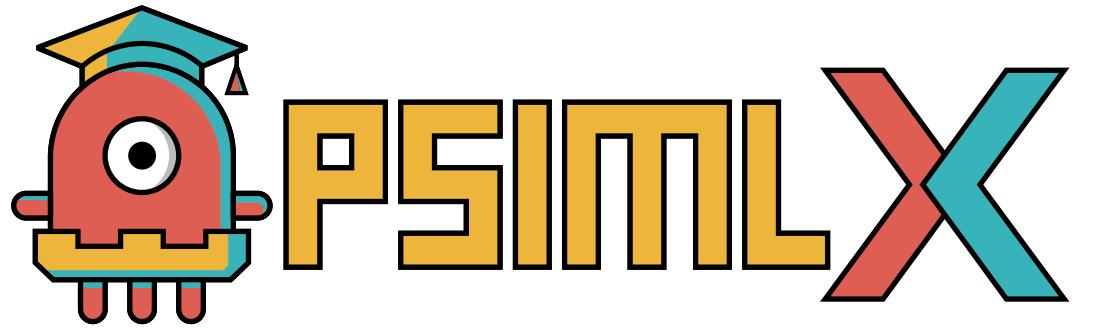
$$S' = (1, 2)$$

x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	0	-1	-2
-3	-3	-1	-2	
3	-40	0	0	-40
-40	0	-40	0	0



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

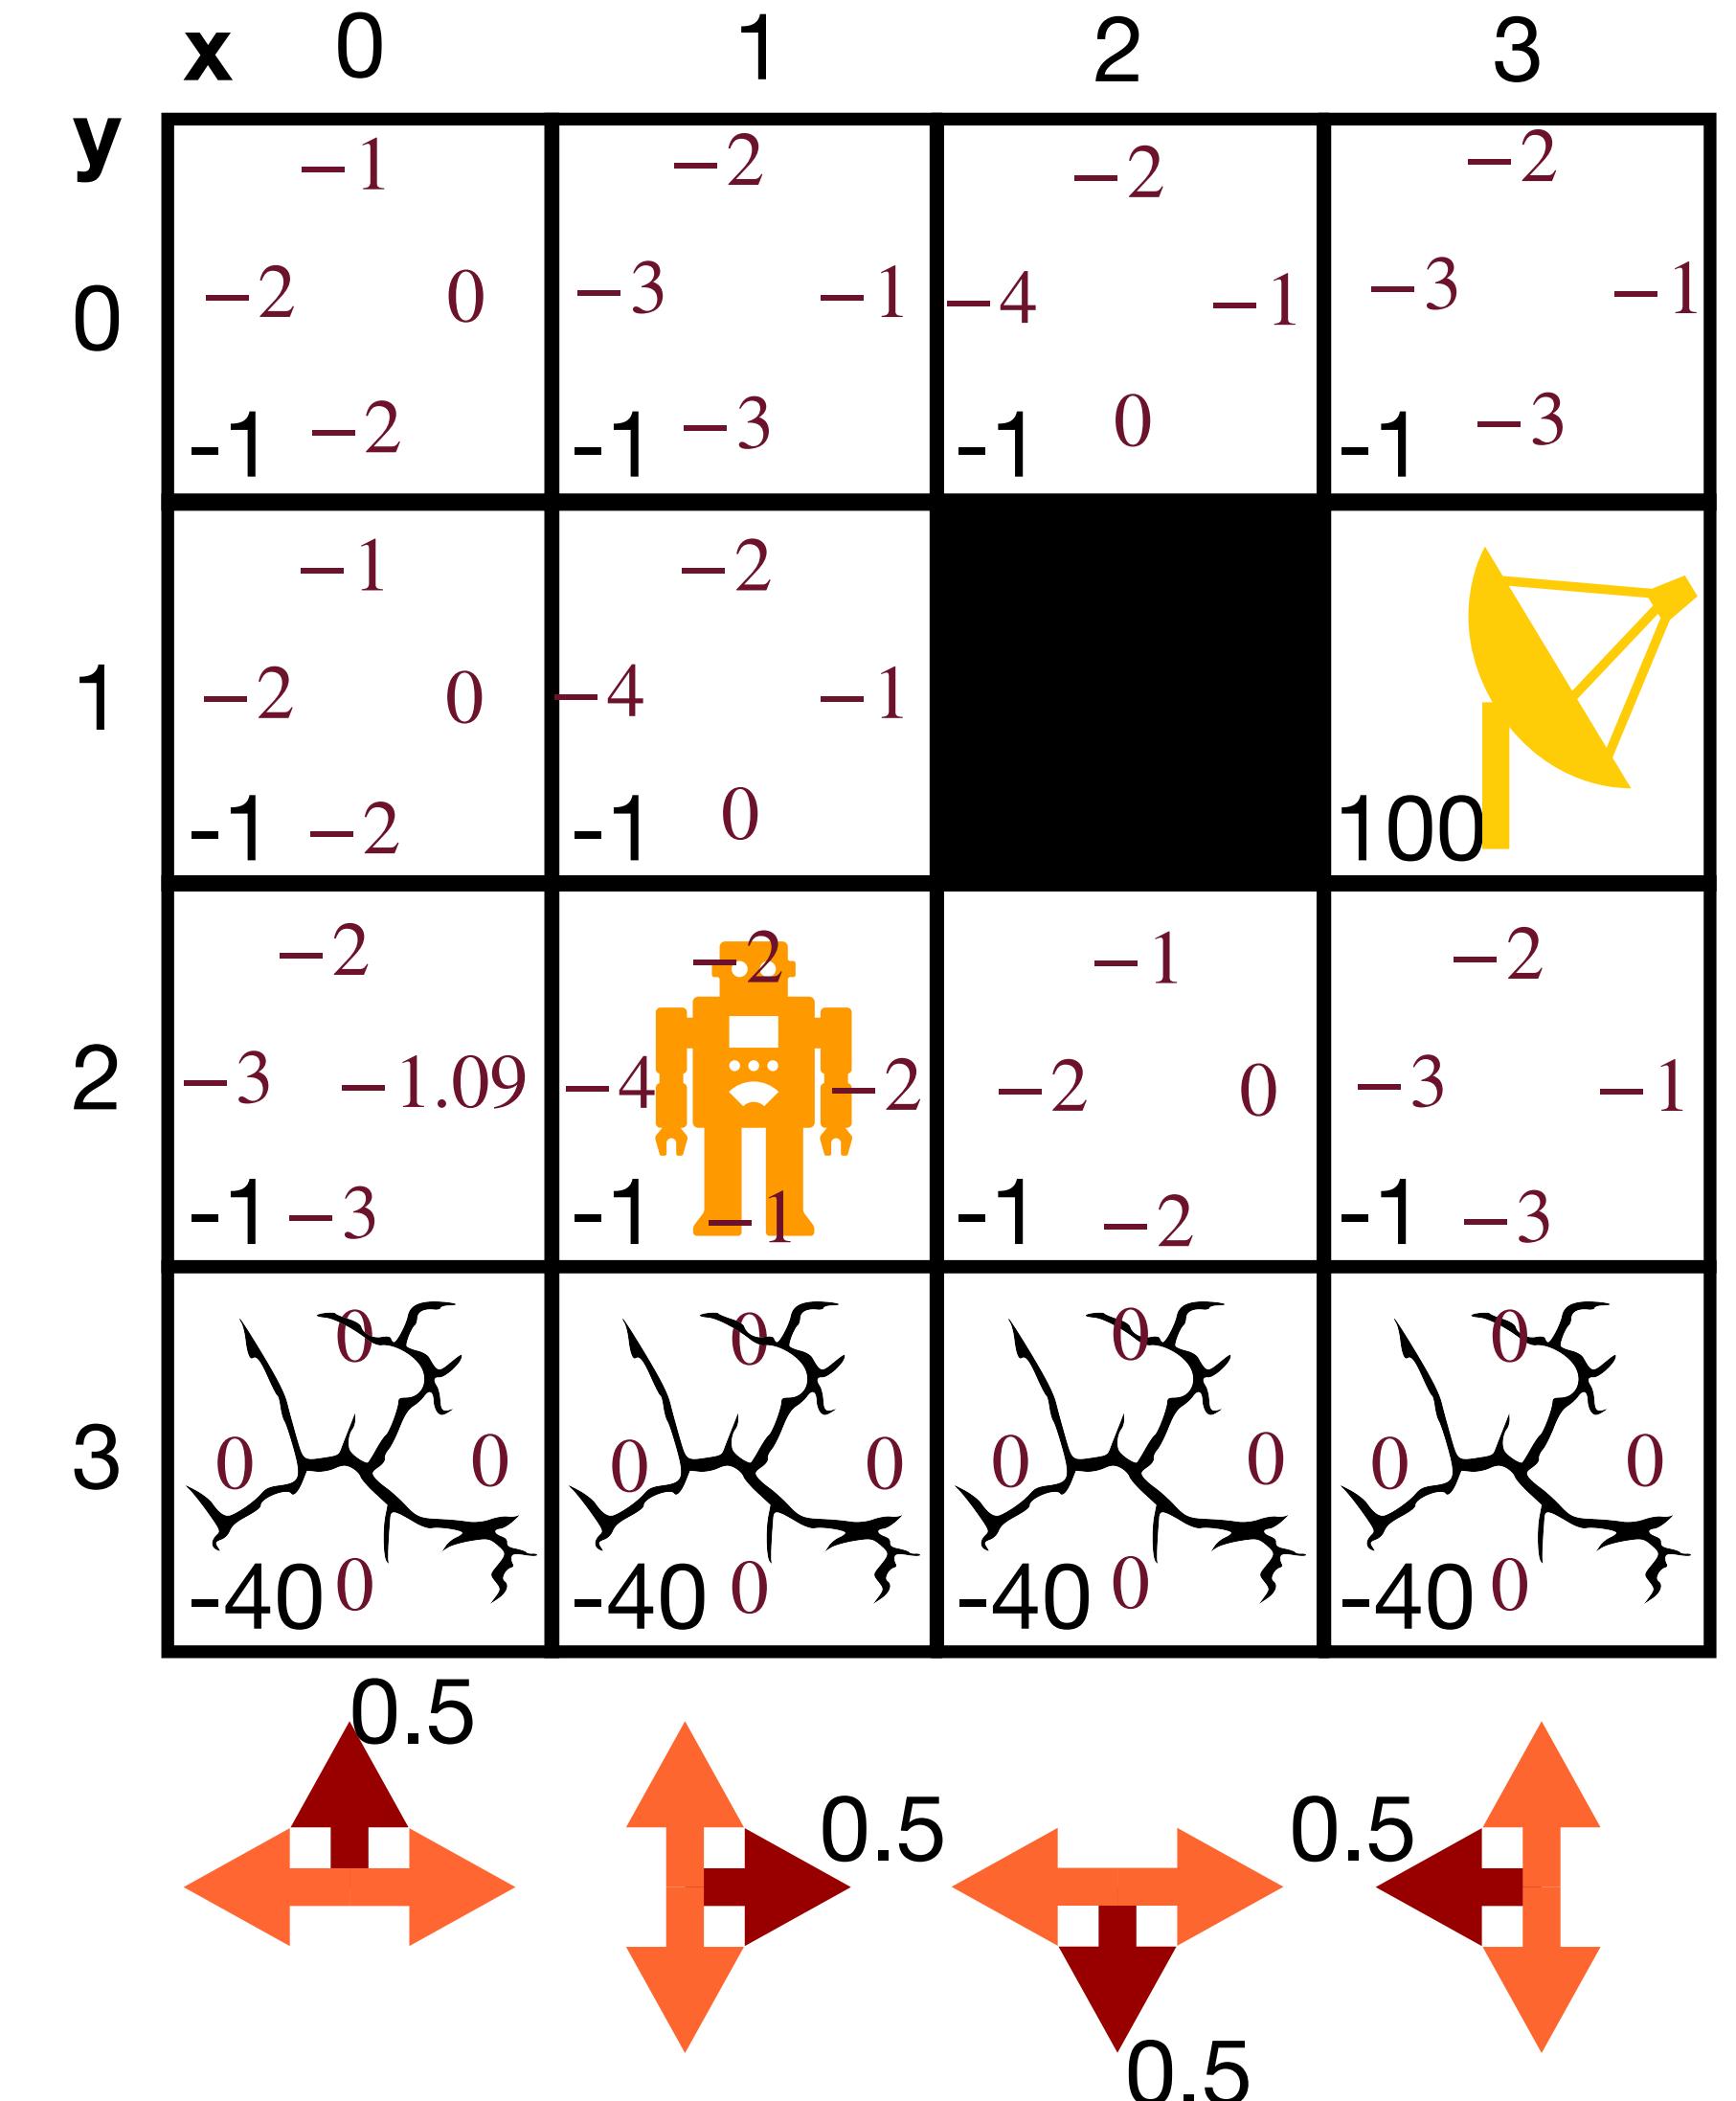
 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

$$S = (0, 2) \quad c \sim U_{[0,1]} = 0.42 > \varepsilon \Rightarrow A = \text{RIGHT} \quad R = -0.1$$

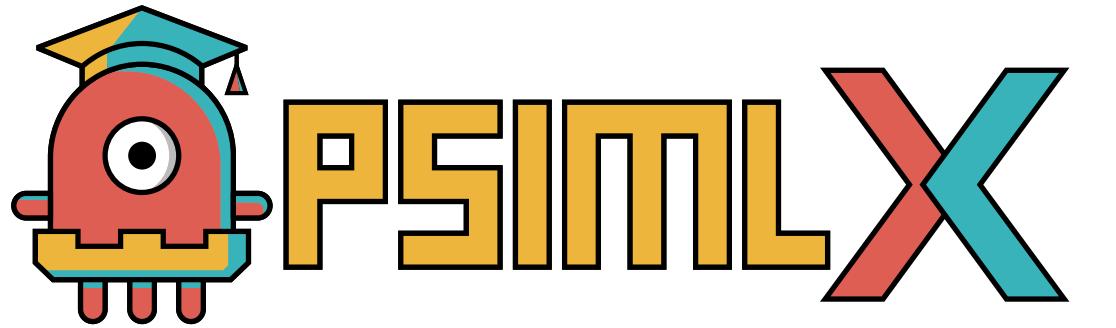
$$S' = (1, 2)$$

$$Q((0, 2), \text{RIGHT}) \leftarrow -1 + 0.1 \cdot [-1 + 0.9 \cdot 0 - (-1)] = -1.09$$



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

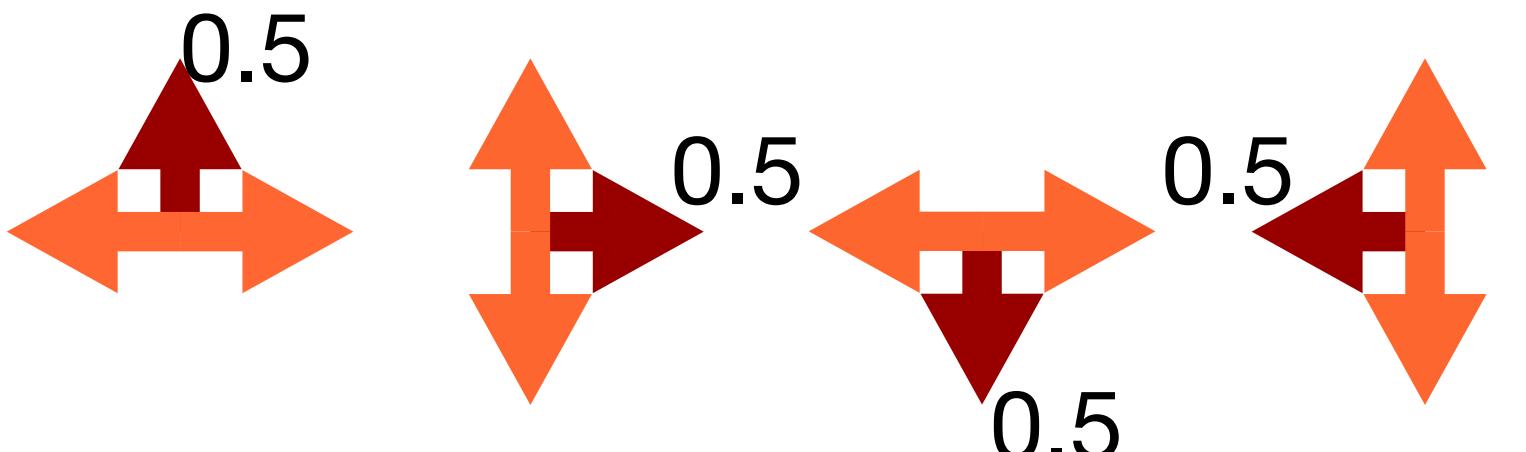
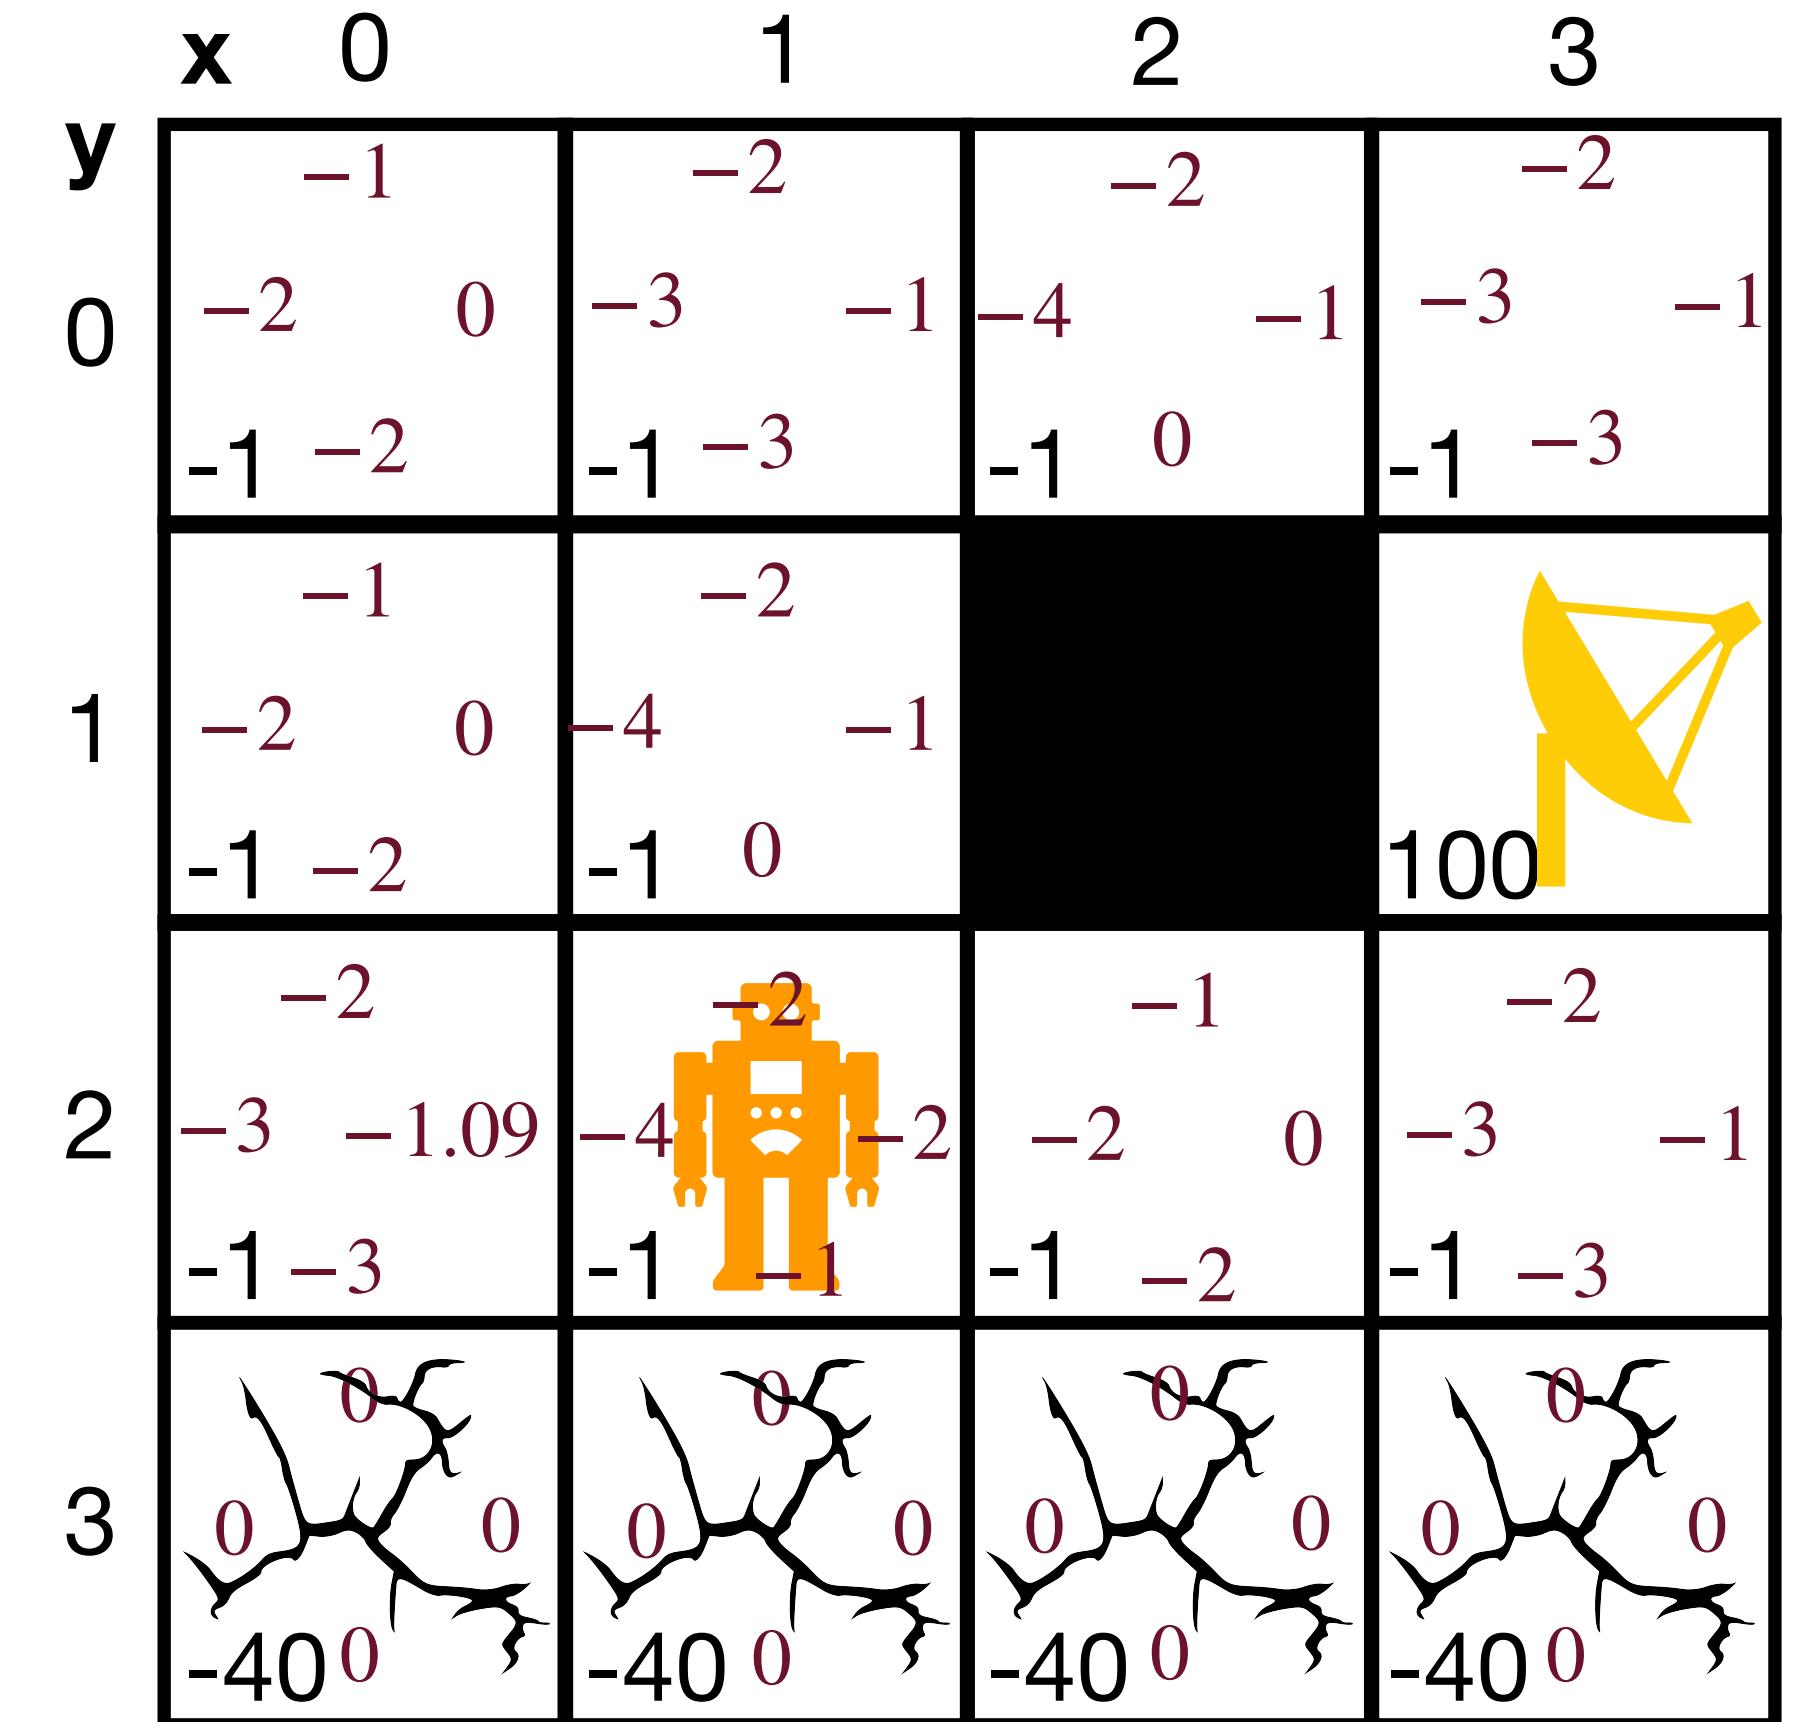
$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

 until S is terminal

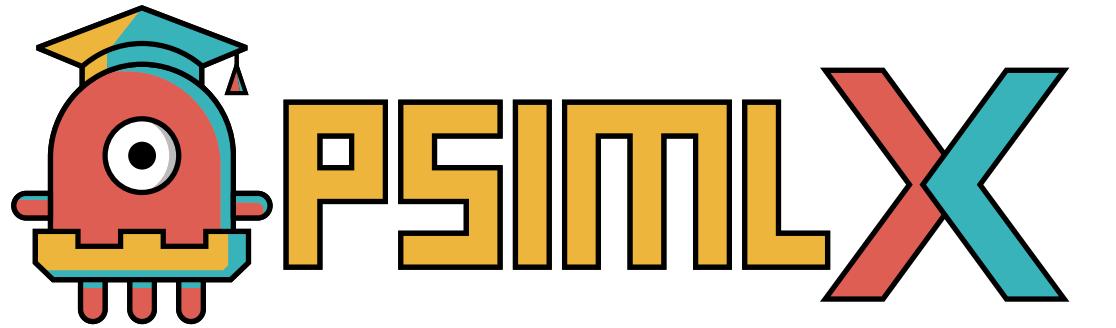
$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

$$S = (1, 2)$$



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

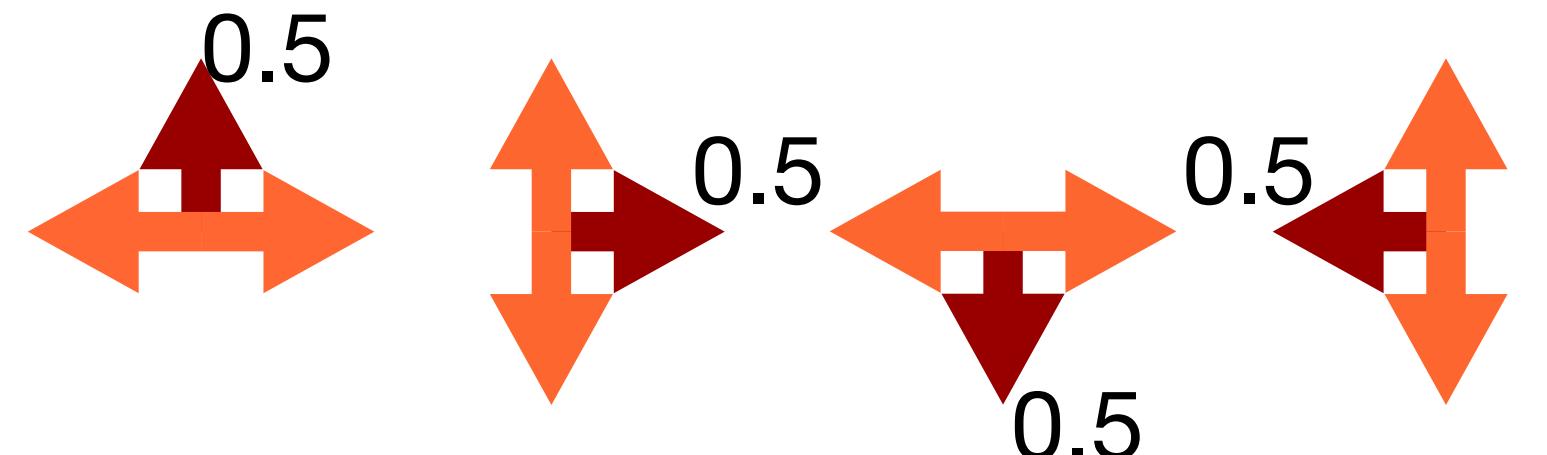
$S \leftarrow S'$

 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

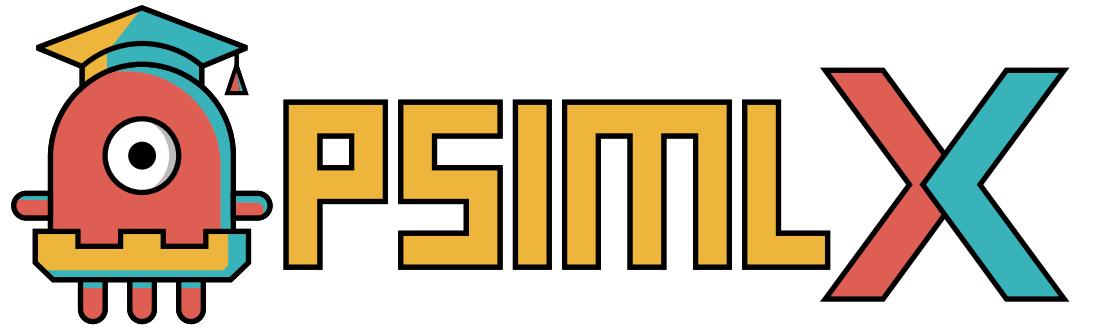
$$S = (1, 2) \quad c \sim U_{[0,1]} = 0.82 > \varepsilon \Rightarrow A = \text{DOWN}$$

x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	0	-1	-2
-3	-3	-1.09	-4	-2
3	0	0	0	0
-40	-40	0	-40	0



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

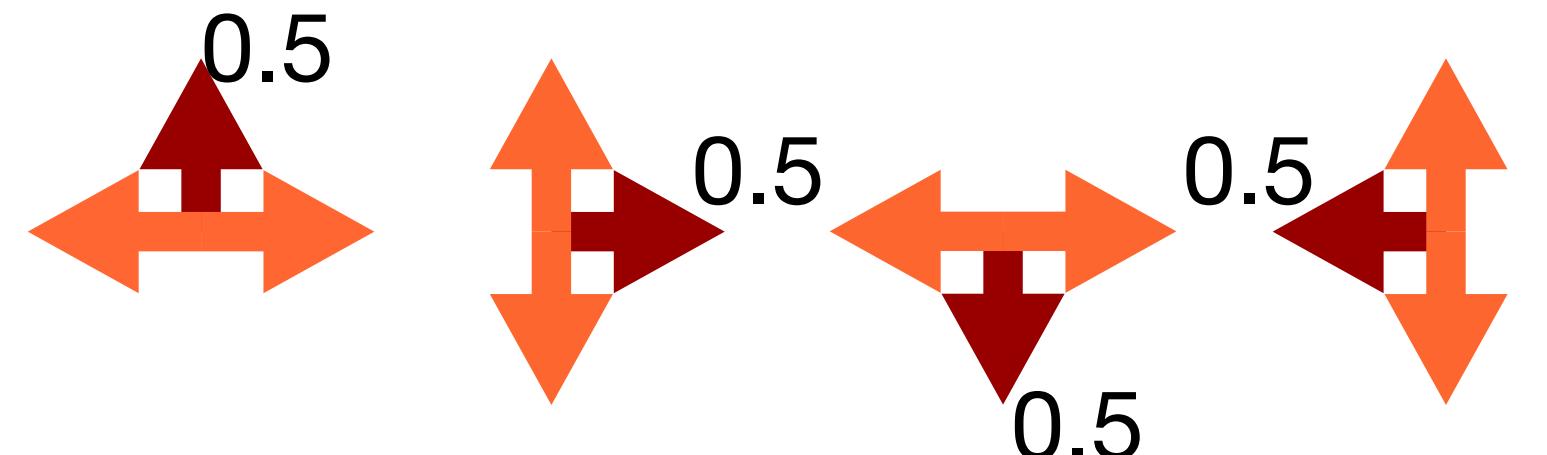
 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

$$S = (1, 2) \quad c \sim U_{[0,1]} = 0.82 < \varepsilon \Rightarrow A = \text{DOWN} \quad R = -40$$

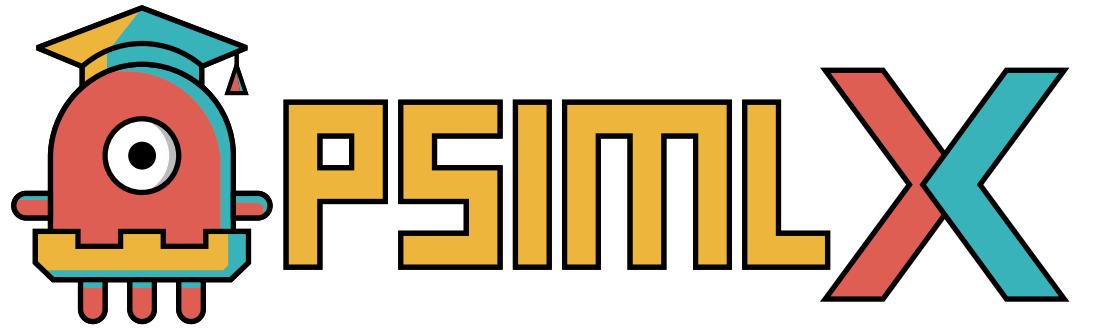
$$S' = (1, 3)$$

x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	-4	-1	
-3	-1	0		
3	-2	-2	-1	-2
-4	-3	-1.09	-4	-2
-1	-1	-3	-1	-2
4	0	0	0	0
-5	-40	0	-40	0
-6	0	0	0	0
-7	-40	0	-40	0
-8	0	0	0	0
-9	-40	0	-40	0
-10	0	0	0	0



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

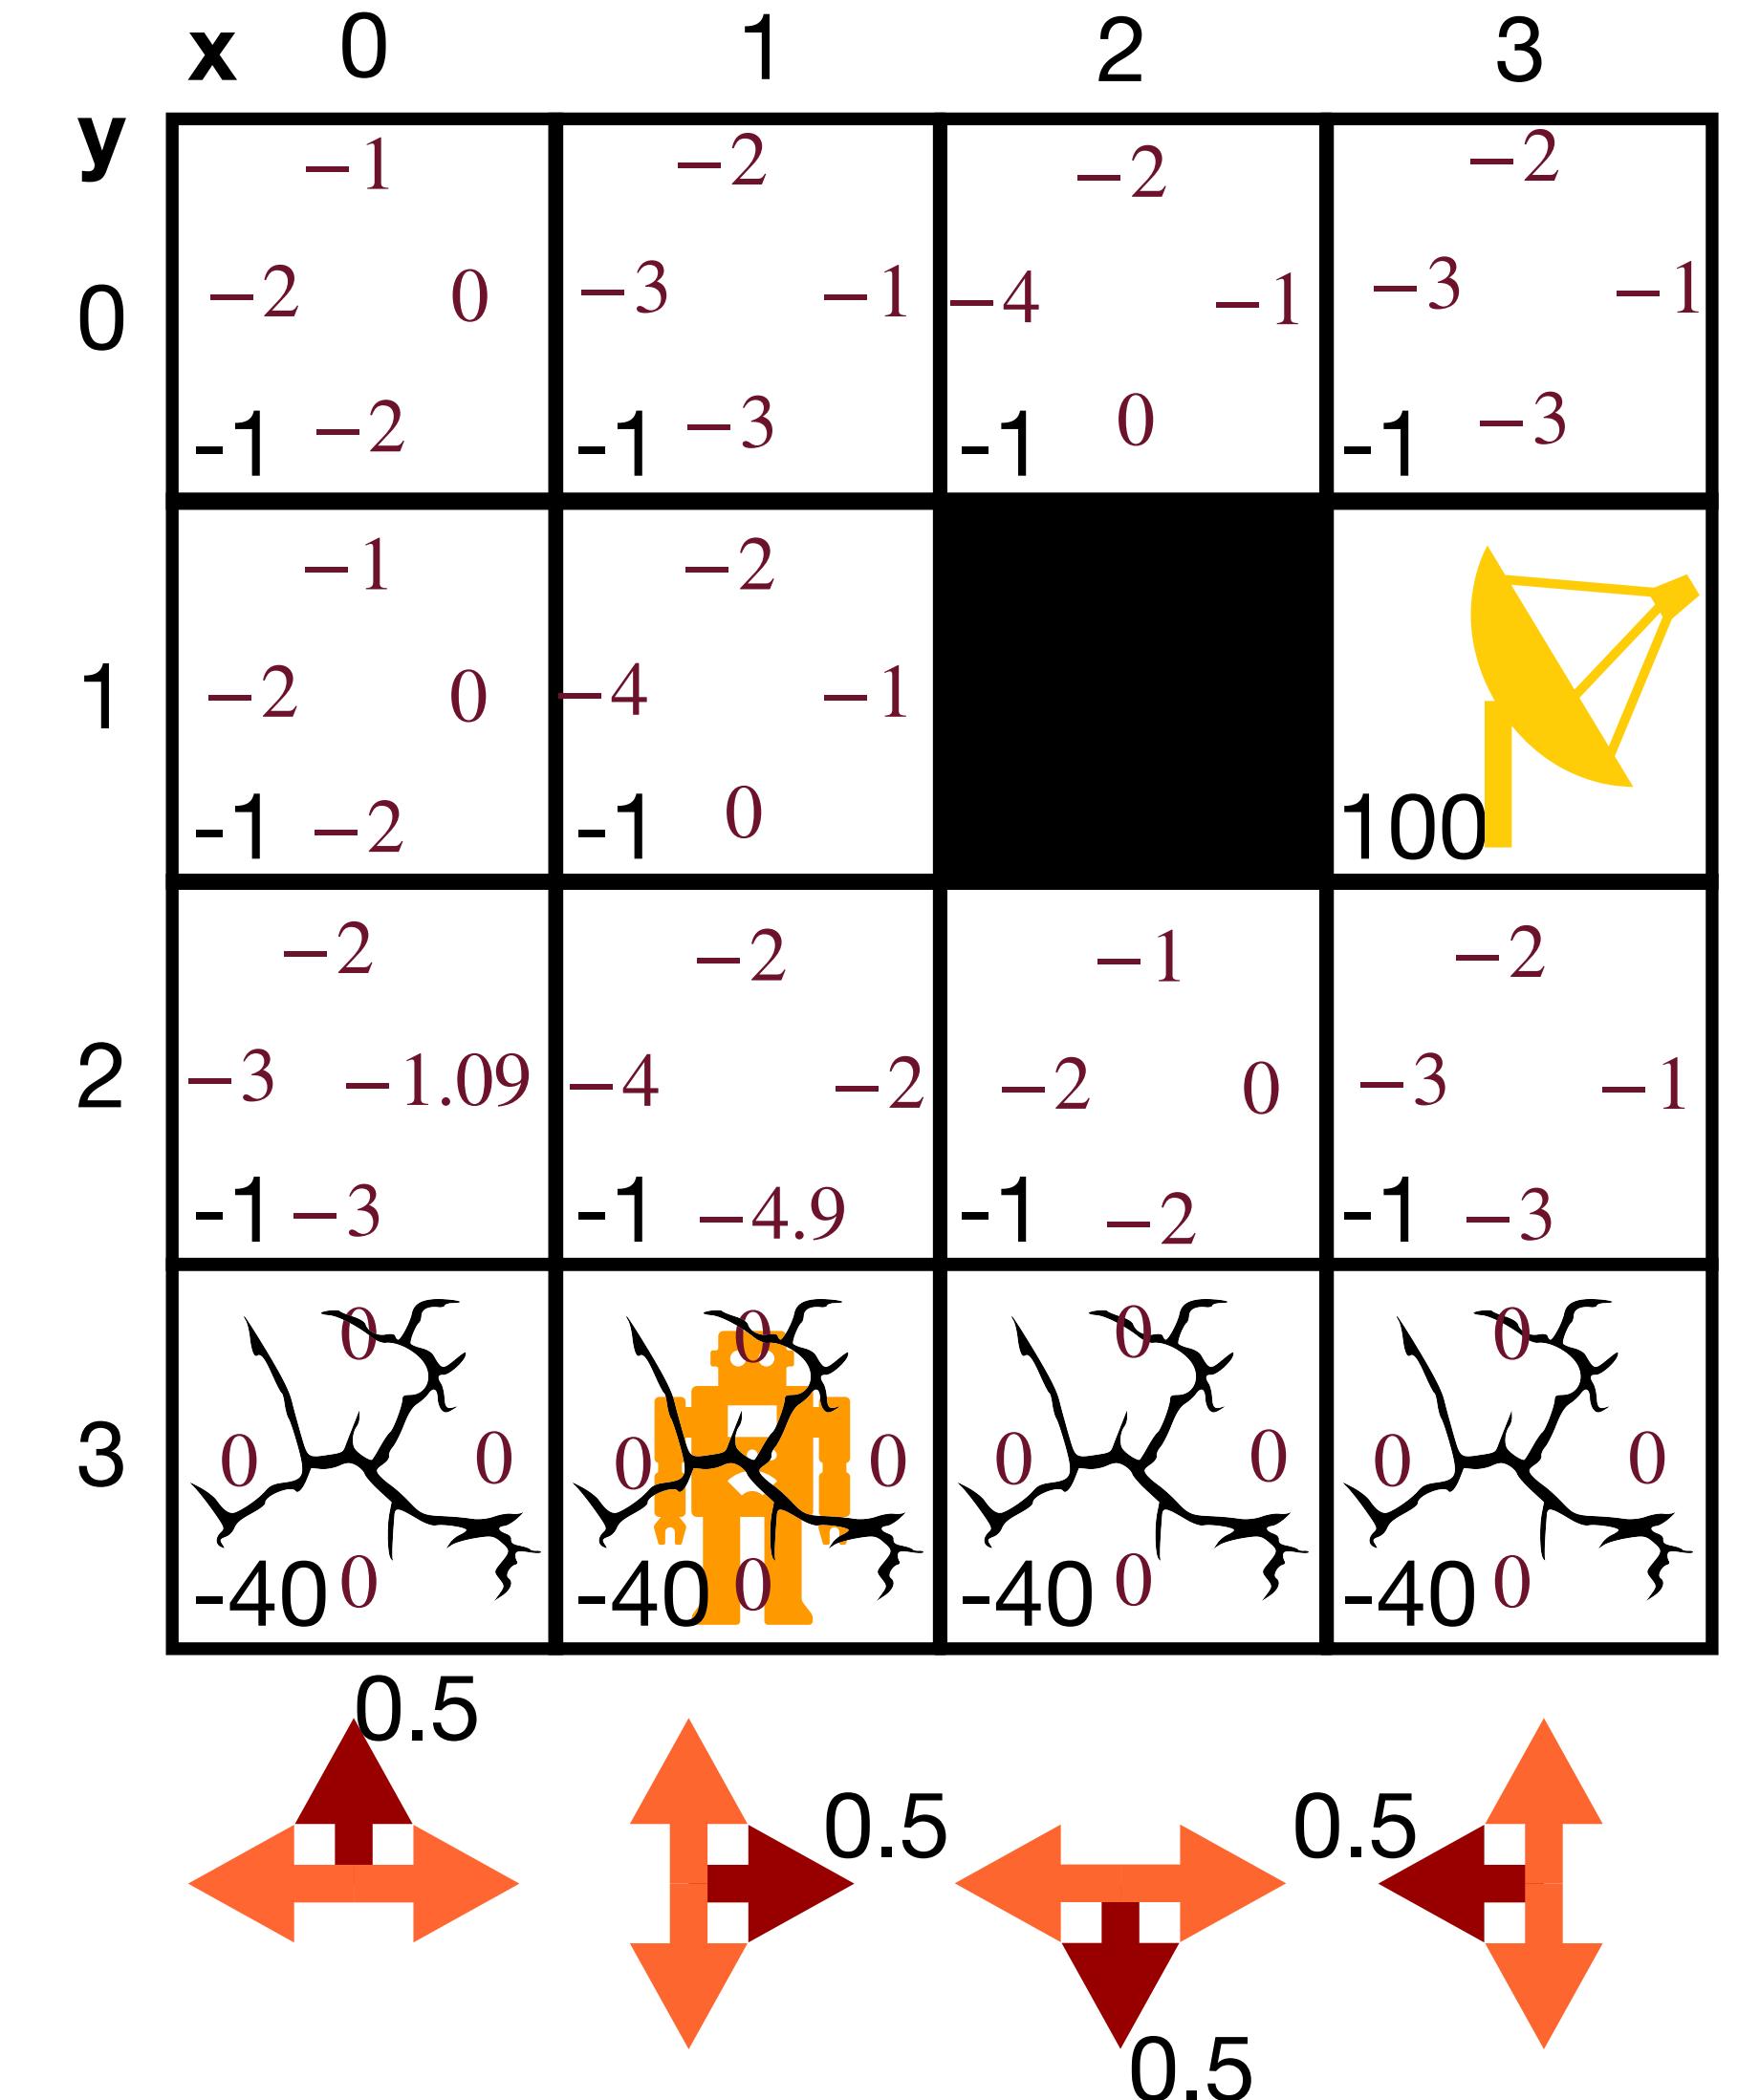
 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

$$S = (1, 2) \quad c \sim U_{[0,1]} = 0.82 < \varepsilon \Rightarrow A = \text{DOWN} \quad R = -40$$

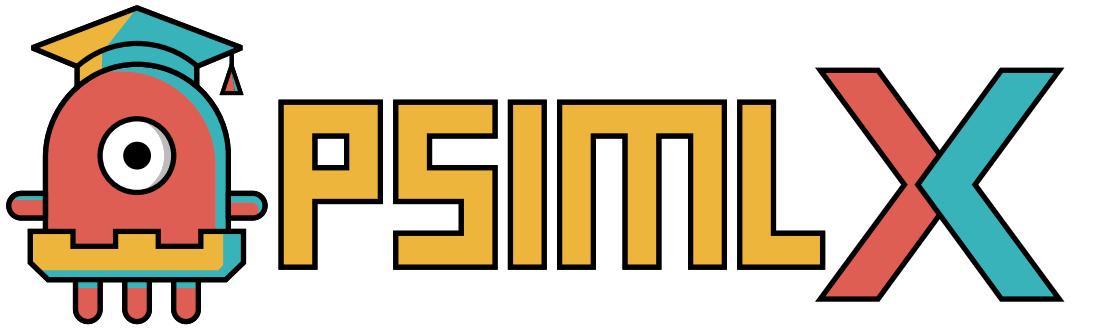
$$S' = (1, 3)$$

$$Q((1, 2), \text{DOWN}) \leftarrow -1 + 0.1 \cdot [-40 + 0.9 \cdot 0 - (-1)] = -4.9$$



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

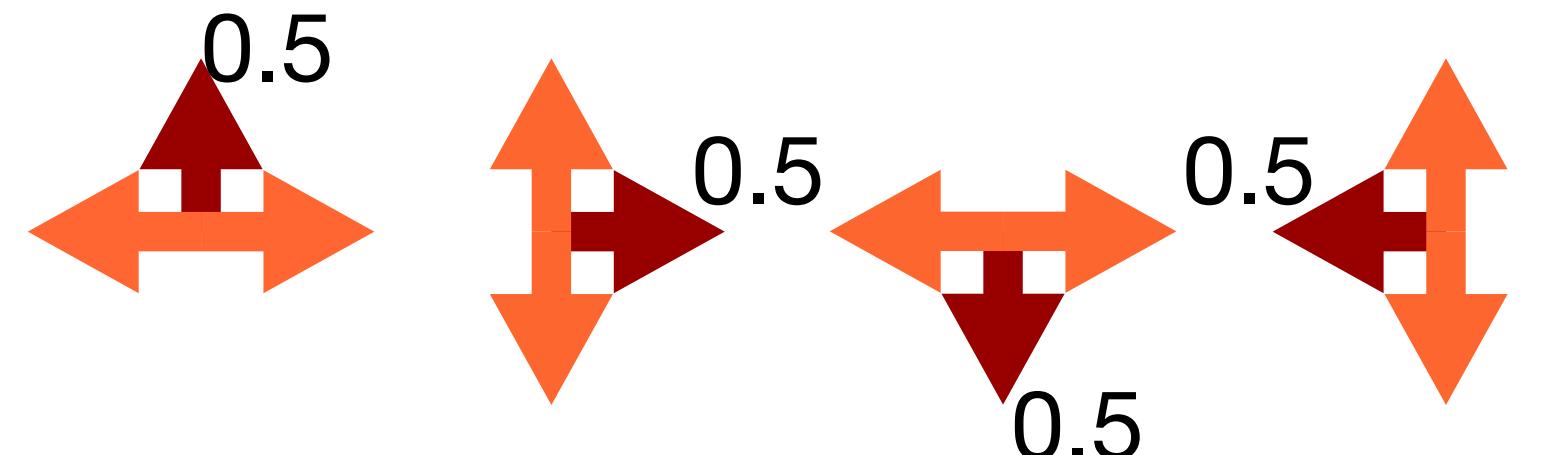
$S \leftarrow S'$

 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

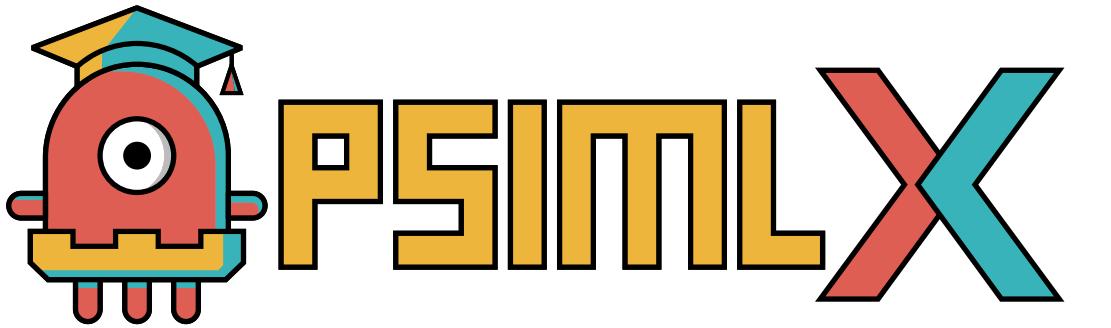
$$S = (1, 3)$$

x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	-4	-1	
-3	-1	0		
3	-2	-2	-1	-2
-4	-3	-1.09	-4	-1
-1	-1	-3	-1	-2
4	0	0	0	0
-5	-40	0	-40	0
-6	0	0	0	0
5	-40	0	-40	0
6	0	0	0	0



Model free RL

Example 2: Off-Policy TD Algorithm



Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

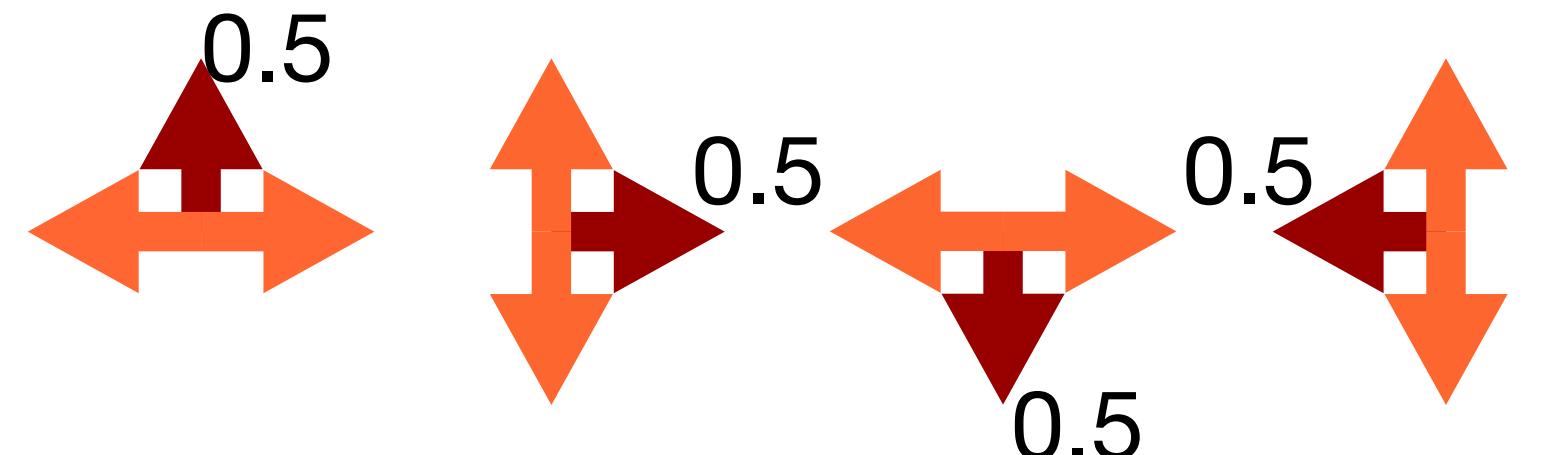
$S \leftarrow S'$

 until S is terminal

$$\alpha = 0.1 \quad \gamma = 0.9 \quad \varepsilon = 0.4$$

$$S = (0, 2)$$

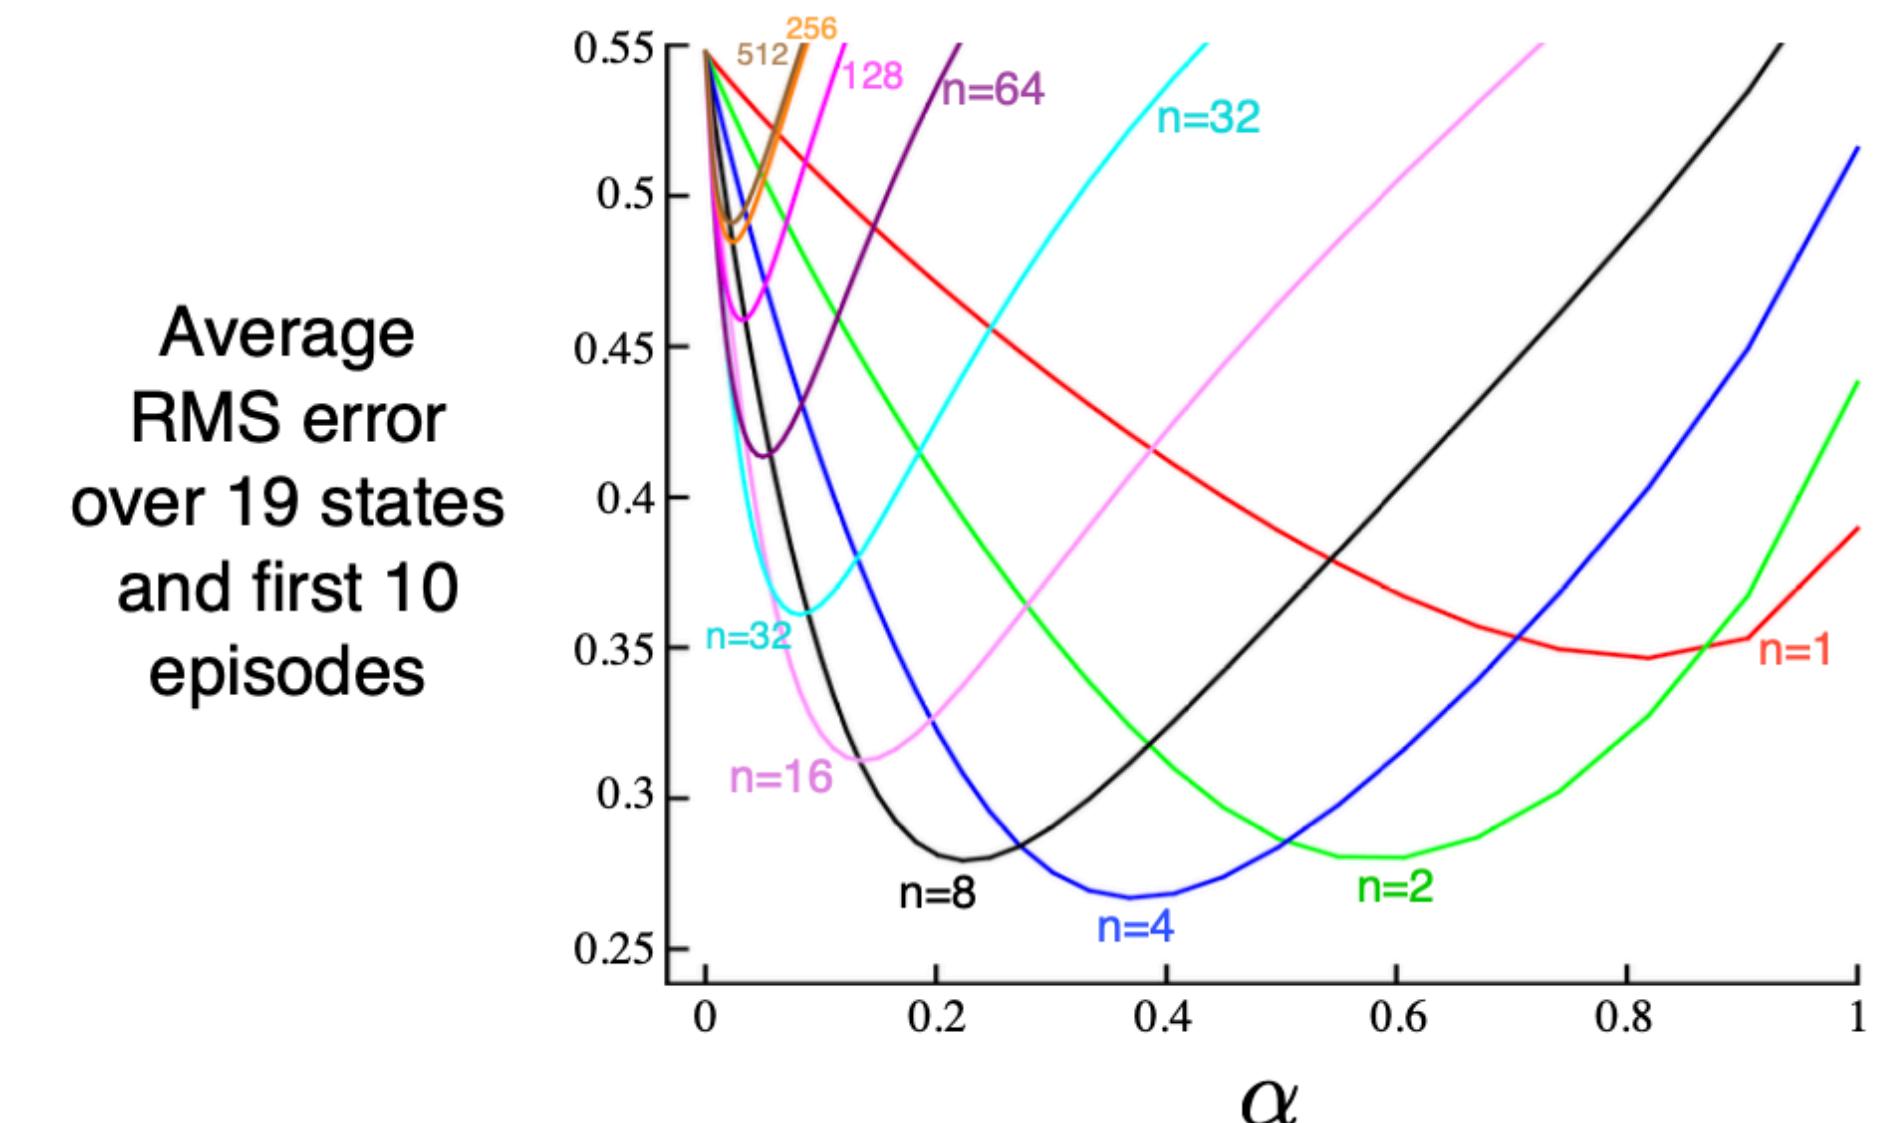
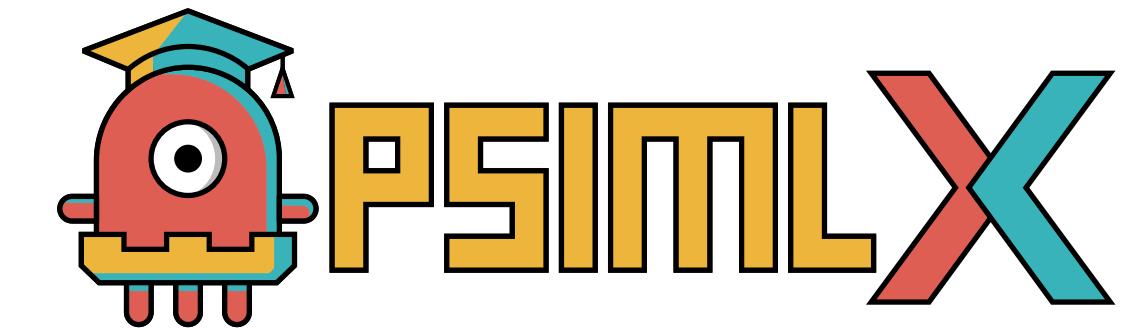
x	0	1	2	3
y	-1	-2	-2	-2
0	-2	0	-3	-1
-1	-1	-2		
1	-2	0	-4	-1
-2	-1	-2		
2	-2	-1	0	
-3	-2	-1	-1	-2
3	-3	-1	-2	-1
-4	-1	0	-2	-3
0	0	0	0	0
-40	0	0	0	0



Model free RL

Conclusions

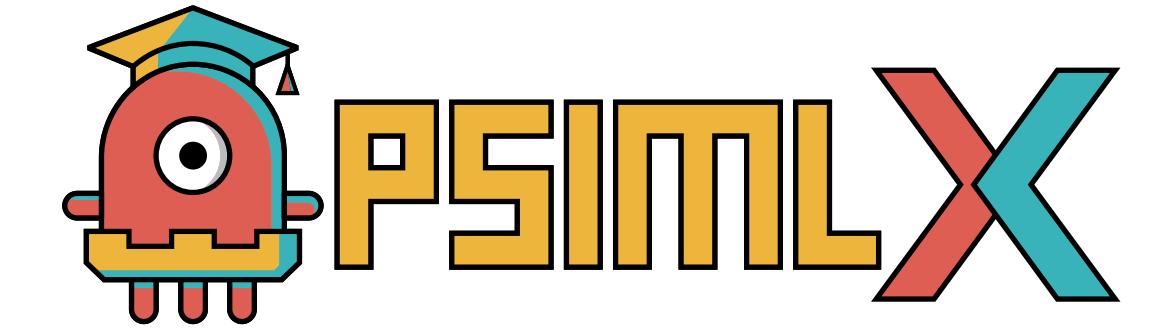
- We have shown on-policy MC methods and off-policy Q-Learning
- This does not mean that all MC methods are on-policy and all TD methods are off-policy
 - Off-policy MC methods: Utilise importance sampling
 - SARSA (Step-Action-Reward-Step-Action): On-policy TD method
- We have seen 1-step TD methods:
 - n-step TD methods bridge the gap between MC and TD paradigms



n-step TD performance with varying n — Intermediate solutions may be the best [Sutton & Barto 2018]

Model free RL

Conclusions

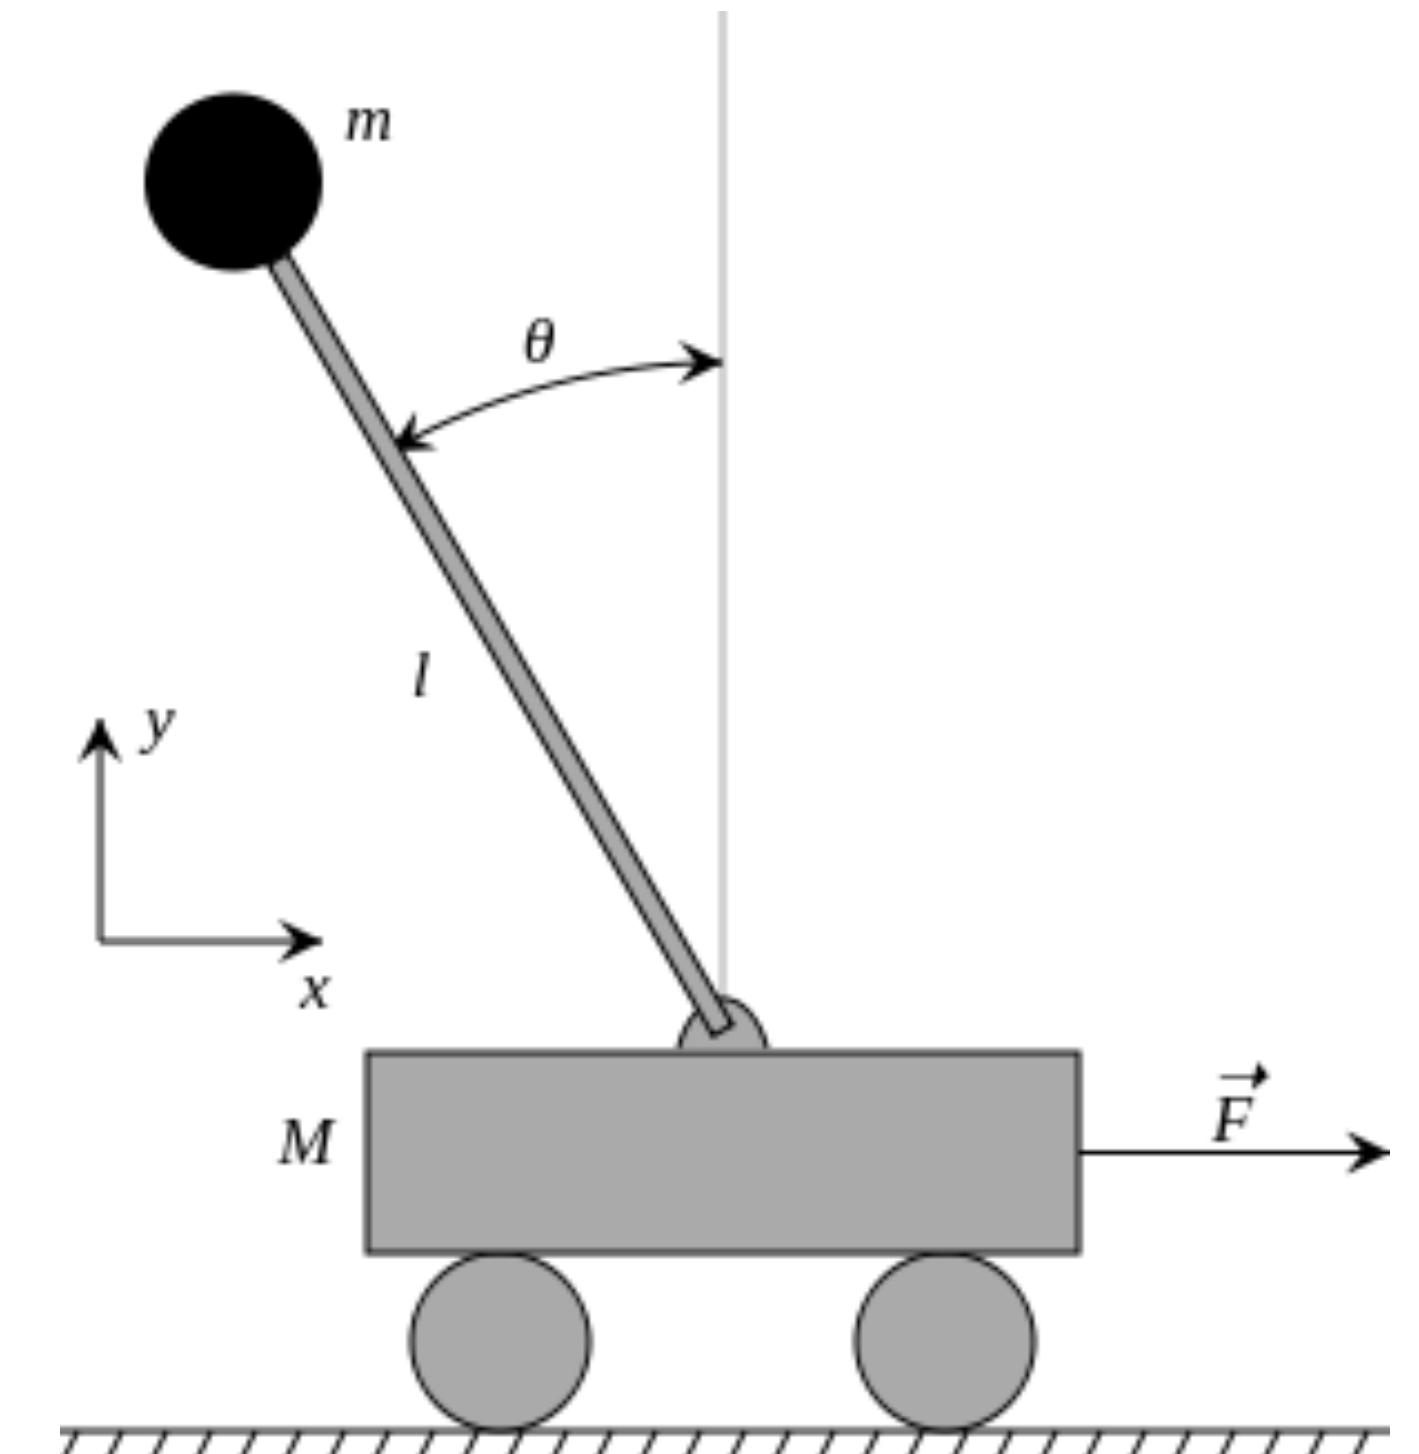
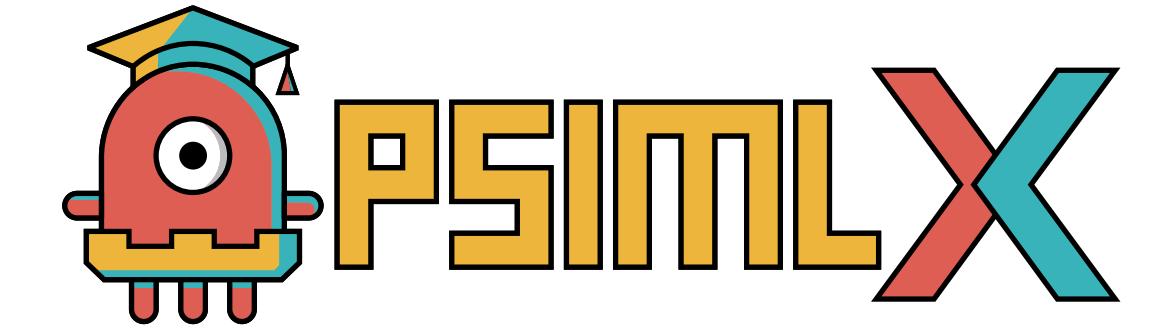


Criteria	Monte Carlo Methods	Temporal Difference Methods
Bias vs Variance	low bias, high variance	high bias, low variance
Online		✓
Bootstrapping		✓ because v_t is based off of v_{t+1}
Estimation	✓	✓
On-Policy	On-policy MC	SARSA
Off-Policy	Off-policy MC with importance sampling	Q-learning
Past vs Future Future data	past experiences	future (models MDP)

Model free RL

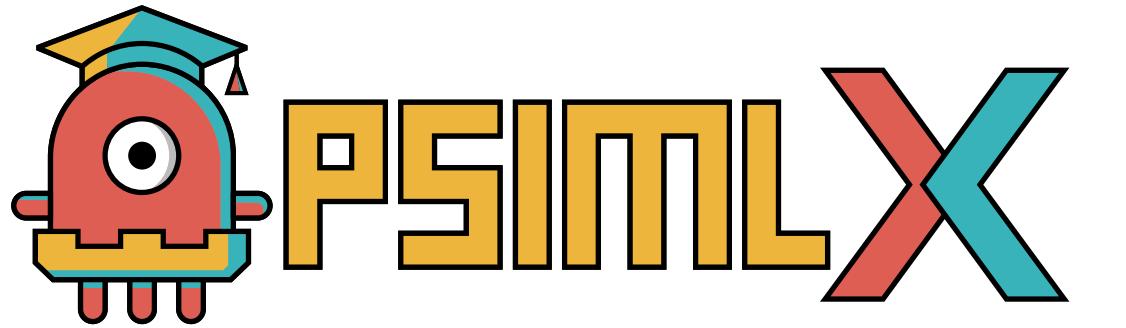
Beyond Tabular Methods

- How to handle large multi-dimensional state-spaces?
- Can we expect similar states in a large state-space to be reasonably similar?
Is interpolation possible?
- How to handle continuous state-spaces?
 - E.g. representing angles
- How to handle continuous actions?
 - E.g. representing force
- Curse of dimensionality
- **Can we utilise (deep) neural networks somehow?**



The cart pole problem

Outline

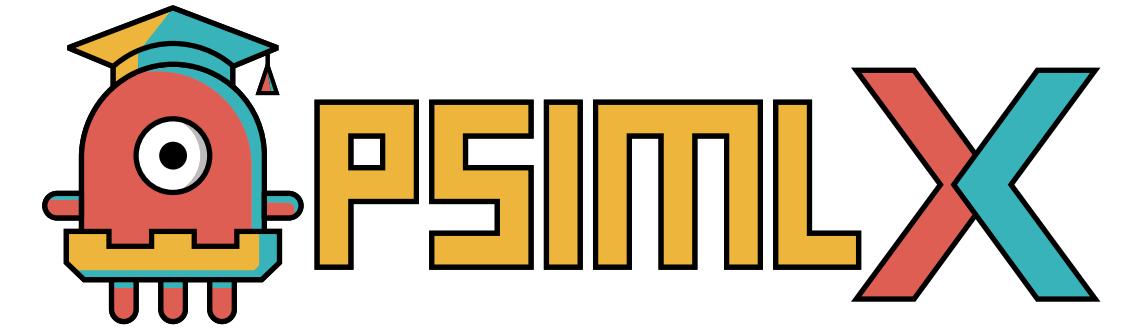


- Introduction
- Reinforcement Learning Formalisation
- Model-Free Reinforcement Learning
- **Value Function Approximation**
 - **Introduction**
 - **Supervised Objective**
 - **Gradient and Semi-Gradient**
 - **Example 1: MC $\hat{v} \approx v_\pi$ Evaluation**
 - **Example 2: TD $\hat{v} \approx v_\pi$ Evaluation**
 - **Conclusions**
- Policy Gradient Methods

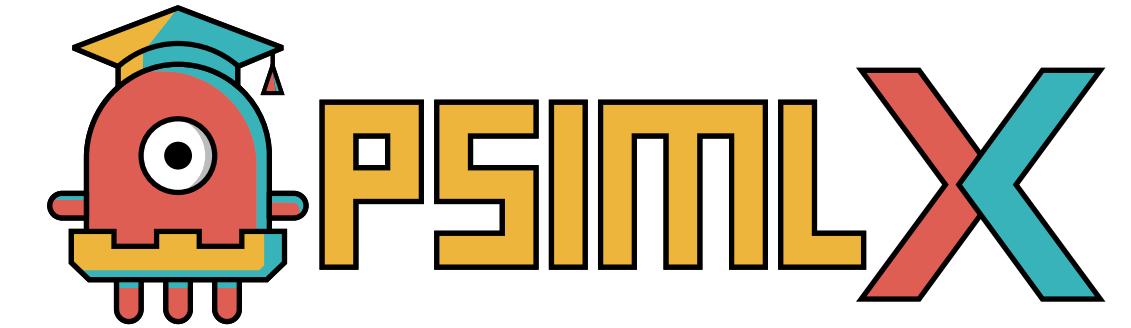
Value Function Approximation

Introduction

- Solutions presented so far were tabular:
 - Every state $s \in S$ has an entry $V(s)$
 - Every state-action pair $s \in S, a \in A$ has an entry $Q(s, a)$ (see slide 67)
- Three main problems:
 - Large (but potentially discrete) state-spaces:
 - Backgammon 10^{20} states
 - Go 10^{70}
 - Continuous state-spaces
 - Physical properties: distances, velocities, angles, ...
 - Robotics applications
 - Continuous actions

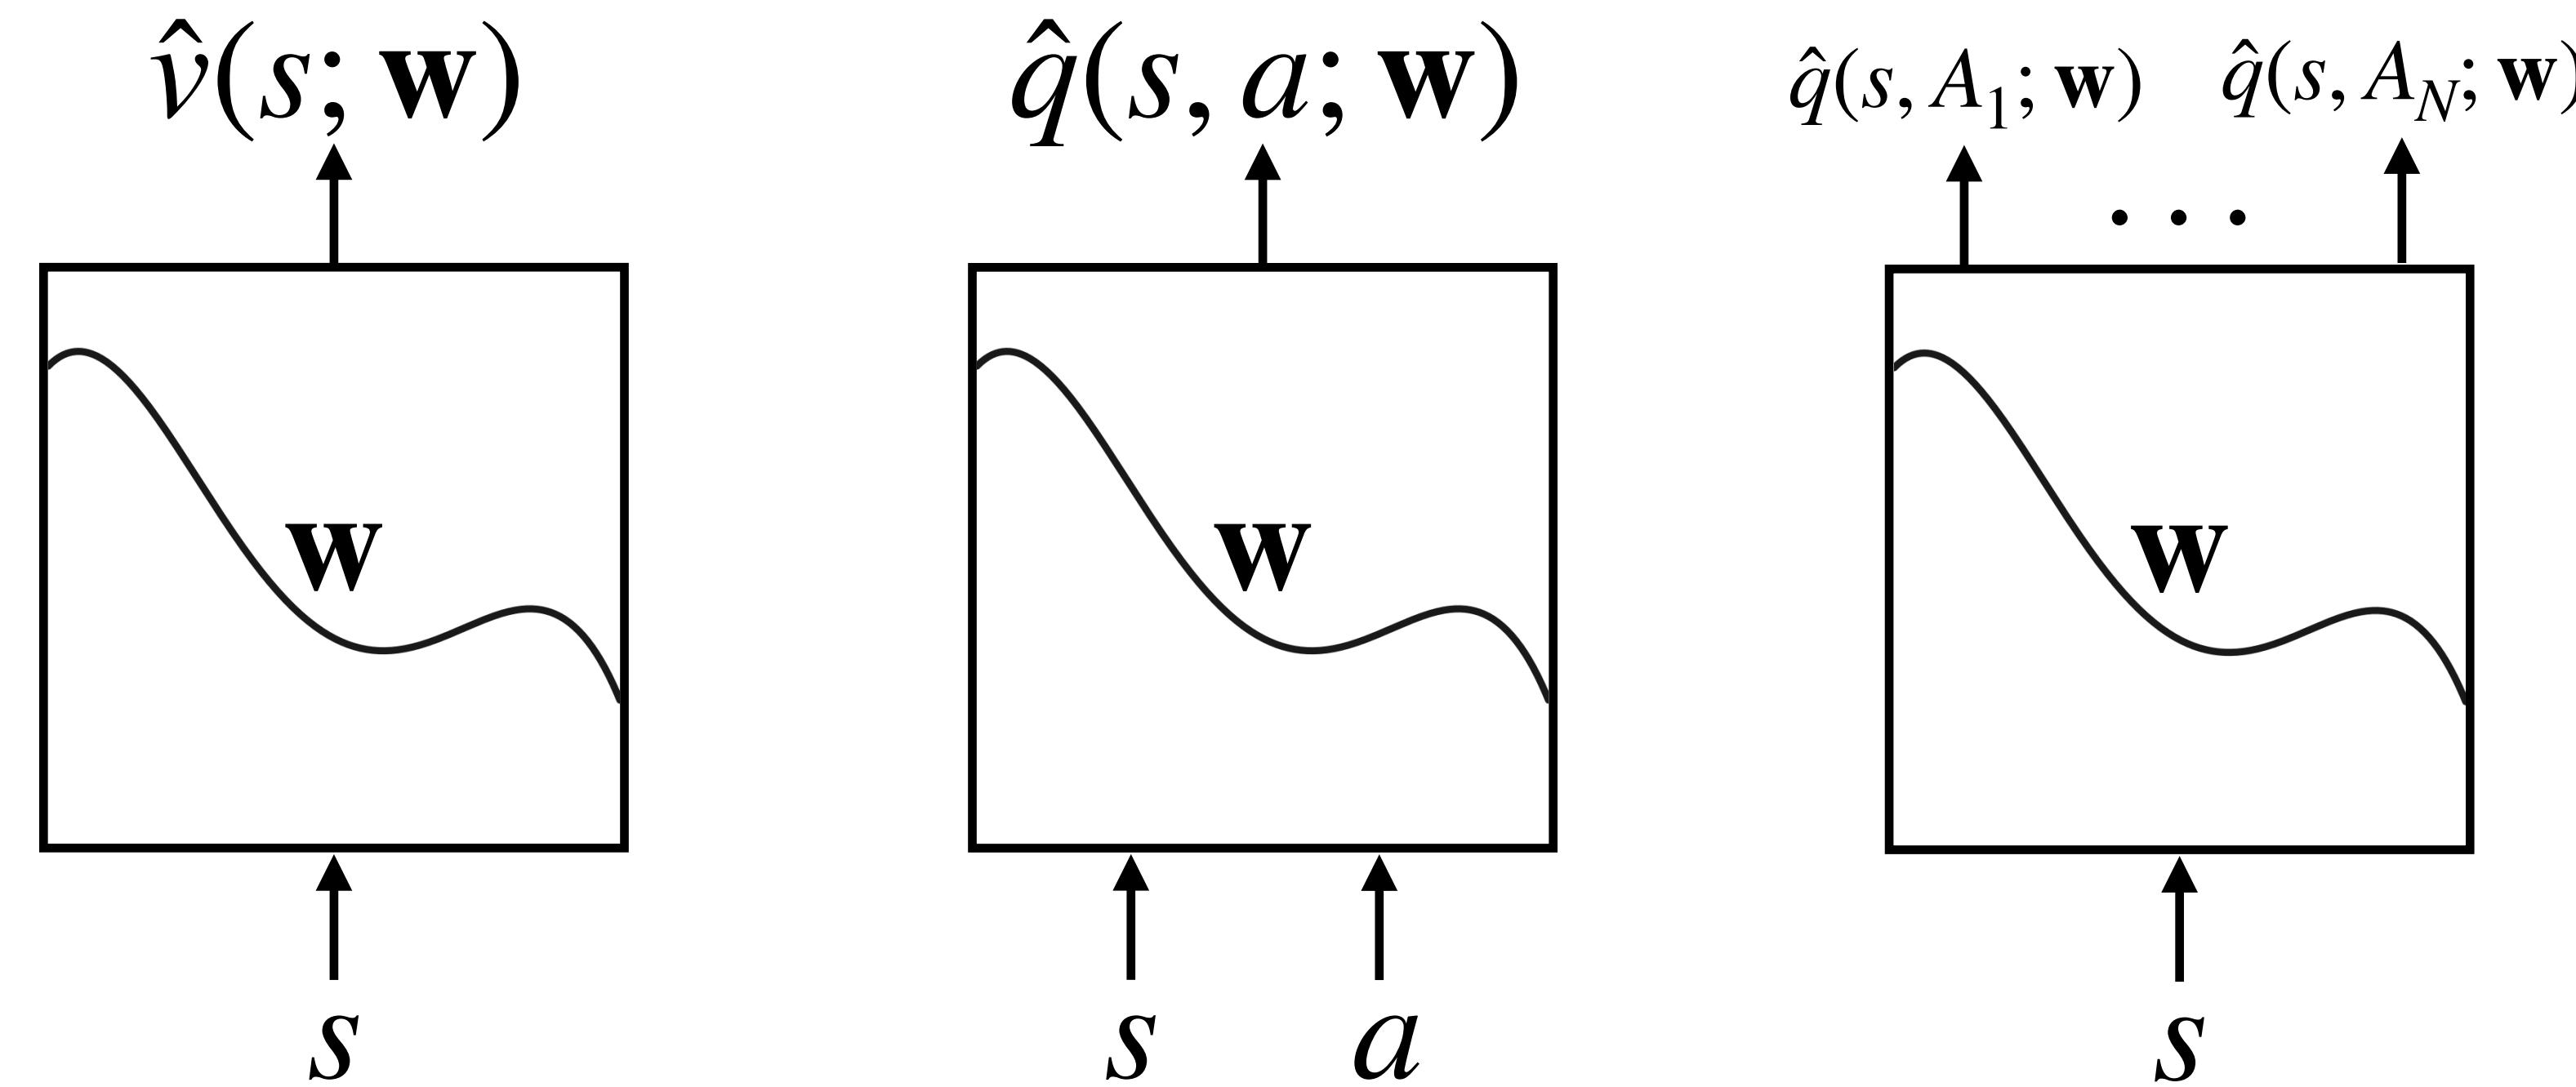


Value Function Approximation



Introduction

- **Idea:** Use supervised learning to train the function approximator
 - Artificial Neural Networks (ANN) + Stochastic Gradient Descent (SGD)



Left: State-Value function approximation for a given state; Middle: Action-Value function approximation for a given state and action pair; Action-Value function approximation for each action for a given state

Value Function Approximation

Supervised Objective

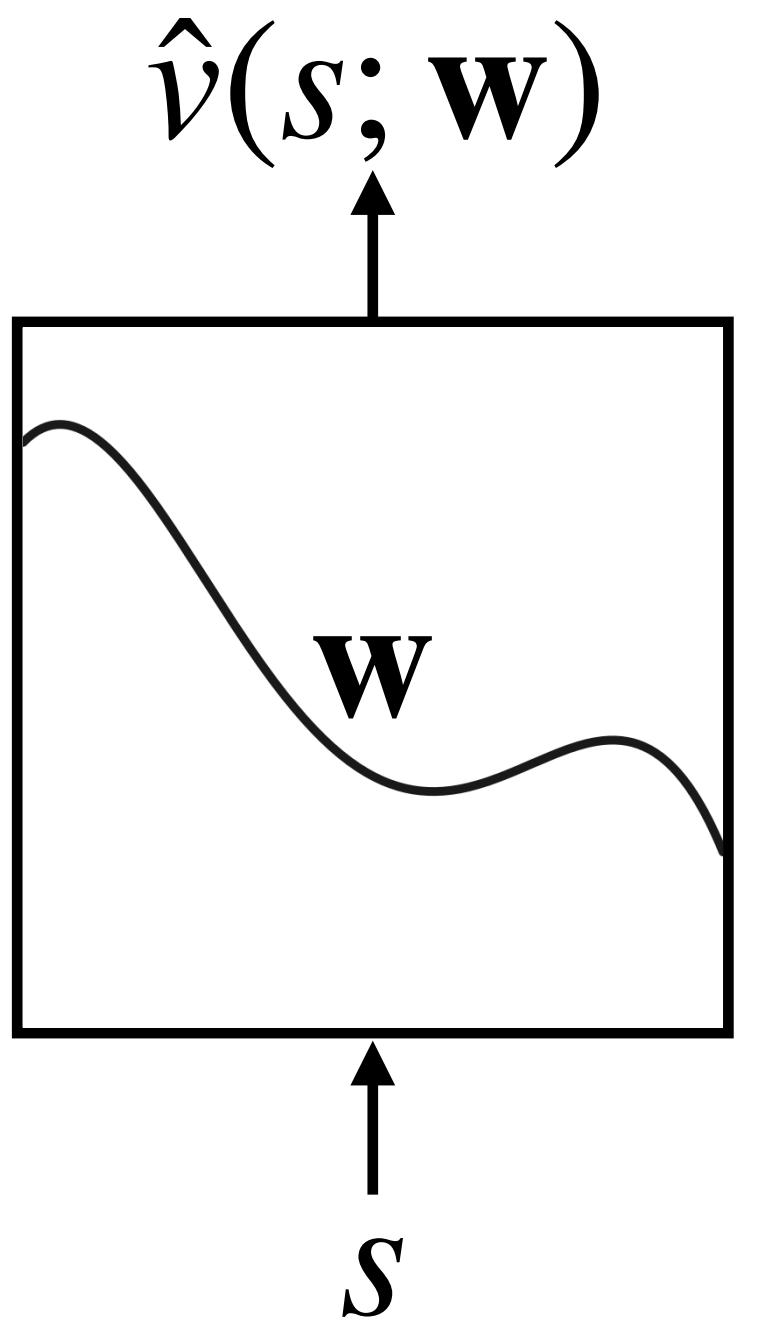
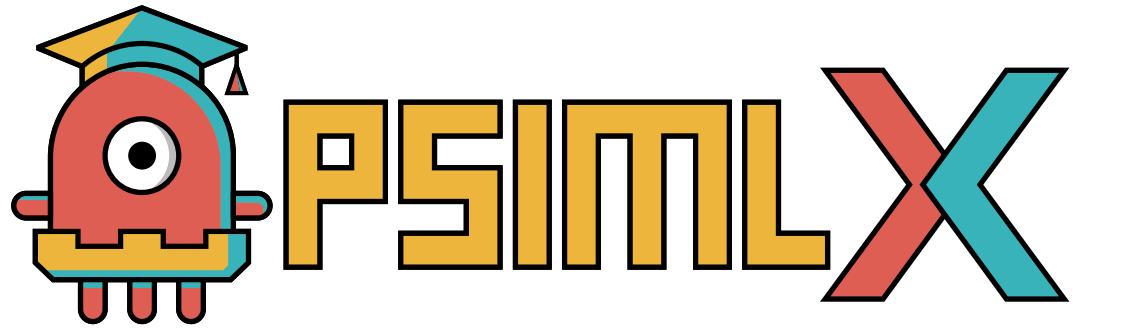
- Use standard Mean Squared Error (MSE) Loss:

$$J(\mathbf{w}) \doteq \sum_{s \in S} \mu(s) [\nu_\pi(s) - \hat{\nu}(s; \mathbf{w})]^2$$

- Scale each error by its importance as captured by the state visitation frequency under policy π :

$$\eta(s) = h(s) + \sum_{\bar{s}} \eta(\bar{s}) \sum_a \pi(a | \bar{s}) p(s | \bar{s}, a)$$

$$\mu(s) = \frac{\eta(s)}{\sum_{s'} \eta(s')} \quad \text{on-policy distribution}$$



Value Function Approximation

Gradient and Semi-Gradient Methods

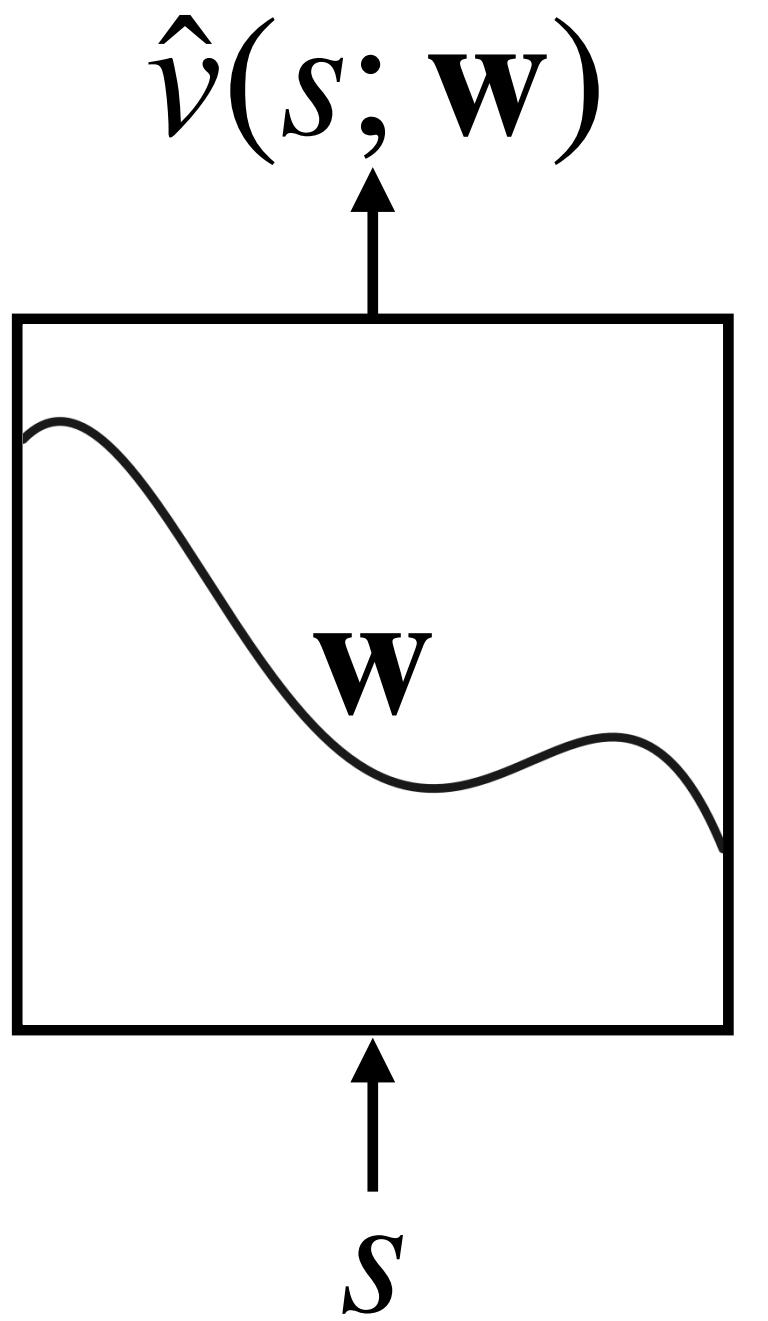
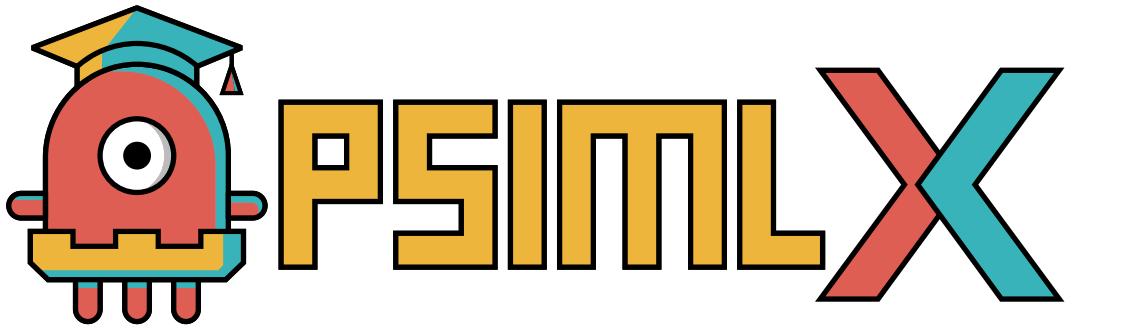
- Use standard Mean Squared Error (MSE) Loss:

$$J(\mathbf{w}) \doteq \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s; \mathbf{w})]^2$$

- Combined with SGD:

$$\begin{aligned} \mathbf{w}_{t+1} &\doteq \mathbf{w}_t - \frac{1}{2}\alpha \nabla [v_\pi(S_t) - \hat{v}(S_t; \mathbf{w}_t)]^2 \\ &= \mathbf{w}_t + \alpha [v_\pi(S_t) - \hat{v}(S_t; \mathbf{w}_t)] \nabla \hat{v}(S_t; \mathbf{w}_t) \end{aligned}$$

$$\nabla f(\mathbf{w}) \doteq \left(\frac{\partial f(\mathbf{w})}{\partial w_1}, \frac{\partial f(\mathbf{w})}{\partial w_2}, \dots, \frac{\partial f(\mathbf{w})}{\partial w_d} \right)^\top$$



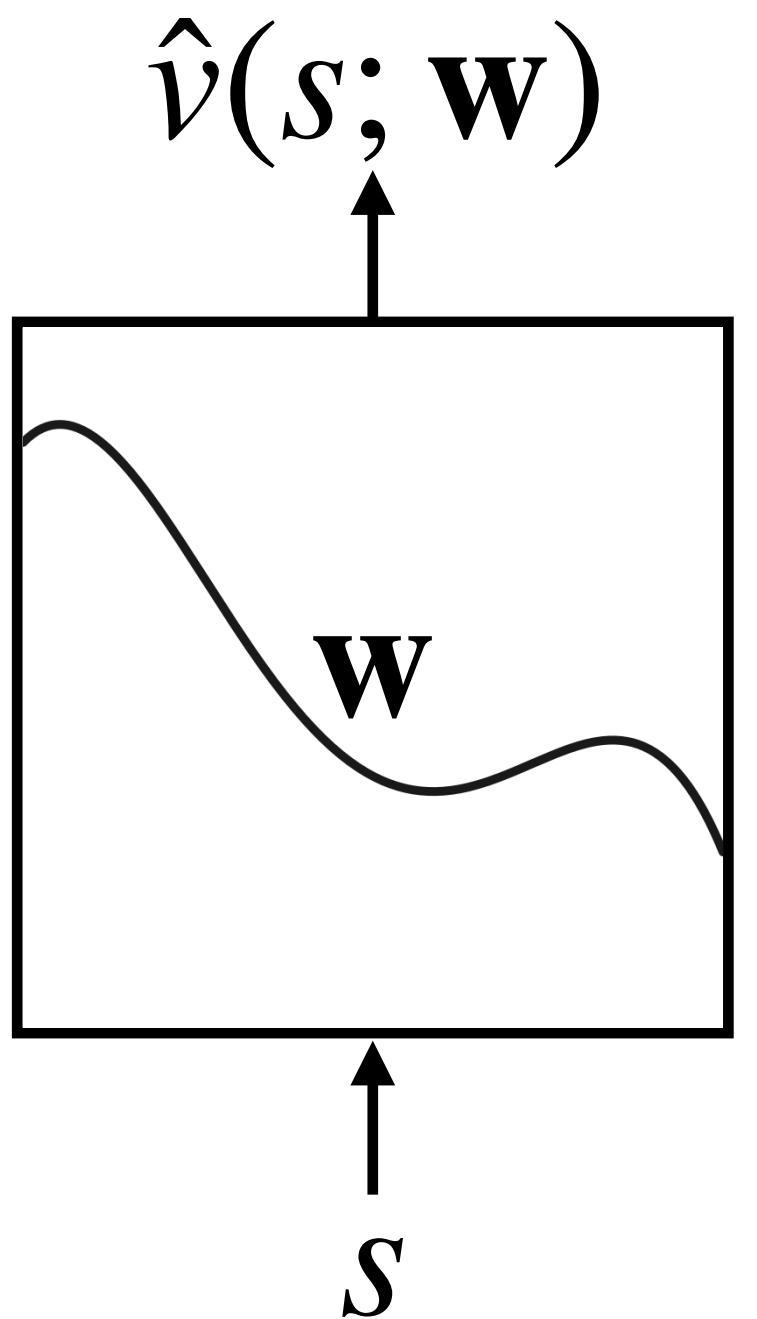
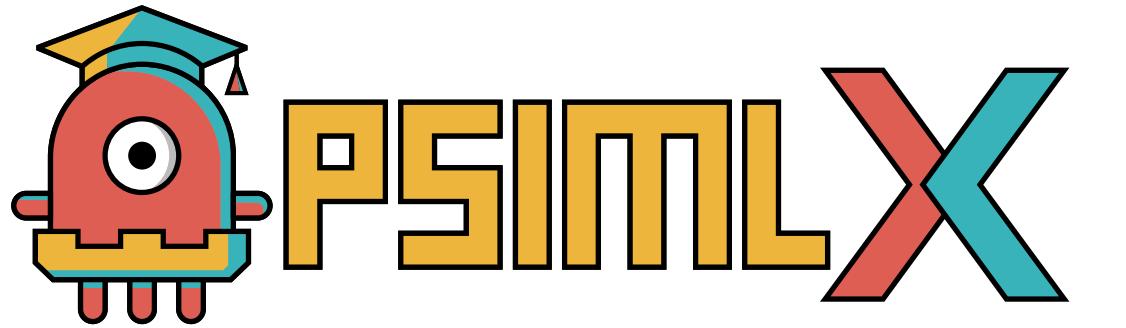
Value Function Approximation

Gradient and Semi-Gradient Methods

- Use standard Mean Squared Error (MSE) Loss:

$$J(\mathbf{w}) \doteq \sum_{s \in S} \mu(s) [\nu_\pi(s) - \hat{\nu}(s; \mathbf{w})]^2$$

- **Problem:** As this is not an actual supervised learning setting, we do not have access to $\nu_\pi(s)$!



Value Function Approximation

Gradient and Semi-Gradient Methods

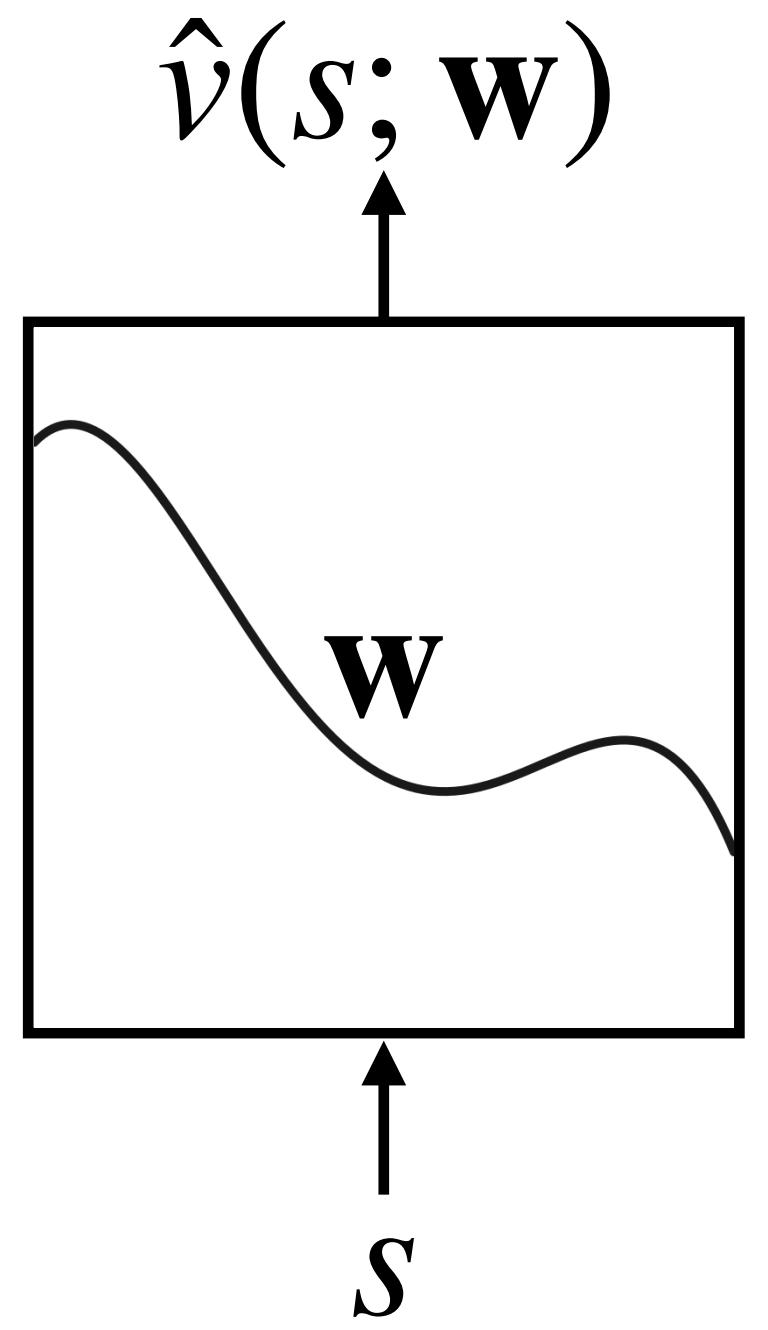
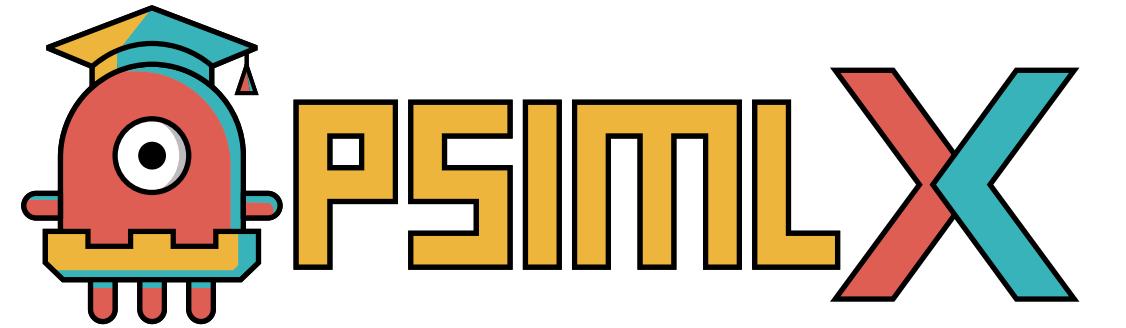
- Use standard Mean Squared Error (MSE) Loss:

$$J(\mathbf{w}) \doteq \sum_{s \in S} \mu(s) [\nu_\pi(s) - \hat{\nu}(s; \mathbf{w})]^2$$

- **Problem:** As this is not an actual supervised learning setting, we do not have access to $\nu_\pi(s)$!

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{\nu}(S_t; \mathbf{w}_t)] \nabla \hat{\nu}(S_t; \mathbf{w}_t)$$

$$U_t = \begin{cases} G_t & \text{MC approach – true gradient} \\ R_{t+1} + \gamma \hat{\nu}(S_{t+1}; \mathbf{w}_t) & \text{TD approach – semi-gradient} \end{cases}$$



Value Function Approximation

Example 1: MC $\hat{v} \approx v_\pi$ Evaluation



Gradient Monte Carlo Algorithm for Estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameter: step size $\alpha > 0$

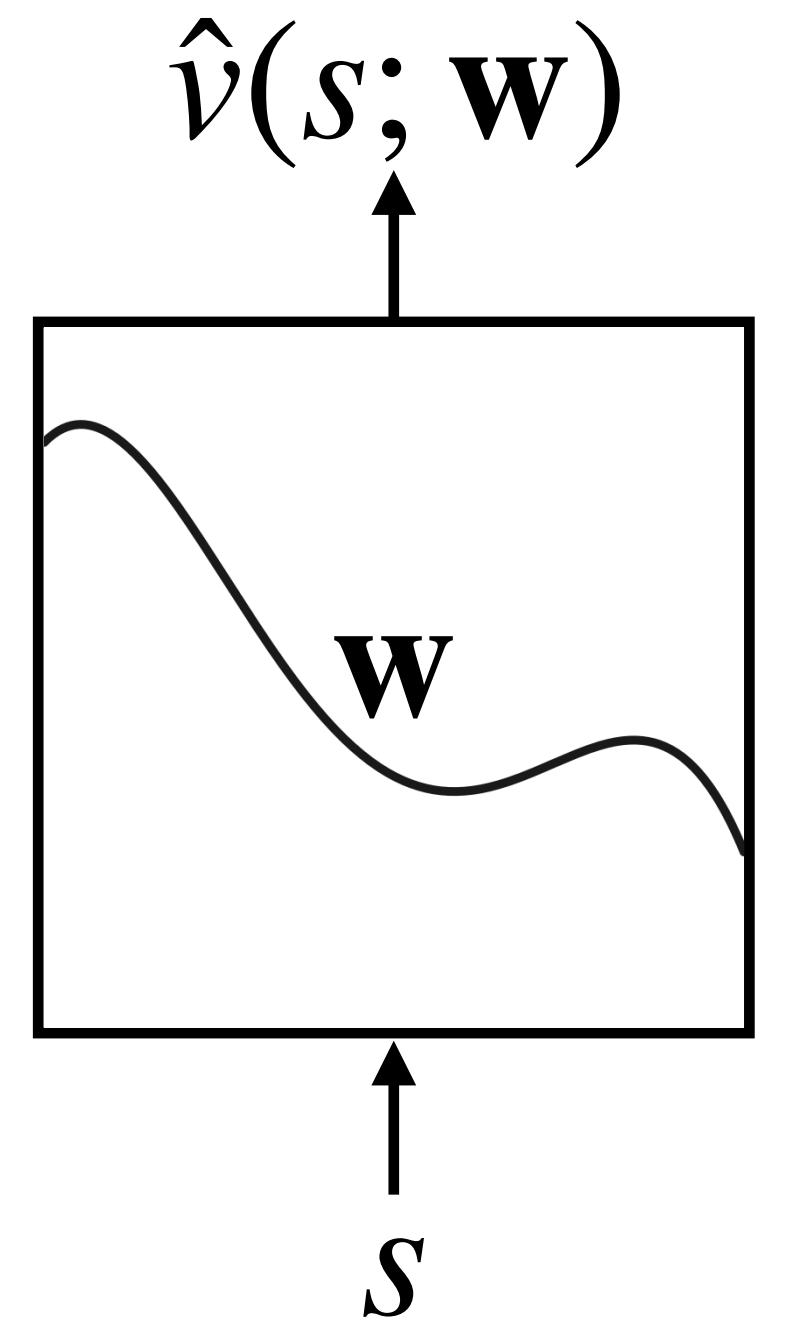
Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop forever (for each episode):

Generate an episode $S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T$ using π

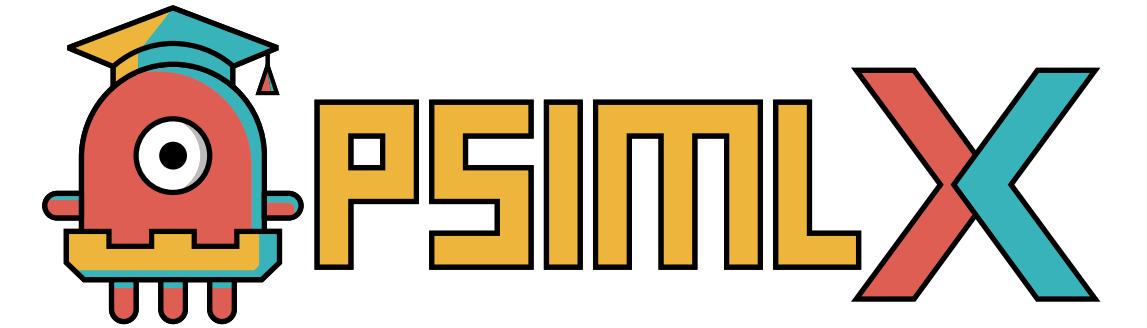
Loop for each step of episode, $t = 0, 1, \dots, T - 1$:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$



Value Function Approximation

Example 2: TD $\hat{v} \approx v_\pi$ Evaluation



Semi-gradient TD(0) for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Algorithm parameter: step size $\alpha > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

 Initialize S

 Loop for each step of episode:

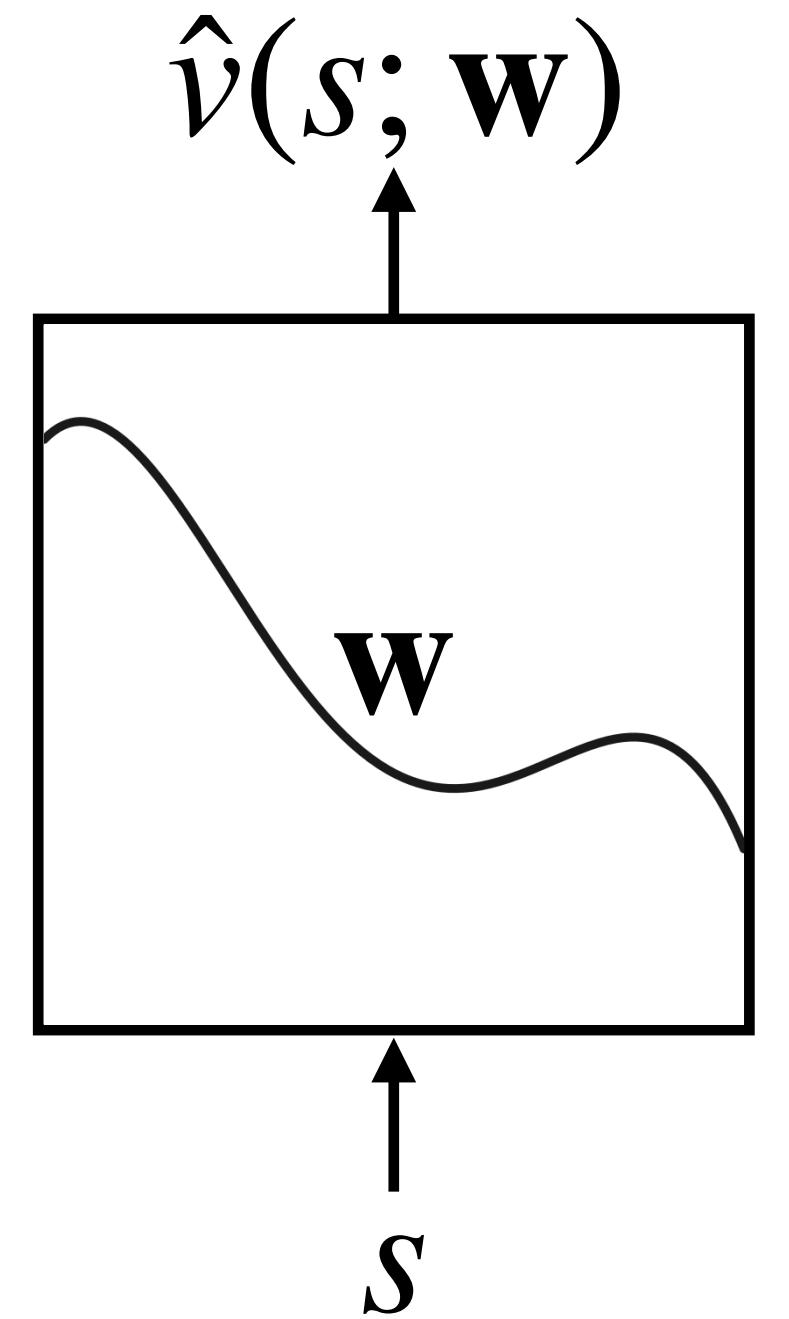
 Choose $A \sim \pi(\cdot | S)$

 Take action A , observe R, S'

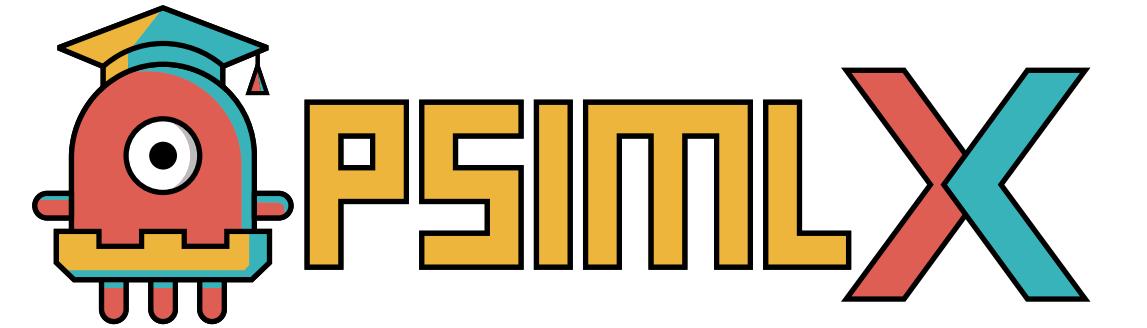
$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$$

$S \leftarrow S'$

 until S is terminal

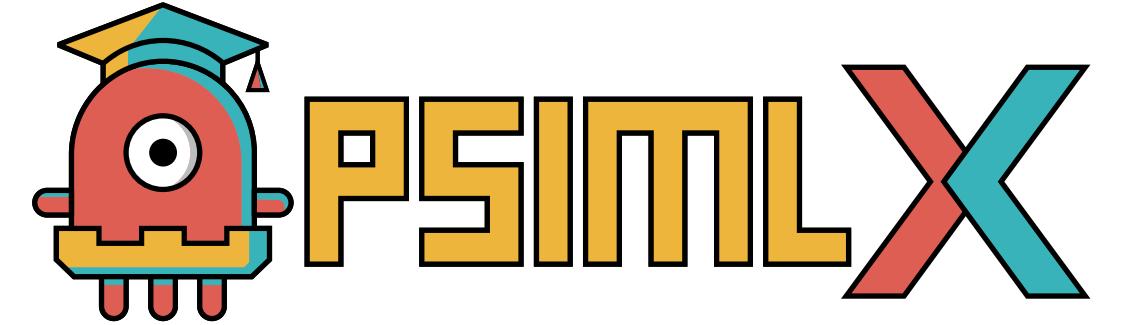


Outline



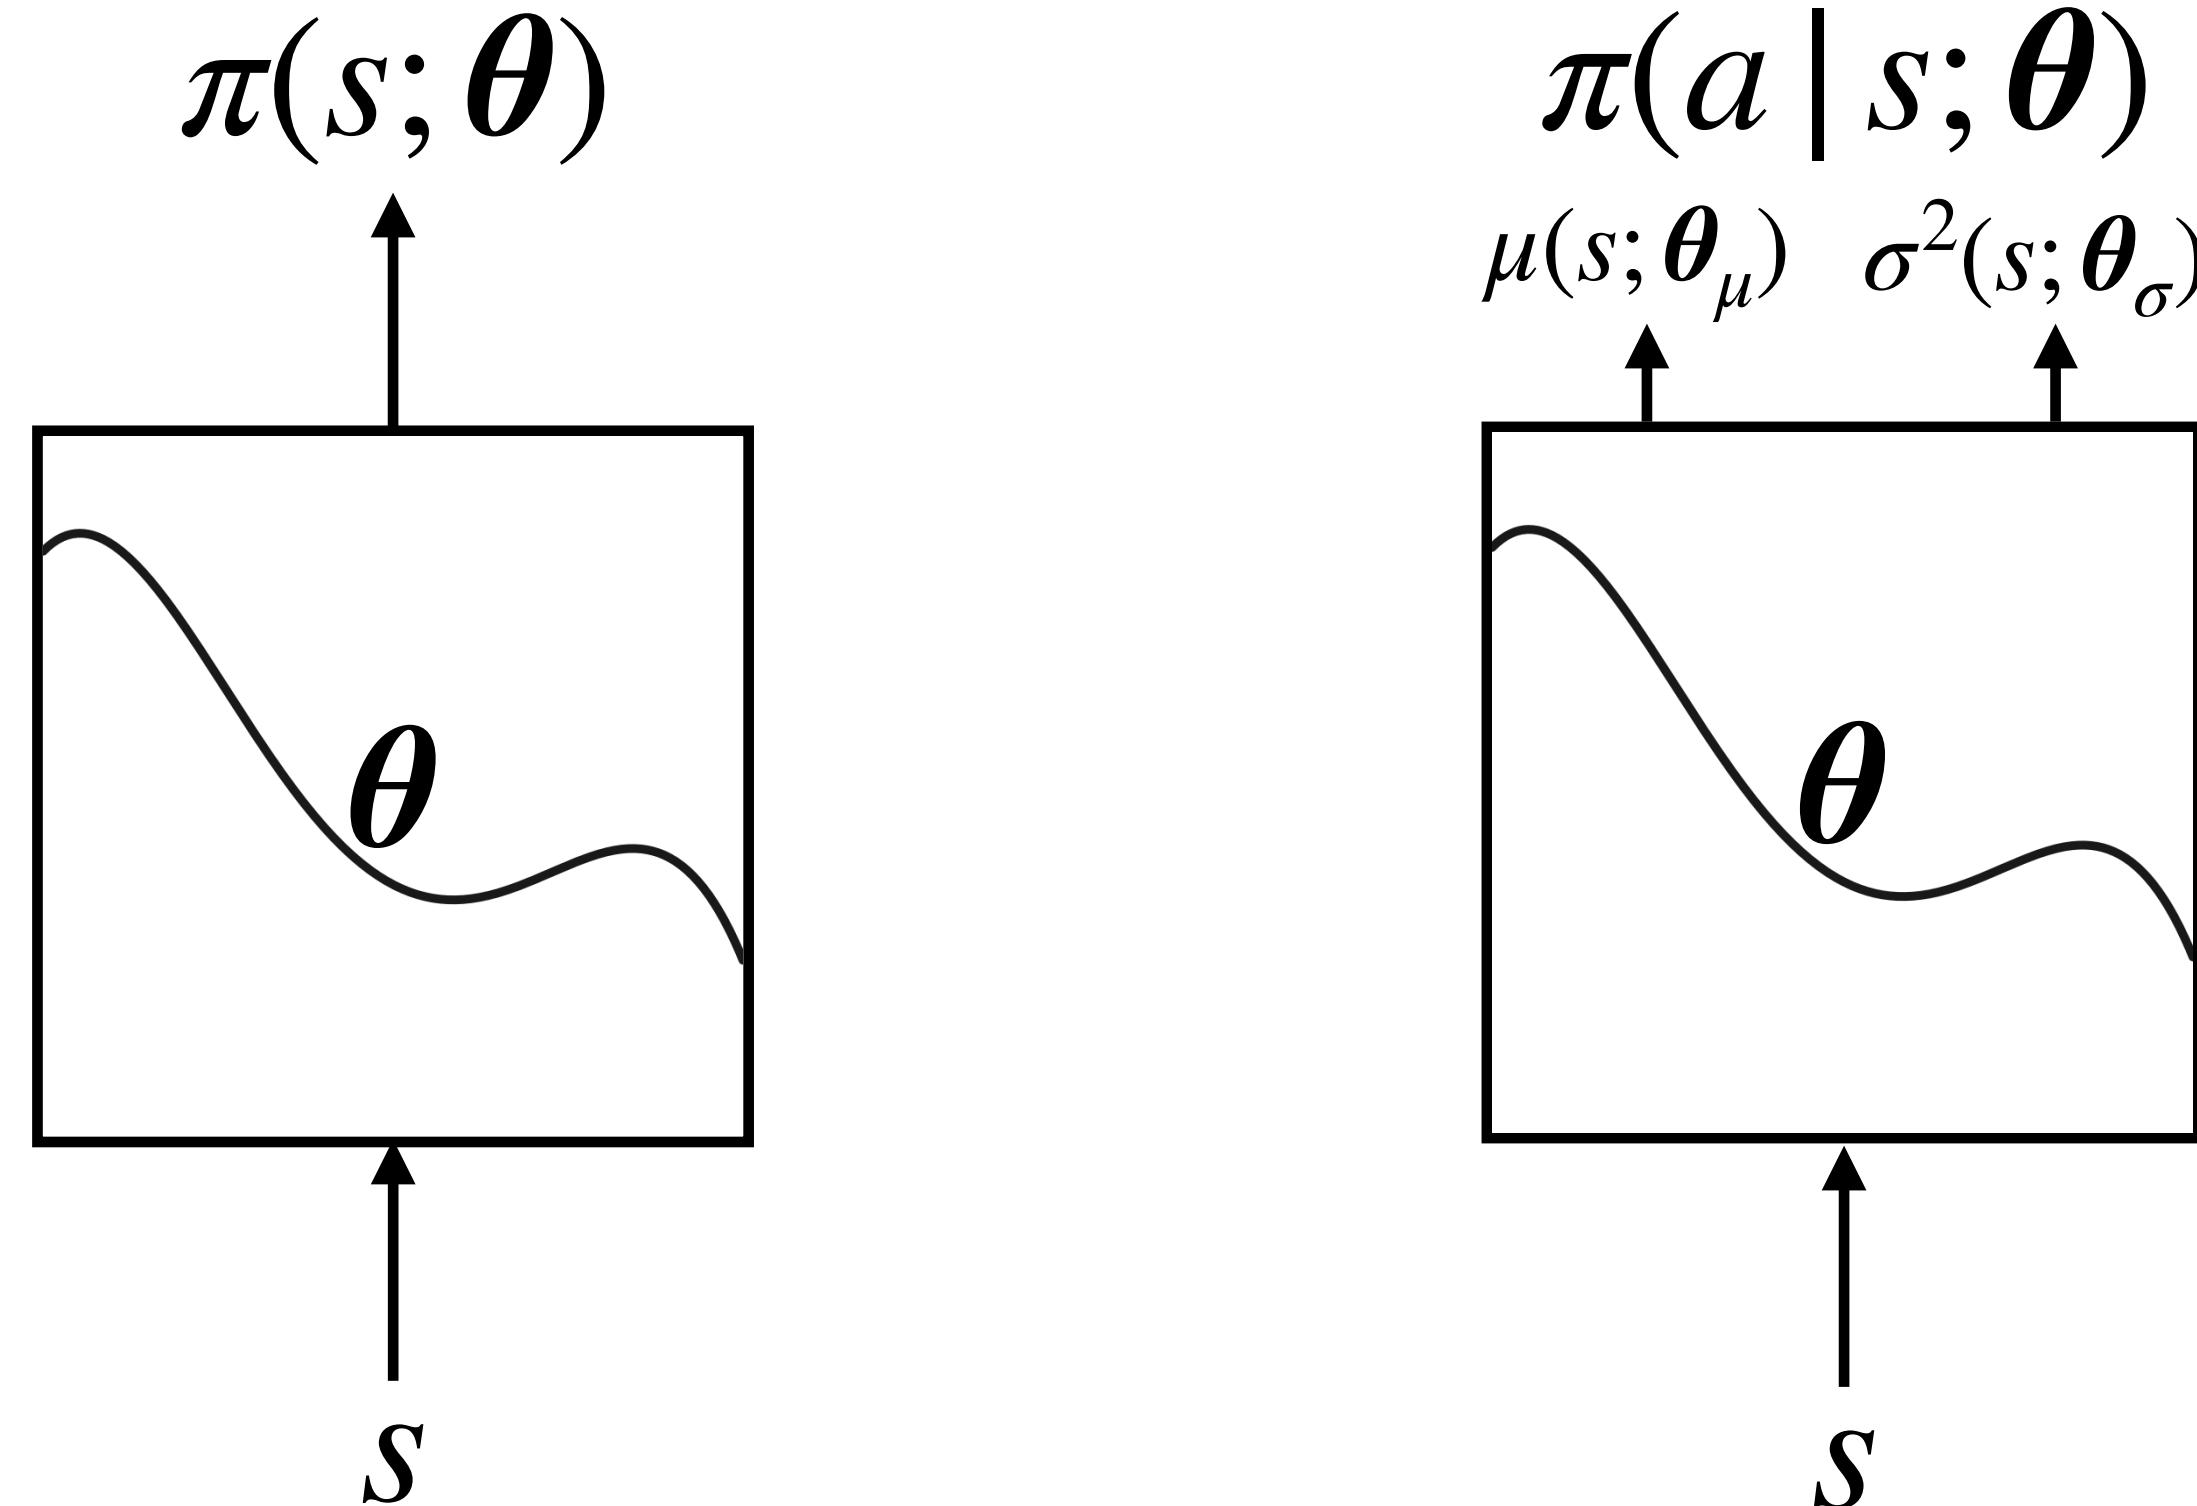
- Introduction
- Reinforcement Learning Formalisation
- Model-Free Reinforcement Learning
- Interlude: RL Taxonomy
- Value Function Approximation
- **Policy Gradient Methods**
 - **Introduction**
 - **The Policy Gradient Theorem: Statement**
 - **The Policy Gradient Theorem: Derivation**
 - **REINFORCE Algorithm**
 - **Actor-Critic Methods**
 - **Policy Parametrisation for Continuous Actions**
 - **Conclusions**

Policy Gradient Methods

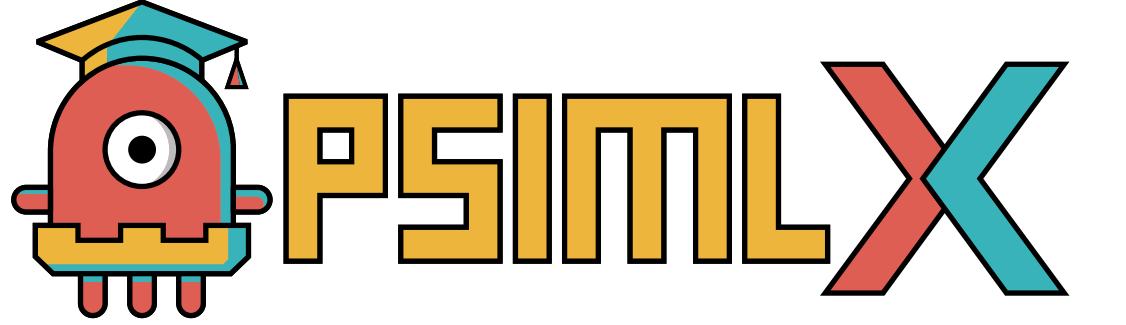


Introduction

- So far policies were **implicit** – we modelled state value functions; policy followed states with high values
- **Idea:** Explicitly model the policy with an ANN



Left: Deterministic policy that produces an action for a given state; Right: Stochastic policy that produces a distribution over actions given the state



Policy Gradient Methods

The Policy Gradient Theorem: Statement

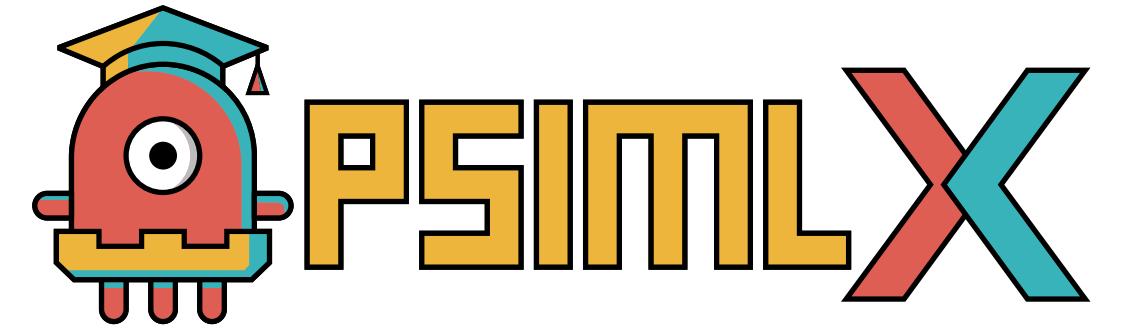
- We wish to maximise the performance under policy π parameterised by θ over an entire episode:

$$J(\theta) \doteq v_{\pi_\theta}(s_0)$$

$$\nabla J(\theta) = \mathbb{E}_{s \sim \mu_\pi, a \sim \pi} [q_\pi(s, a) \nabla \ln \pi(a | s; \theta)]$$

- We can now improve performance using the gradient of the policy represented as an ANN, but we still have $q_\pi(s, a)$:
 - We can explicitly model $q_\pi(s, a) \approx \hat{q}(s, a; \mathbf{w})$
 - We can replace it with the return G_t as $\mathbb{E}_\pi[G_t | S_t = s, A_t = a] = q_\pi(S_t, A_t)$:

$$\nabla J(\theta) = \mathbb{E}_\pi [G_t \nabla \ln \pi(A_t | S_t; \theta)] \quad \text{REINFORCE}$$



Policy Gradient Methods

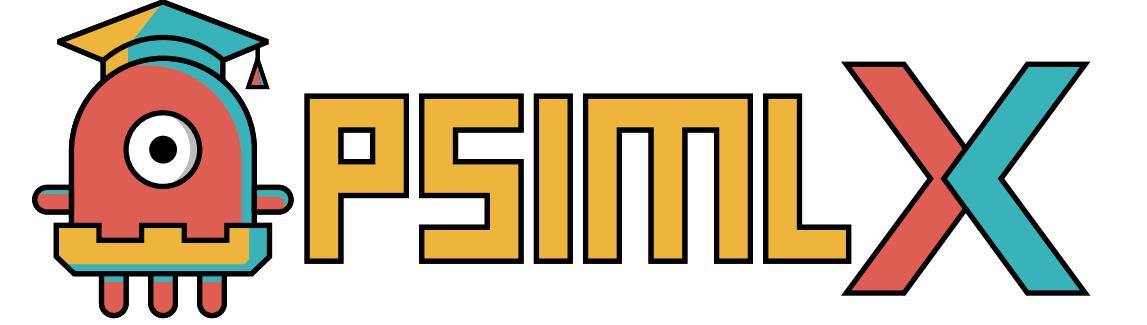
The Policy Gradient Theorem: Derivation

$$\nabla v_\pi(s) = \nabla \left[\sum_a \pi(a | s; \theta) q_\pi(s, a) \right] \text{ (slide 23)}$$

$$= \sum_a [\nabla \pi(a | s; \theta) q_\pi(s, a) + \pi(a | s; \theta) \nabla q_\pi(s, a)] \text{ (derivative product rule)}$$

$$= \sum_a \left[\nabla \pi(a | s; \theta) q_\pi(s, a) + \pi(a | s; \theta) \nabla \left[r(s, a) + \sum_{s'} p(s' | s, a) v_\pi(s') \right] \right] \text{ (slide 23)}$$

$$= \sum_a \left[\nabla \pi(a | s; \theta) q_\pi(s, a) + \pi(a | s; \theta) \sum_{s'} p(s' | s, a) \nabla v_\pi(s') \right] \text{ (recursion)}$$



Policy Gradient Methods

The Policy Gradient Theorem: Derivation

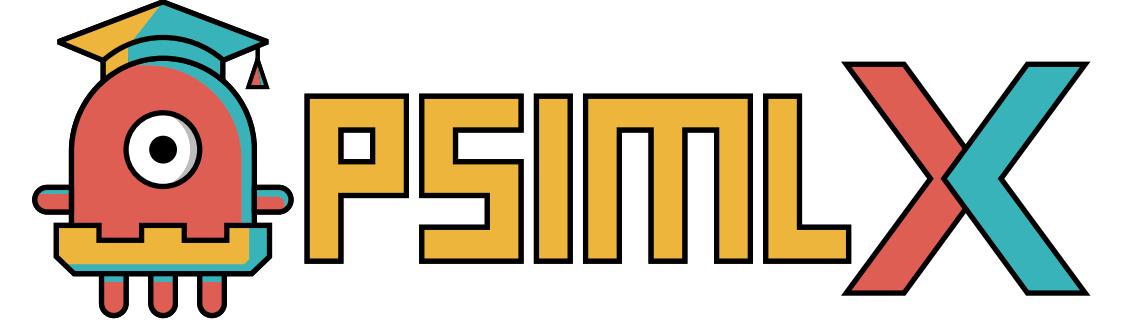
$$\nabla v_\pi(s) = \nabla \left[\sum_a \pi(a | s; \theta) q_\pi(s, a) \right] \text{ (slide 22)}$$

$$= \sum_a [\nabla \pi(a | s; \theta) q_\pi(s, a) + \pi(a | s; \theta) \nabla q_\pi(s, a)] \text{ (derivative product rule)}$$

$$= \sum_a \left[\nabla \pi(a | s; \theta) q_\pi(s, a) + \pi(a | s; \theta) \nabla \left[r(s, a) + \sum_{s'} p(s' | s, a) v_\pi(s') \right] \right] \text{ (slide 22)}$$

$$= \sum_a \left[\nabla \pi(a | s; \theta) q_\pi(s, a) + \pi(a | s; \theta) \sum_{s'} p(s' | s, a) \boxed{\nabla v_\pi(s')} \right] \text{ (recursion)}$$

$$= \sum_{x \in S} \sum_{k=0}^{\infty} P(s \rightarrow x, k, \pi) \sum_a \nabla \pi(a | x; \theta) q_\pi(x, a) \text{ (unroll recursion)}$$



Policy Gradient Methods

The Policy Gradient Theorem: Derivation

$$\nabla J(\theta) = \nabla v_\pi(s_0)$$

$$= \sum_s \left(\sum_{k=0}^{\infty} P(s_0 \rightarrow s, k, \pi) \right) \sum_a \nabla \pi(a | s; \theta) q_\pi(s, a)$$

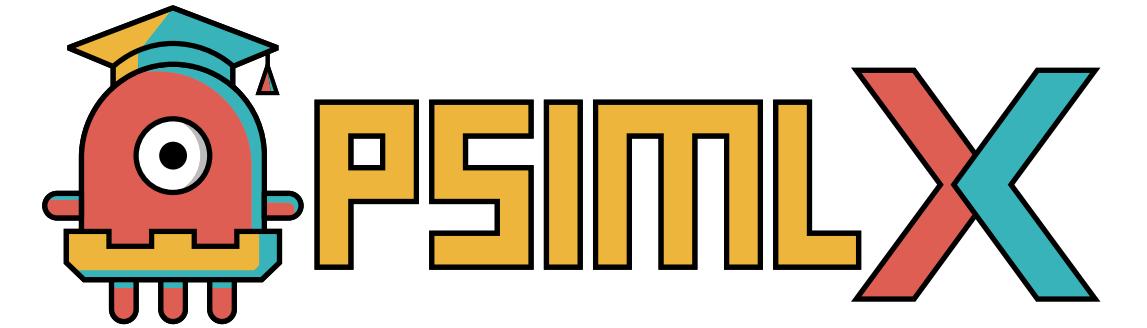
$$= \sum_s \eta_\pi(s) \sum_a \nabla \pi(a | s; \theta) q_\pi(s, a) \quad (\text{slide 86})$$

$$= \sum_{s'} \eta_\pi(s') \sum_s \mu_\pi(s) \sum_a \nabla \pi(a | s; \theta) q_\pi(s, a) \quad (\text{slide 86})$$

$$\propto \sum_s \mu_\pi(s) \sum_a \nabla \pi(a | s; \theta) q_\pi(s, a)$$

Policy Gradient Methods

The Policy Gradient Theorem: Derivation



$$\nabla J(\theta) = \nabla v_{\pi}(s_0)$$

$$\propto \sum_s \mu_{\pi}(s) \sum_a \nabla \pi(a | s; \theta) q_{\pi}(s, a)$$

$$= \sum_s \mu_{\pi}(s) \sum_a \pi(a | s; \theta) q_{\pi}(s, a) \frac{\nabla \pi(a | s; \theta)}{\pi(a | s; \theta)} \text{ and } \nabla \ln \pi(a | s; \theta) = \frac{1}{\pi(a | s; \theta)} \cdot \nabla \pi(a | s; \theta)$$

$$= \sum_s \mu_{\pi}(s) \sum_a \pi(a | s) q_{\pi}(s, a) \nabla \ln \pi(a | s; \theta)$$

$$= \mathbb{E}_{s \sim \mu_{\pi}, a \sim \pi} [q_{\pi}(s, a) \nabla \ln \pi(a | s; \theta)]$$

*grad log derivative trick
or
eligibility vector*

$$J(\theta) \doteq v_{\pi_{\theta}}(s_0)$$

$$\nabla J(\theta) = \mathbb{E}_{s \sim \mu_{\pi}, a \sim \pi} [q_{\pi}(s, a) \nabla \ln \pi(a | s; \theta)]$$

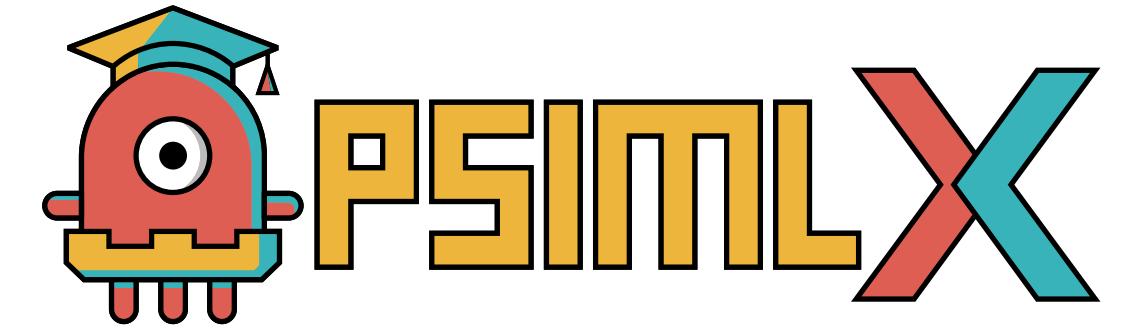
Policy Gradient Methods

REINFORCE Algorithm

$$\nabla J(\boldsymbol{\theta}) = \mathbb{E}_{\pi} [q_{\pi}(s, a) \nabla \ln \pi(a | s; \boldsymbol{\theta})]$$

$$= \mathbb{E}_{\pi} [G_t \nabla \ln \pi(A_t | S_t; \boldsymbol{\theta})] \text{ as } q_{\pi} \doteq \mathbb{E}_{\pi} [G_t | A_t = a, S_t = s] \text{ (Slide 22)}$$

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha G_t \nabla \ln \pi(A_t | S_t; \boldsymbol{\theta}_t) \quad \text{Gradient Ascent Step}$$



REINFORCE: Monte-Carlo Policy-Gradient Control (episodic) for π_*

Input: a differentiable policy parameterization $\pi(a|s, \boldsymbol{\theta})$

Algorithm parameter: step size $\alpha > 0$

Initialize policy parameter $\boldsymbol{\theta} \in \mathbb{R}^{d'}$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|s, \boldsymbol{\theta})$

Loop for each step of the episode $t = 0, 1, \dots, T - 1$:

$$\begin{aligned} G &\leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k & (G_t) \\ \boldsymbol{\theta} &\leftarrow \boldsymbol{\theta} + \alpha \gamma^t G \nabla \ln \pi(A_t | S_t, \boldsymbol{\theta}) \end{aligned}$$

Policy Gradient Methods

Actor-Critic Methods

- Speed-up learning and reduce variance by utilising bootstrapping
- Use $\hat{q}(s, a; \mathbf{w})$ or $\hat{v}(s; \mathbf{w})$ to estimate the TD residual δ_t

$$\begin{aligned}\theta_{t+1} &\doteq \theta_t + \alpha \left(R_{t+1} + \gamma \hat{v}(S_{t+1}; \mathbf{w}) - \hat{v}(S_t; \mathbf{w}) \right) \nabla \ln \pi(A_t | S_t; \theta_t) \\ &= \theta_t + \alpha \delta_t \ln \pi(A_t | S_t; \theta_t)\end{aligned}$$

One-step Actor–Critic (episodic), for estimating $\pi_\theta \approx \pi_*$

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization $\hat{v}(s, \mathbf{w})$

Parameters: step sizes $\alpha^\theta > 0$, $\alpha^\mathbf{w} > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ and state-value weights $\mathbf{w} \in \mathbb{R}^d$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

 Initialize S (first state of episode)

$I \leftarrow 1$

 Loop while S is not terminal (for each time step):

$A \sim \pi(\cdot | S, \theta)$

 Take action A , observe S', R

$\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$ (if S' is terminal, then $\hat{v}(S', \mathbf{w}) \doteq 0$)

$\mathbf{w} \leftarrow \mathbf{w} + \alpha^\mathbf{w} \delta \nabla \hat{v}(S, \mathbf{w})$

$\theta \leftarrow \theta + \alpha^\theta I \delta \nabla \ln \pi(A | S, \theta)$

$I \leftarrow \gamma I$

$S \leftarrow S'$



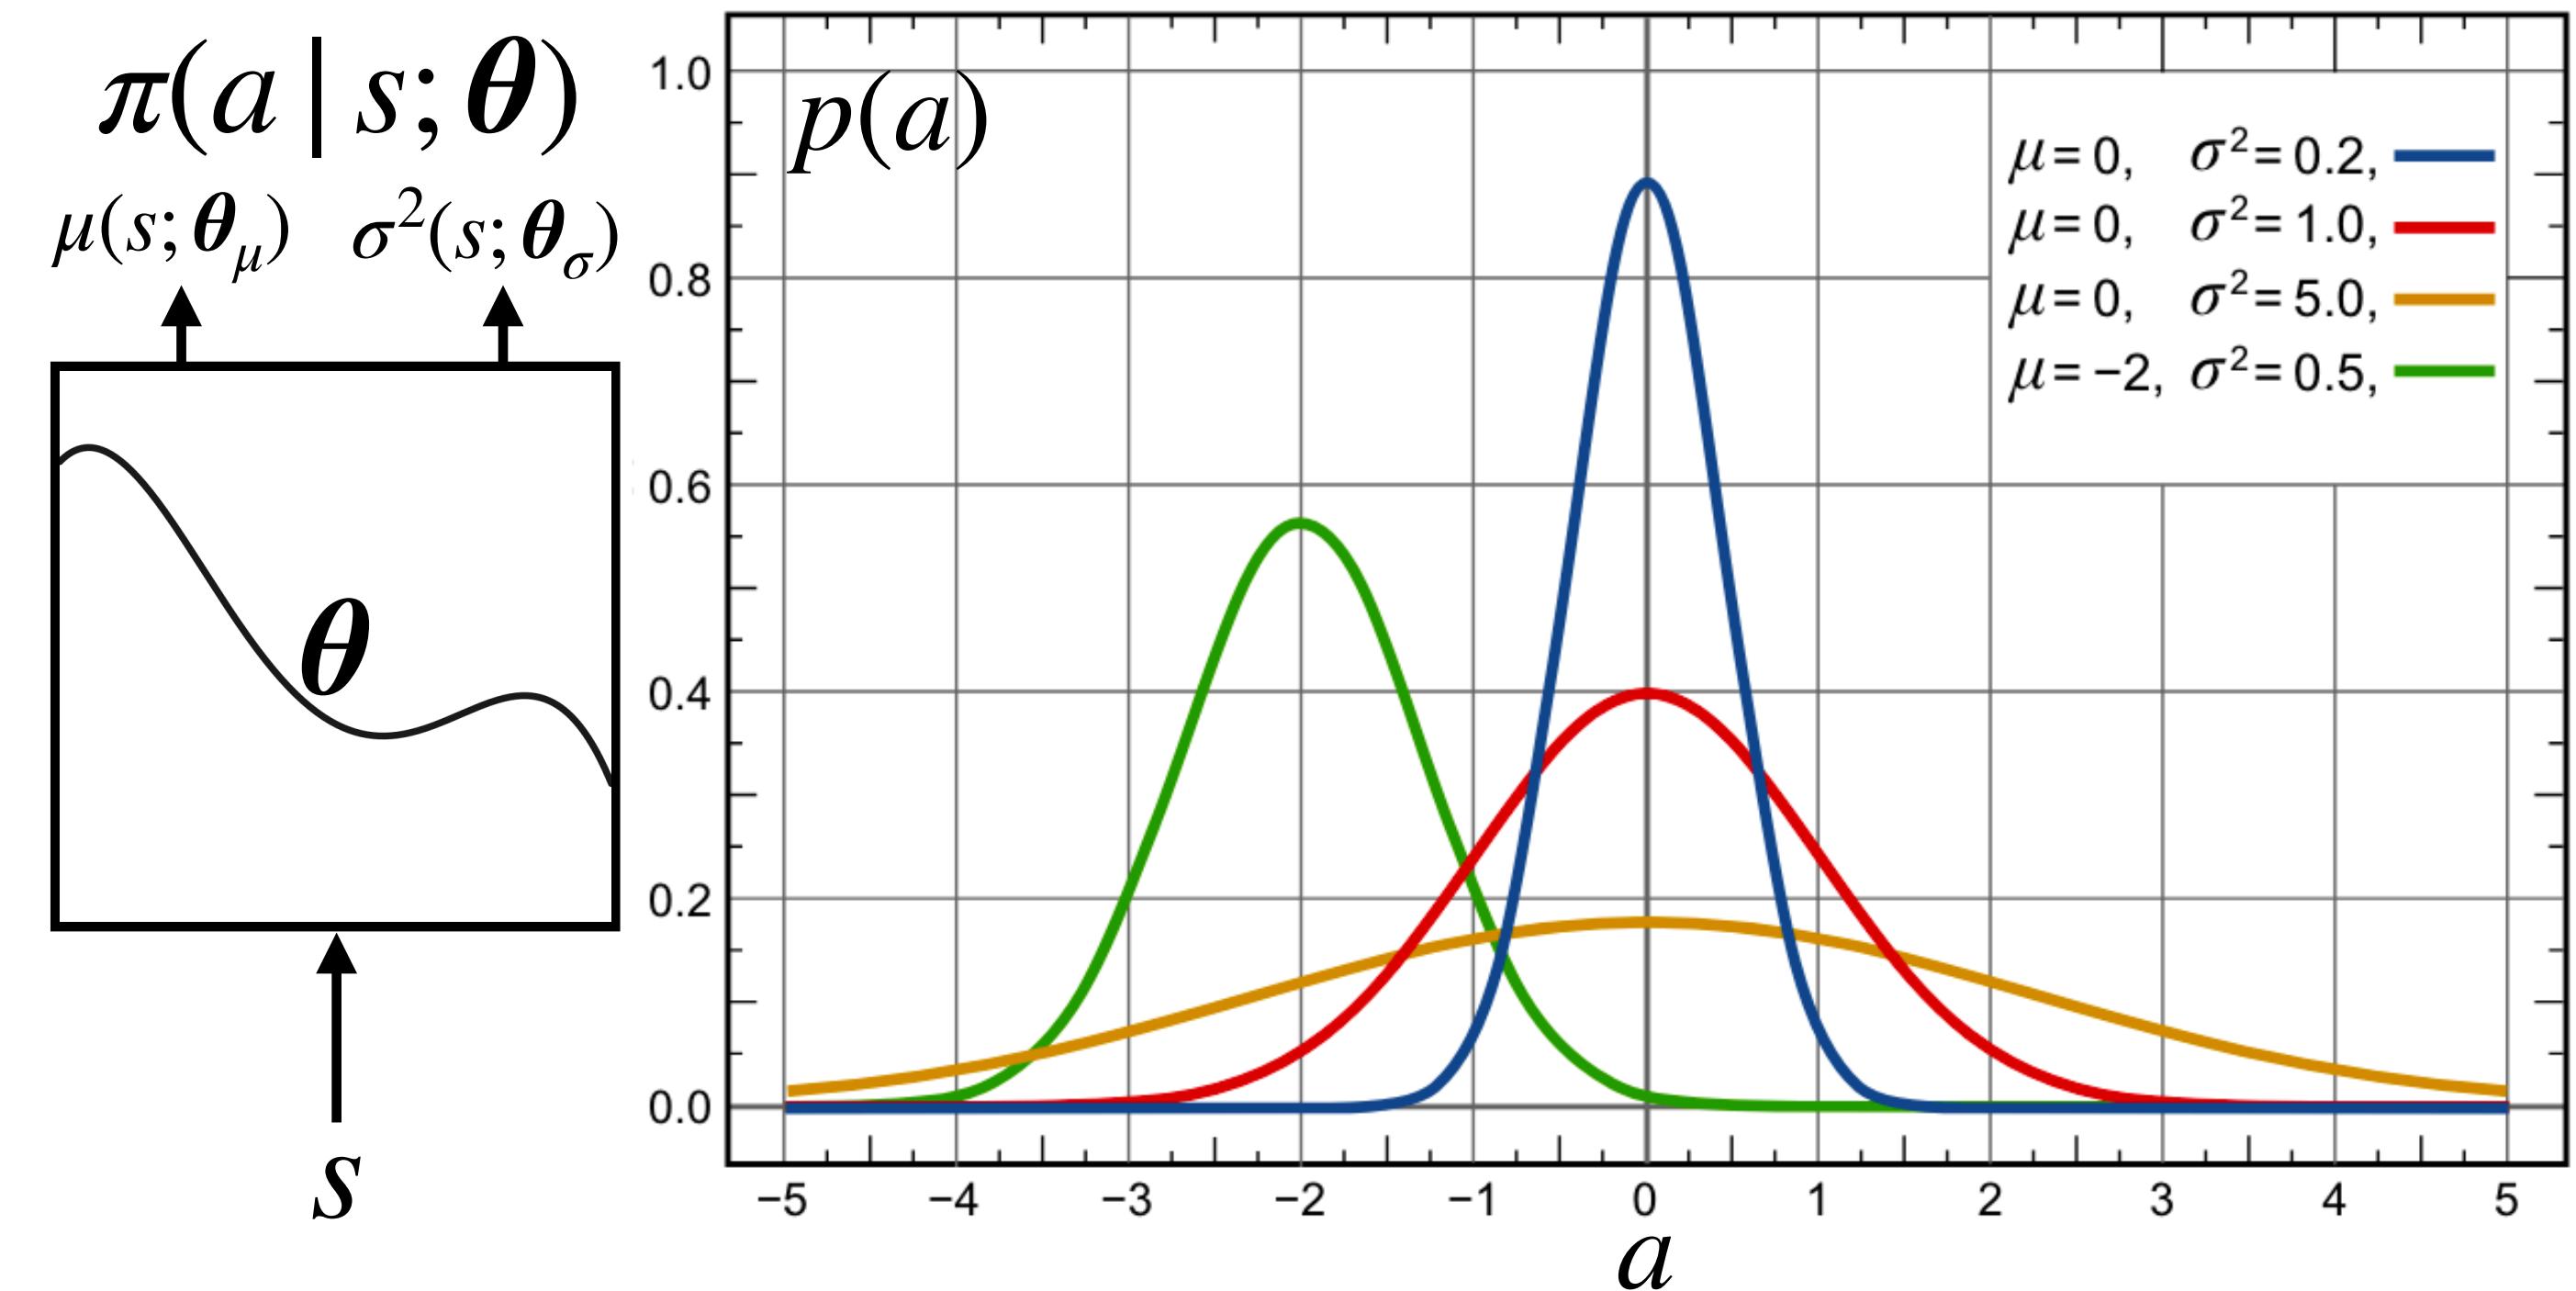
Policy Gradient Methods

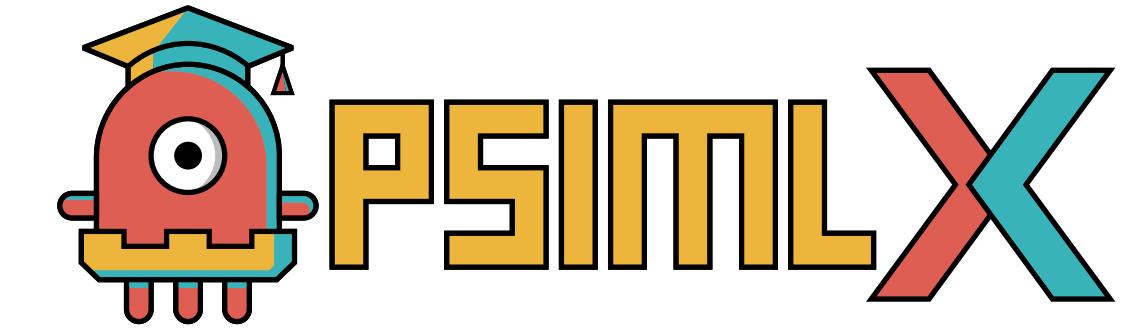
Policy Parametrisation for Continuous Actions

- Define policy as the Gaussian probability density over the real-valued actions
- Use function approximation for $\mu(s; \theta_\mu)$ and $\sigma^2(s; \theta_\sigma)$, with potentially the same feature extractor base $\mathbf{x}(s)$
- We can either learn the variance, or keep it fixed to ensure sufficient exploration throughout learning

$$\pi(a | s; \theta) \doteq \frac{1}{\sigma(s; \theta_\sigma)\sqrt{2\pi}} e^{\left(-\frac{(a - \mu(s; \theta_\mu))^2}{2\sigma(s; \theta_\sigma)^2}\right)}, \theta = [\theta_\mu, \theta_\sigma]^\top$$

$$\mu(s, \theta_\mu) \doteq \theta_\mu^\top \mathbf{x}(s), \sigma(s, \theta_\sigma) \doteq e^{(\theta_\sigma^\top \mathbf{x}(s))}$$





Policy Gradient Methods

Policy Gradient (Monte Carlo) vs TD Learning

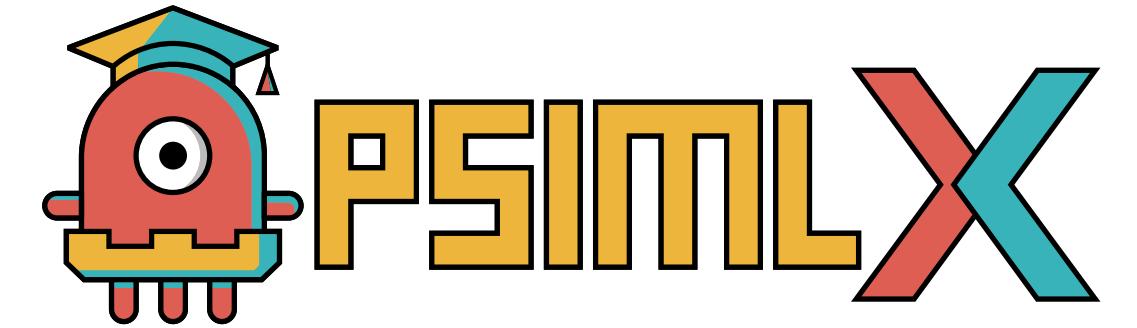
Criteria	Policy Gradient	Temporal Difference Methods
Bias vs Variance	low bias, high variance	high bias, low variance
Online		✓
Bootstrapping		✓ because v_t is based off of v_{t+1}
Estimation	✓	✓
On-Policy	✓	✓
Off-Policy	✓	✓
Exploration vs Exploitation	may be naturally handled	not naturally handled
Past vs Future Future data	past experiences	future (models MDP)
Convergence speed	✓	
Convergence guarantees & stability*	✓	
Sample efficiency		✓
Stochastic policy representation	✓	
Applicability to continuous action spaces	✓	

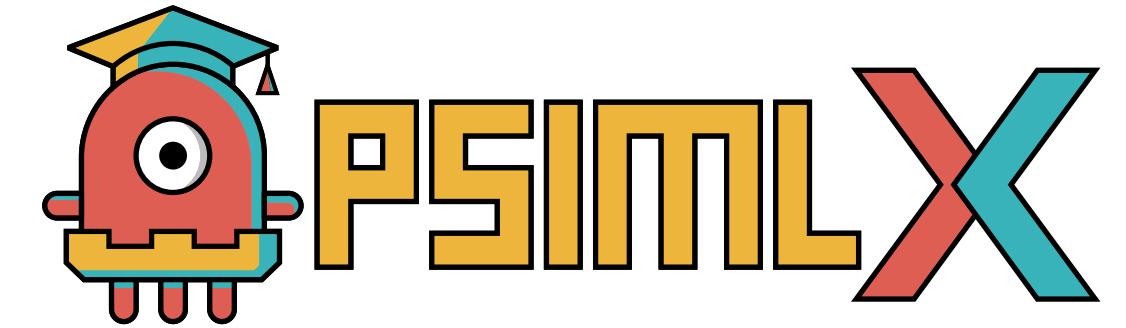
- **The Deadly Triad:** Bootstrapping & Function approximation & Off-policy

Policy Gradient Methods

Conclusions

- Stronger convergence guarantees compared to TD function approximation methods due to the Policy Gradient Theorem
- Naturally applicable on continuous action spaces
- Can represent stochastic policies and approach deterministic policies asymptotically
- Most modern state of the art algorithms belong to either Actor-Critic methods **which combine both Monte Carlo and TD approaches**

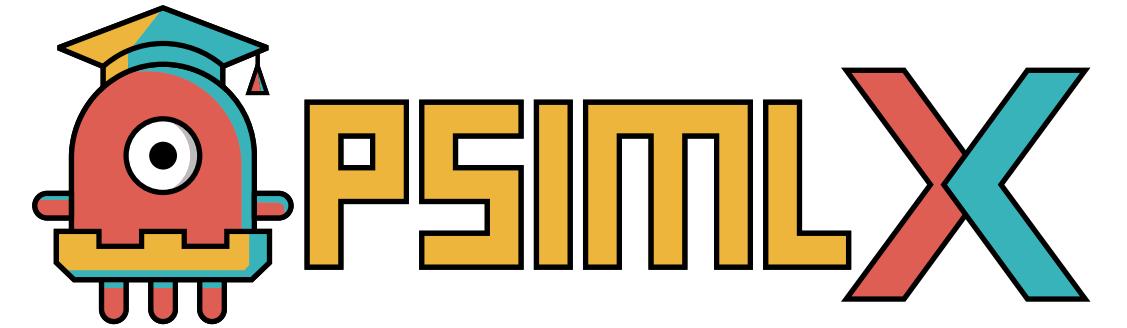




Next Steps

RL Areas and Papers

- State of the art **model-free** algorithms:
 - Proximal Policy Optimisation (Schulman et al, 2017), Soft Actor Critic (Haarnoja et al, 2018)
- **Model-based** approaches:
 - Dyna-Q (Sutton and Barto 2018), Monte Carlo Tree Search (Sutton and Barto, 2018), World Models (Ha and Schmidhuber, 2018)
- **Hierarchical** reinforcement learning:
 - Between MDPs and Semi-MDPs (The options framework) (Sutton et al. 1999)
- **Intrinsically motivated** reinforcement learning:
 - Intrinsic Motivation and Reinforcement Learning (Barto, 2013), Curiosity-driven Exploration by Self-supervised Prediction (Pathak et al, 2017)



(~ˇˇ)~ Thanks! (°՞՞)

